

# Rendu différé

Source cours de Steve Marschner

# Light reflection physics

# Radiometry redux

**Power**

**Intensity** power per unit solid angle

**Irradiance** power per unit area

**Radiance** power per unit (solid angle  $\times$  area)

# Sources of light

## Point sources

- intensity
- can be directionally varying—spotlights

## Area sources

- radiance
- can be spatially varying

## Directional sources

- irradiance (normal irradiance)

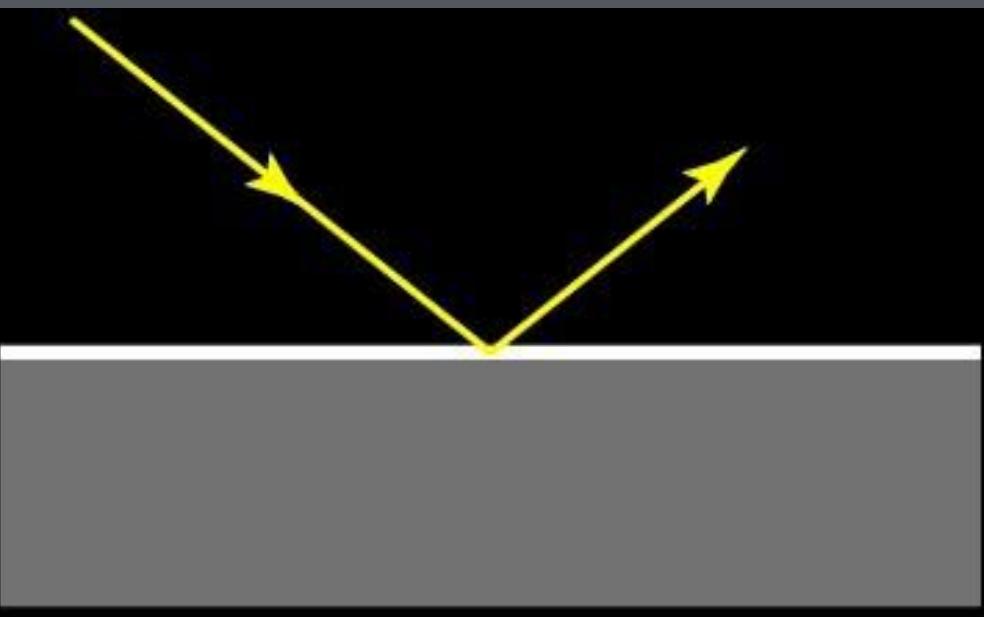
## Environment lighting

- radiance (usually spatially varying)
- sun-sky models

# Simple kinds of scattering

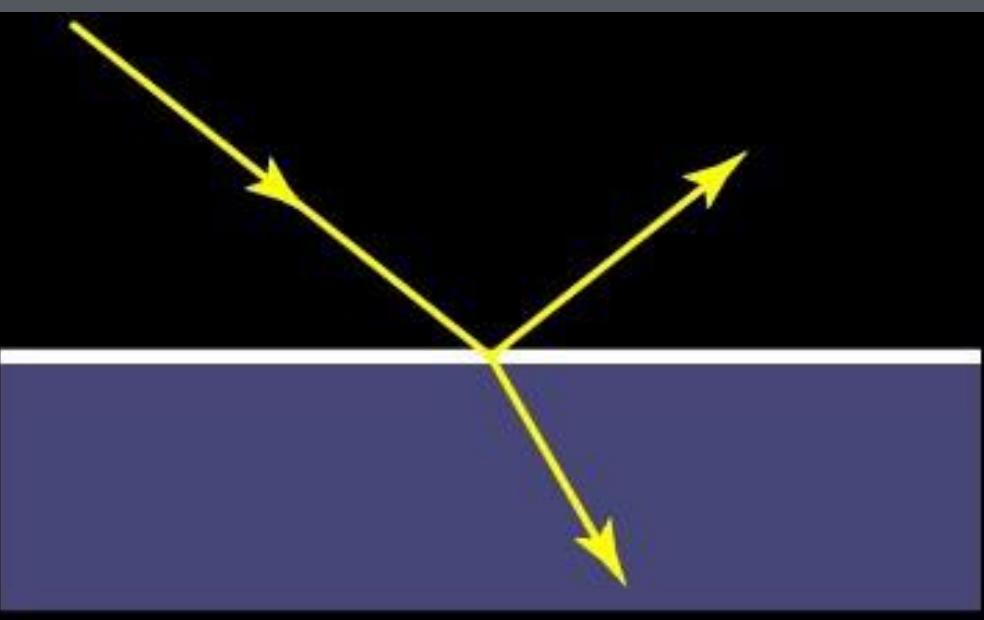
## Ideal specular reflection

- incoming ray reflected to a single direction
- mirror-like behavior
- arises at smooth surfaces



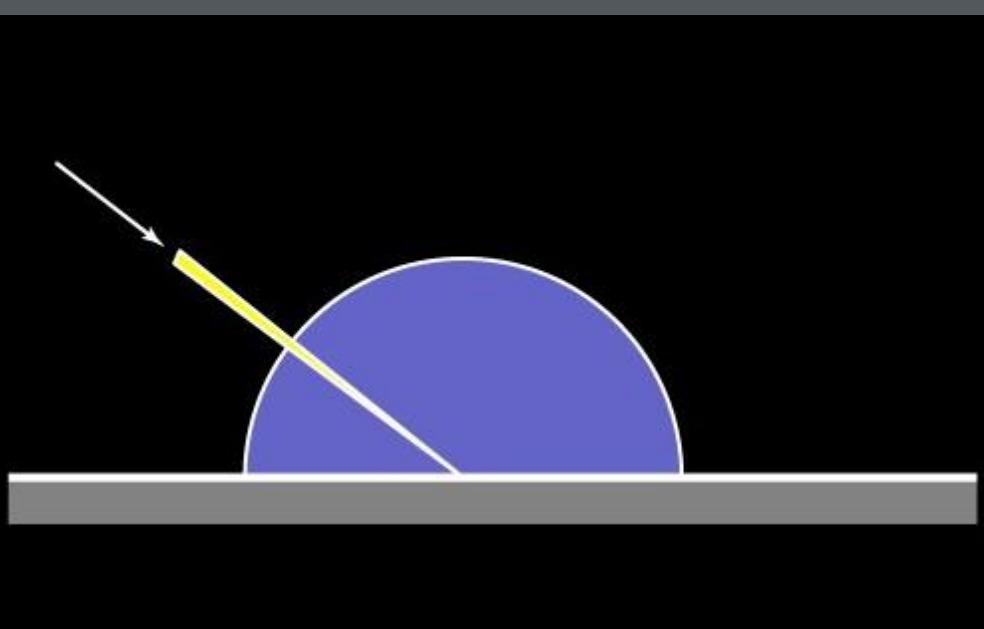
## Ideal specular transmission

- incoming ray refracted to a single direction
- glass-like behavior
- arises at smooth dielectric (nonmetal) surfaces

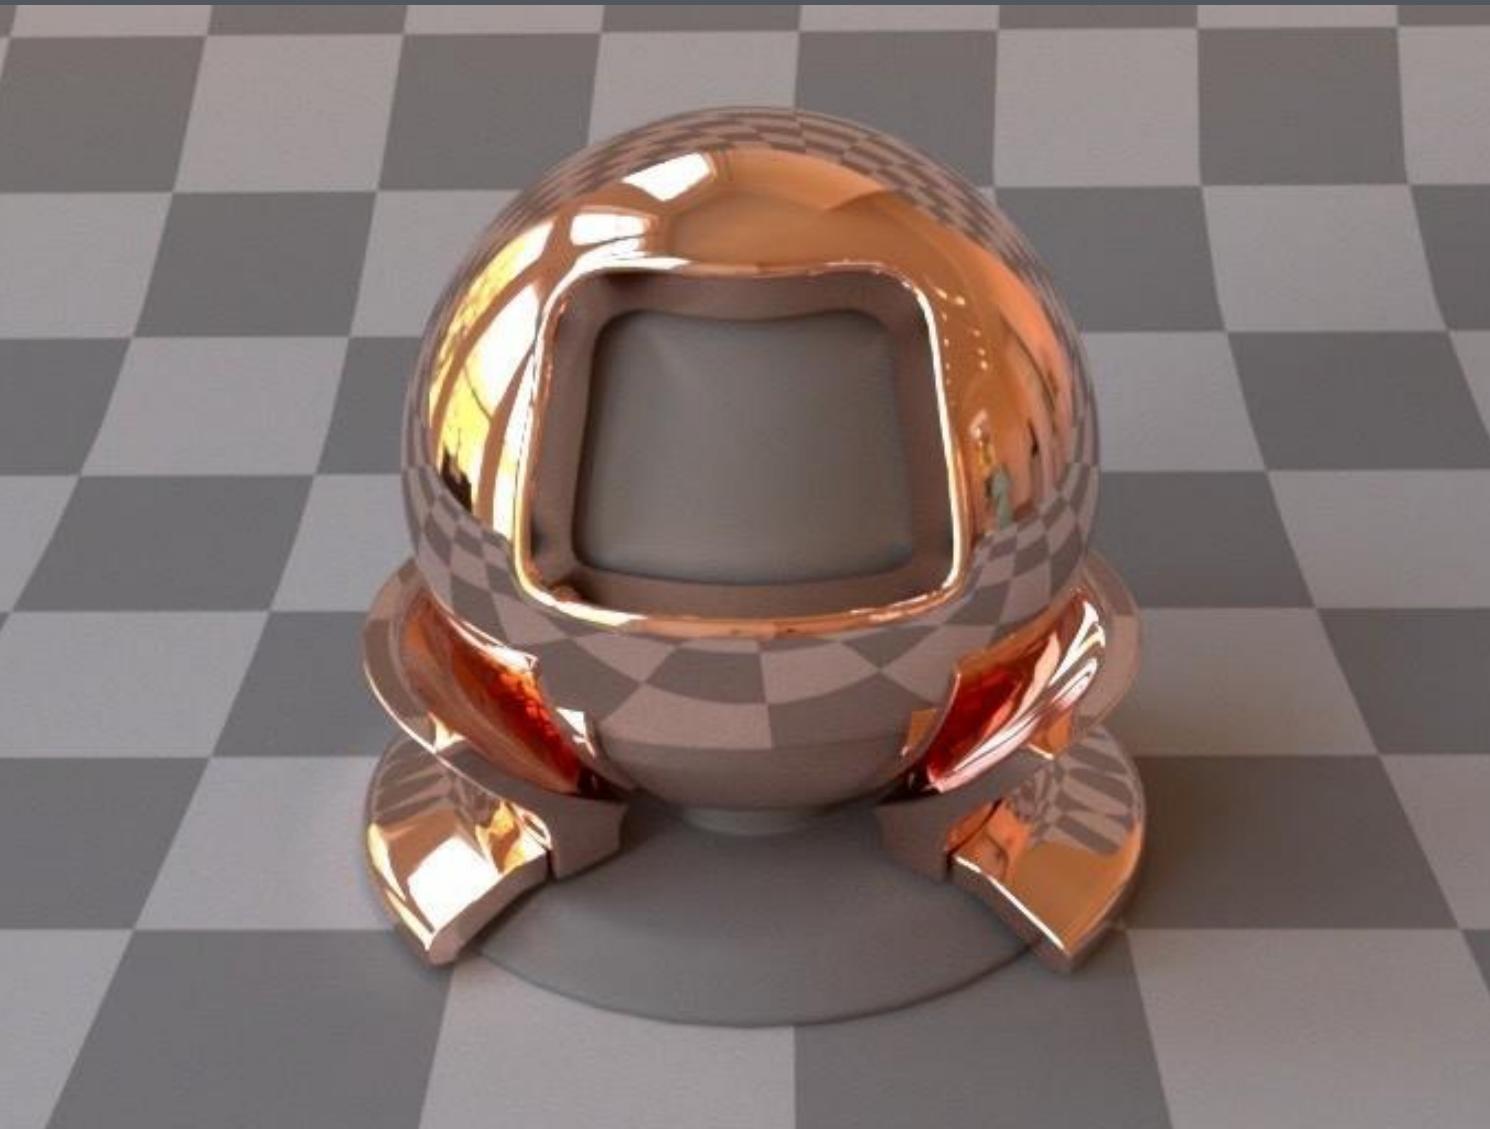


## Ideal diffuse reflection or transmission

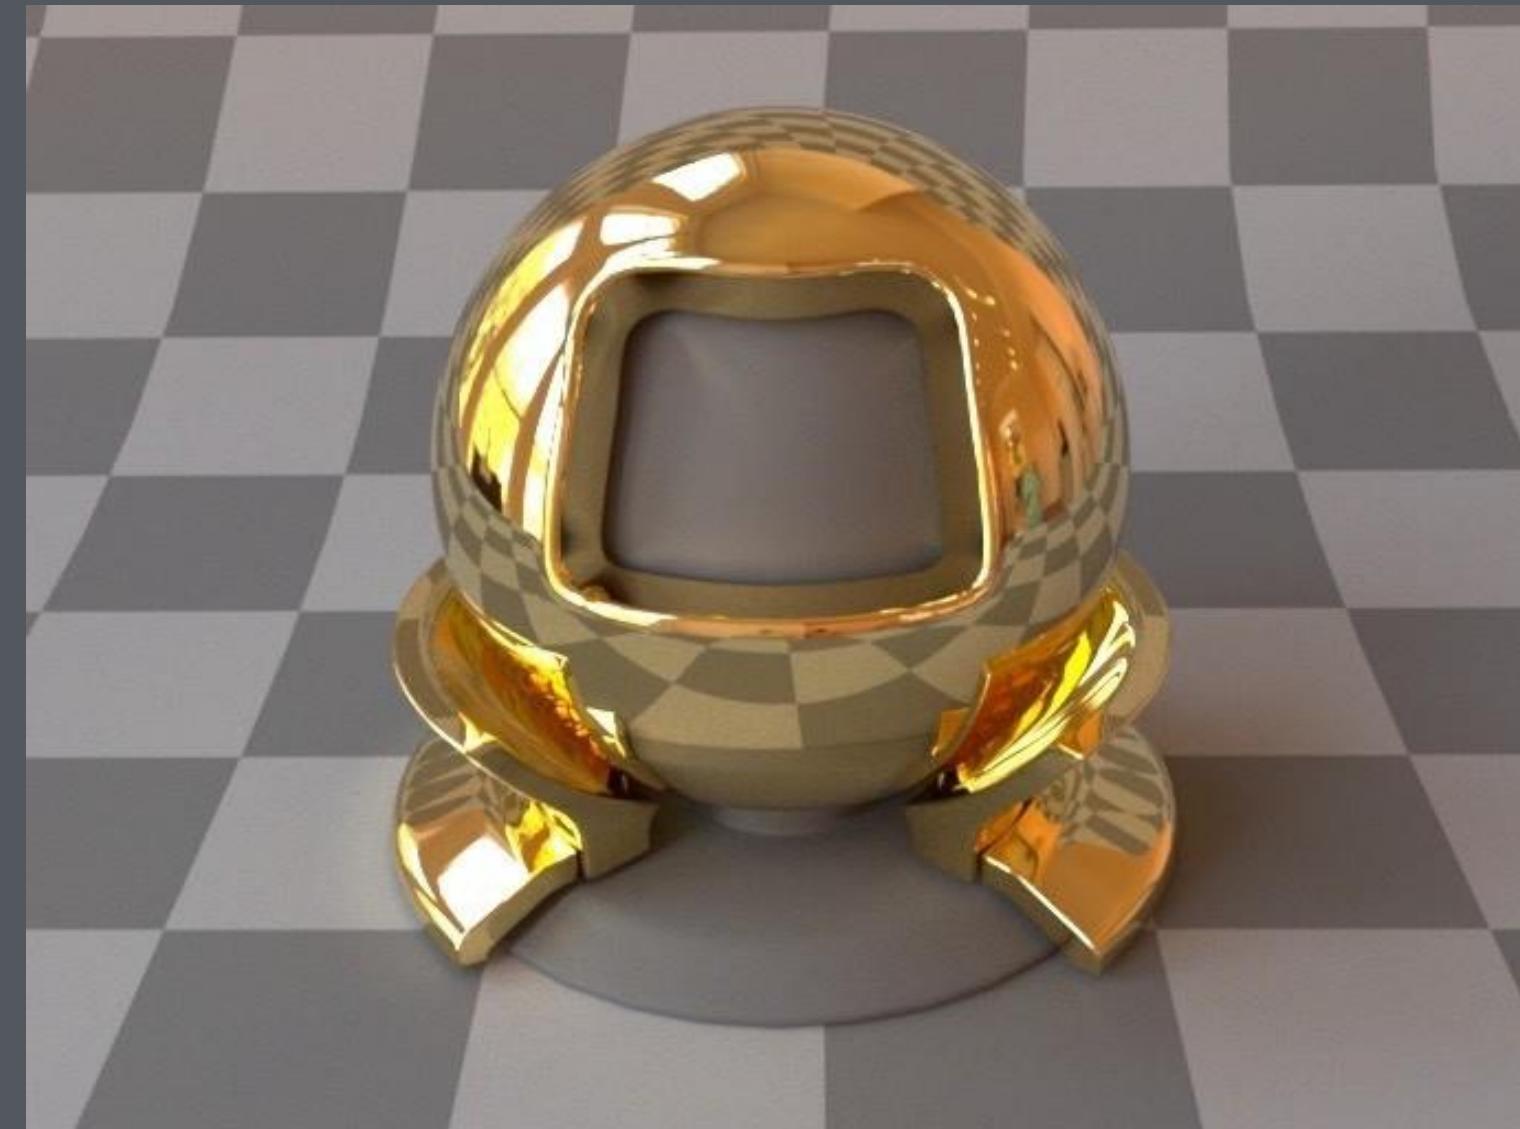
- outgoing radiance independent of direction
- arises from subsurface multiple scattering



# Ideal specular reflection from metals



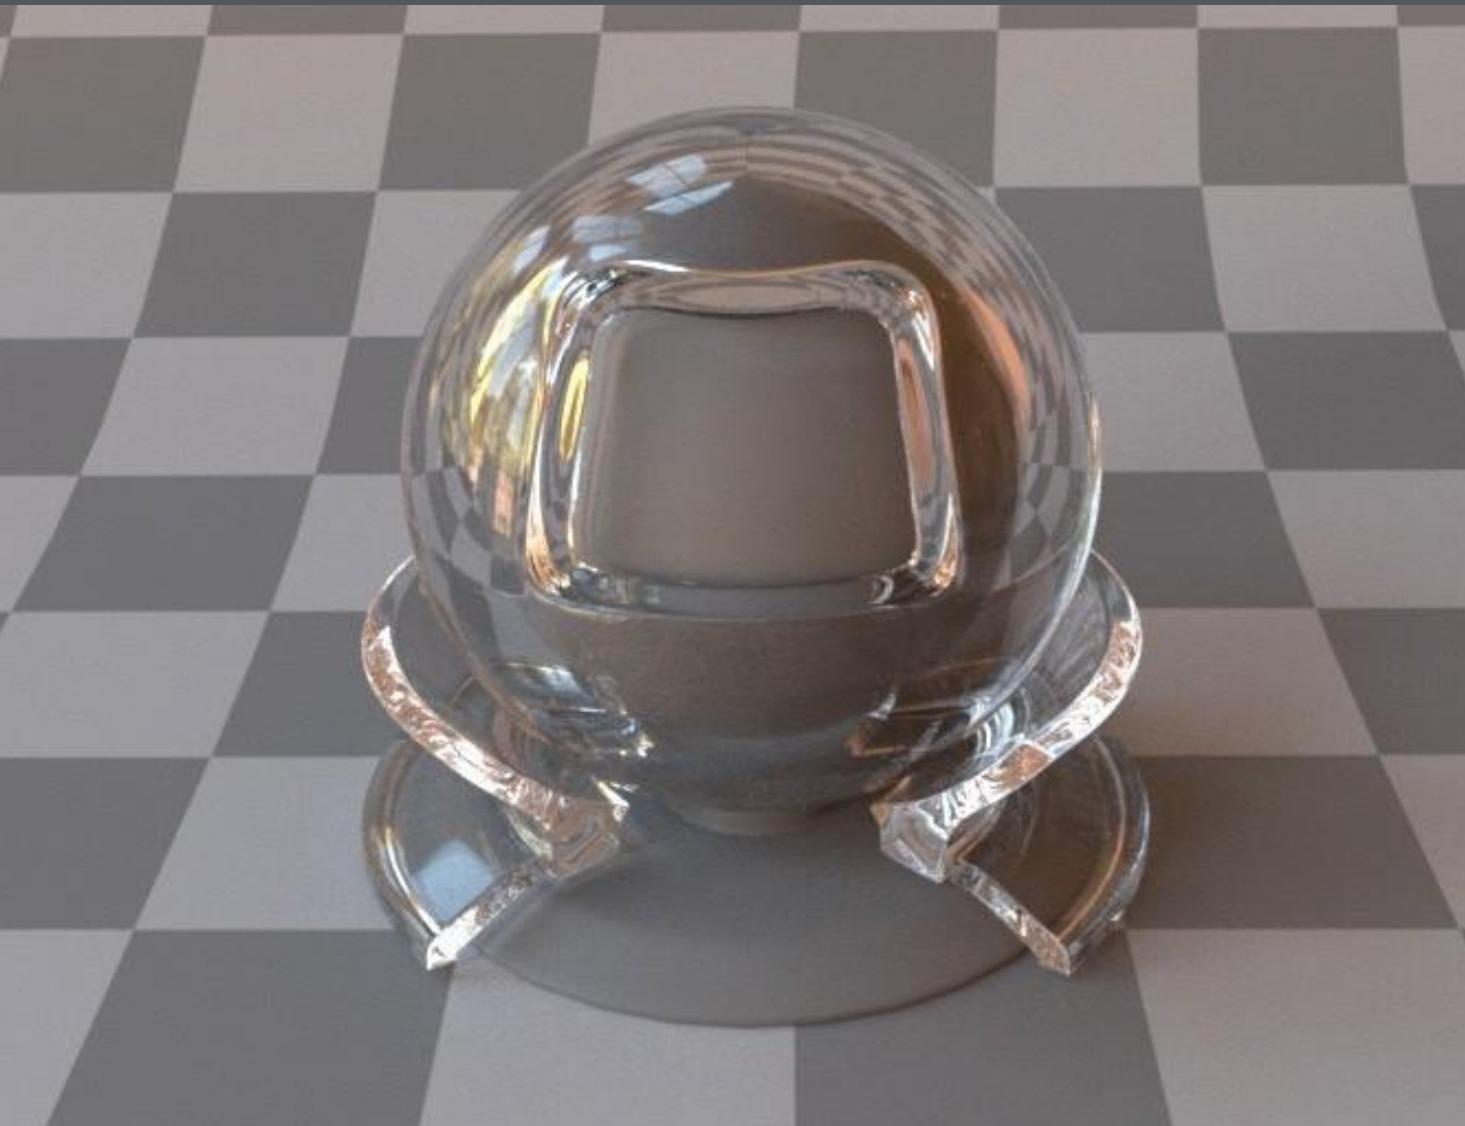
Cu



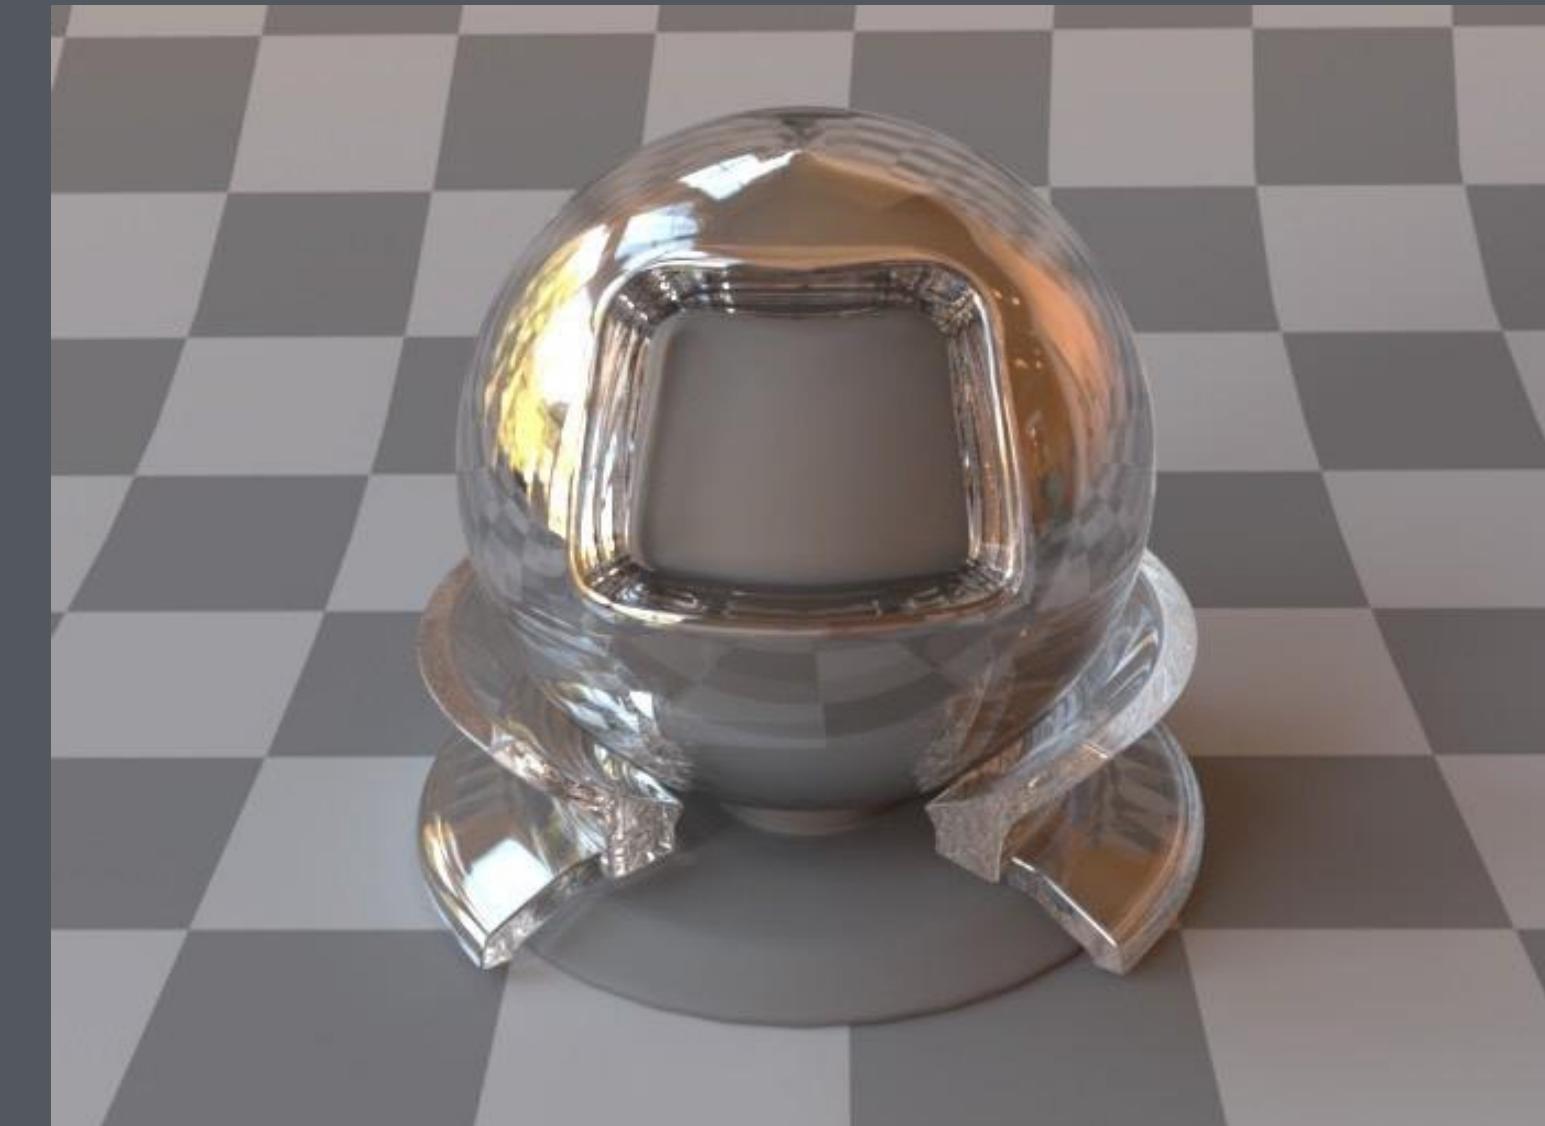
Au

Wenzel Jakob / Mistuba

# Ideal reflection and transmission from smooth dielectrics



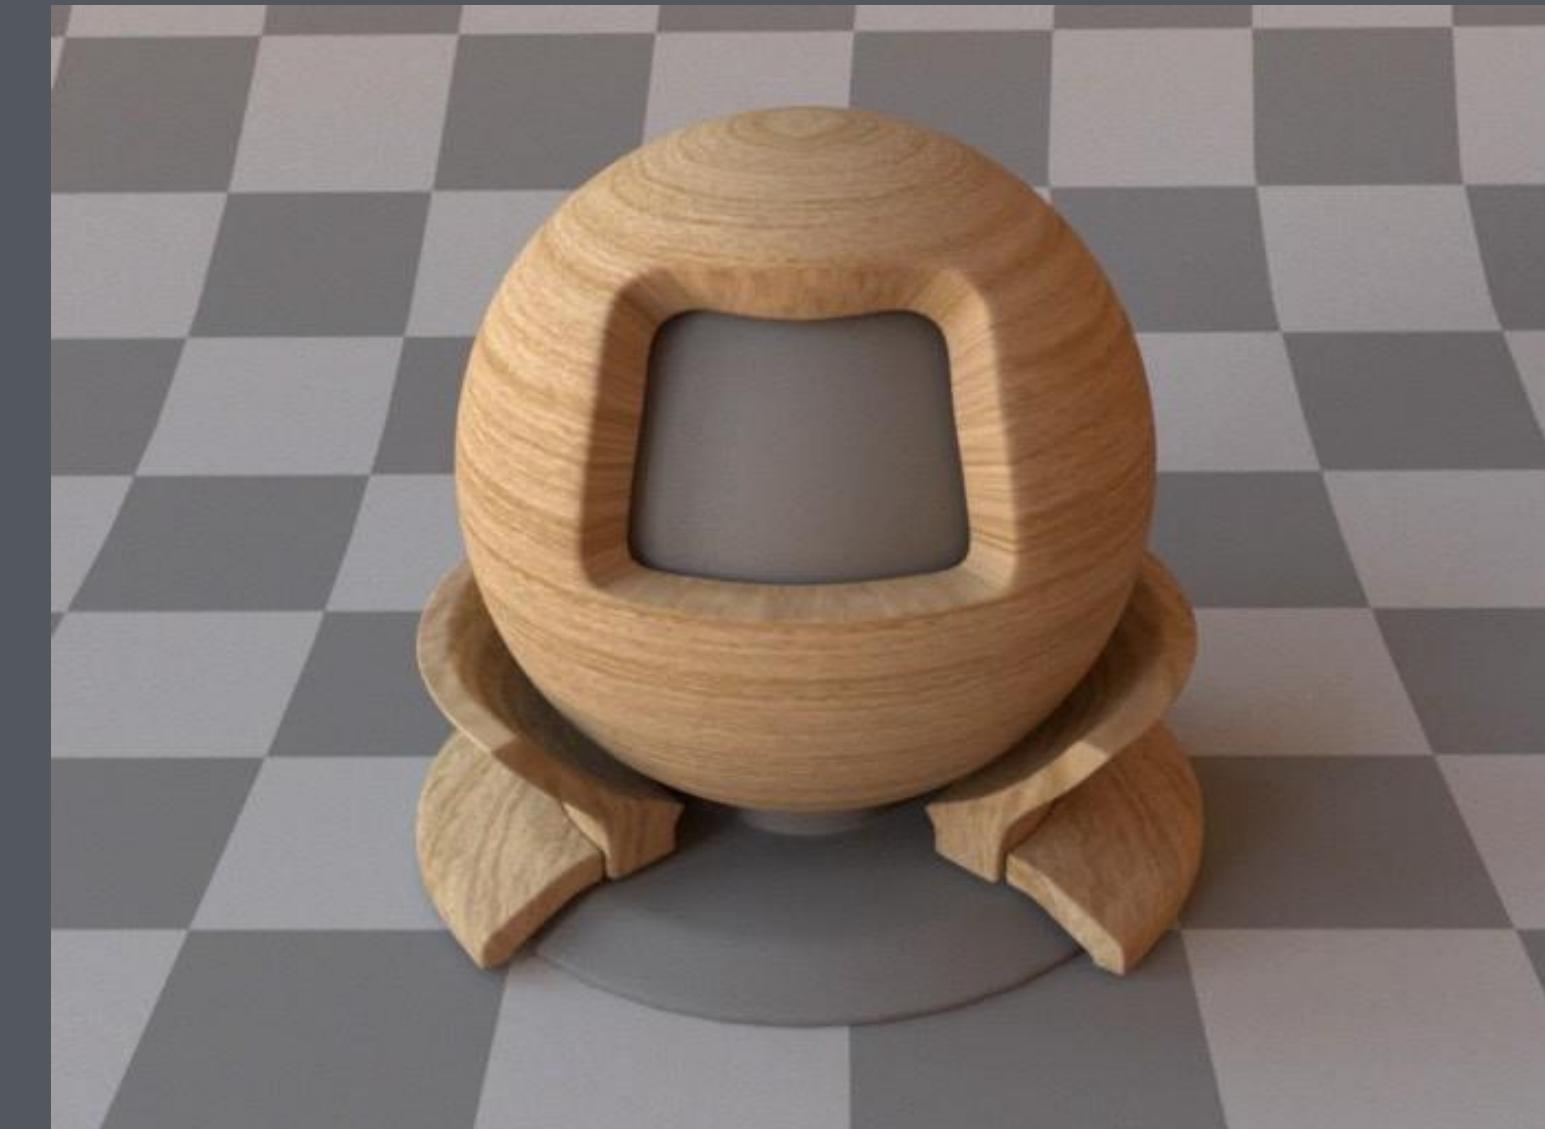
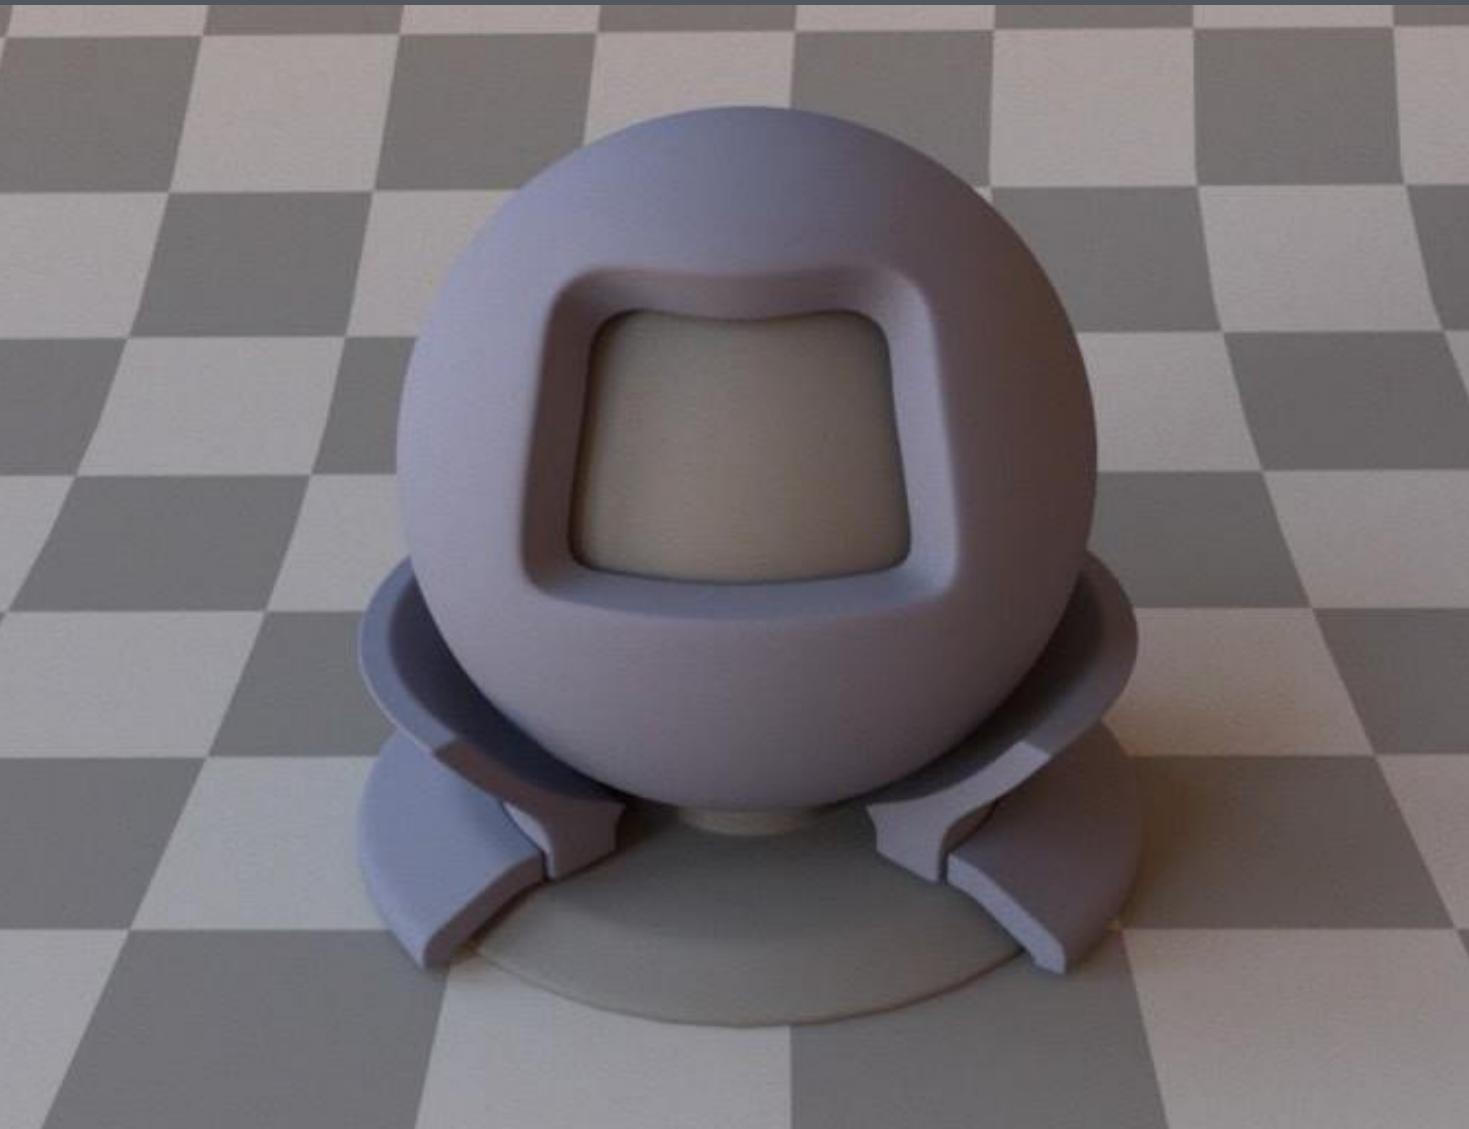
Water (ior = 1.33)



Diamond (ior = 2.4)

Wenzel Jakob / Mistuba

# Two diffuse surfaces

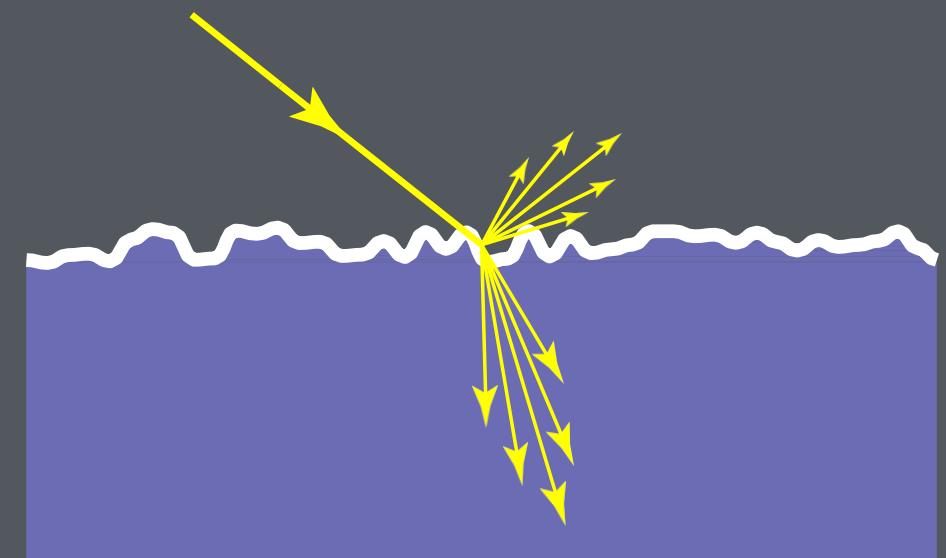
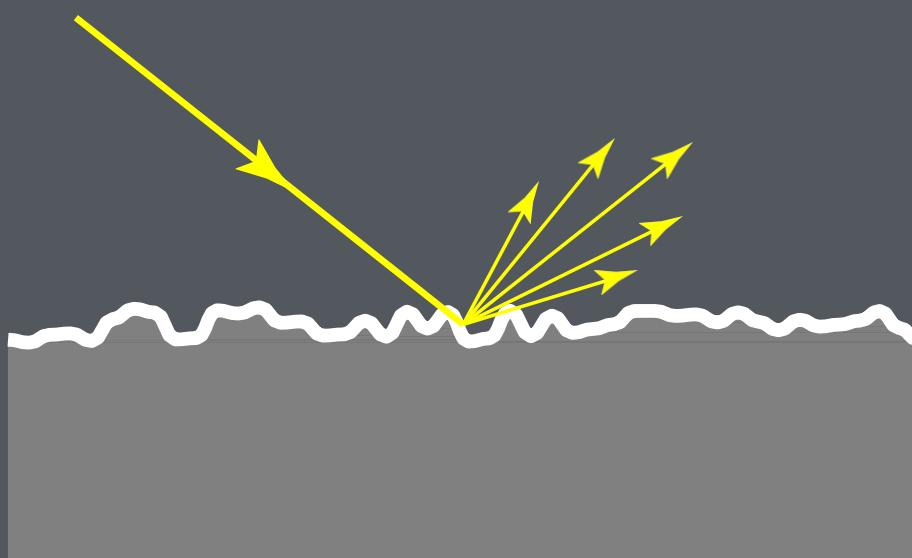


Wenzel Jakob / Mistuba

# More complex scattering

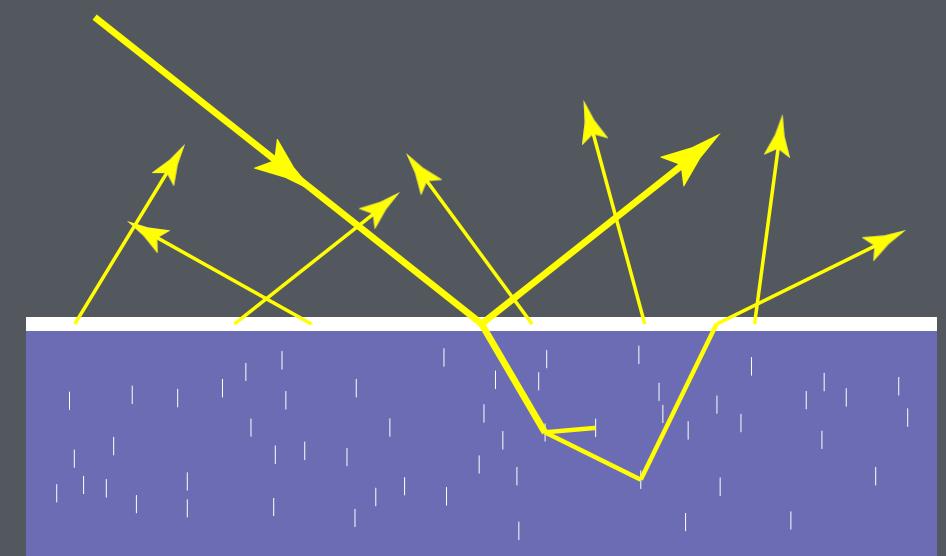
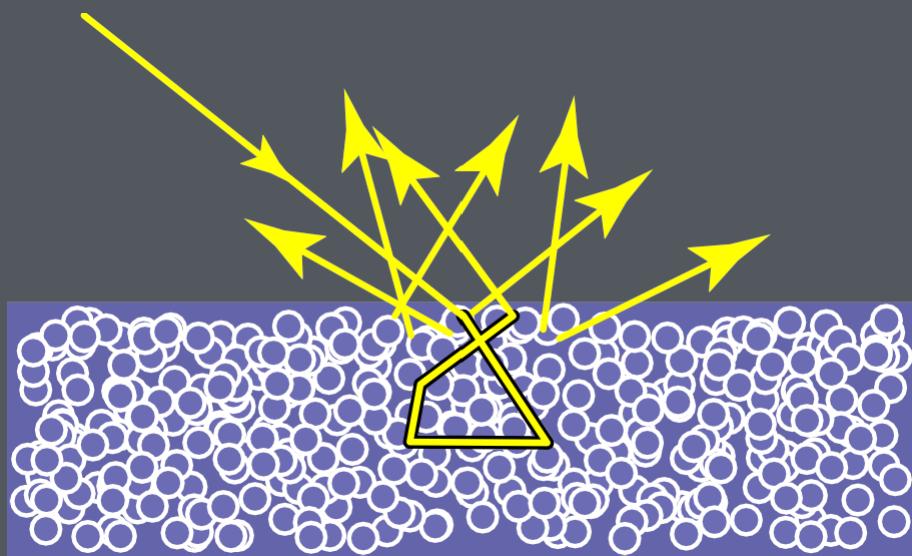
## Rough interfaces

- metal interfaces: blurred reflection
- dielectric interfaces: blurred transmission

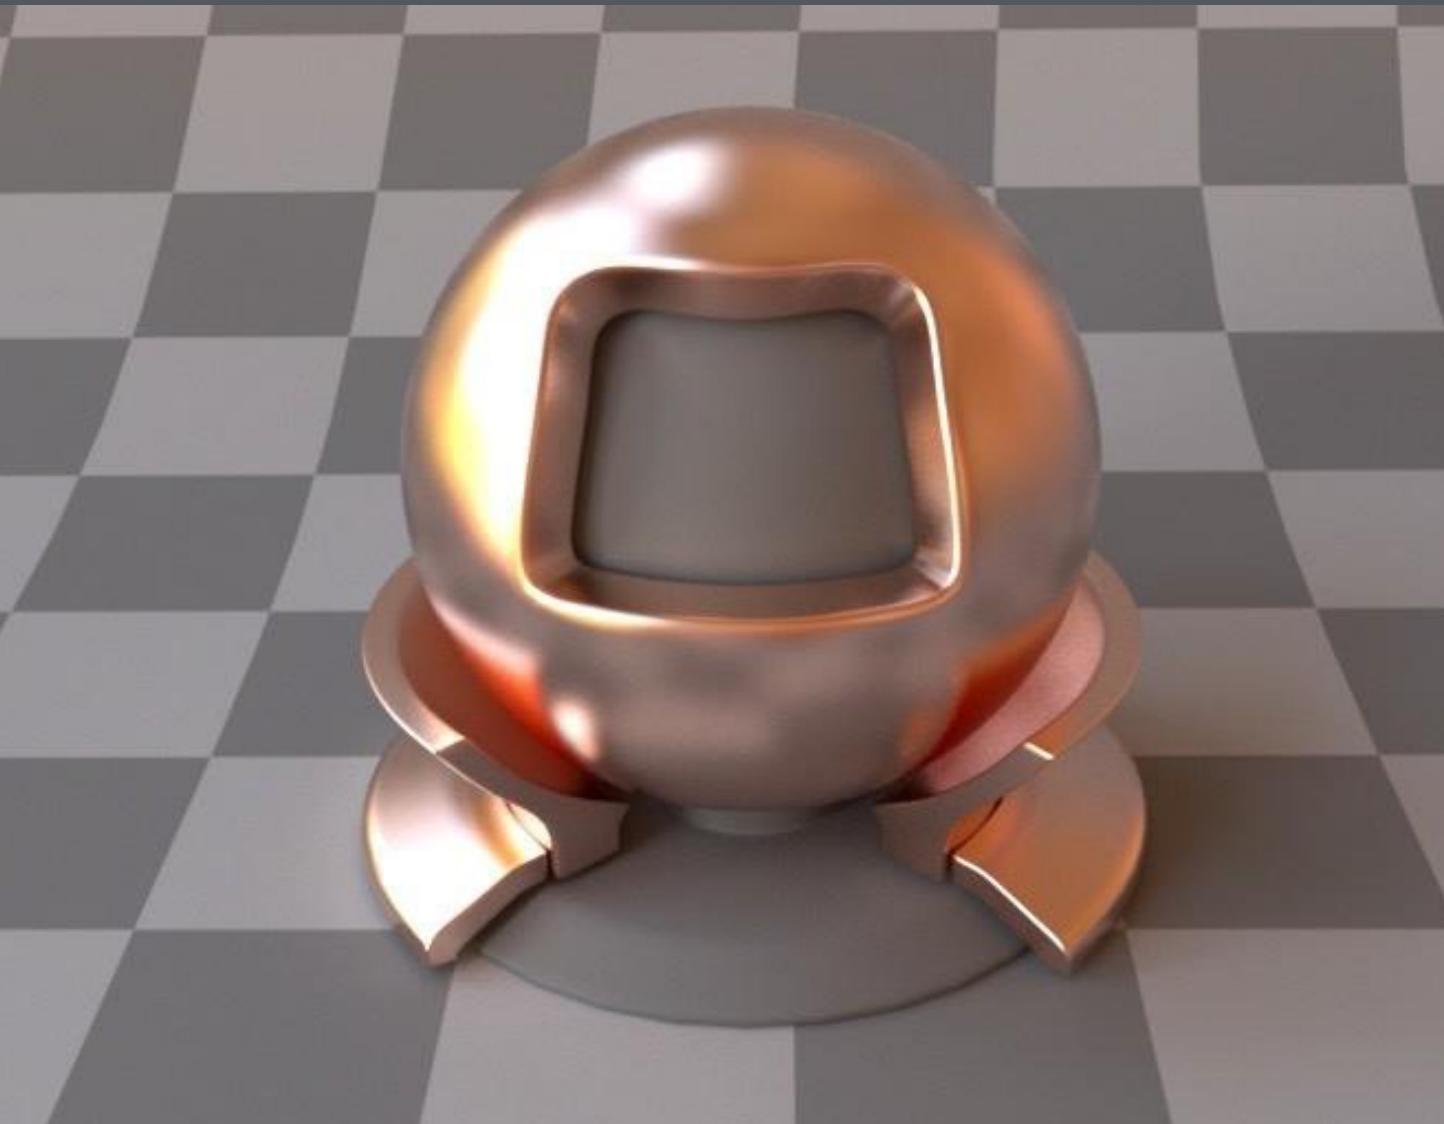


## Subsurface scattering

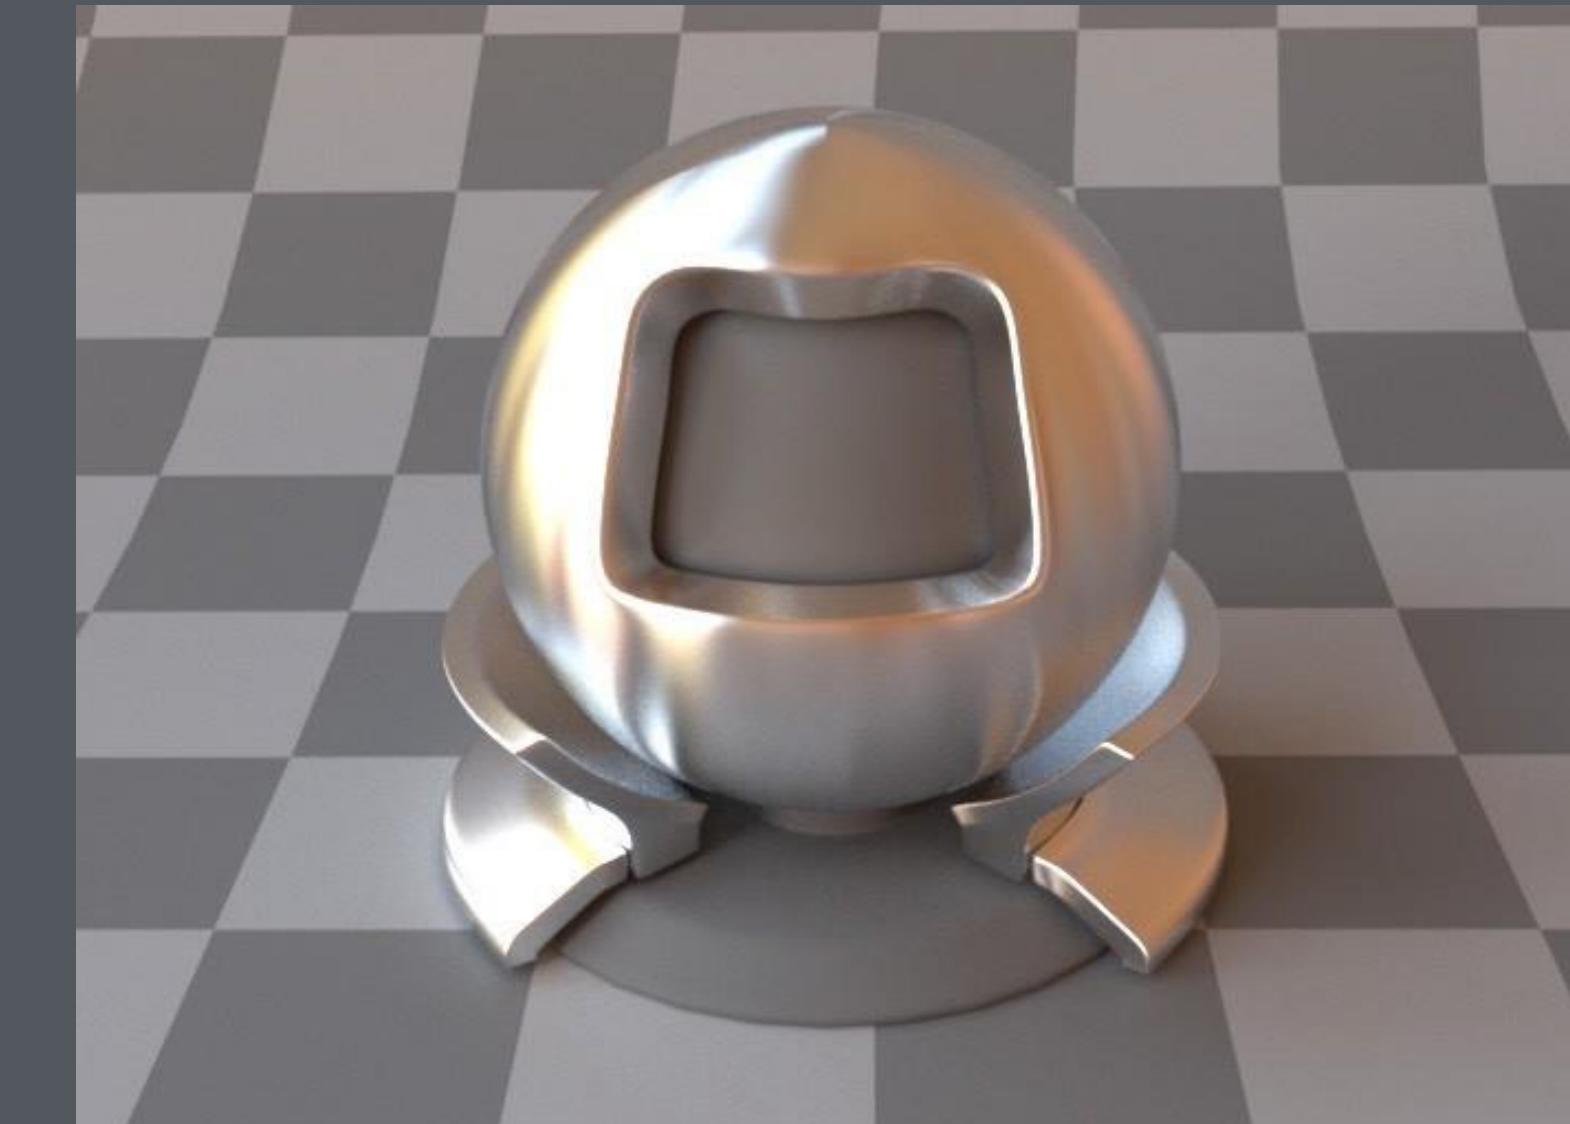
- liquids—milk, juice, beer, ...
- coatings—paint, glaze, varnish, ...
- natural materials—wood, marble, ...
- biological materials—skin, plants, ...
- low optical density leads to *translucency*



# Reflection from rough metal interfaces



Cu ( $\alpha = 0.1$ )



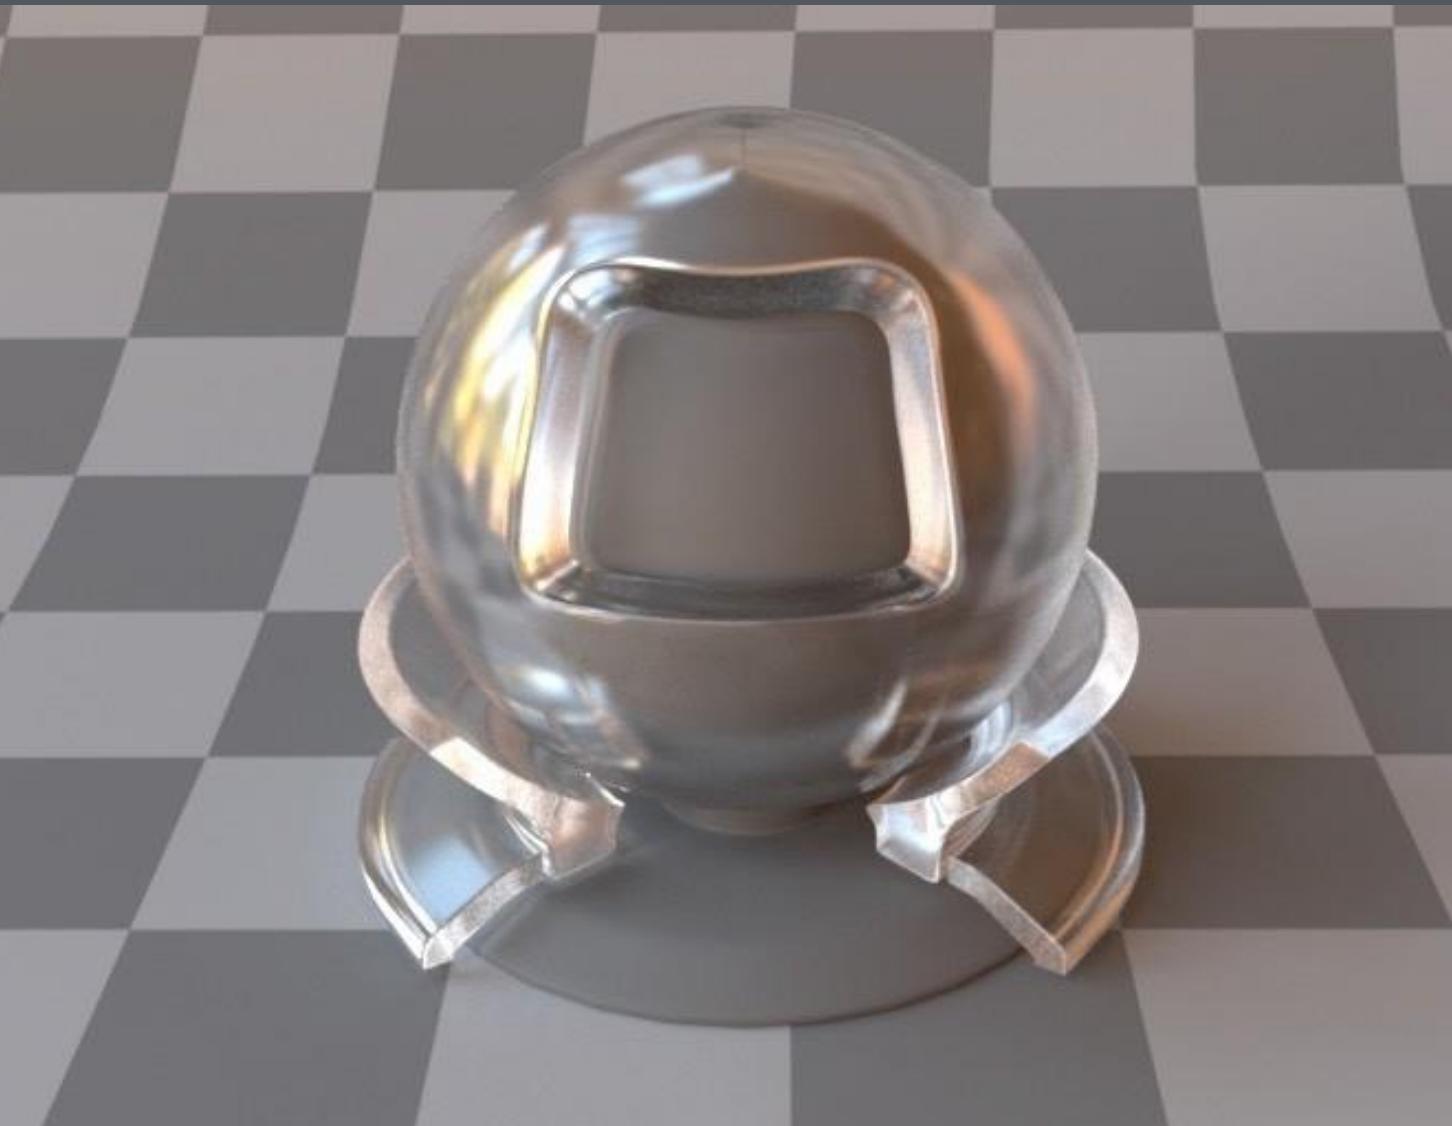
Al (anisotropic)

Wenzel Jakob / Mistuba

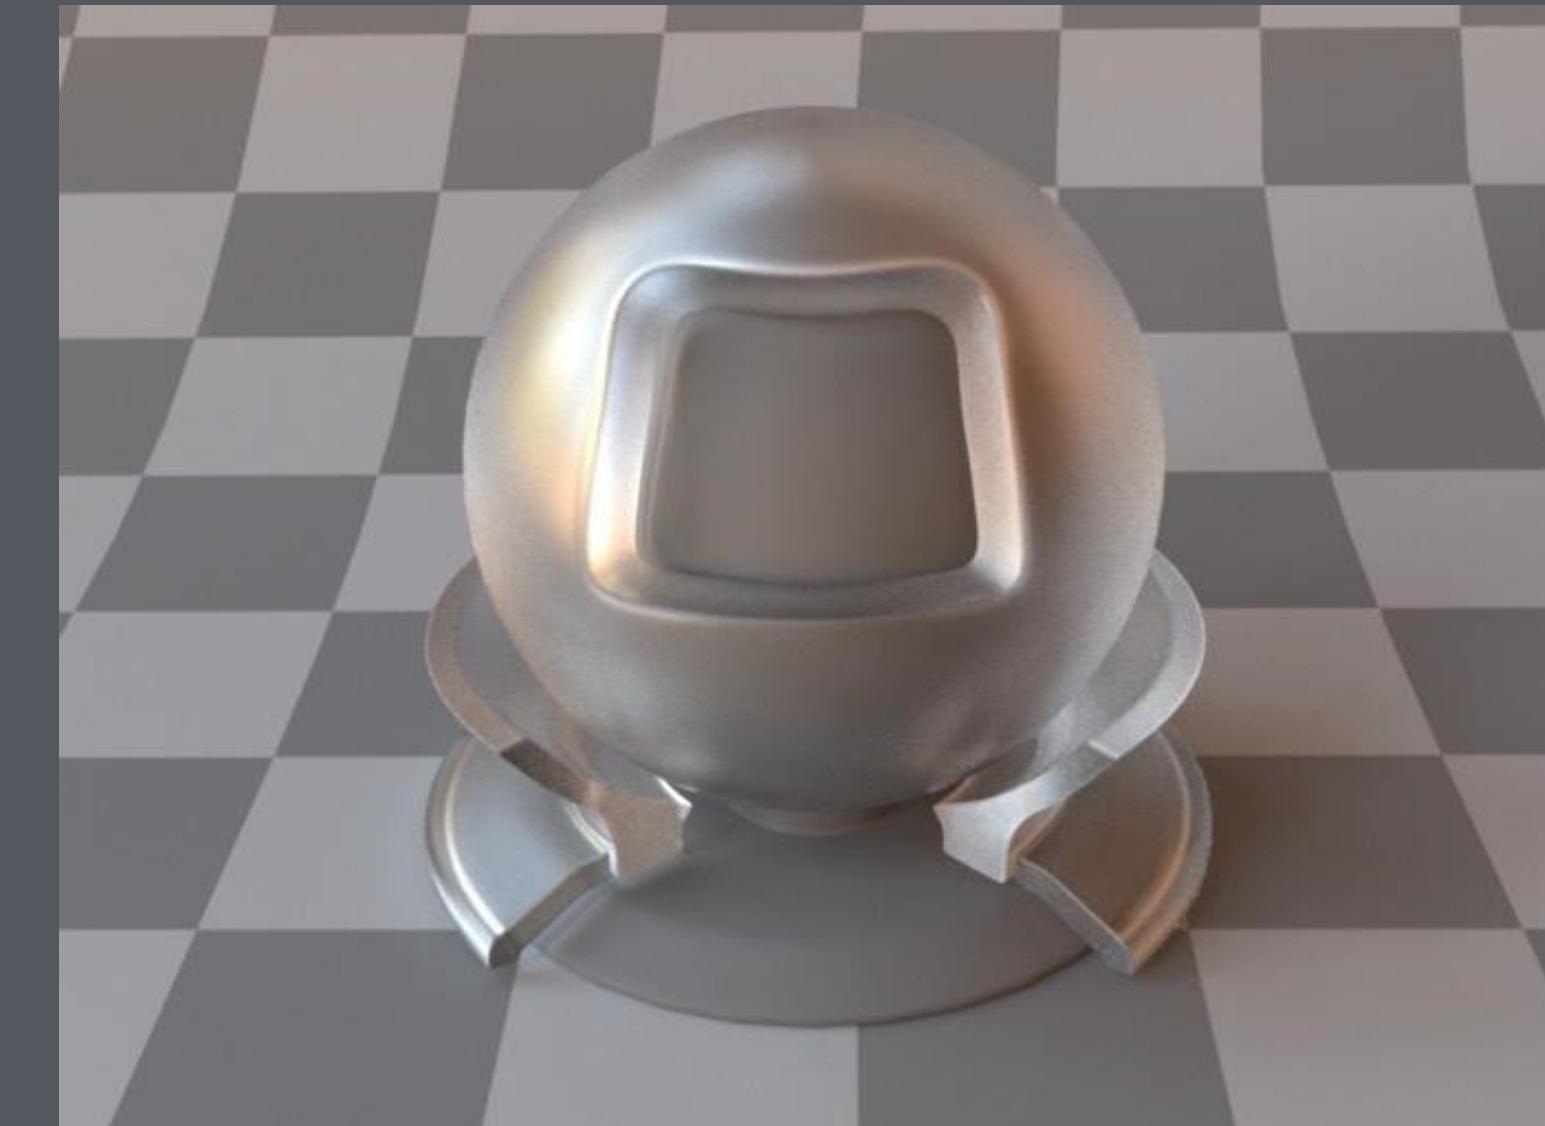




# Reflection and refraction at rough dielectric interfaces



Anti-glare glass ( $\alpha = 0.02$ )



Etched glass ( $\alpha = 0.1$ )

Wenzel Jakob / Mistuba

# Translucent materials



“skim milk”

Wenzel Jakob / Mistuba

Wenzel Jakob / Mistuba



low  
optical  
density



high  
optical  
density

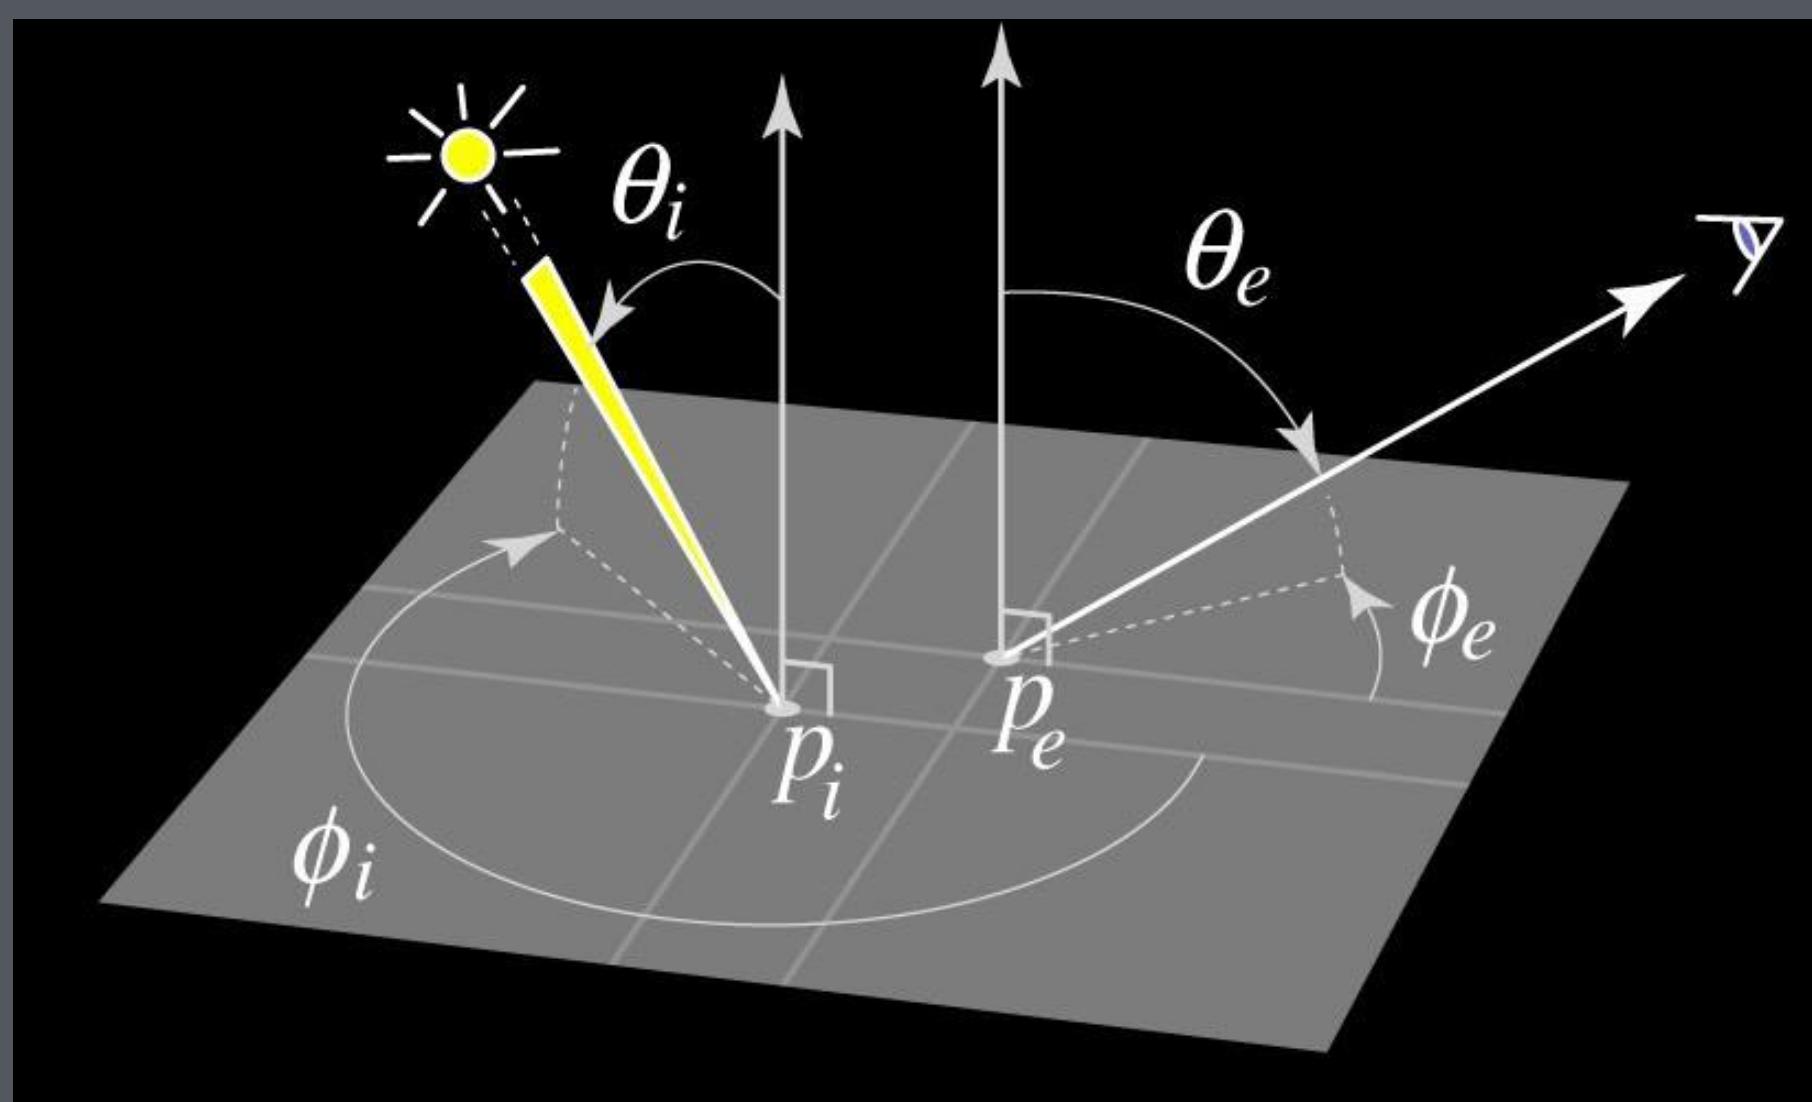
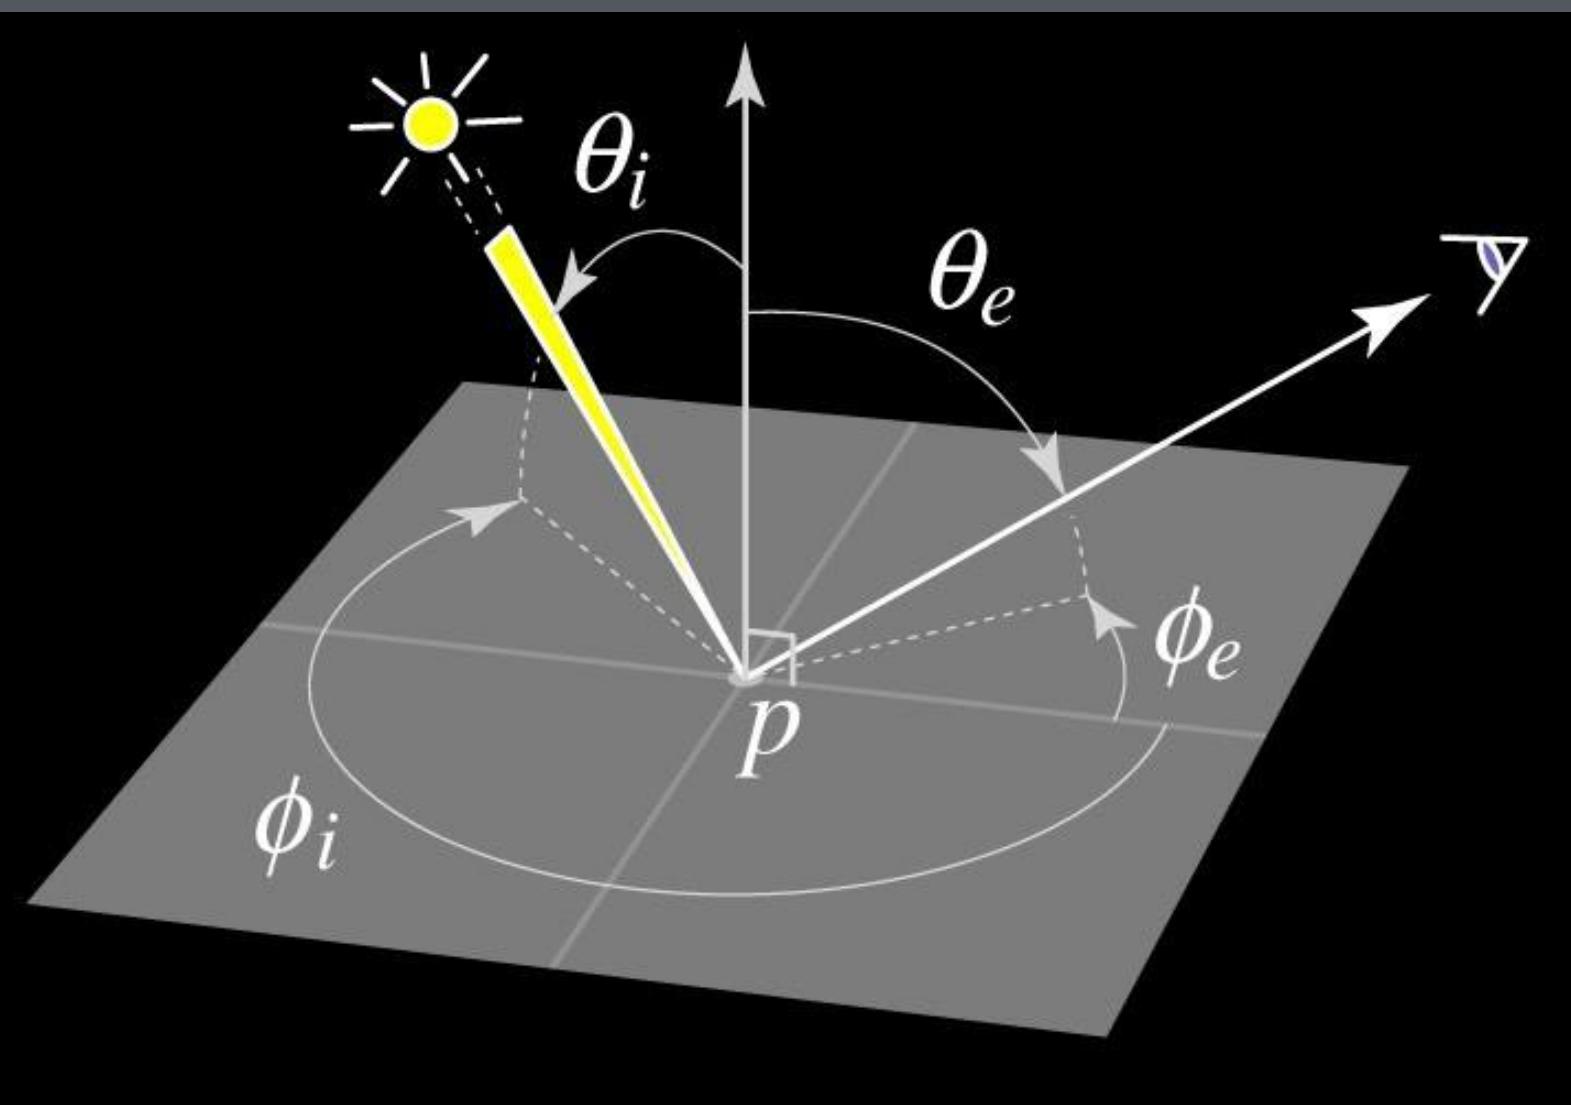
# Modeling complex scattering

## Opaque materials

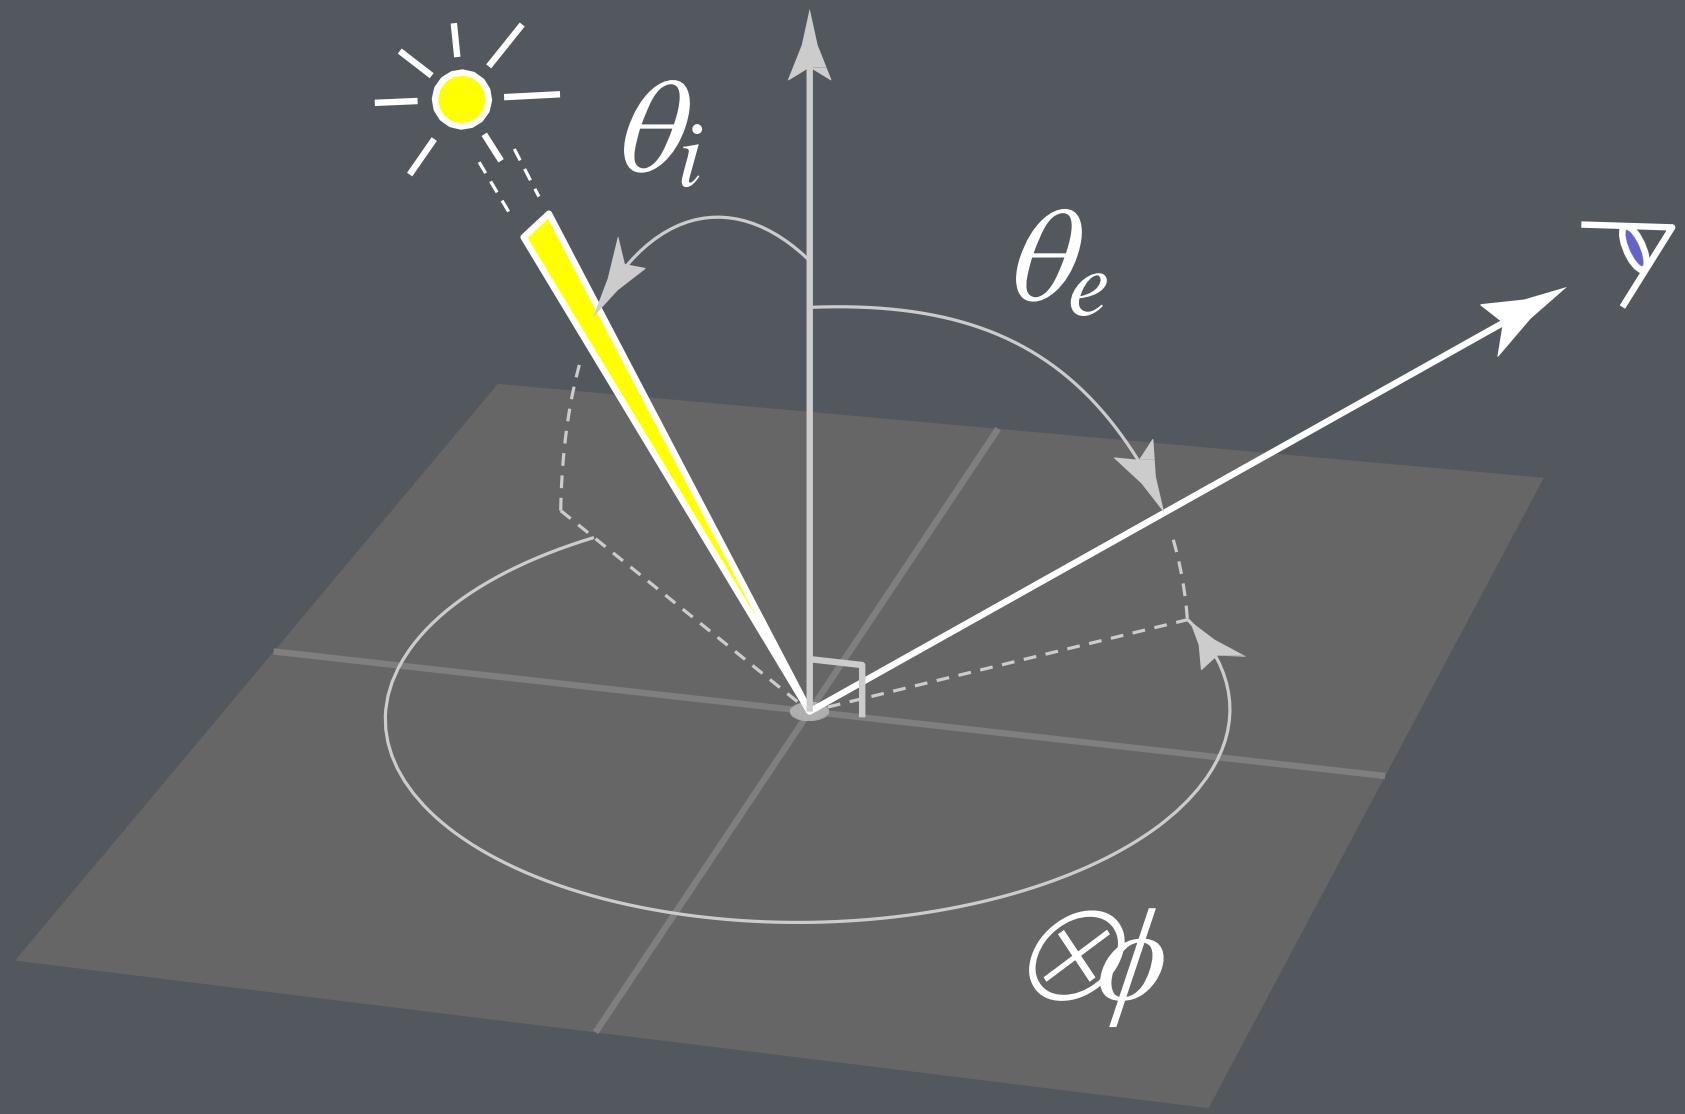
- reflection: bidirectional reflectance distribution function (BRDF)
- transmission: bidirectional transmittance distribution function (BTDF)
- both: bidirectional scattering distribution function (BSDF)

## Translucent materials

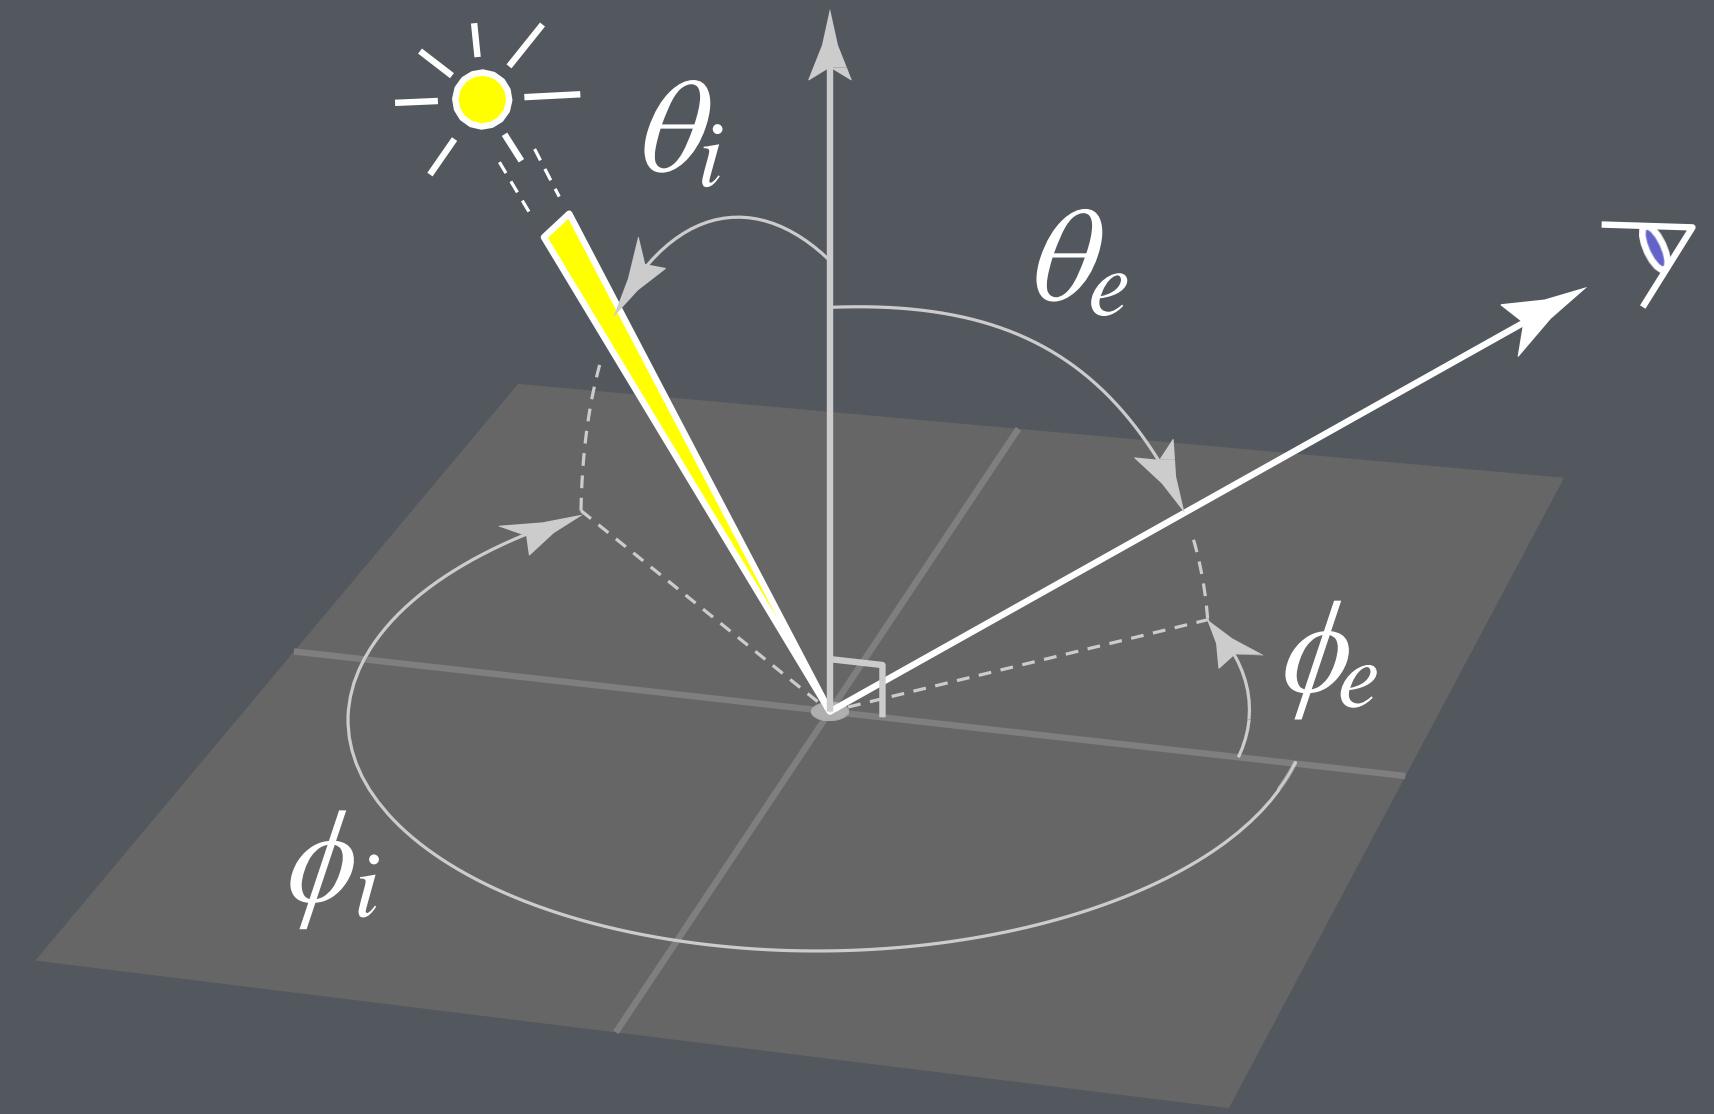
- bidirectional subsurface scattering reflectance distribution function (BSSRDF)



# Isotropy vs. anisotropy



isotropic

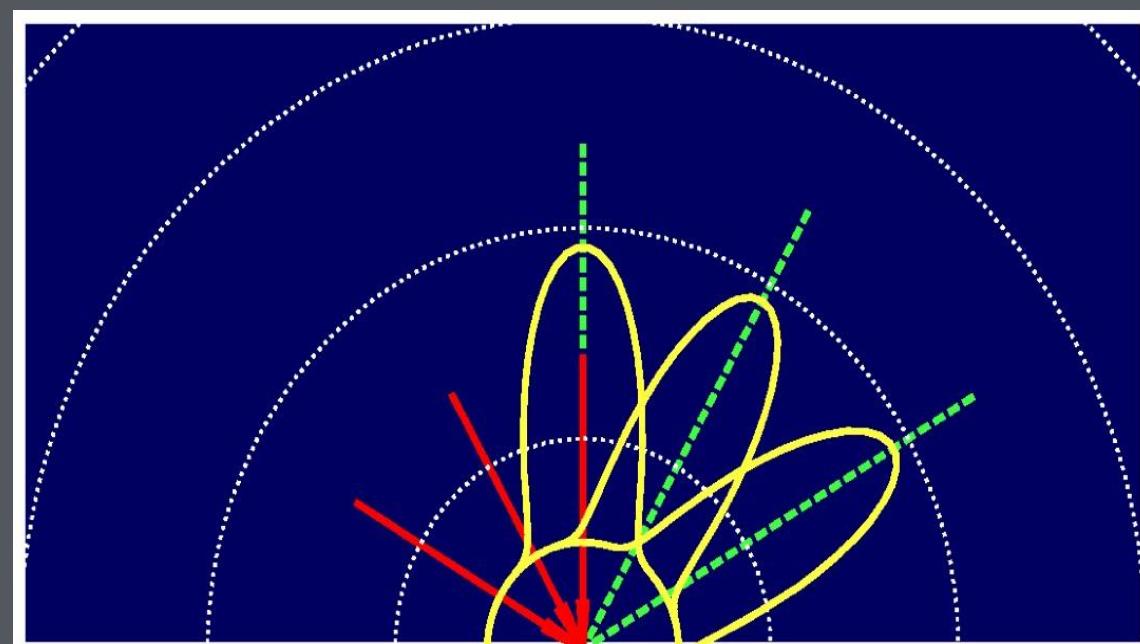


anisotropic

# Types of BRDF/BSDF models

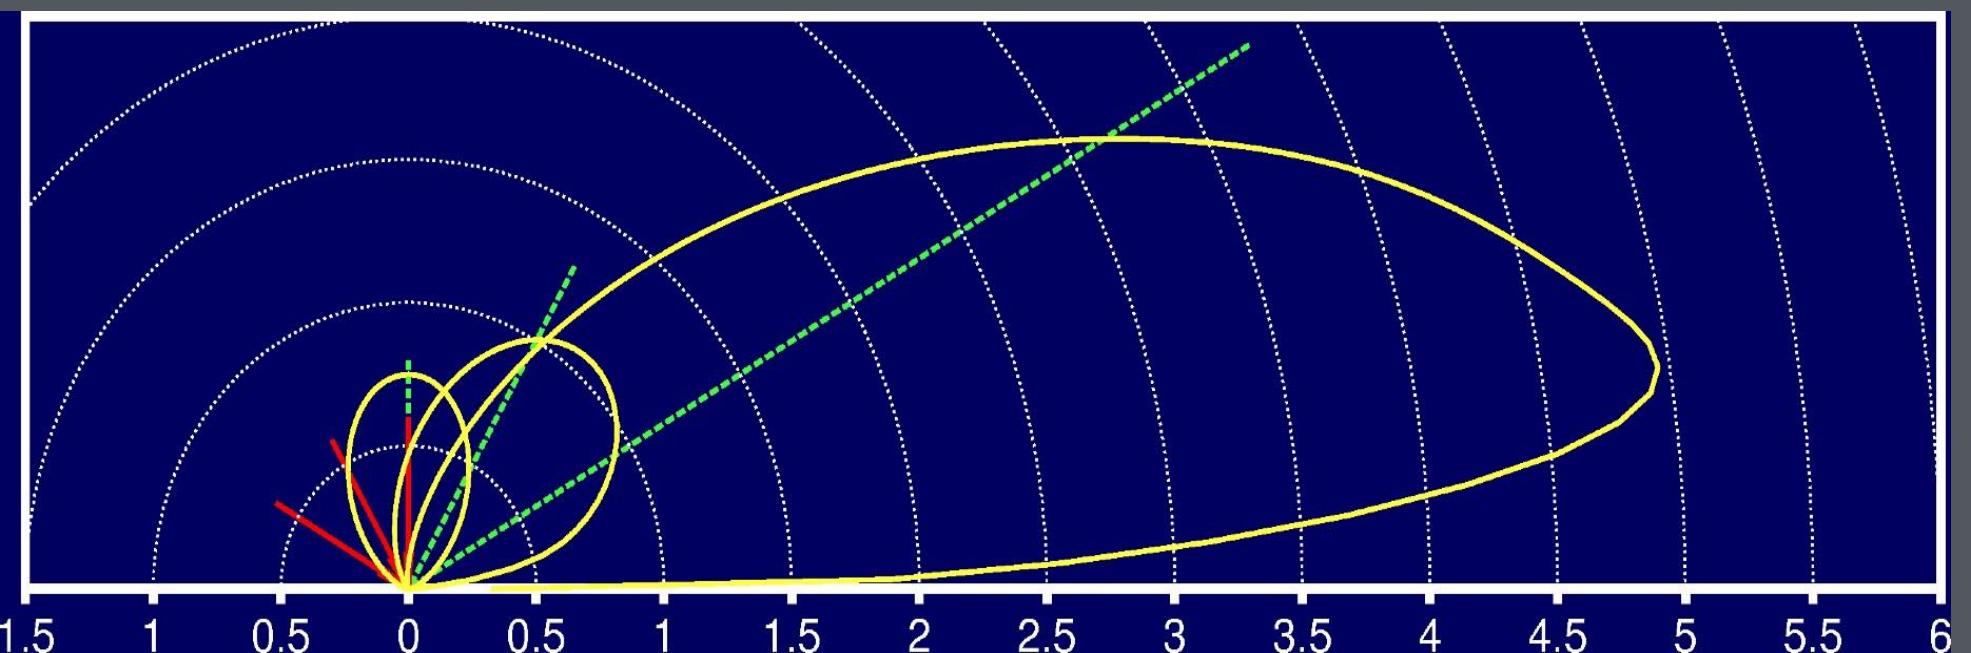
## Ad hoc formulas

- e.g. Blinn-Phong



## Physics-based analytical models

- Lambertian
- Microfacet-based models
- Kirchhoff-based models



## Measured data

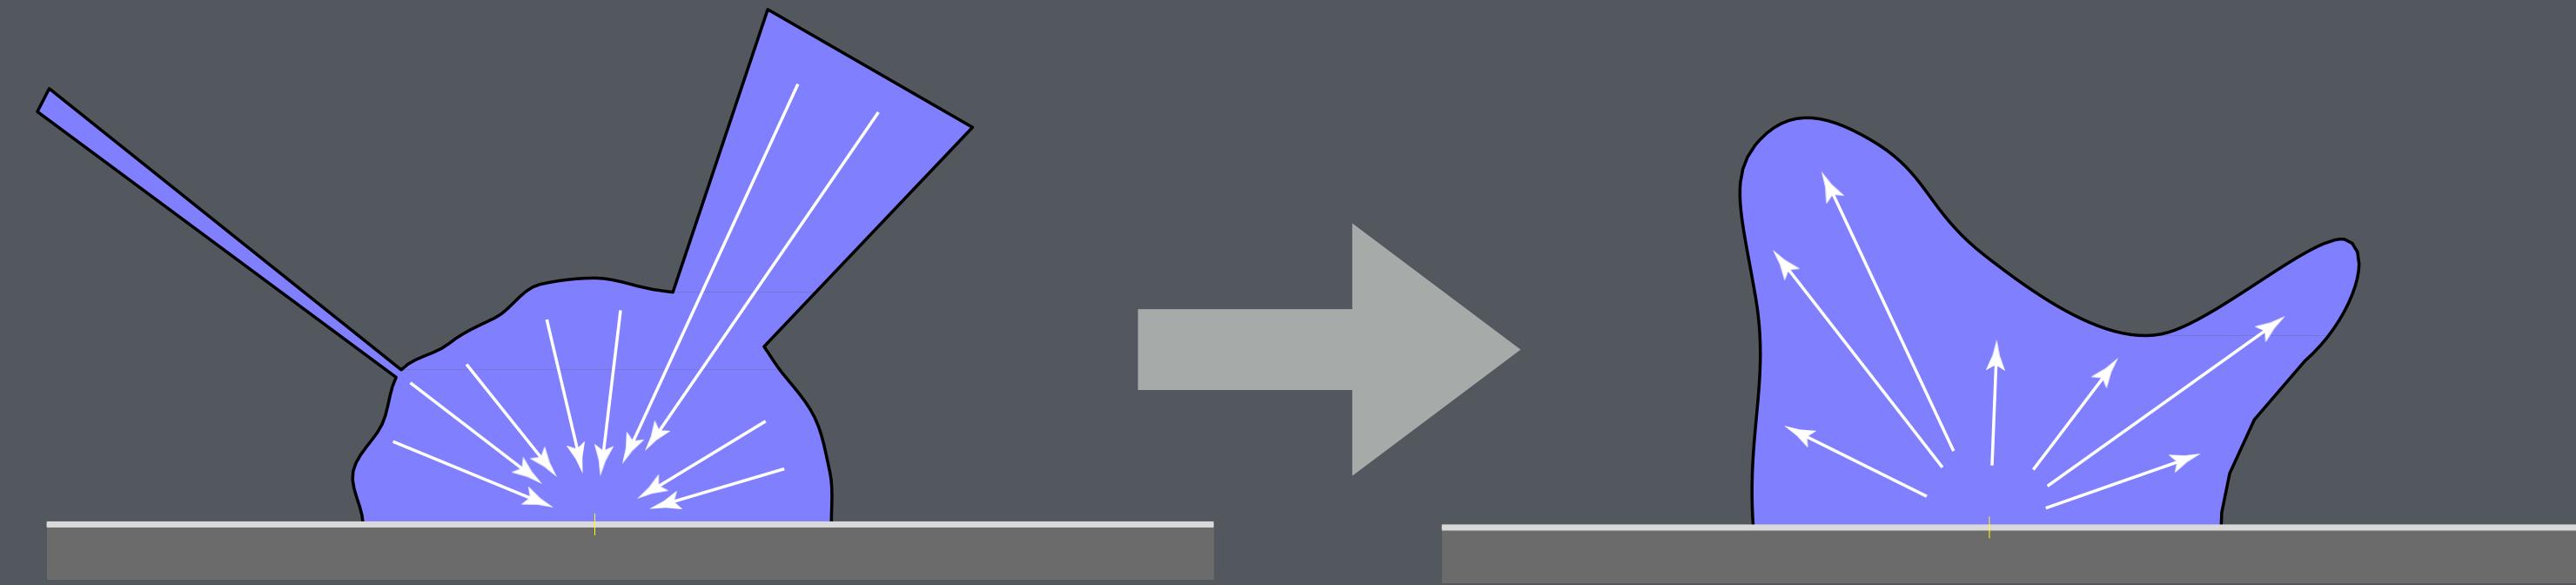
- tables of data from pointwise BRDF measurements
- image-based BRDF measurements

# Light reflection in shaders

# Light reflection: full picture

**all types of reflection reflect all types of illumination**

- diffuse, glossy, mirror reflection
- environment, area, point illumination



incident distribution  
(function of direction)

reflected distribution  
(function of direction)

# Categories of illumination

	diffuse	glossy	mirror
indirect	soft indirect illumination	blurry reflections of other objects	reflected images of other objects
environment	soft shadows	blurry reflection of environment	reflected image of environment
area	soft shadows	shaped specular highlight	reflected image of source
point/directional	hard shadows	simple specular highlight	point reflections



= easy to compute using standard shaders

# How to compute shading

**Basic case: point or directional lights; diffuse or glossy BRDF**

**Type in BRDF model, plug in illumination and view direction**

- can write down model in world space, use world-space vectors
  - can write down model in surface frame, transform vectors
    - really not different

**Subtleties are all about what frame to use for shading**

# Interpolated shading

**Coarse triangle meshes are fast**

**Discontinuities are bad**

**Therefore: interpolate geometric quantities across triangles**

- goal: shading is smooth across edges

**What do we interpolate?**

- what do we need to compute shading?

# Shading frames

**When we carry around a normal, we are defining a tangent plane**

- interpolated normal defines an approximate, smoothly varying tangent plane

**For some purposes, the tangent plane is enough**

- e.g. computing shading for isotropic BRDFs
- any coordinate system conforming to the normal is equally good

**In other cases, need a complete frame**

- whenever directions within the plane are inequivalent
- e.g. anisotropic BRDFs
- e.g. tangent-frame normal maps

# What to interpolate

**Need plane: can just interpolate a normal**

**Need frame: interpolate enough data to define a tangent frame**

**One and a half vectors rounds up to two**

- normal and one tangent vector
- two tangent vectors

**Rebuilding a frame from the vectors**

- worry about handedness matching texture coordinates (or not)
- orthonormality gets broken by interpolation (when does that matter?)

# What you need for shading

## When/why you need full frames

- when you care (or not) what the orientation is
- when you care (or not) about orthonormality

## What to interpolate

- underlying math question: representation of frames
- representations that behave well under interpolation

## How to author orientation

- with maps
- by following a parameterization

## How to deal with corner cases

# Rendering: forward shading

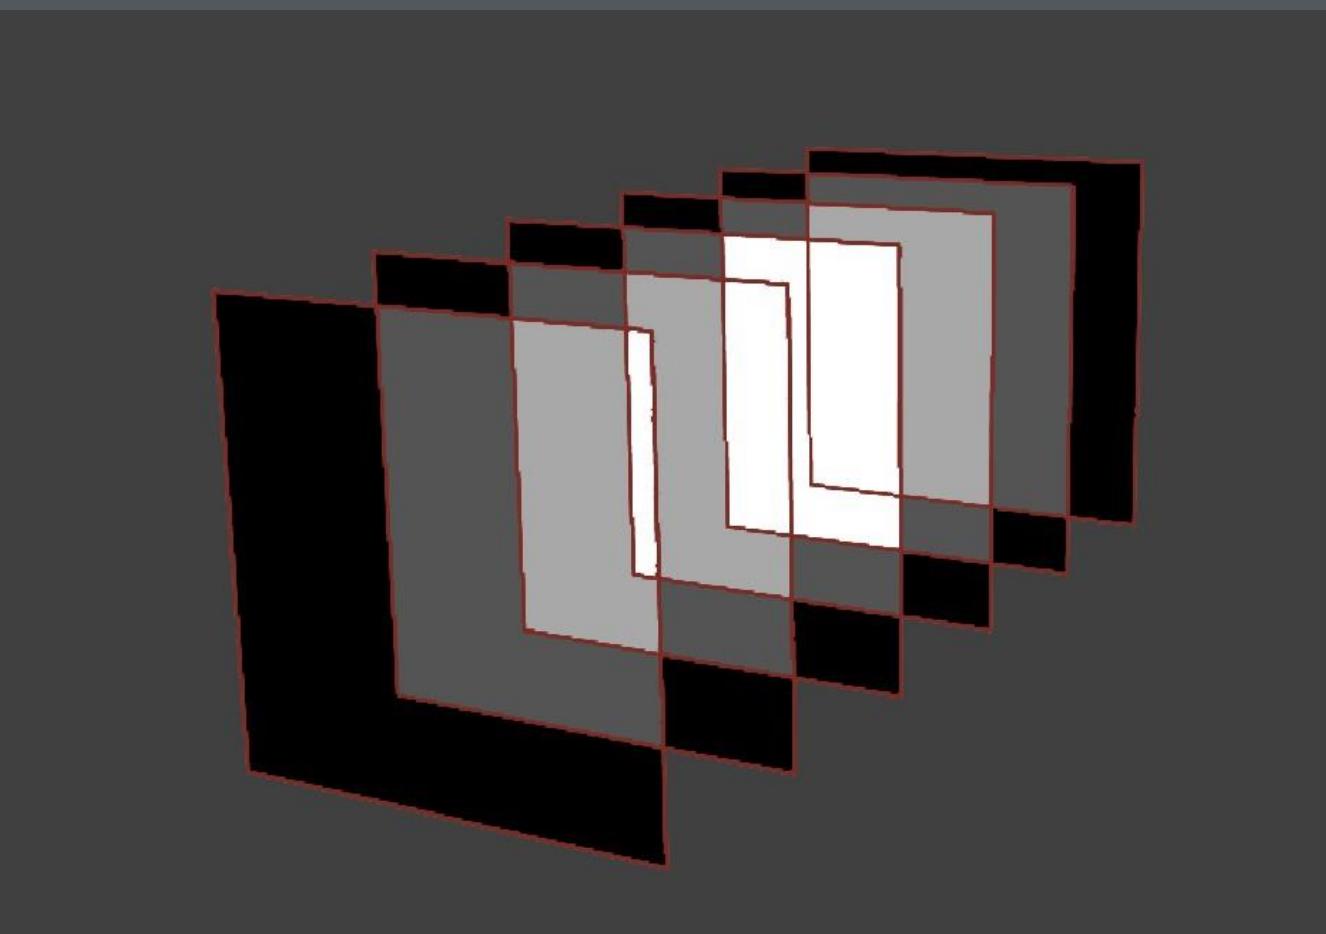
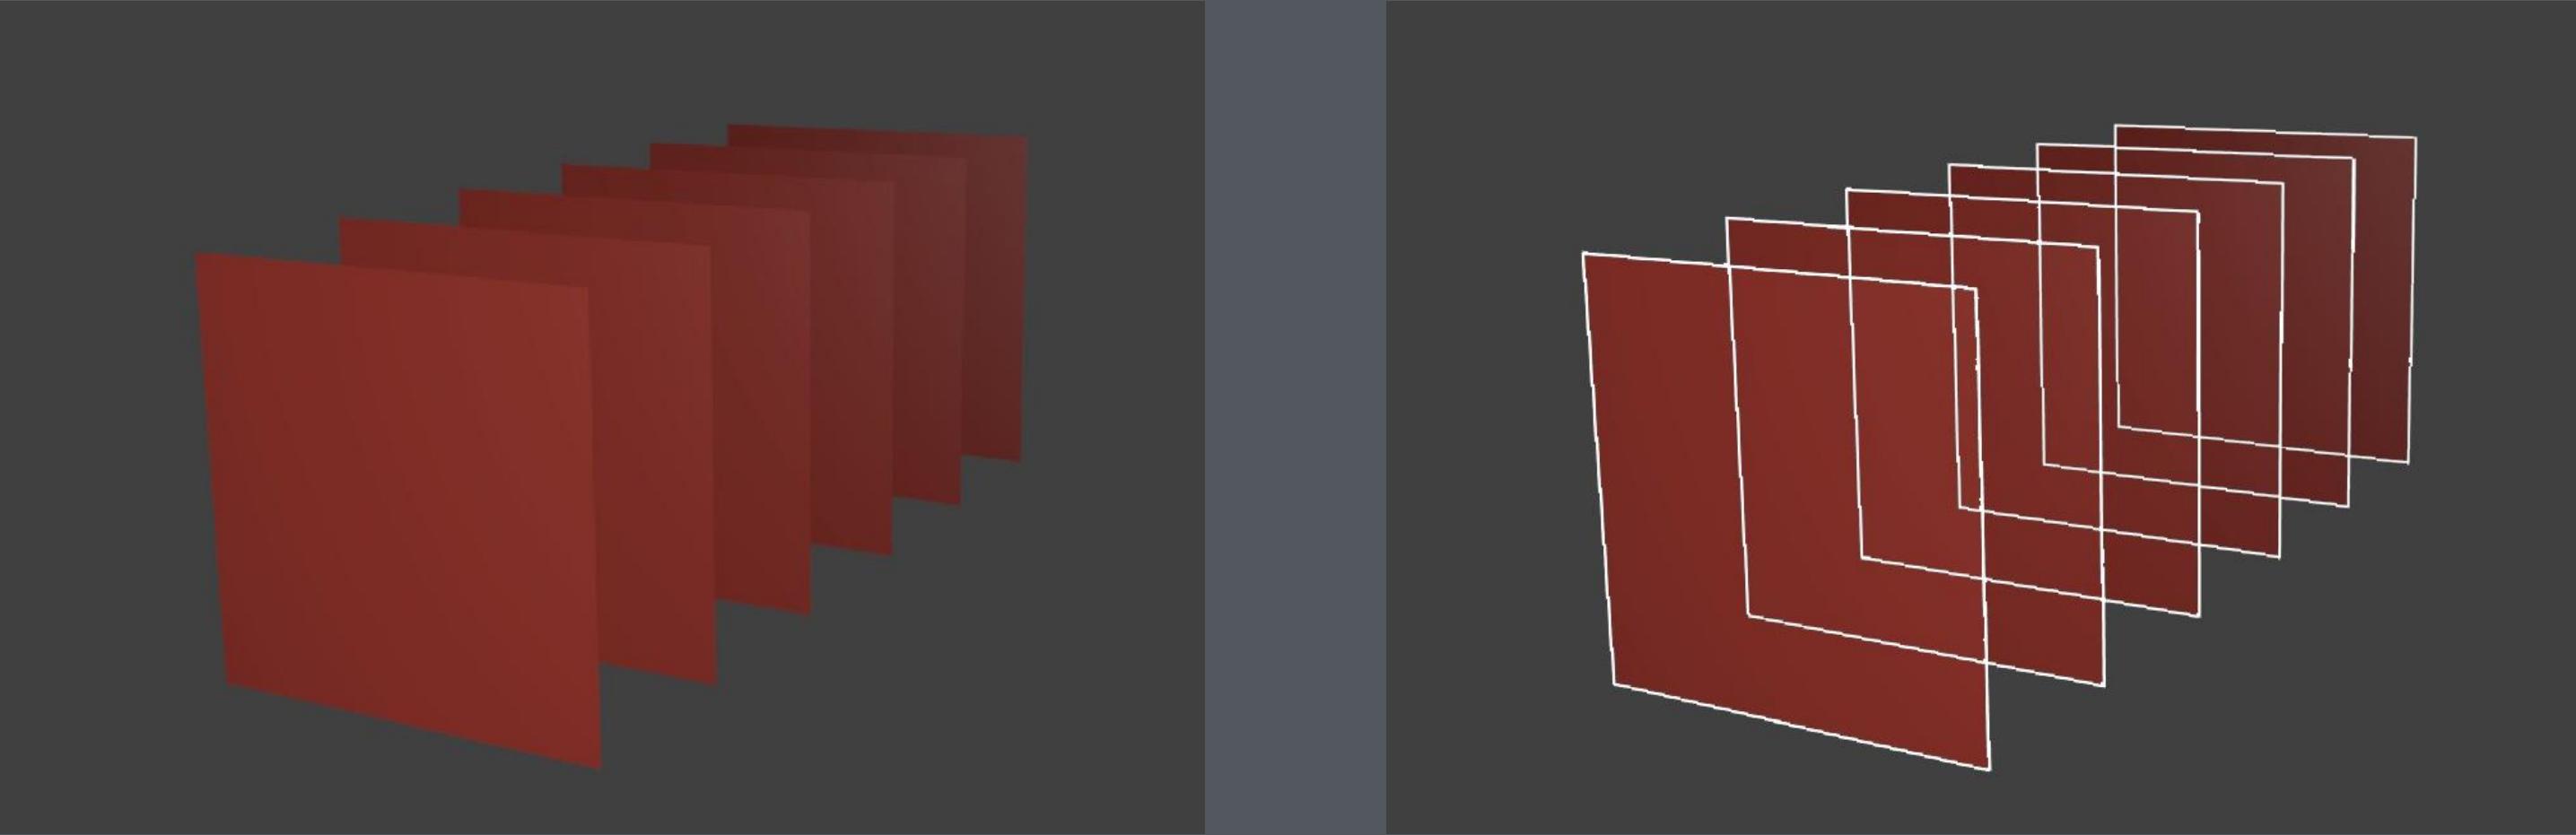
```
for each object in the scene
    for each triangle in this object
        for each fragment f in this triangle

            gl_FragColor = shade(f)

            if (depth of f < depthbuffer[x, y])
                framebuffer[x, y] = gl_FragColor
                depthbuffer[x, y] = depth of f
            end if

        end for
    end for
end for
```

# Problem: Overdraw



# Problem: lighting complexity



# Missed opportunity: spatial processing

**Fragments cannot talk to each other**

- a fundamental constraint for performance

**Many interesting effects depend on neighborhood and geometry**

- bloom
- ambient occlusion
- motion blur
- depth of field
- edge-related rendering effects

# Deferred shading approach

## First render pass

- draw all geometry
- compute material- and geometry-related inputs to lighting model
- don't compute lighting
- write shading inputs to intermediate buffer

## Second render pass

- don't draw geometry
- use stored inputs to compute shading
- write to output

## Post-processing pass (optional, can also be used with fwd. rendering)

- process final image to produce output pixels

# Deferred shading step 1

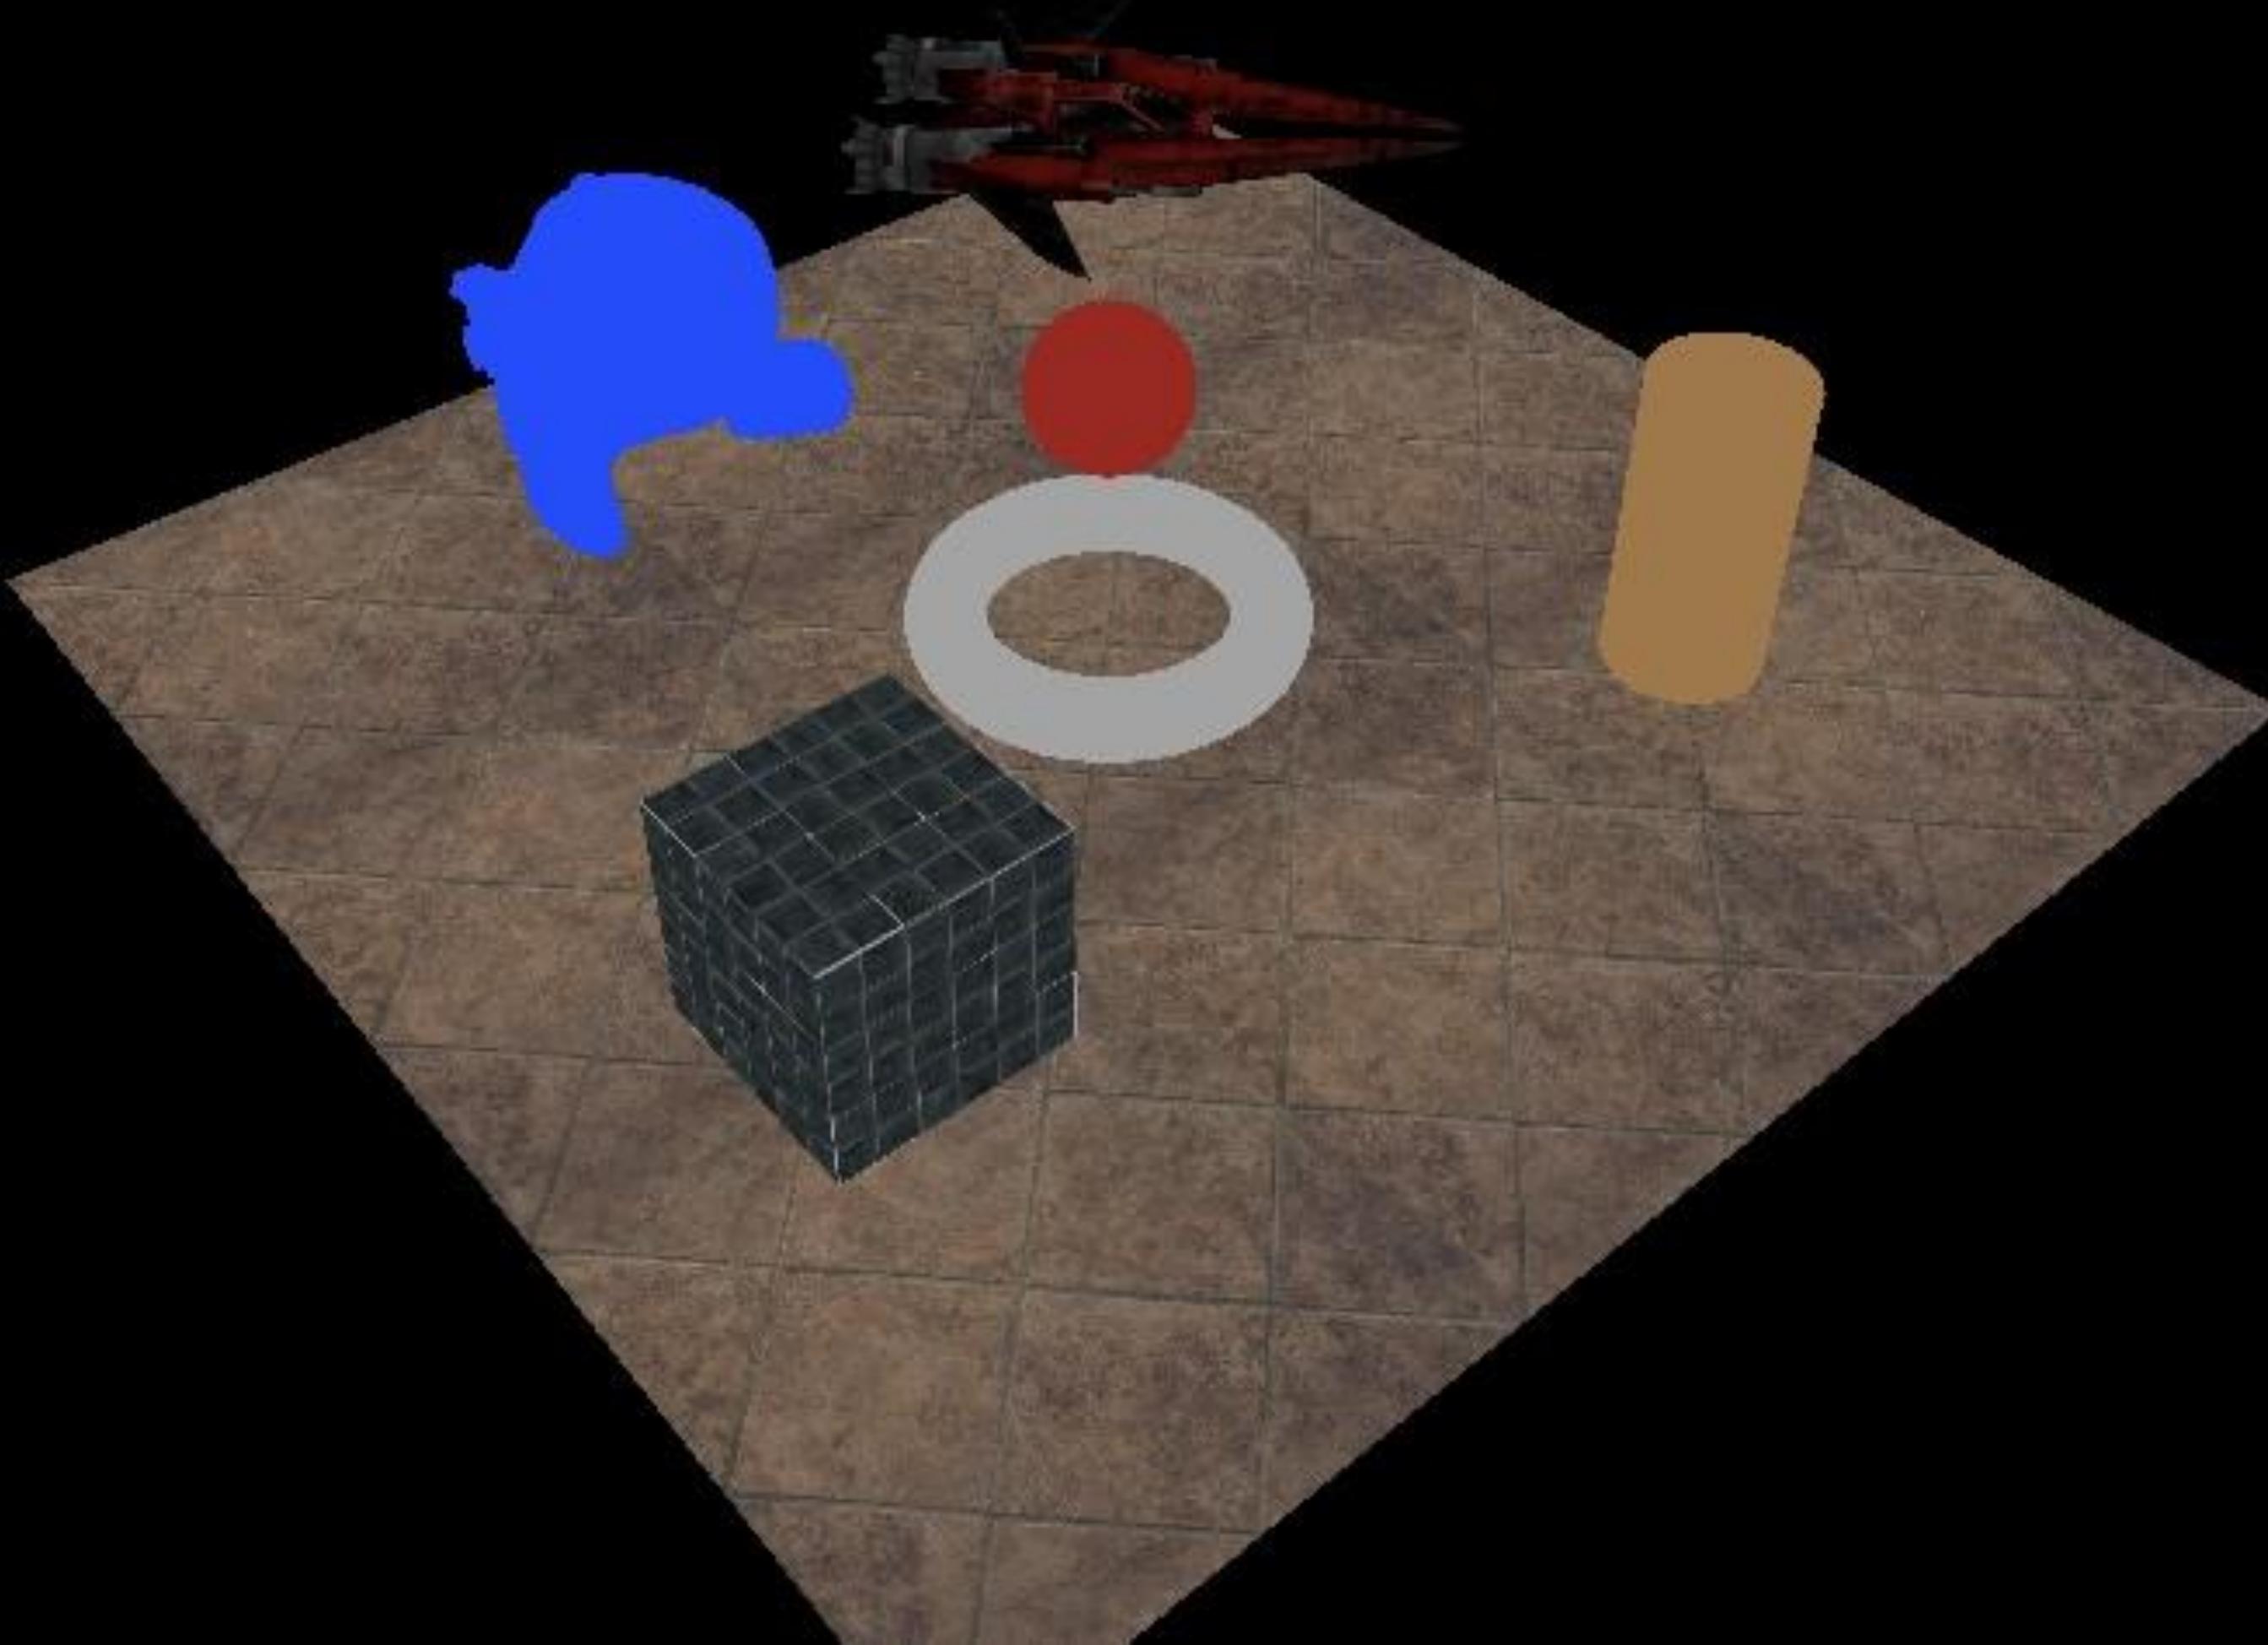
```
for each object in the scene
    for each triangle in this object
        for each fragment f in this triangle

            gl_FragData[...] = material properties of f
            if (depth of f < depthbuffer[x, y])
                gbuffer[...][x, y] = gl_FragData[...]
                depthbuffer[x, y] = depth of f
            end if

        end for
    end for
end for
```

Here we're making use of an OpenGL feature called "Multiple Render Targets" in which the familiar `gl_FragColor` is replaced by an array of values, each of which is written to a different buffer.

# First pass: output just the materials



# Deferred Shading Step 2

```
for each fragment f in the gbuffer  
    framebuffer[x, y] = shade (f)  
end for
```

Key improvement: **shade (f)** only executed for **visible** fragments.

Output is the same →



```
for each object in the scene
    for each triangle in this object
        for each fragment f in this triangle

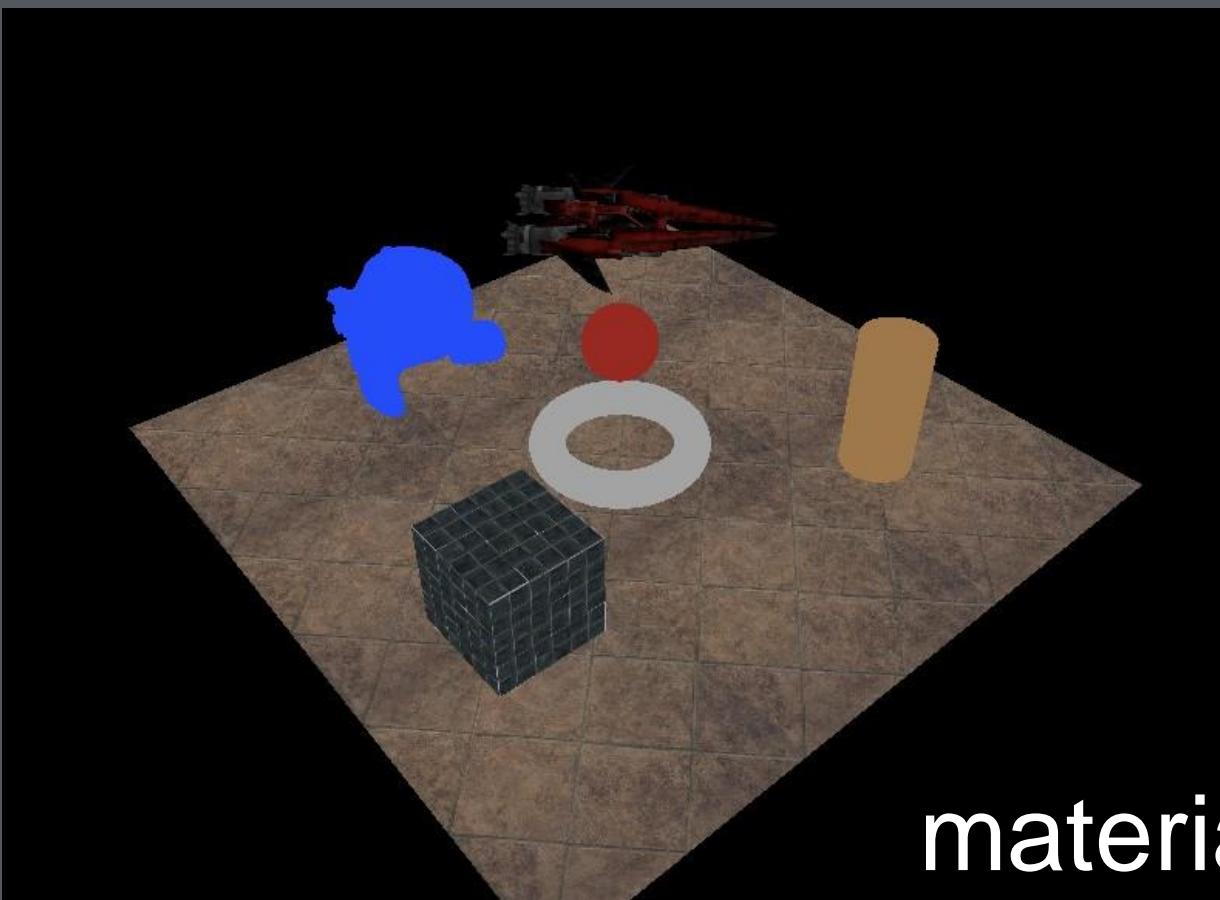
            gl_FragData[...] = material properties of f
            if (depth of f < depthbuffer[x, y])
                gbuffer[...][x, y] = gl_FragData[...]
                depthbuffer[x, y] = depth of f
            end if

        end for
    end for
end for
```

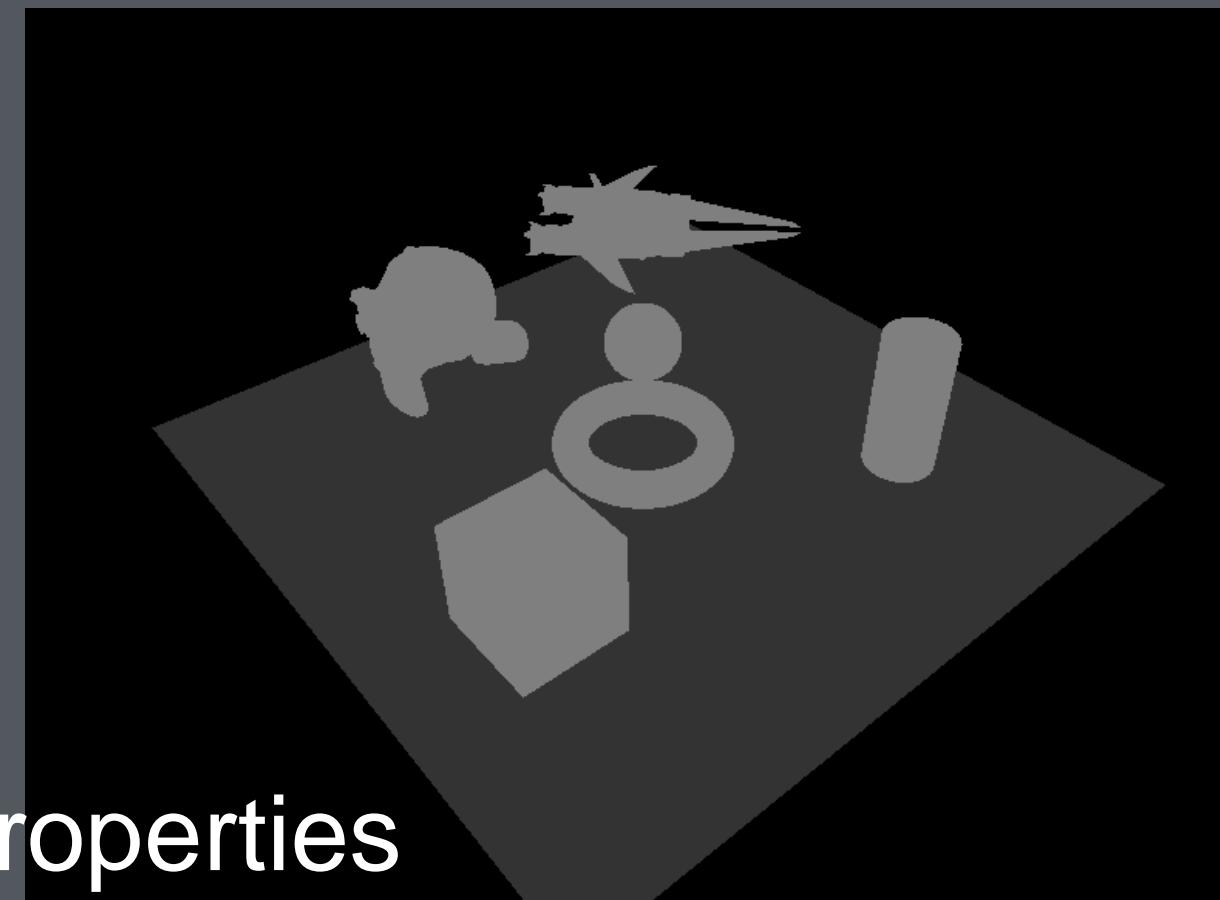
Here we're making use of an OpenGL feature called "Multiple Render Targets" in which the familiar `gl_FragColor` is replaced by an array of values, each of which is written to a different buffer.

```
for each fragment f in the gbuffer
    framebuffer[x, y] = shade (f)
end for
```

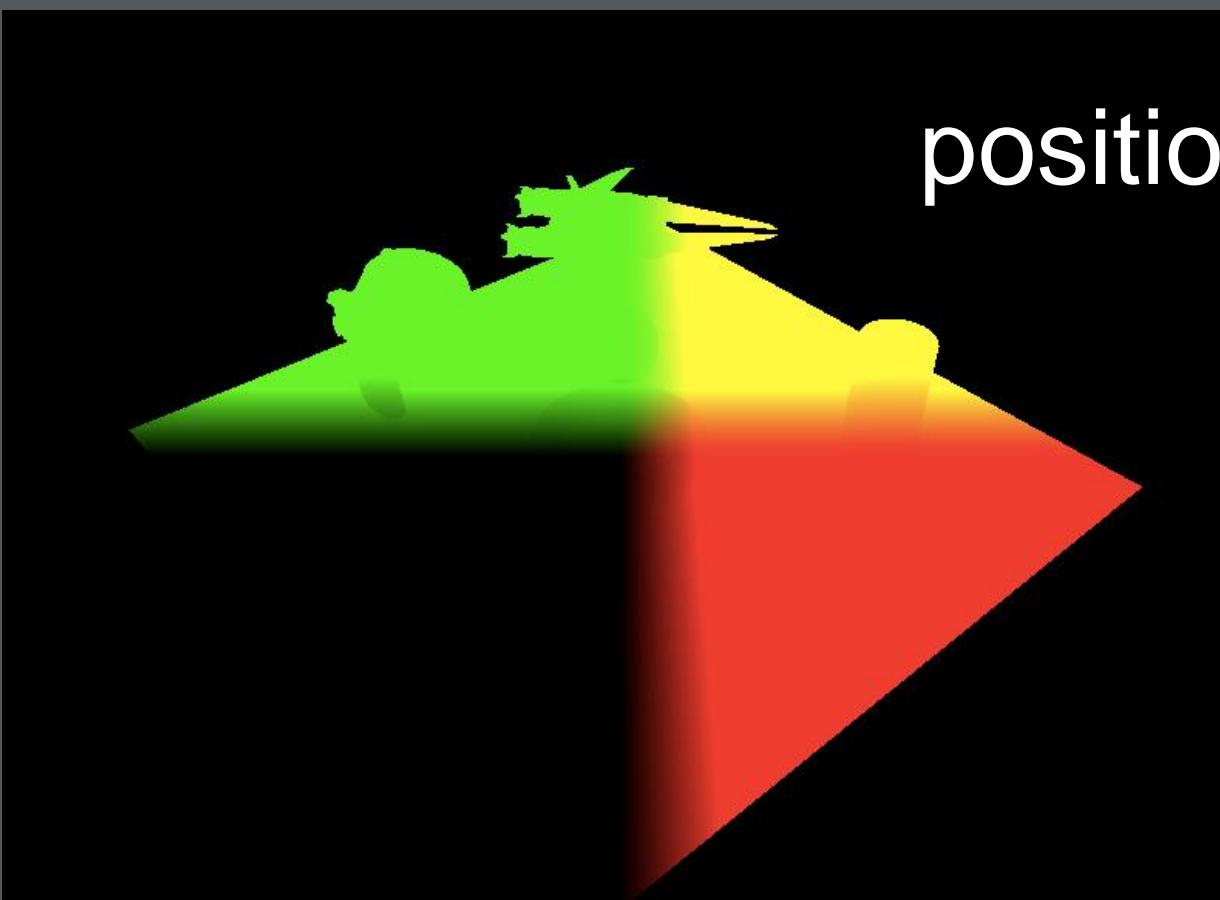
# G-buffer: multiple textures



material properties



position



normal

# The Übershader

**Shader which computes lighting based on g-buffer: has code for all material/lighting models in a single huge shader.**

```
shade (f) {  
    result = 0;  
    if (f is Lambertian) {  
        for each light  
            result += (n . l) * diffuse;  
    end for  
    } else if (f is Blinn-Phong) {  
        ...  
    } else if (f is ...) {  
        ...  
    }  
    return result;  
}
```

# Übershader inputs

**Need access to all parameters of the material for the current fragment:**

- Blinn-Phong: kd, ks, n
- Microfacet: kd, ks, alpha
- etc.

**Also need fragment position and surface normal**

**Solution: write all that out from the material shaders:**

```
{outputs} = {f.material, f.position, f.normal}
if (depth of f < depthbuffer[x, y])
    gbuffer[x, y] = {outputs}
    depthbuffer[x, y] = depth of f
end if
```

# Deferred lighting

## **Single-pass render has to consider all lights for every fragment**

- much wasted effort since only a few lights probably contribute
- batching geometry by which lights affect it is awkward
- straight 2-pass deferred has same problem

## **With deferred shading, fragments can be visited in subsets**

- for each light, draw bounds of (significantly) affected volume
- only compute shading for fragments covered by that
- with depth/stencil games, can only shade fragments inside the volume

# Power of Deferred Shading

**Can do any image processing between step 1 and step 2!**

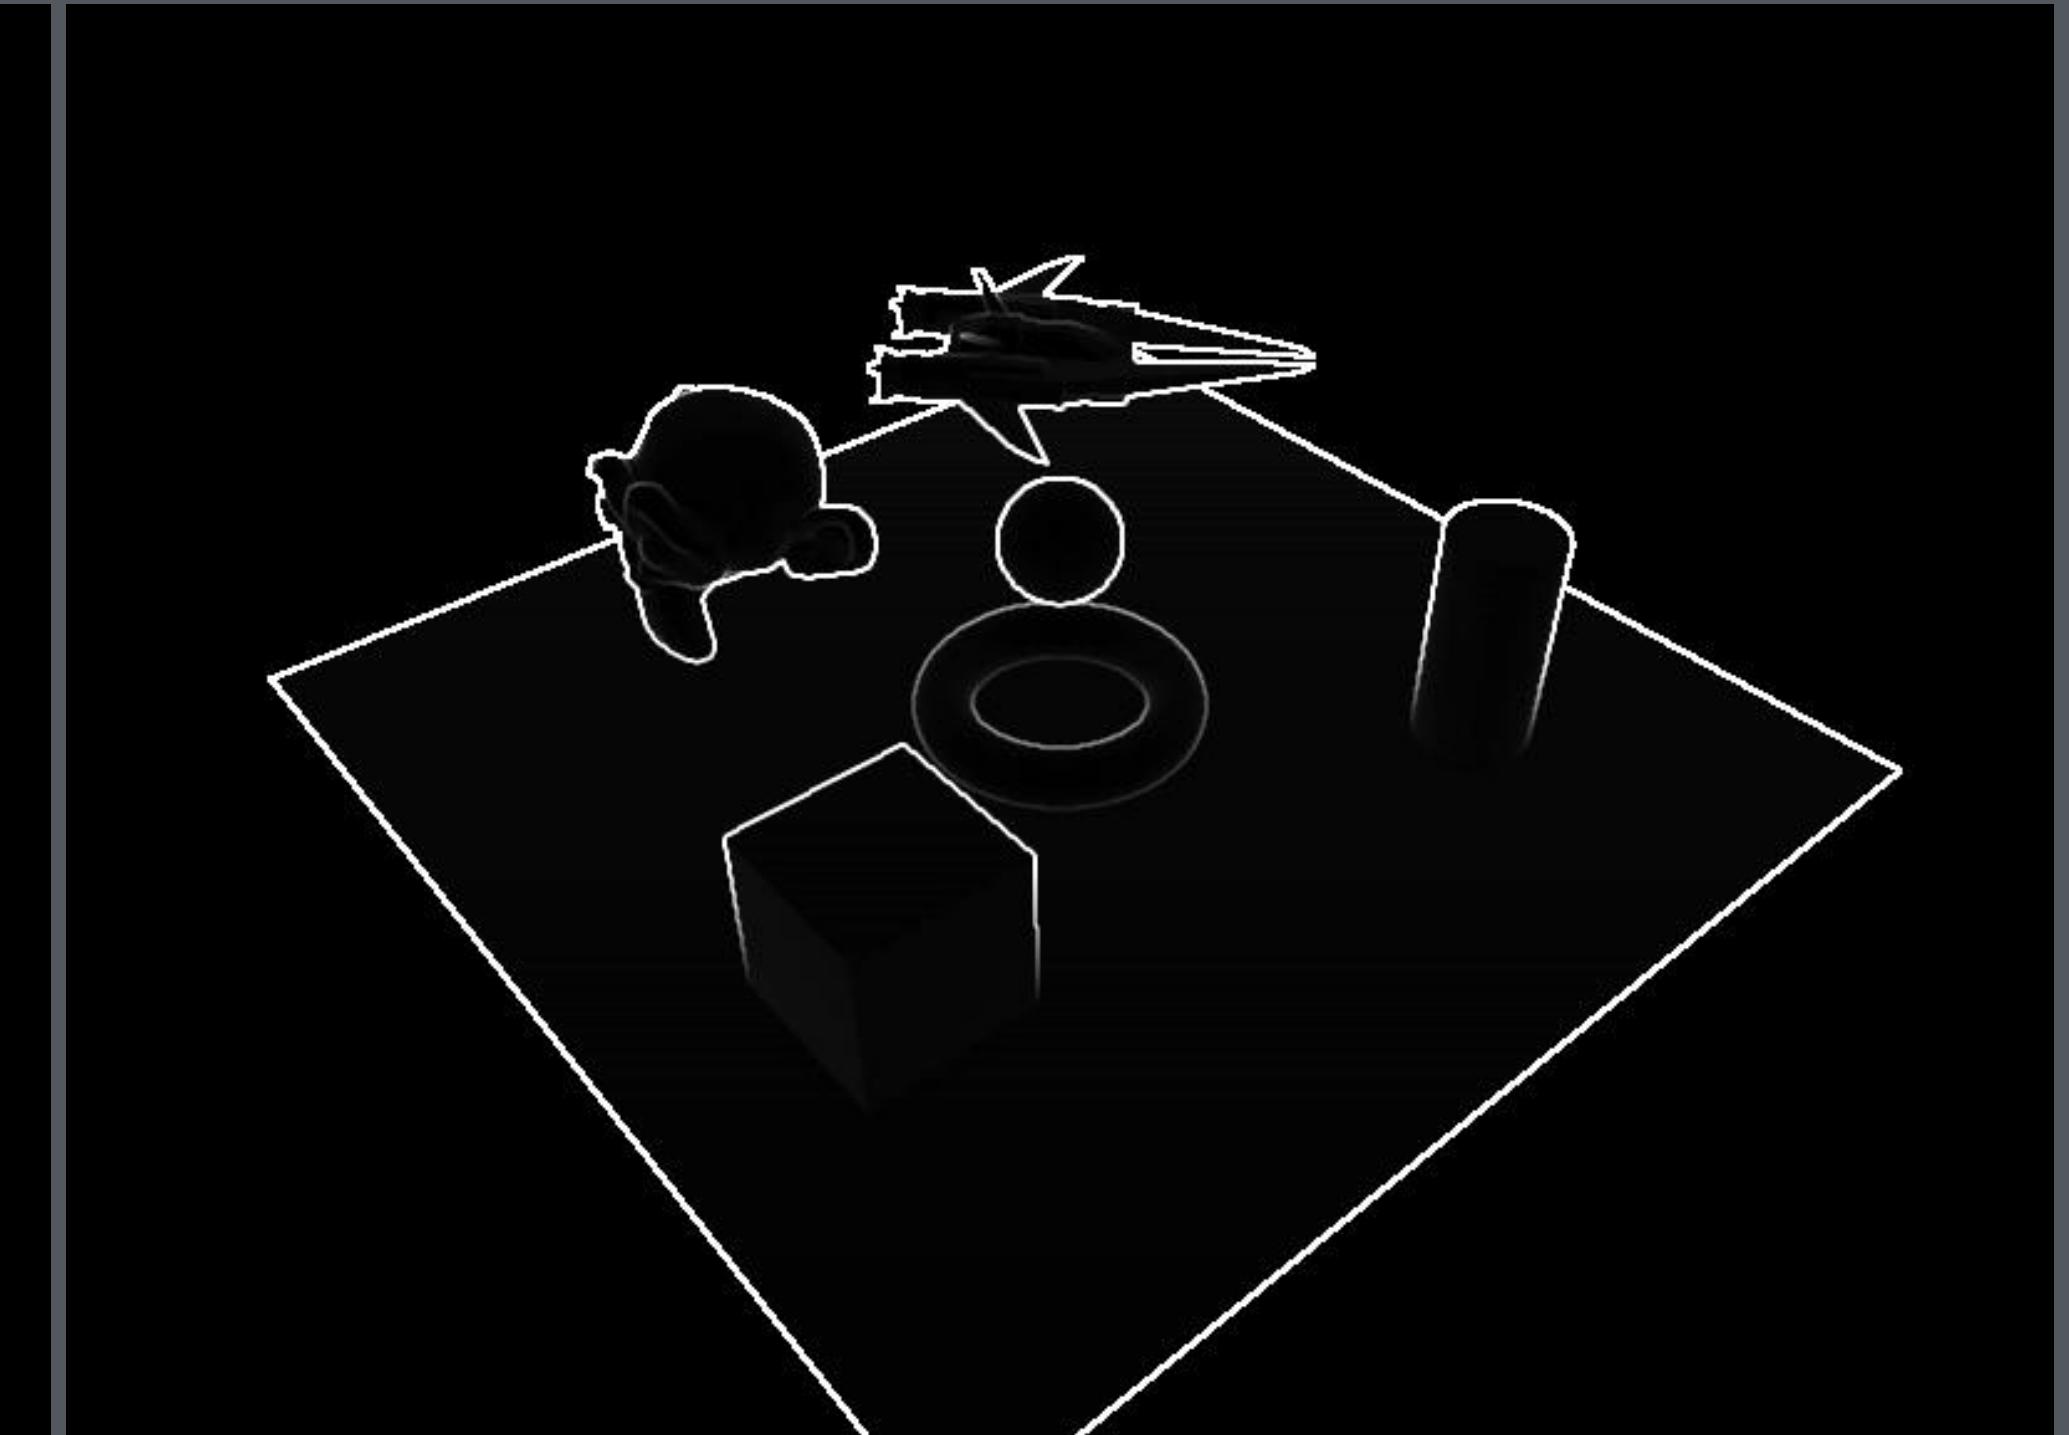
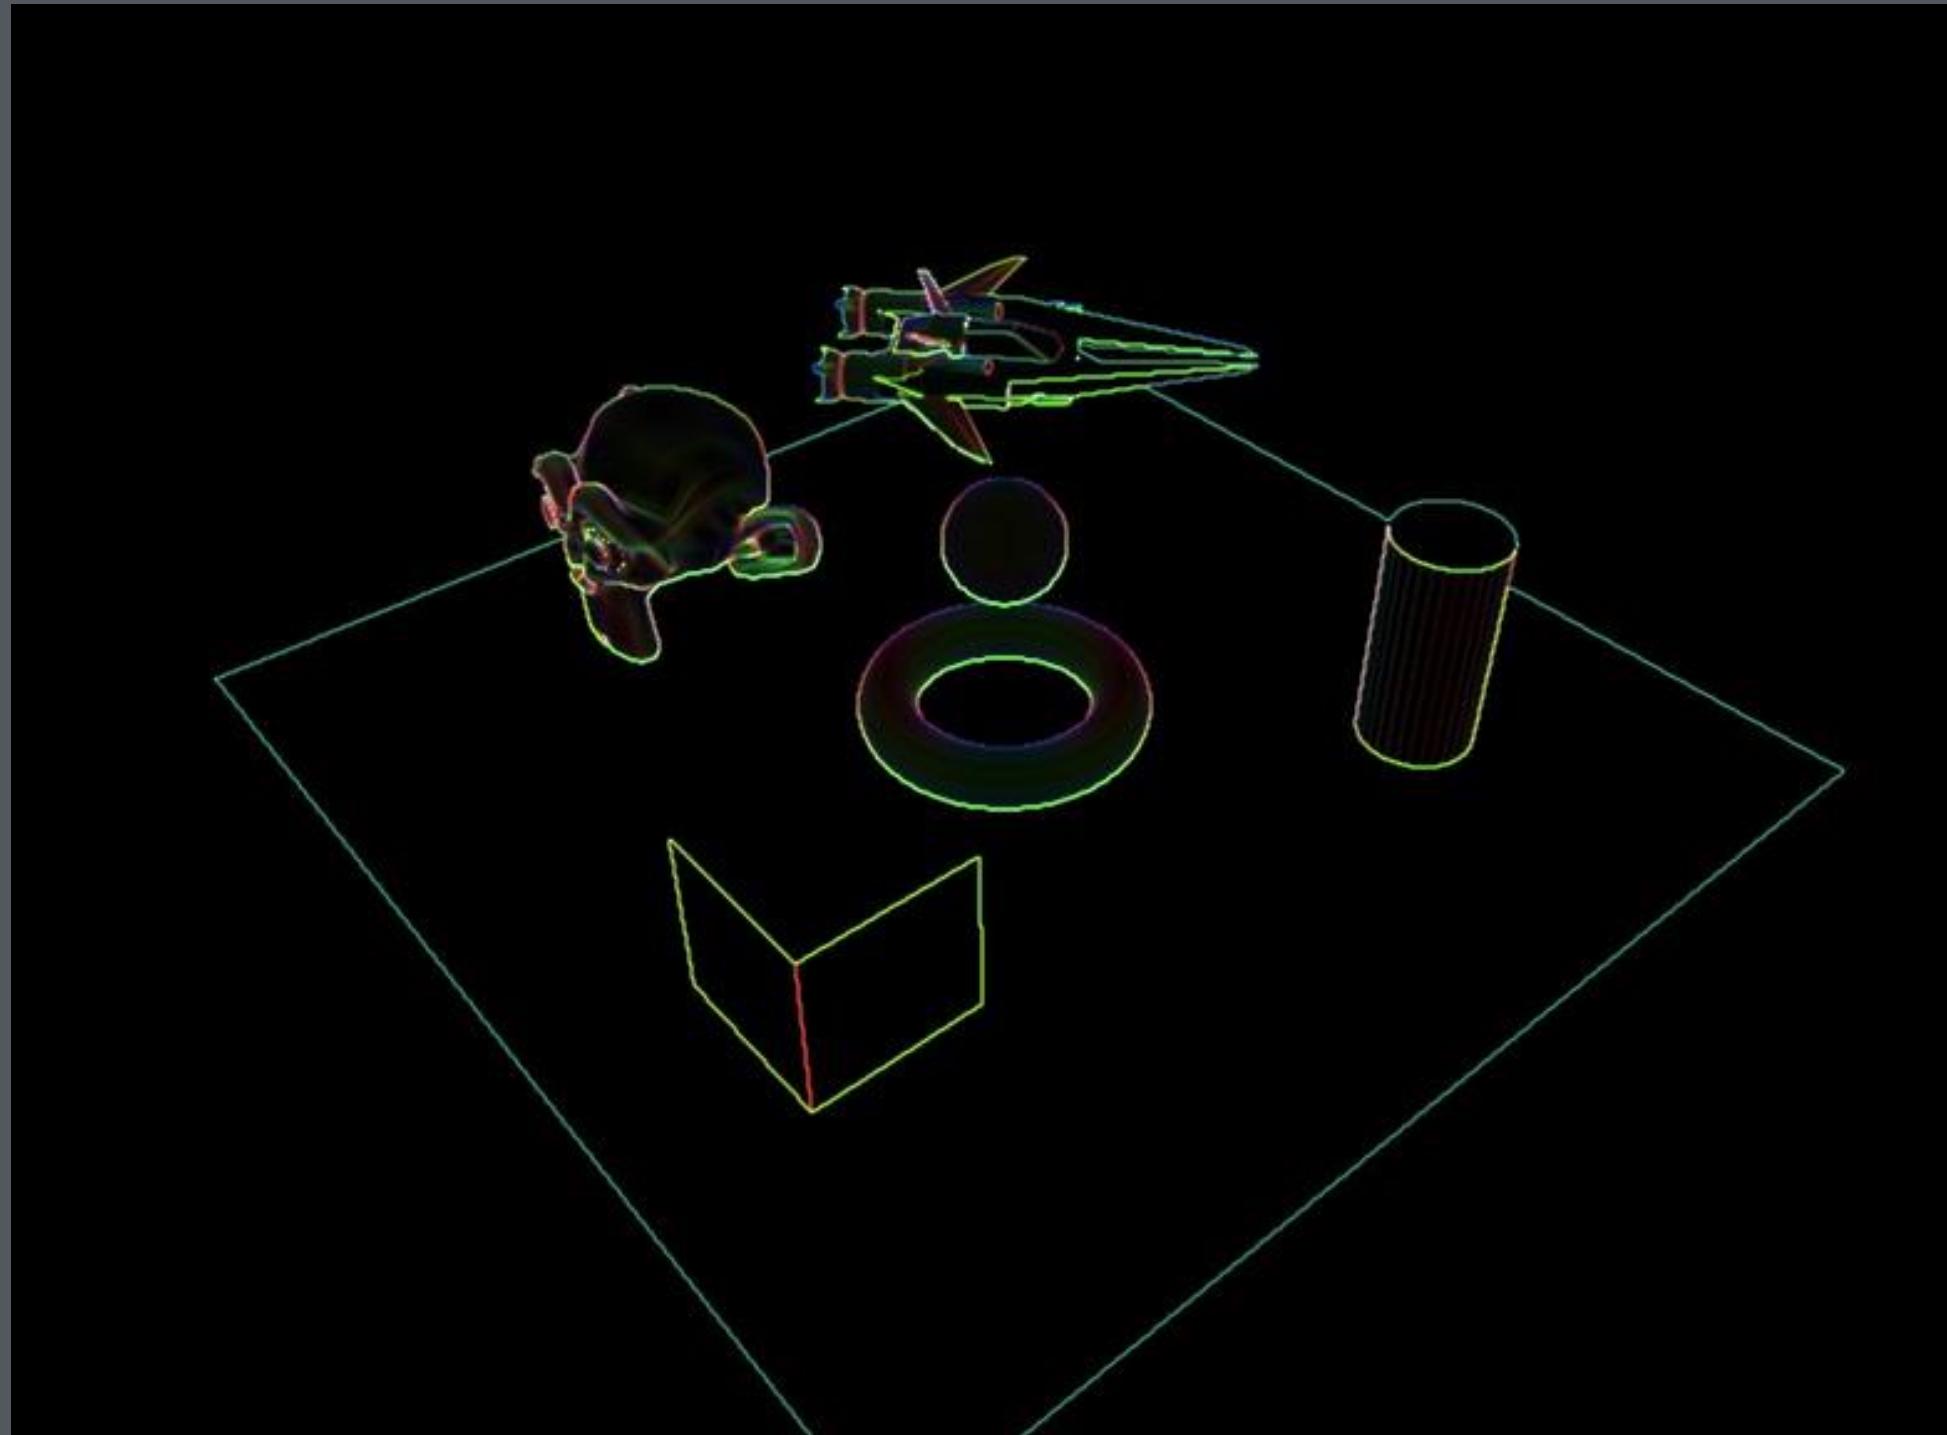
- Recall: step 1 = fill g-buffer, step 2 = light/shade
- Could add a step 1.5 to filter the g-buffer

**Examples:**

- silhouette detection for artistic rendering
- screen-space ambient occlusion
- denoising based on bilateral filter using geometric info

# Silhouette detection

By differentiating the depth buffer, can locate silhouettes and creases





# Ambient occlusion



# Denoising



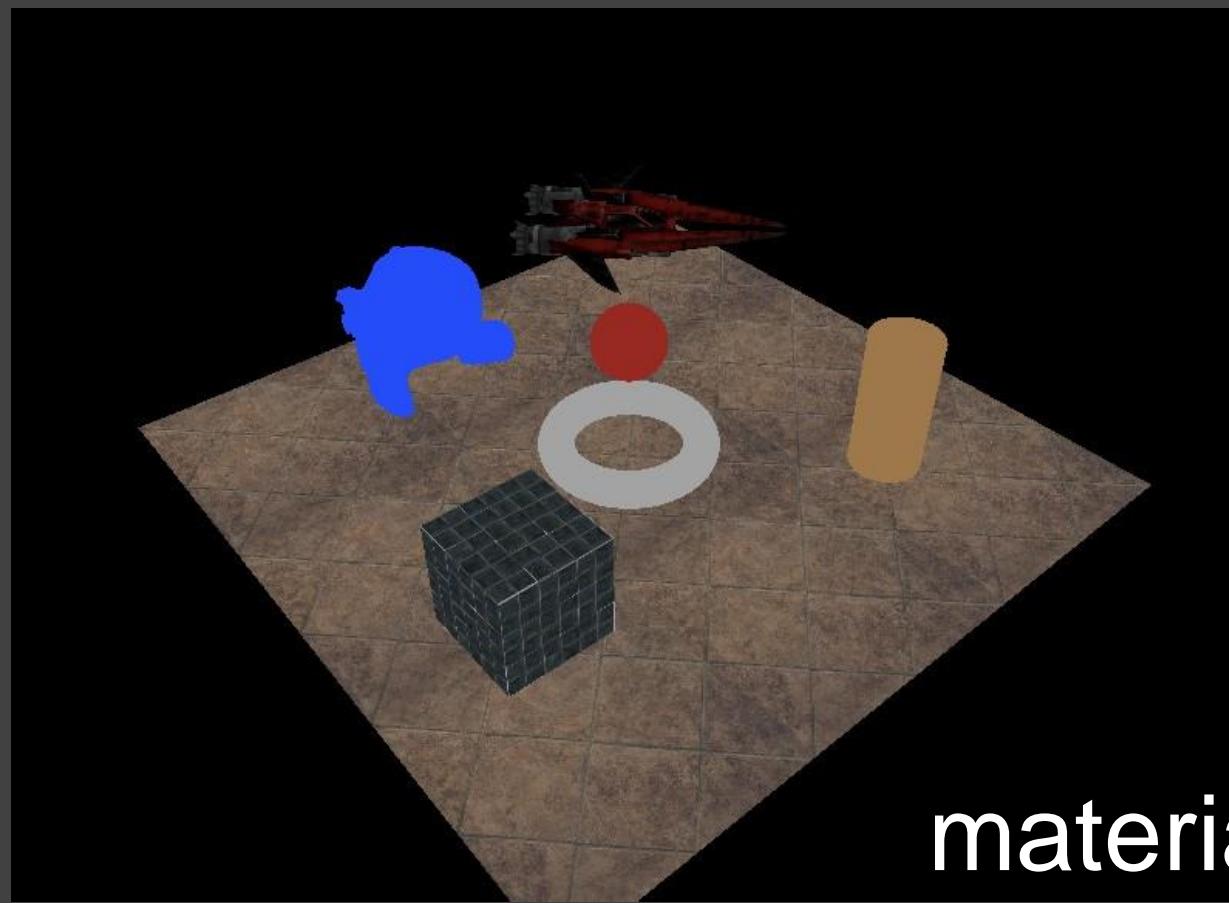
[Mara et al. HPG 2017]

# Denoising

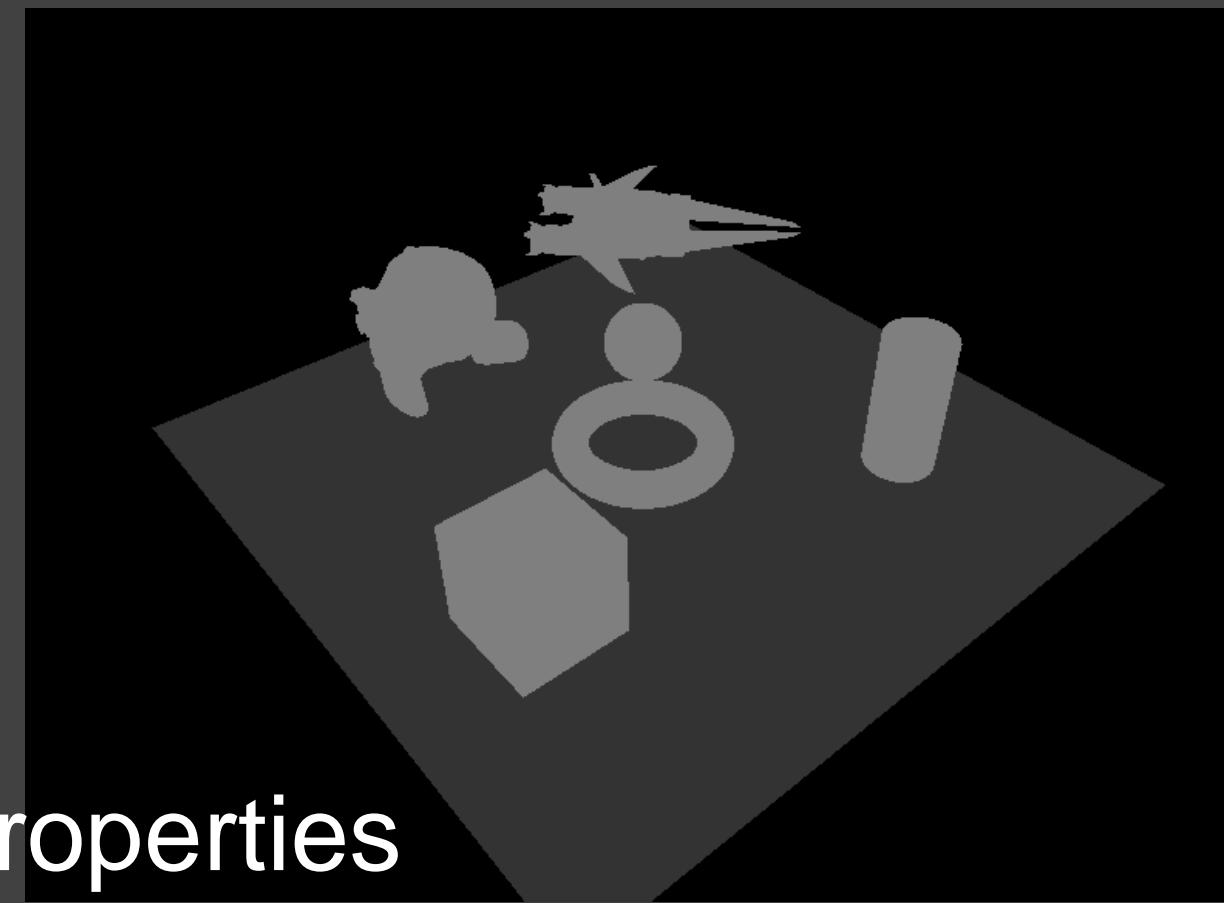


[Mara et al. HPG 2017]

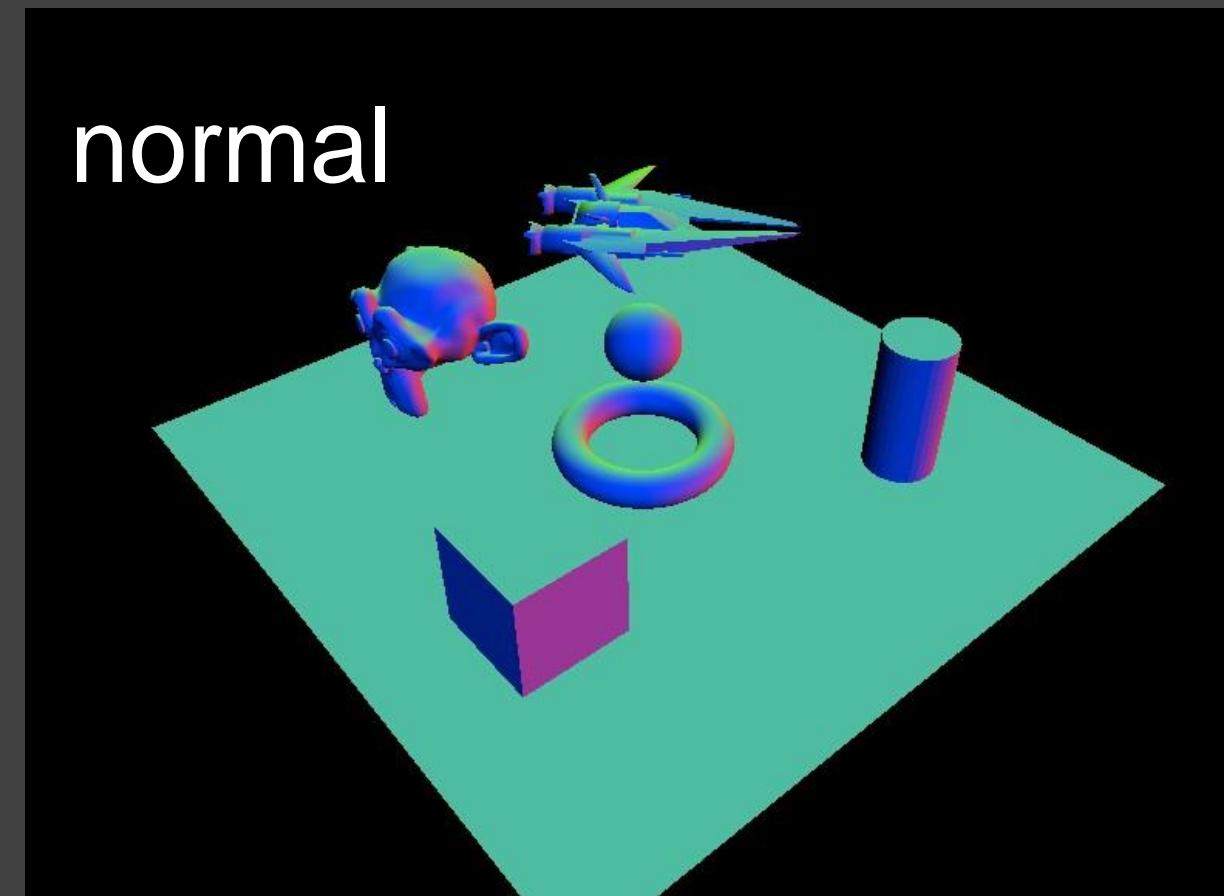
# How to Fill the Buffers?



material properties



position



normal

# Limitations of Deferred Shading

**Each pixel in the g-buffer can only store material and surface info for a single surface.**

- blending/transparency is difficult
- antialiasing is a different ball game

**For transparency: a “hybrid” renderer**

- deferred shading for opaque objects, forward shading for translucent objects
- allows translucent geometry to know about opaque geometry behind it

**For antialiasing: smart blurring**

- use what is in the g-buffers to blur along but not across edges

# Hybrid Rendering



Note how the water effect  
fades out in the shallows —  
access to depth of ground.

Image credit: random guy on  
internet: <http://vimeo.com/14337919>

# Antialiasing

**Single shading sample per pixel**

**Reconstruct by blending nearby samples**

- select them by looking for edges  
(Morphological AA [Reshetov 09])
- learn about edges using multisample depth  
(Subpixel Reconstruction AA [Chajdas et al. 11])

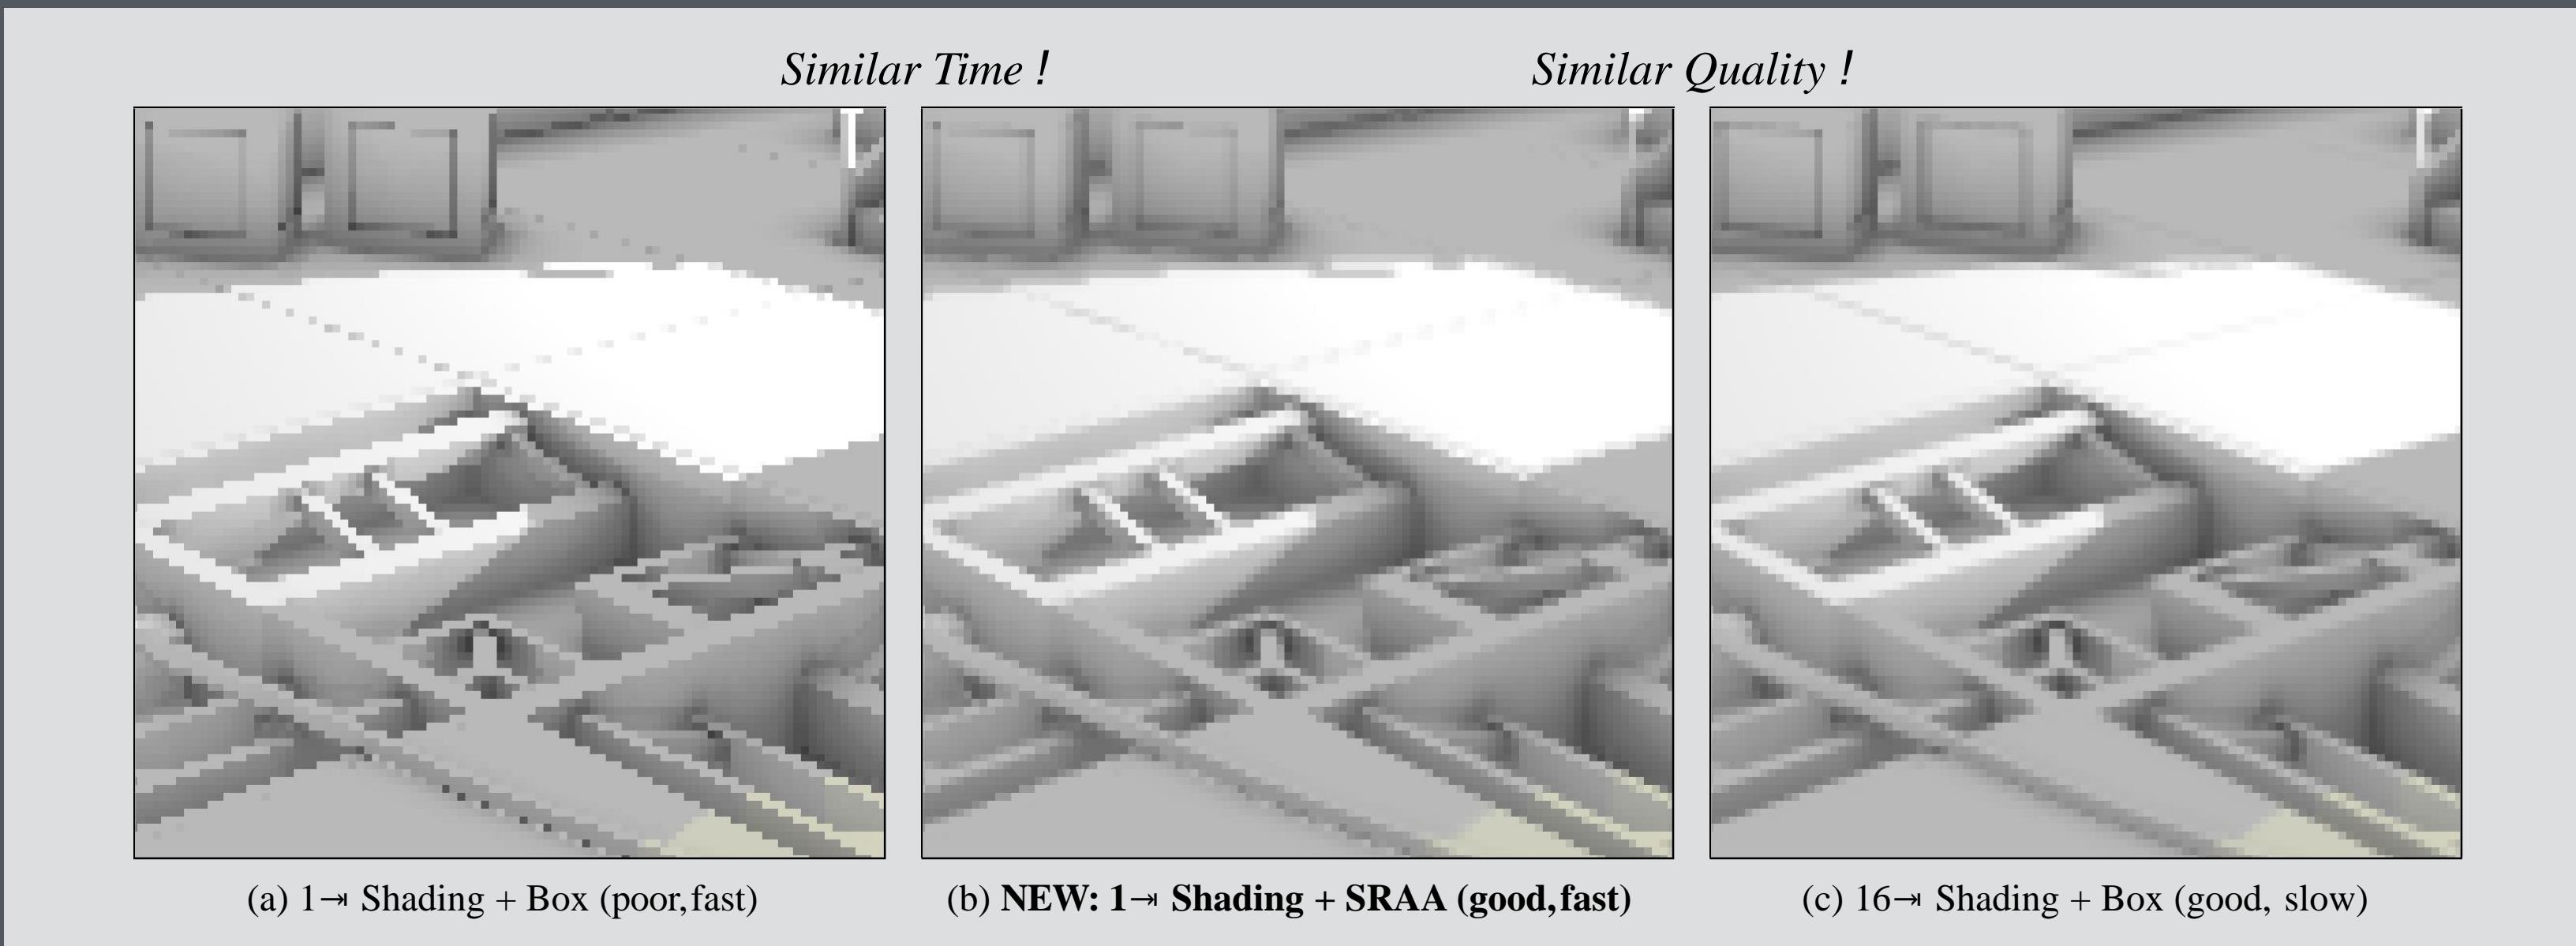
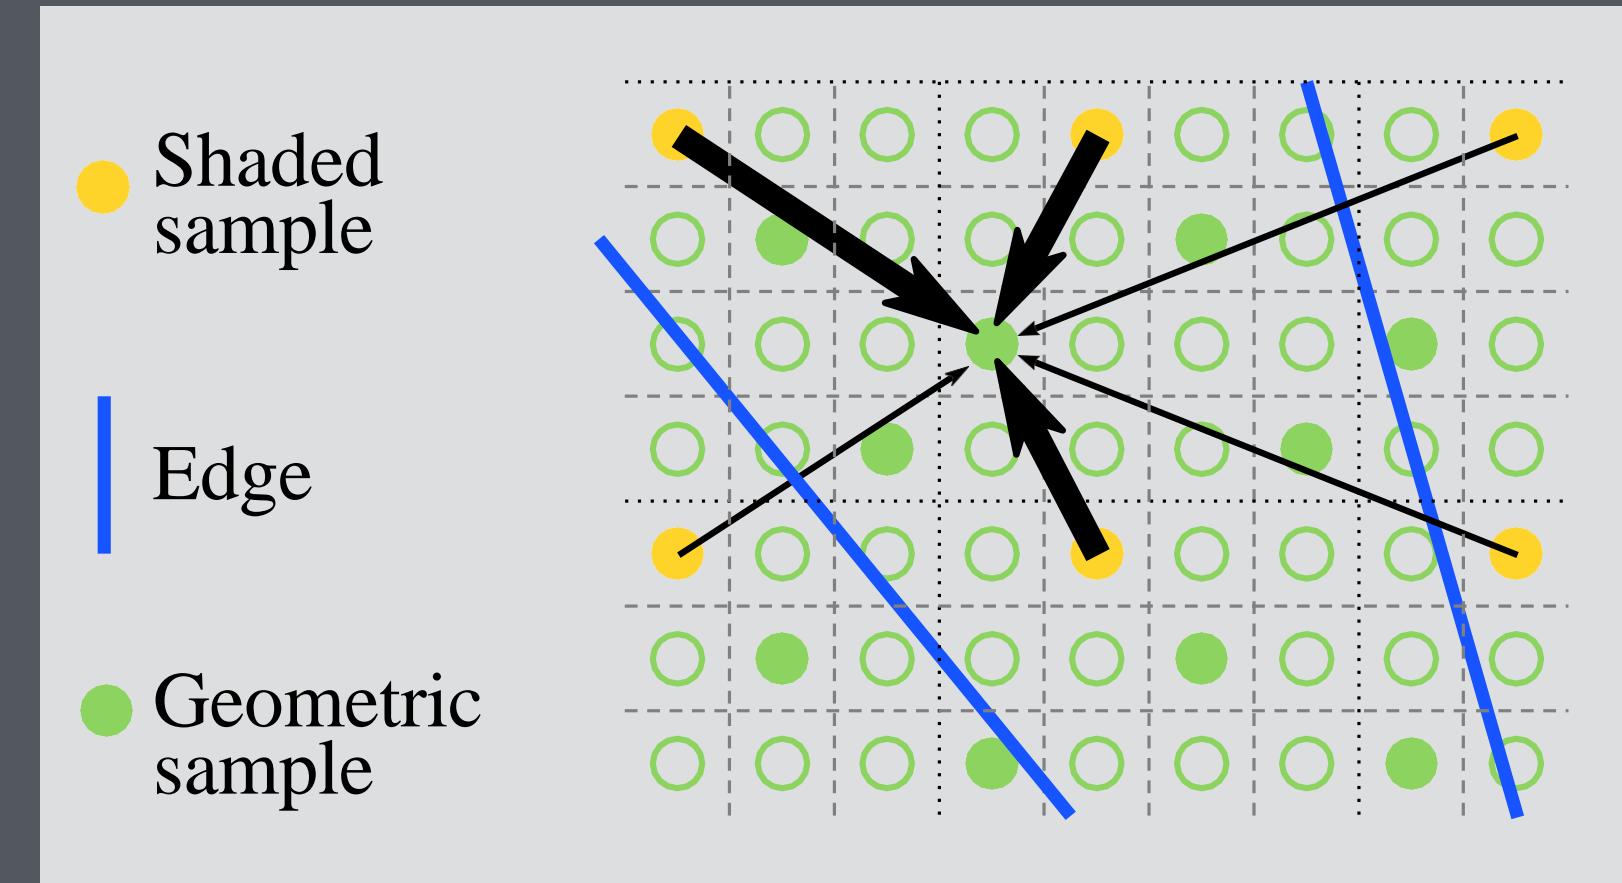
# MLAA

[Reshetov 09]



Detects borders in the resulting image and then finds specific patterns in these.

**Anti-aliasing** : blending pixels in these borders, according to the pattern and their position within the pattern



# Summary: Deferred Shading

## Pros

- Store everything you need in 1st pass
  - normals, diffuse, specular, positions,...
  - G-buffer
- After z-buffer, can shade only what is visible

## Cons:

- transparency (only get one fragment per pixel)
- antialiasing (multisample AA not easy to adapt)

**Standard game engines provide both forward and deferred paths**

# How to do all this in OpenGL

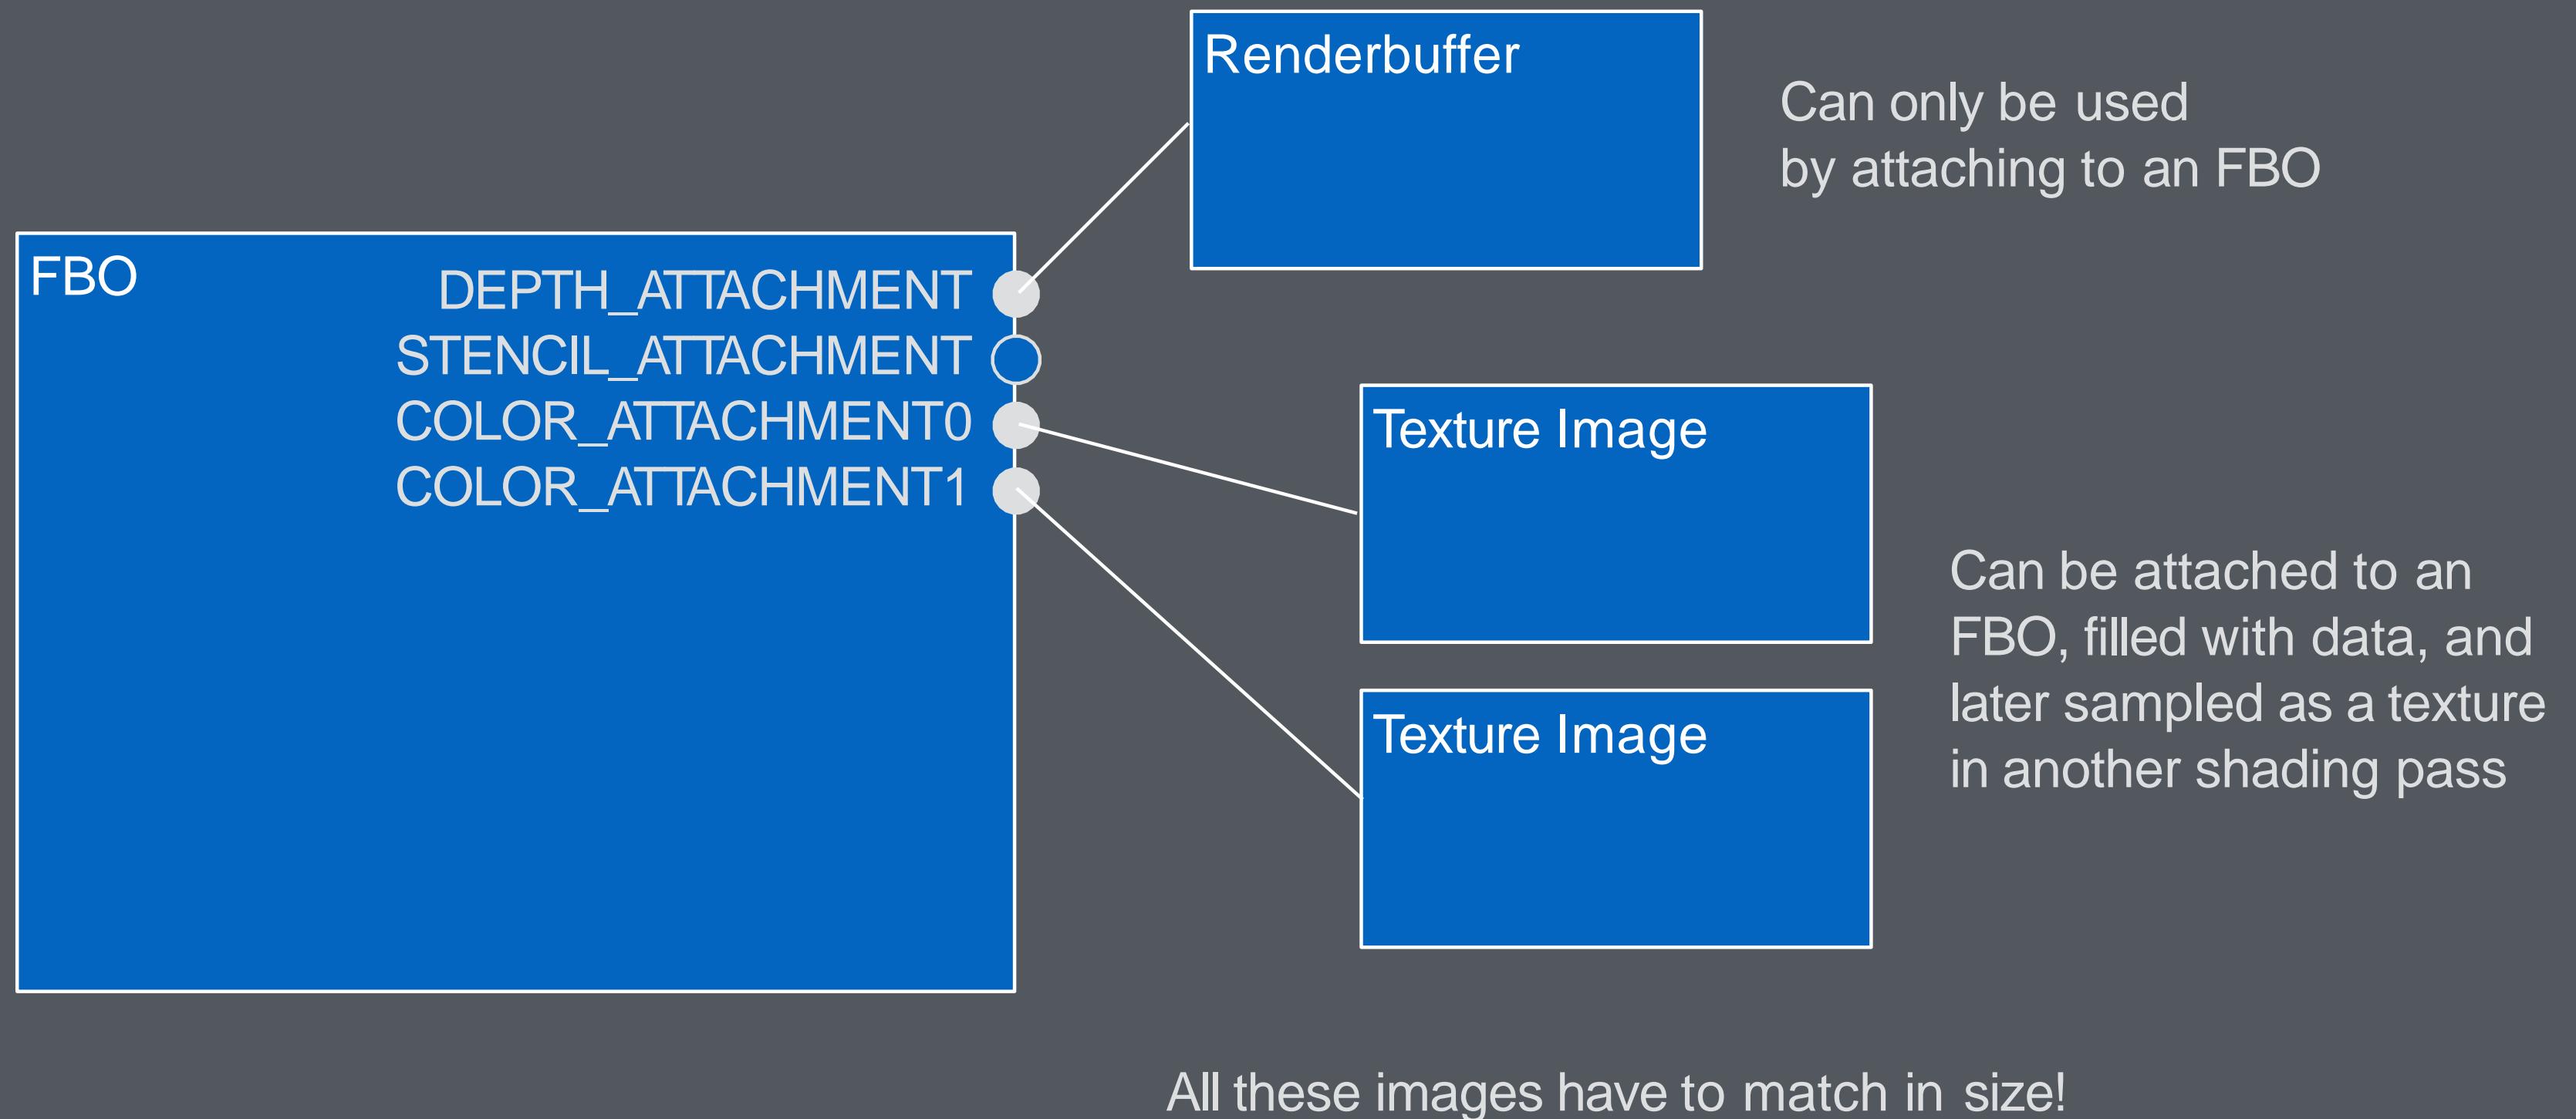
**When you first fire up OpenGL, all fragment shader output is written to the framebuffer that shows up in your window**

- this is the default framebuffer
- it actually can contain multiple buffers: front and back for double-buffering; left and right for stereo/HMD devices. You can control where fragment shader output goes using `glDrawBuffer()`

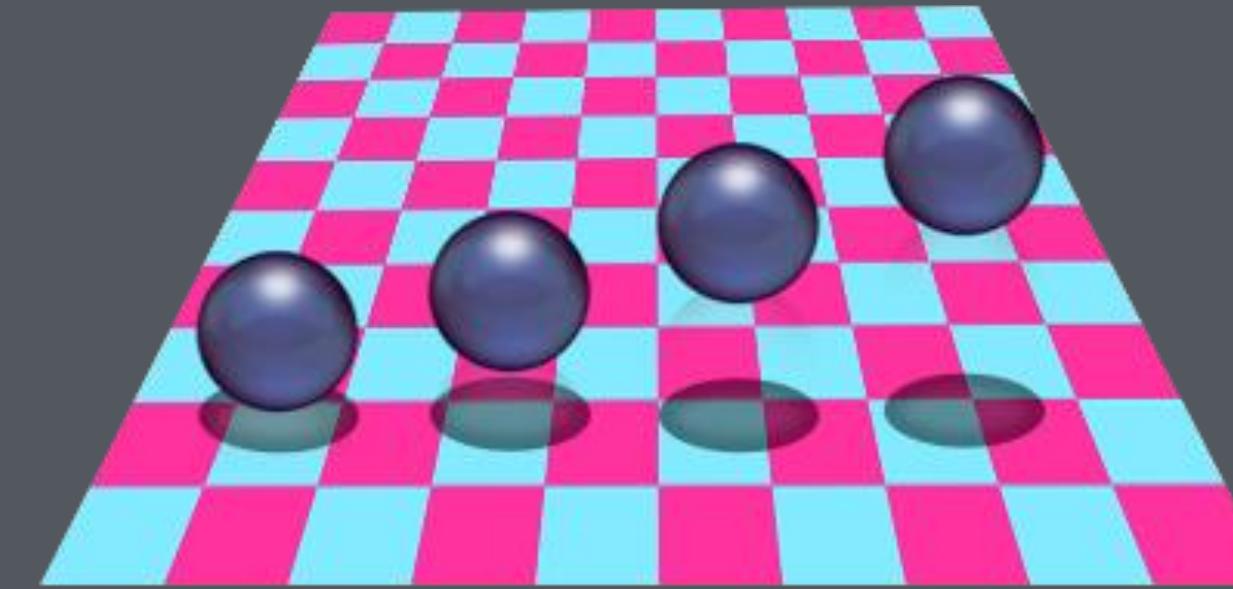
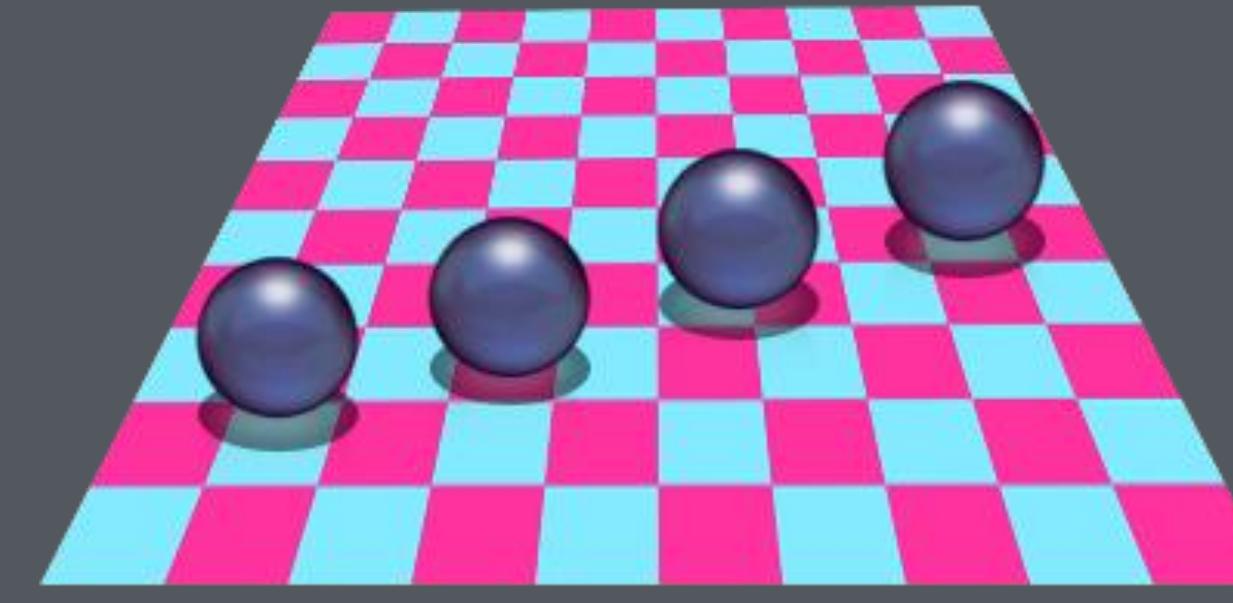
**For deferred shading and other multi-pass methods, you instead create a (non-default) *Framebuffer Object* (FBO)**

- you *attach* images to the FBO to receive fragment shader output
- color attachments (variable number) receive color data from `gl_FragData[...]` (`gl_FragColor` is just an alias for `gl_FragData[0]`)
- a depth attachment is required for z-buffering to function; stencil attachments are also possible

# Framebuffer Object



# Shadows as depth cue



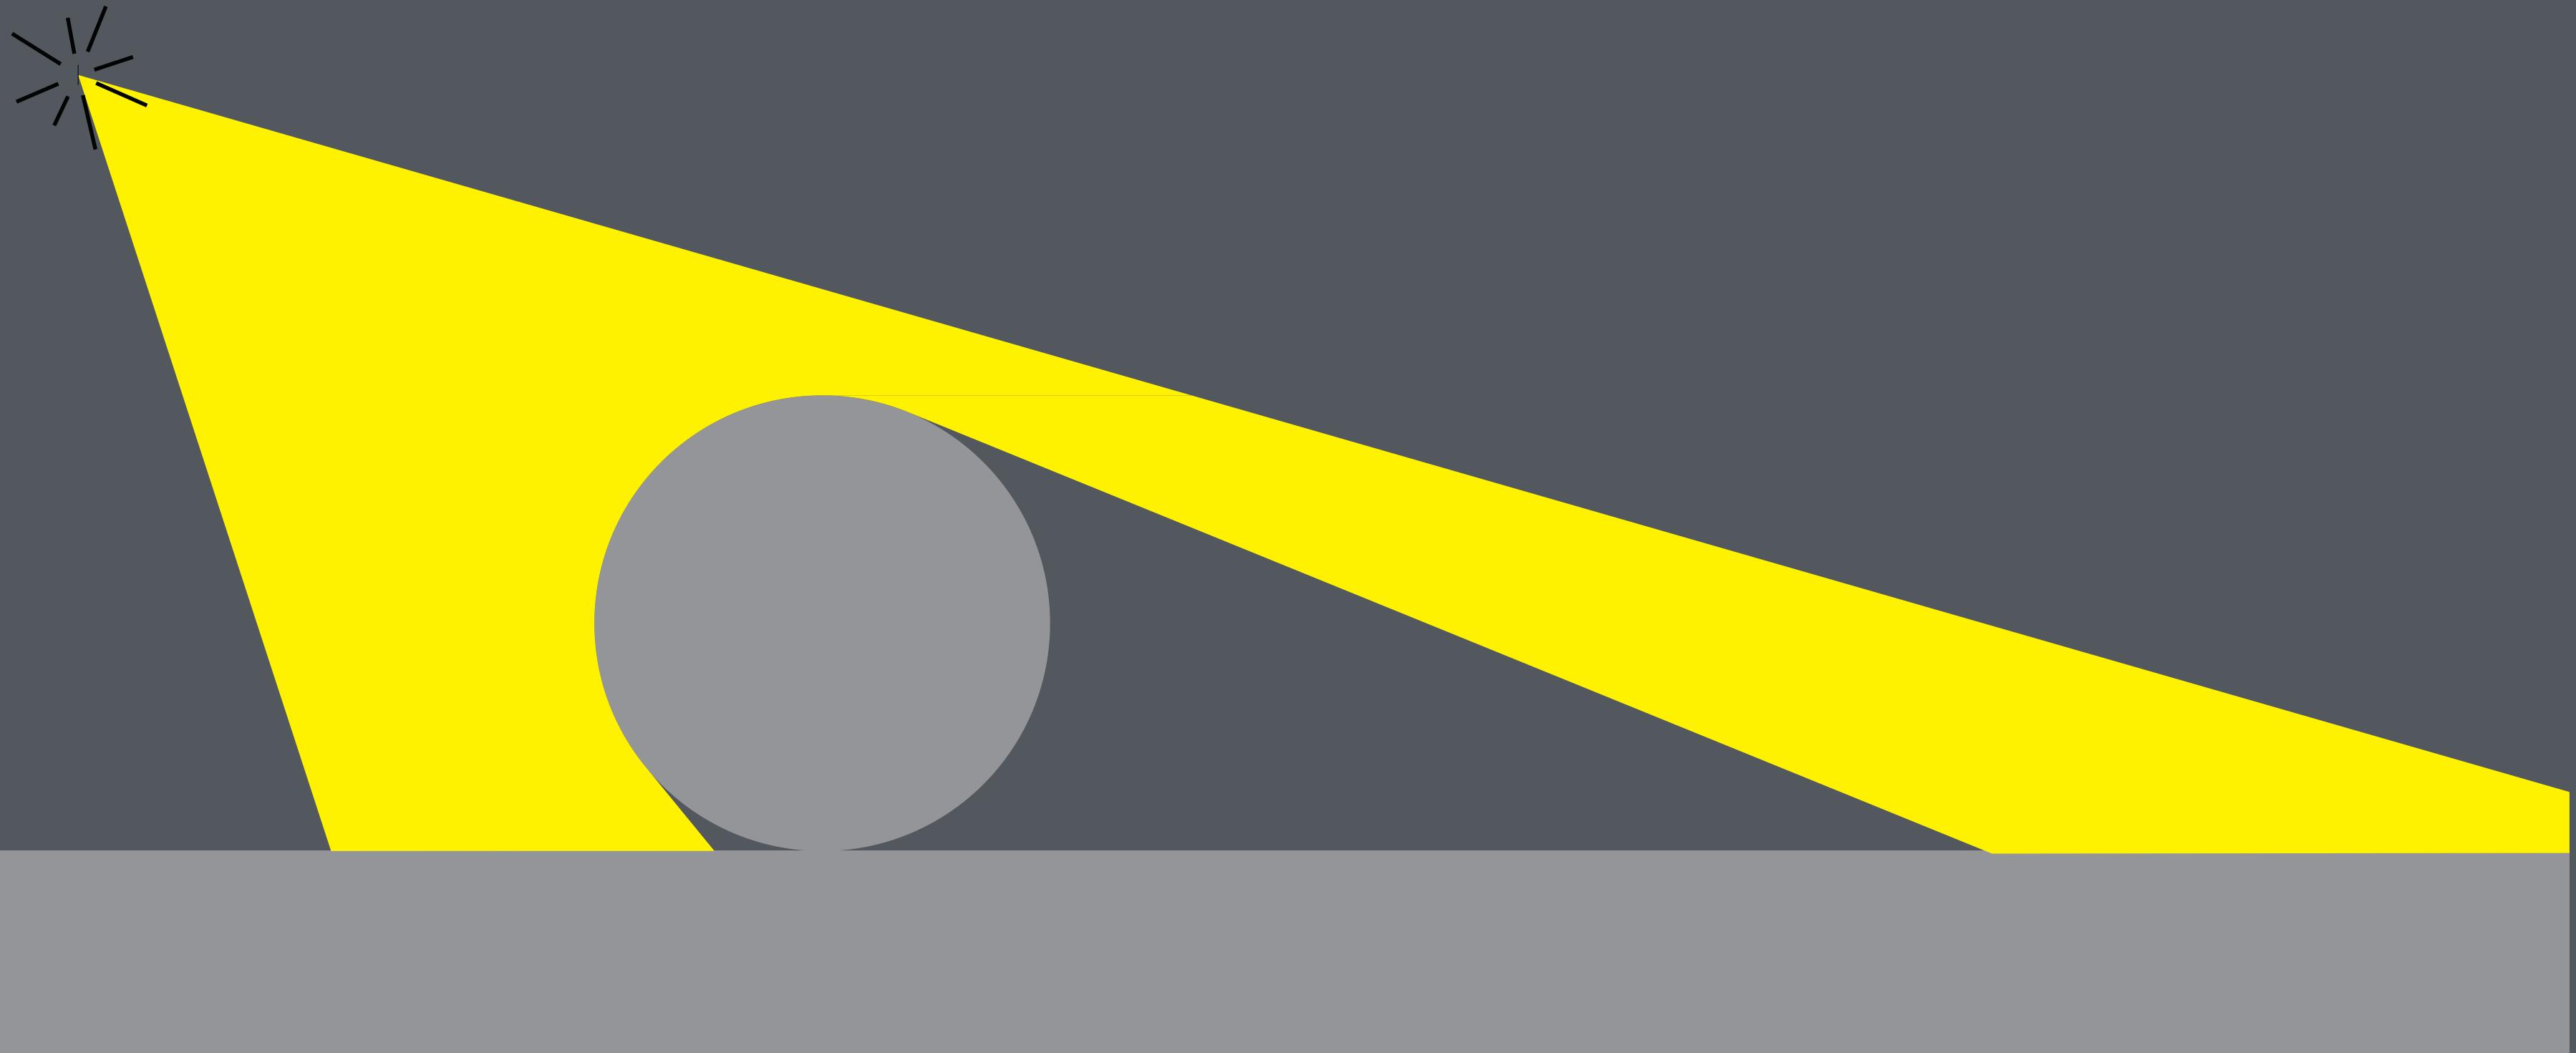
[\[tricks-and-  
illusions.com\]](http://tricks-and-illusions.com)

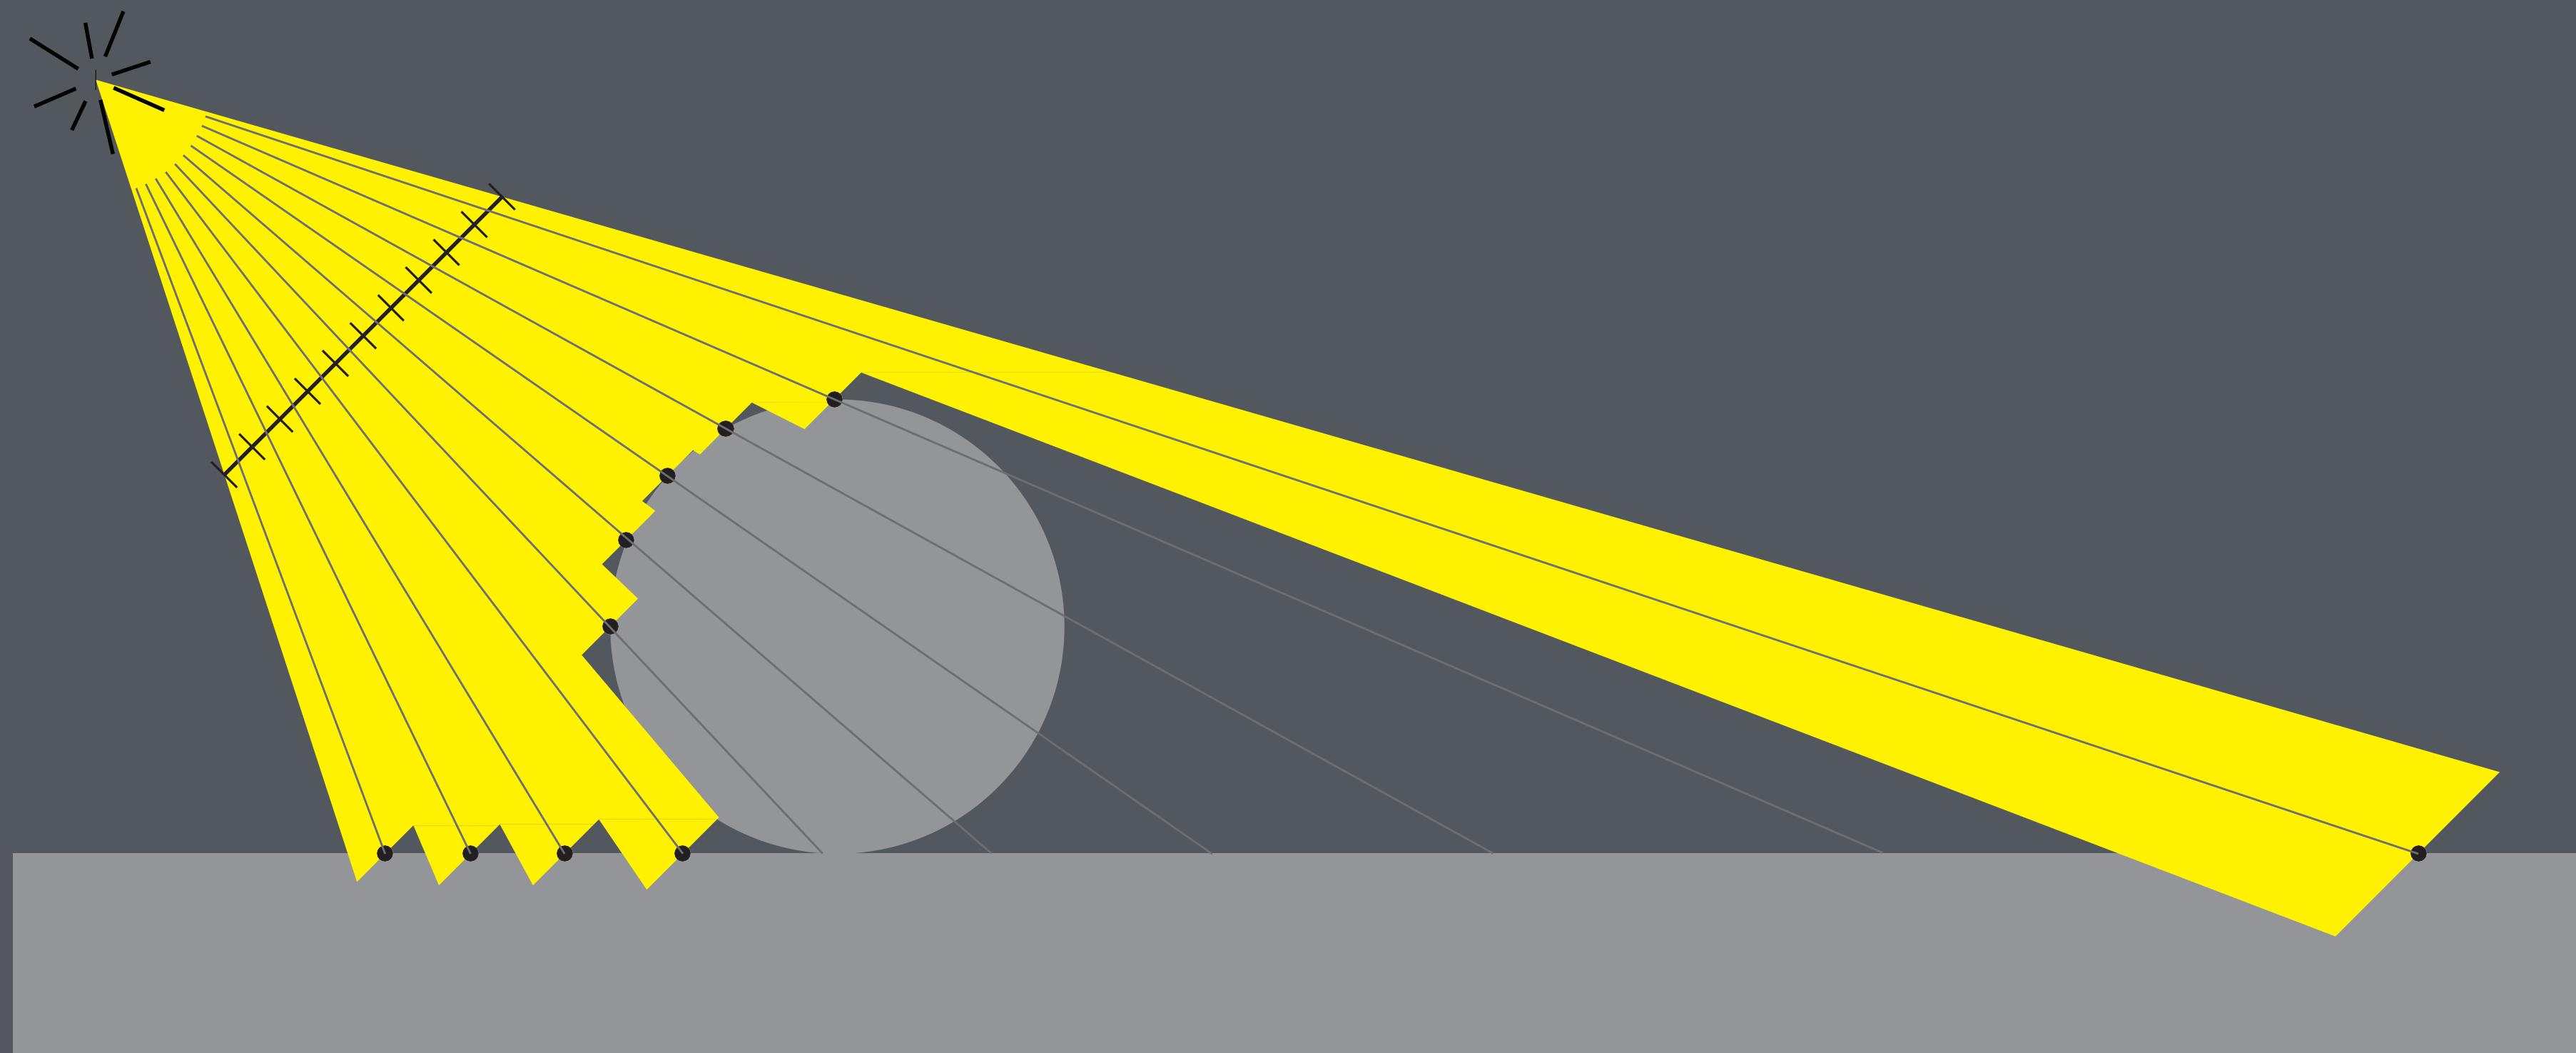
# Shadows as anchors

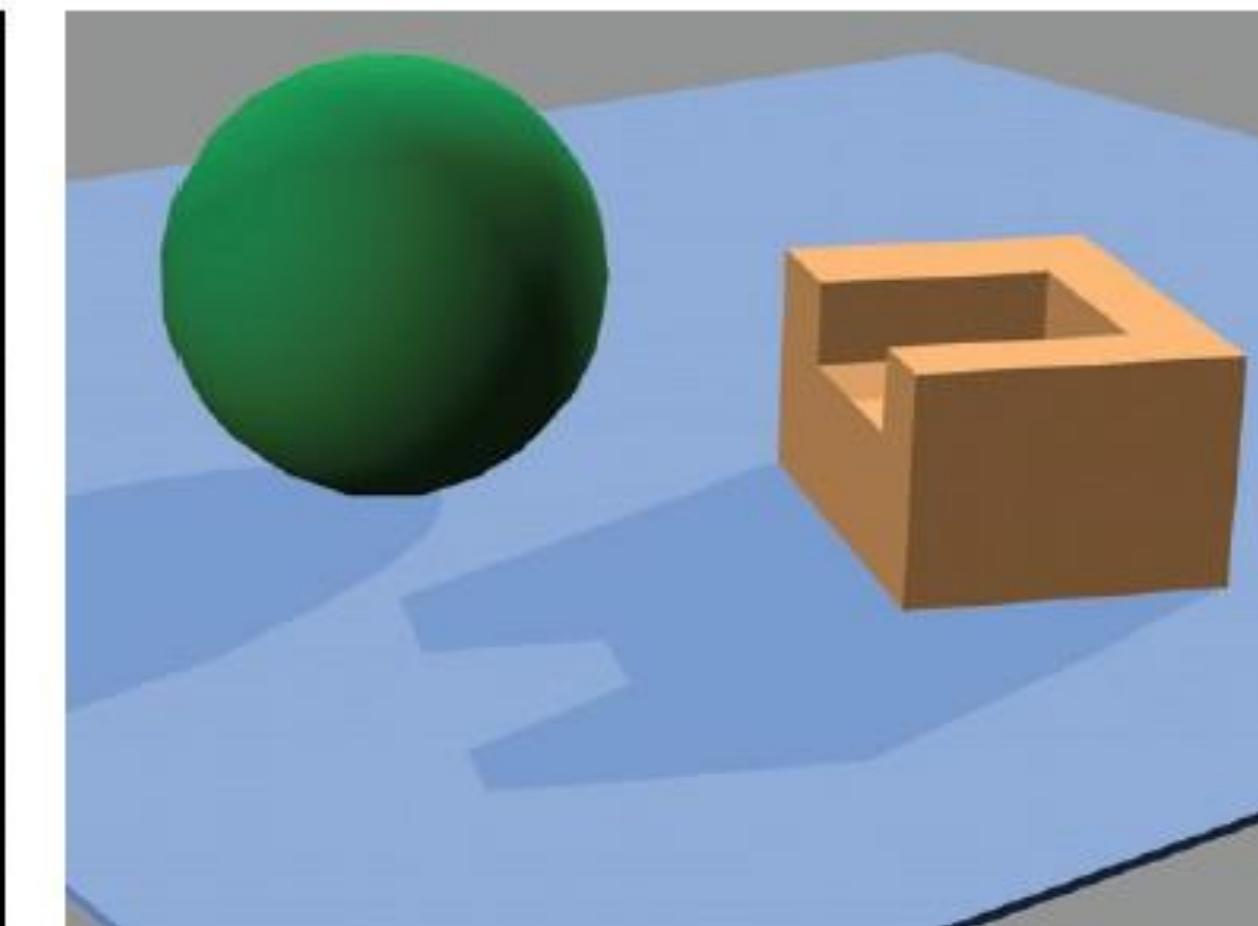
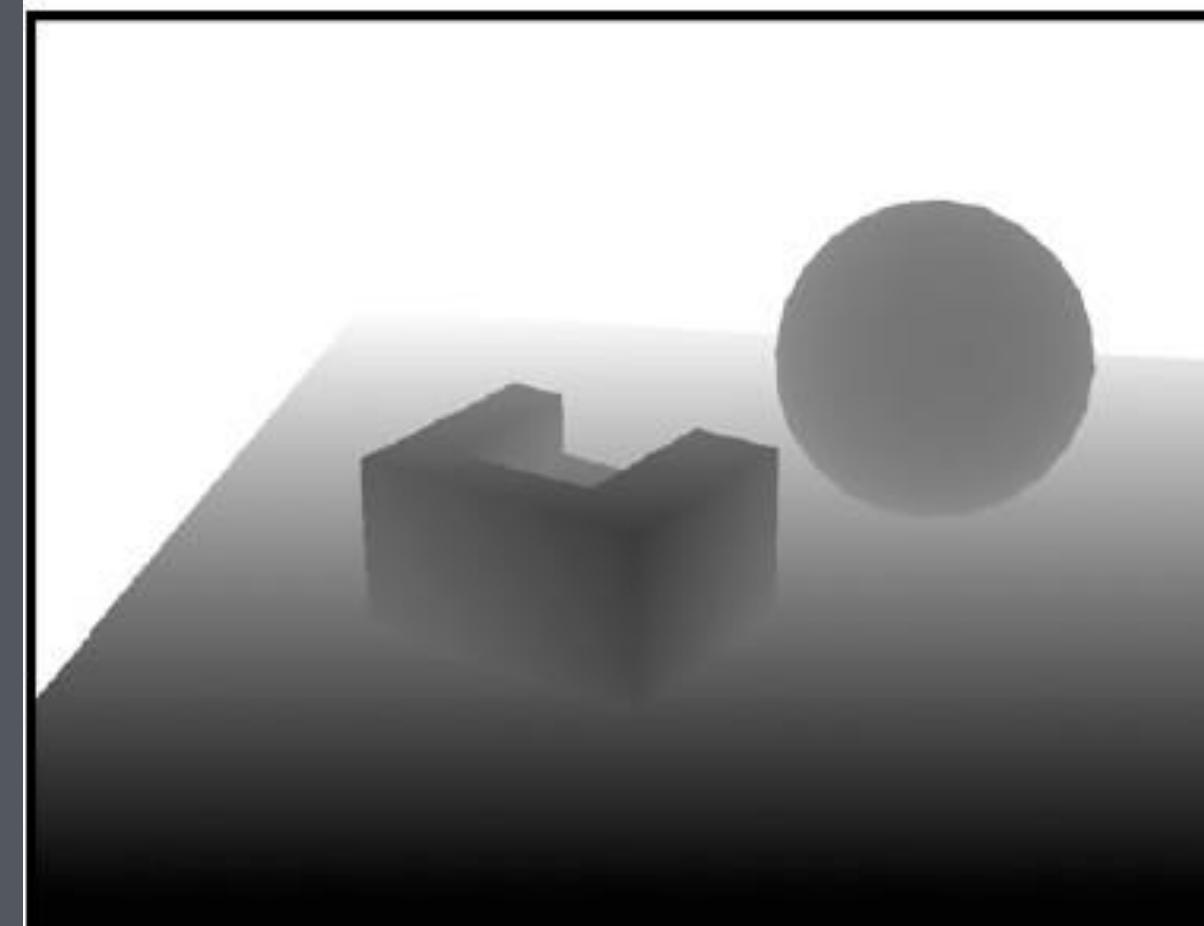
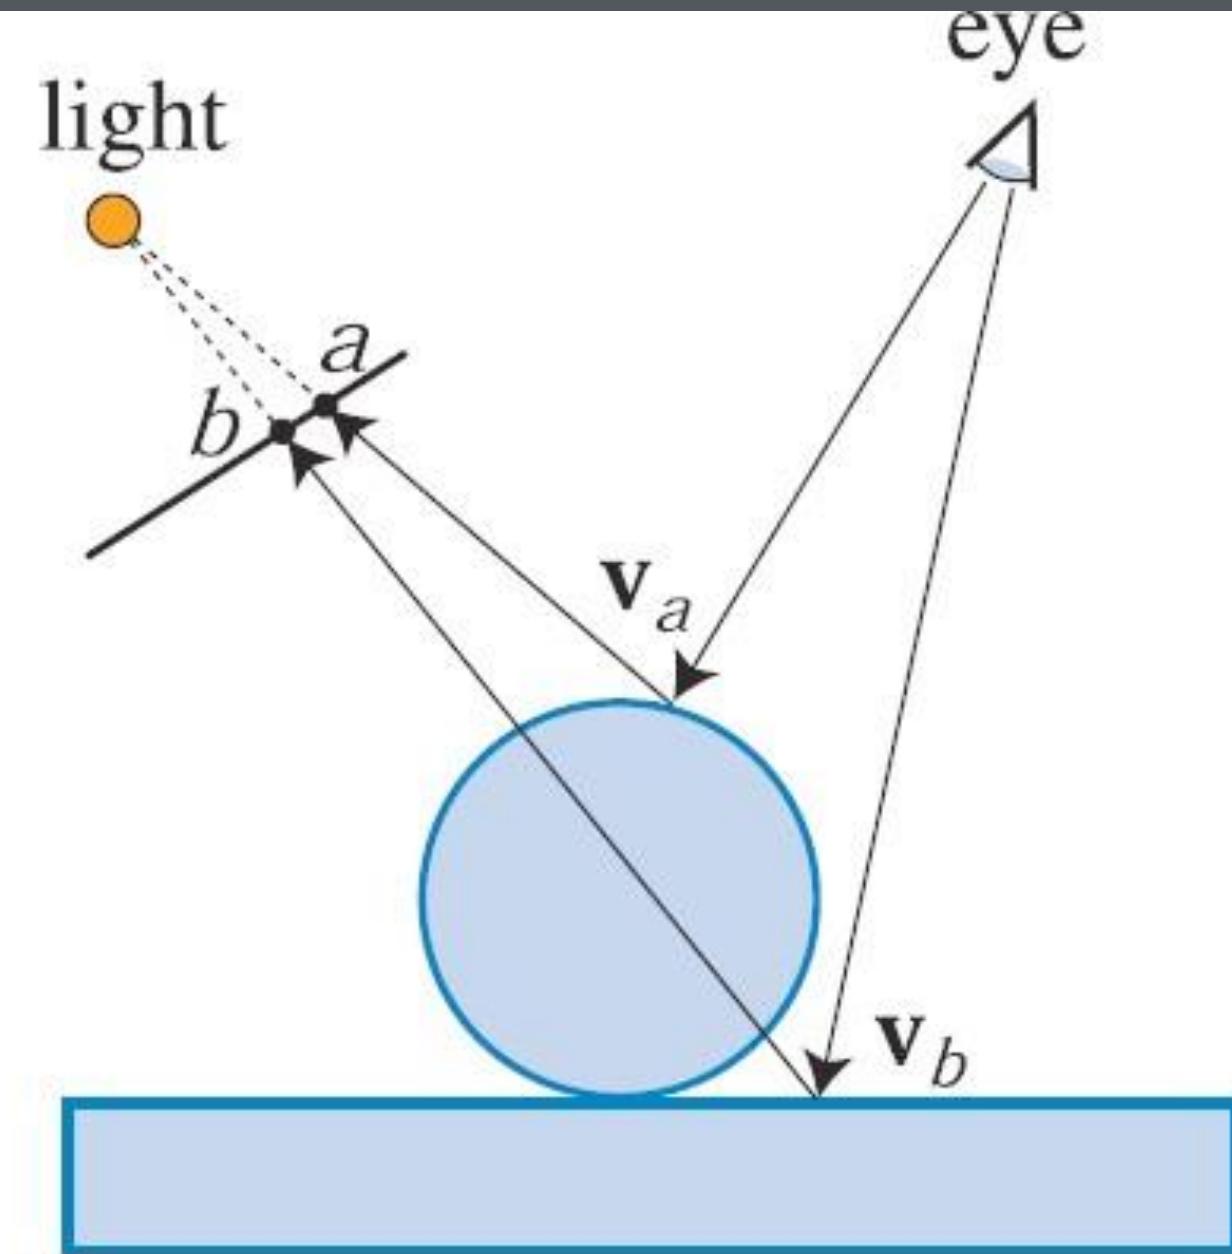
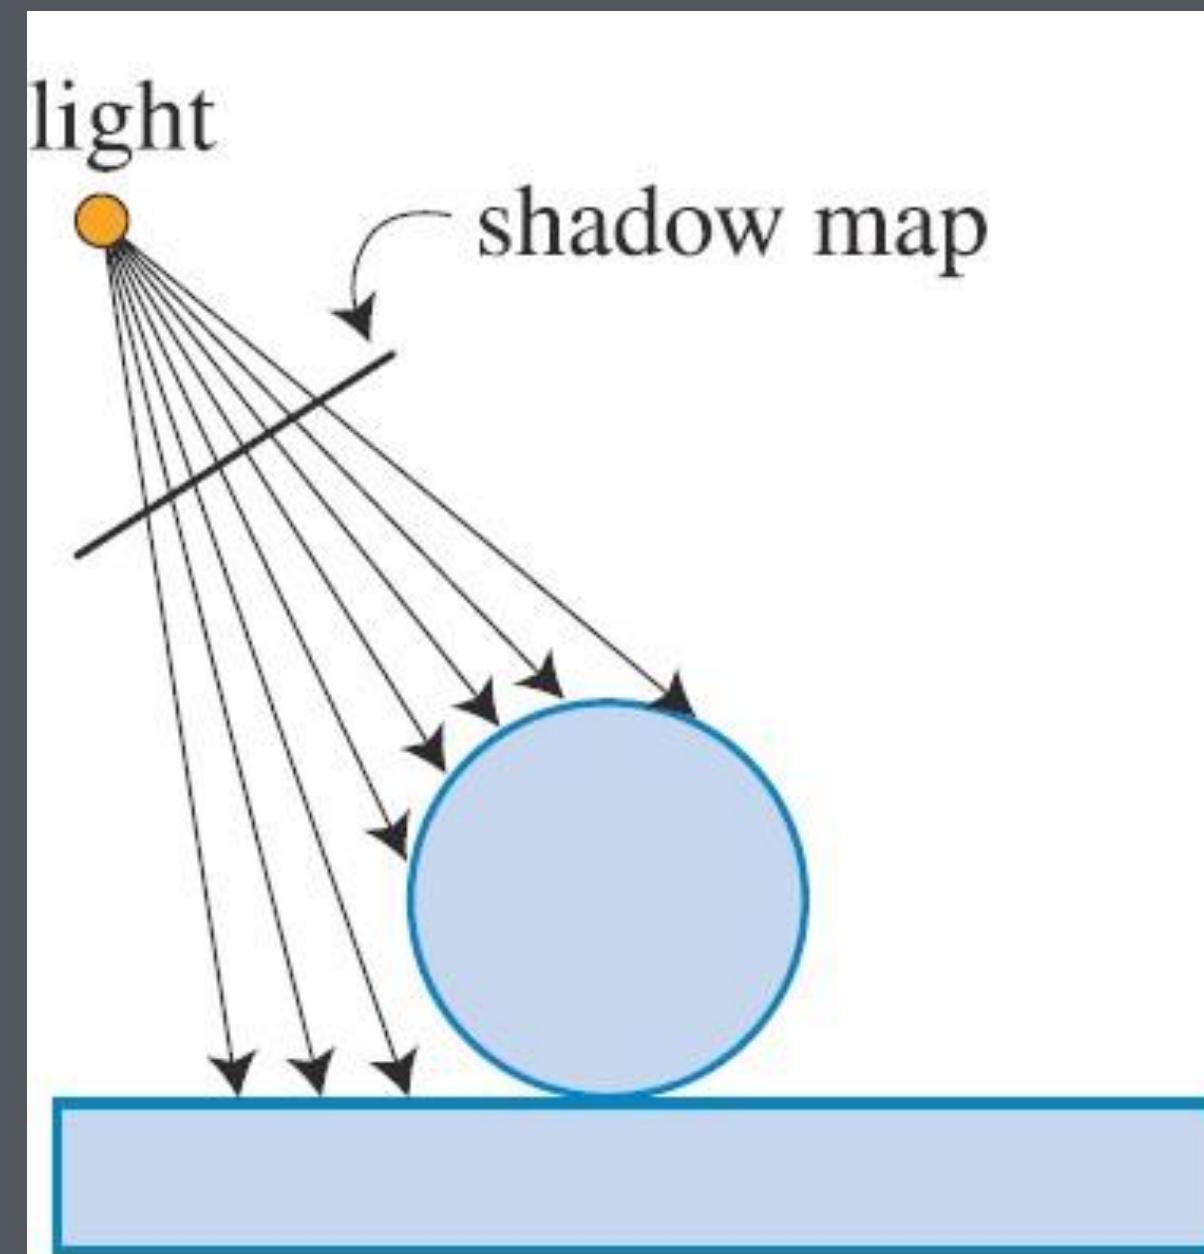


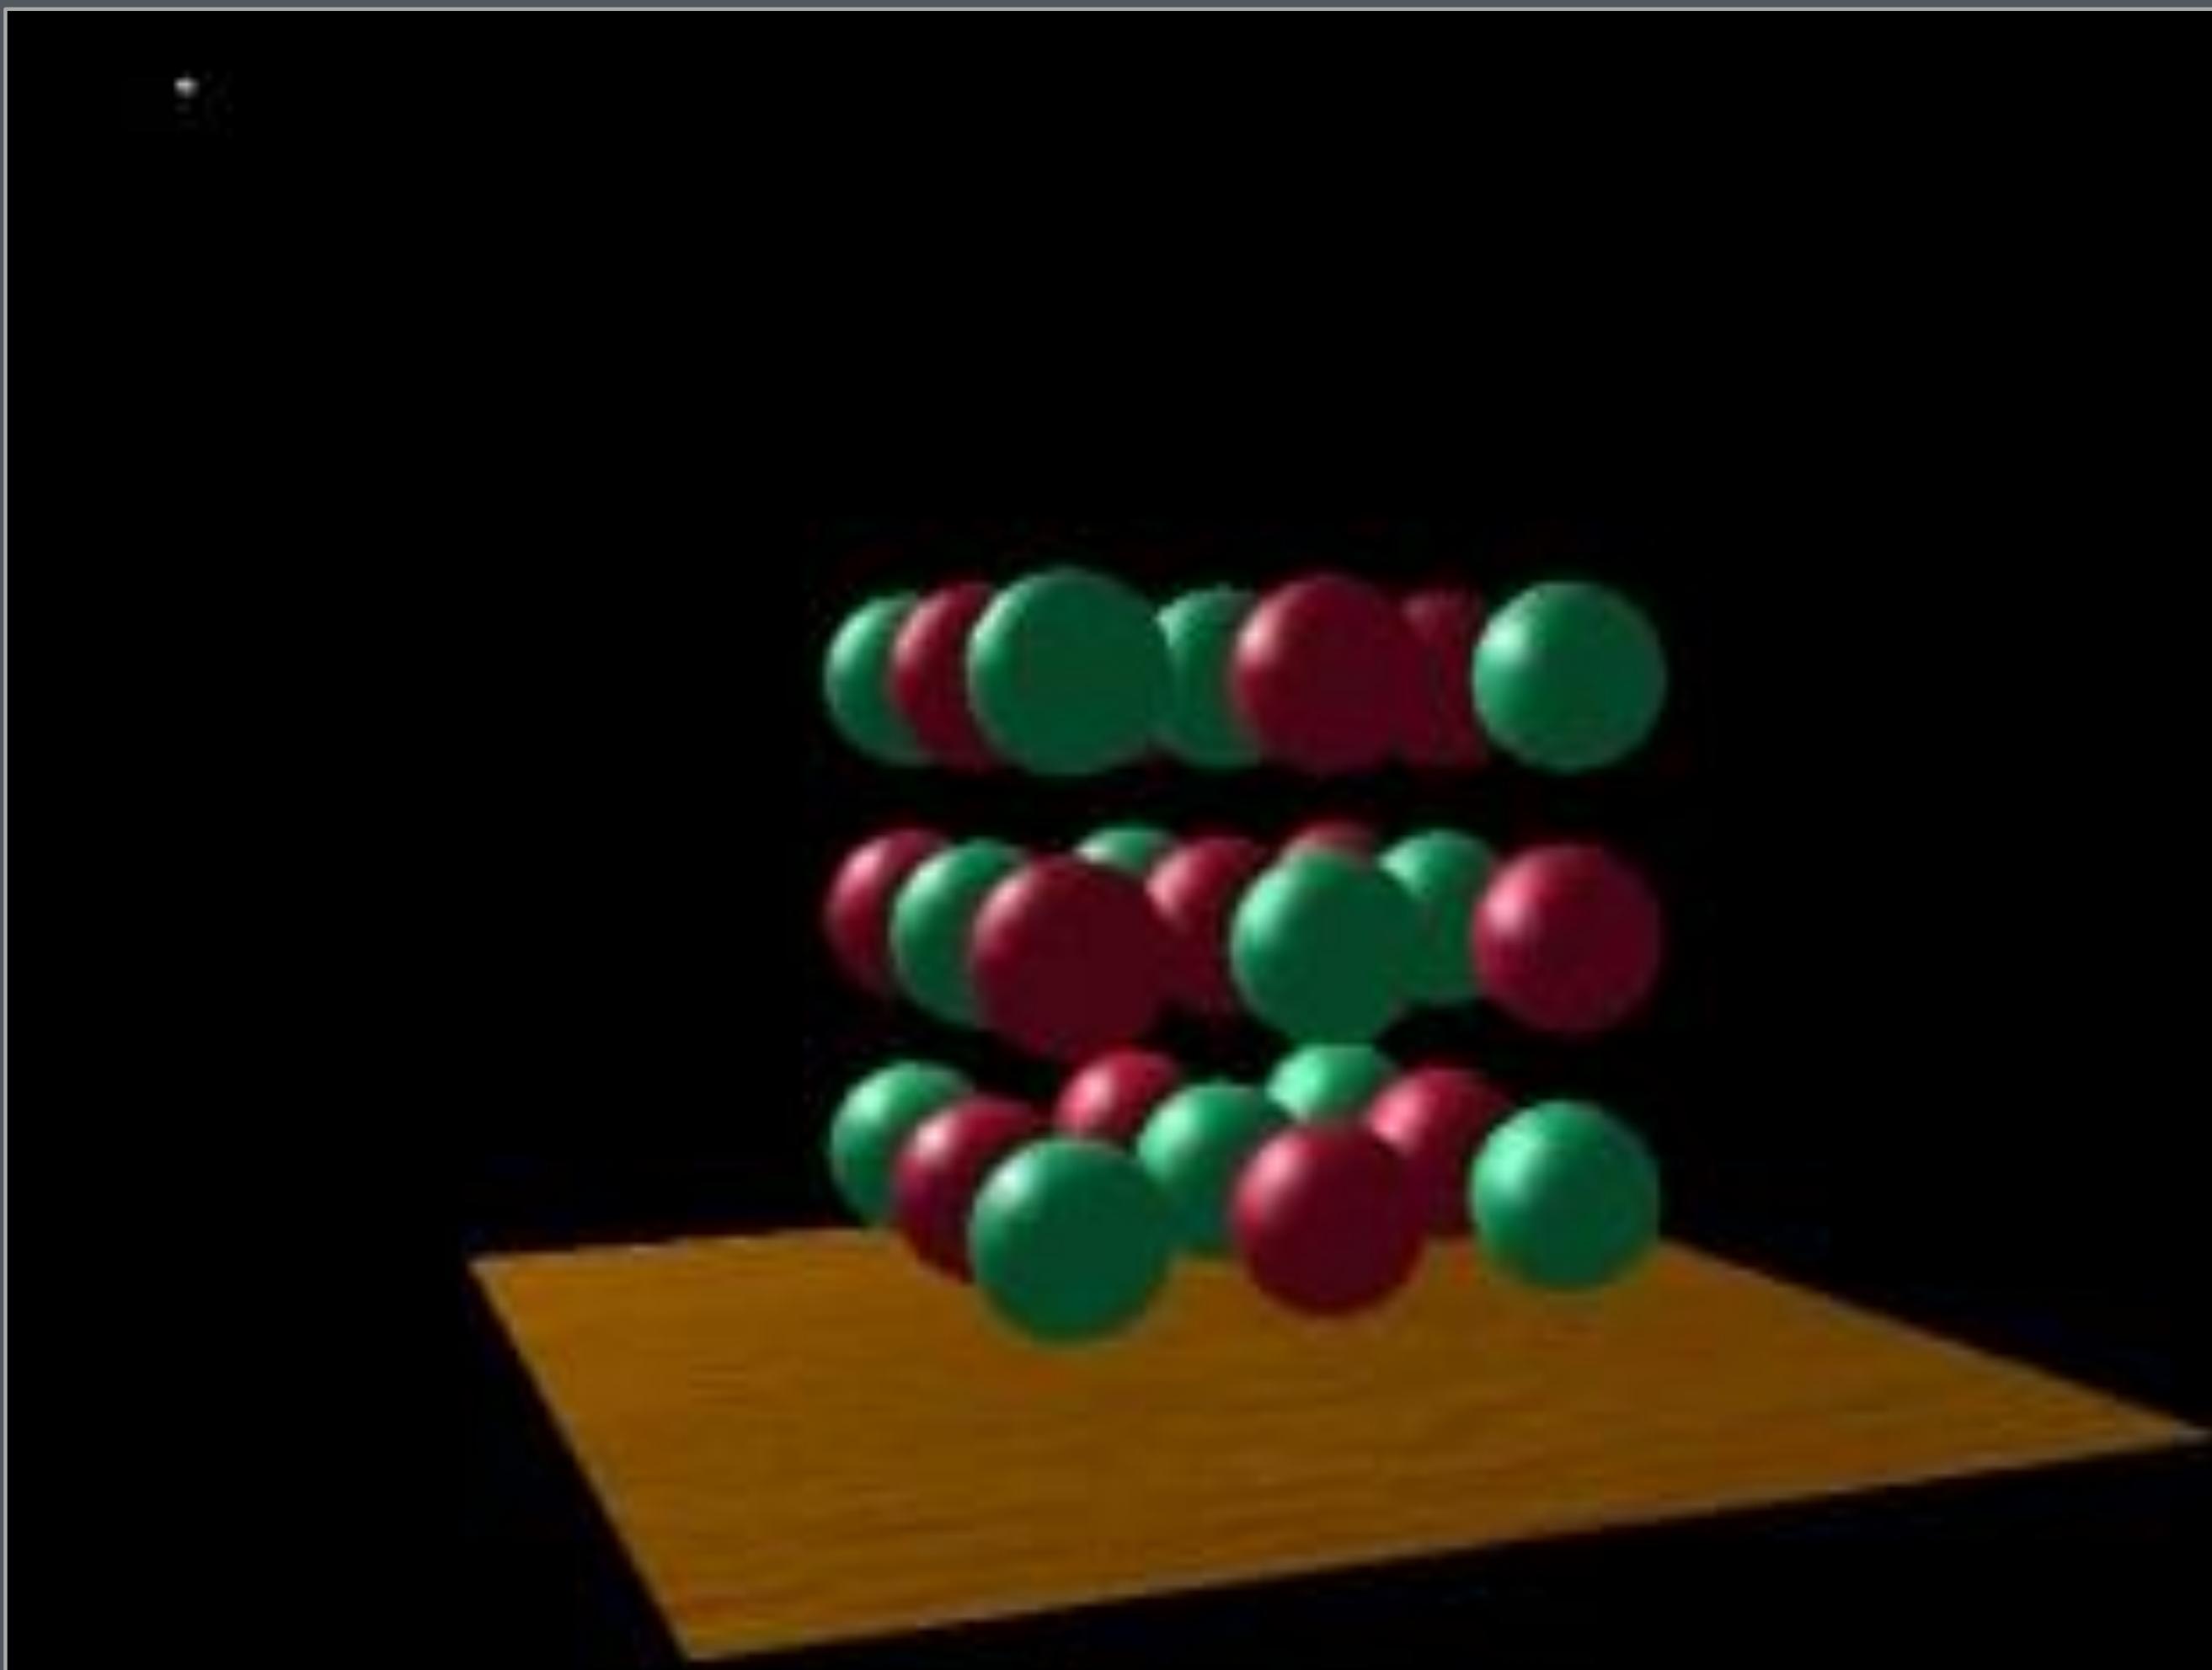
# Shadows as anchors



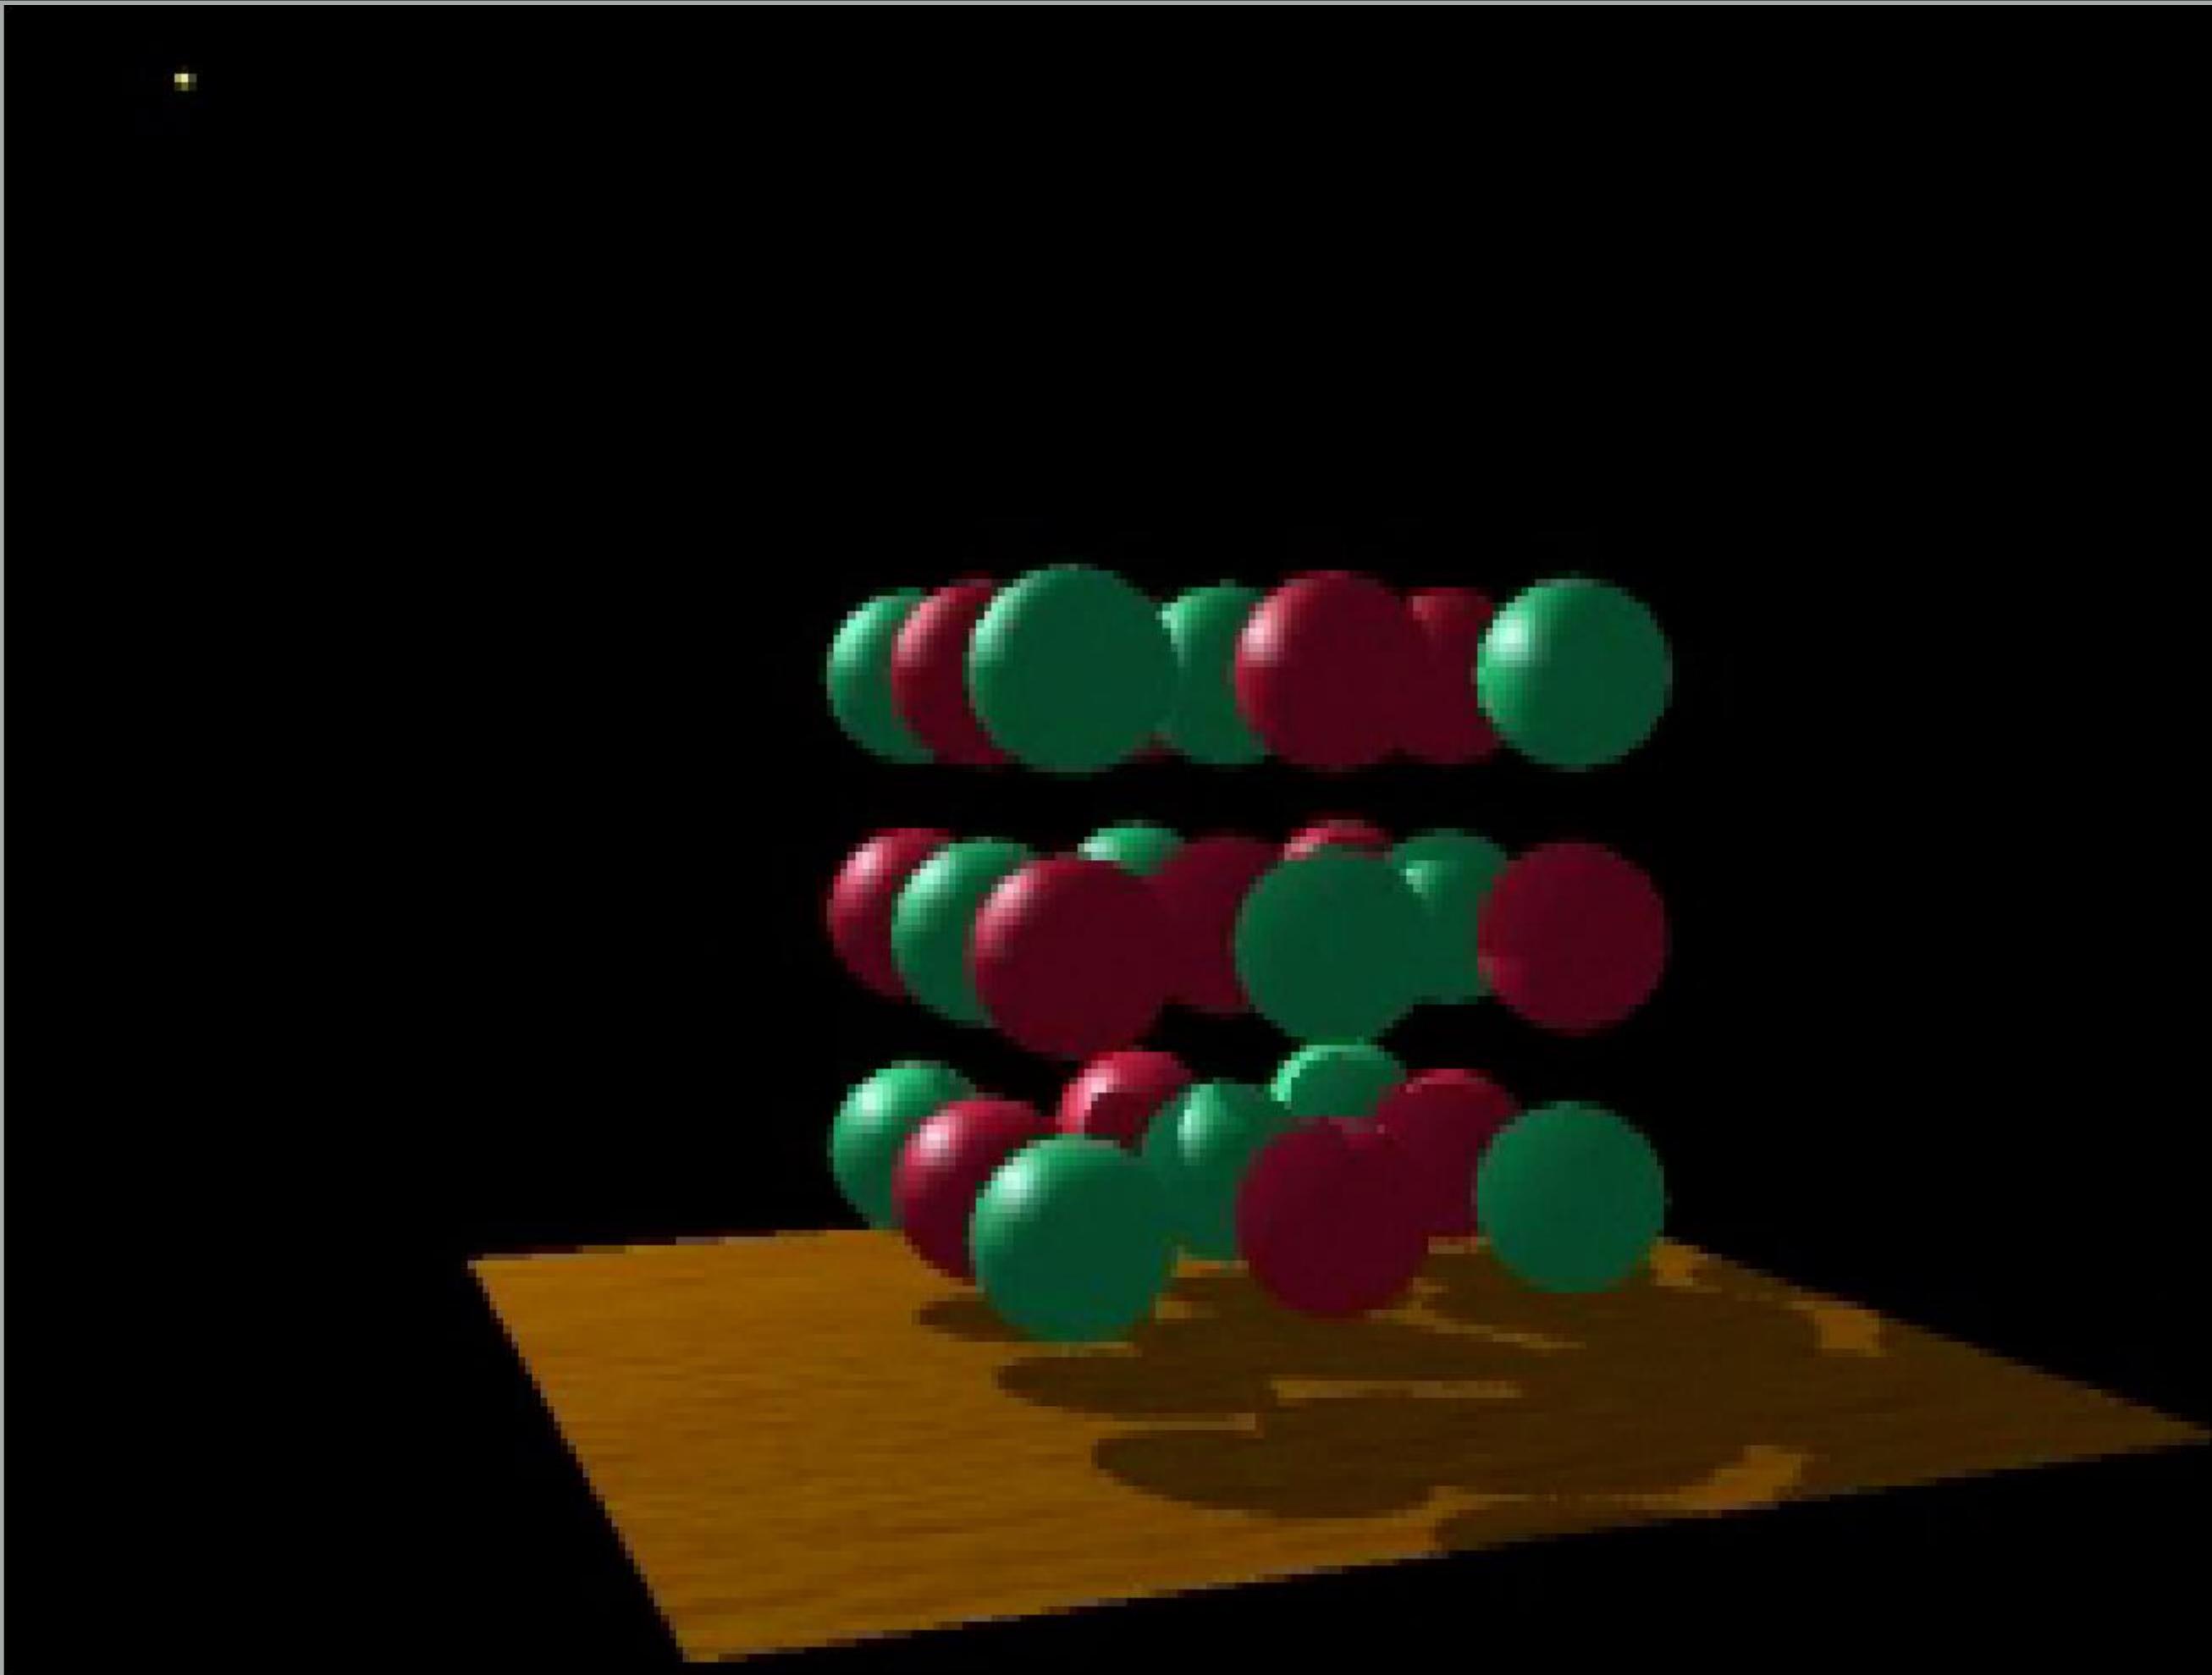




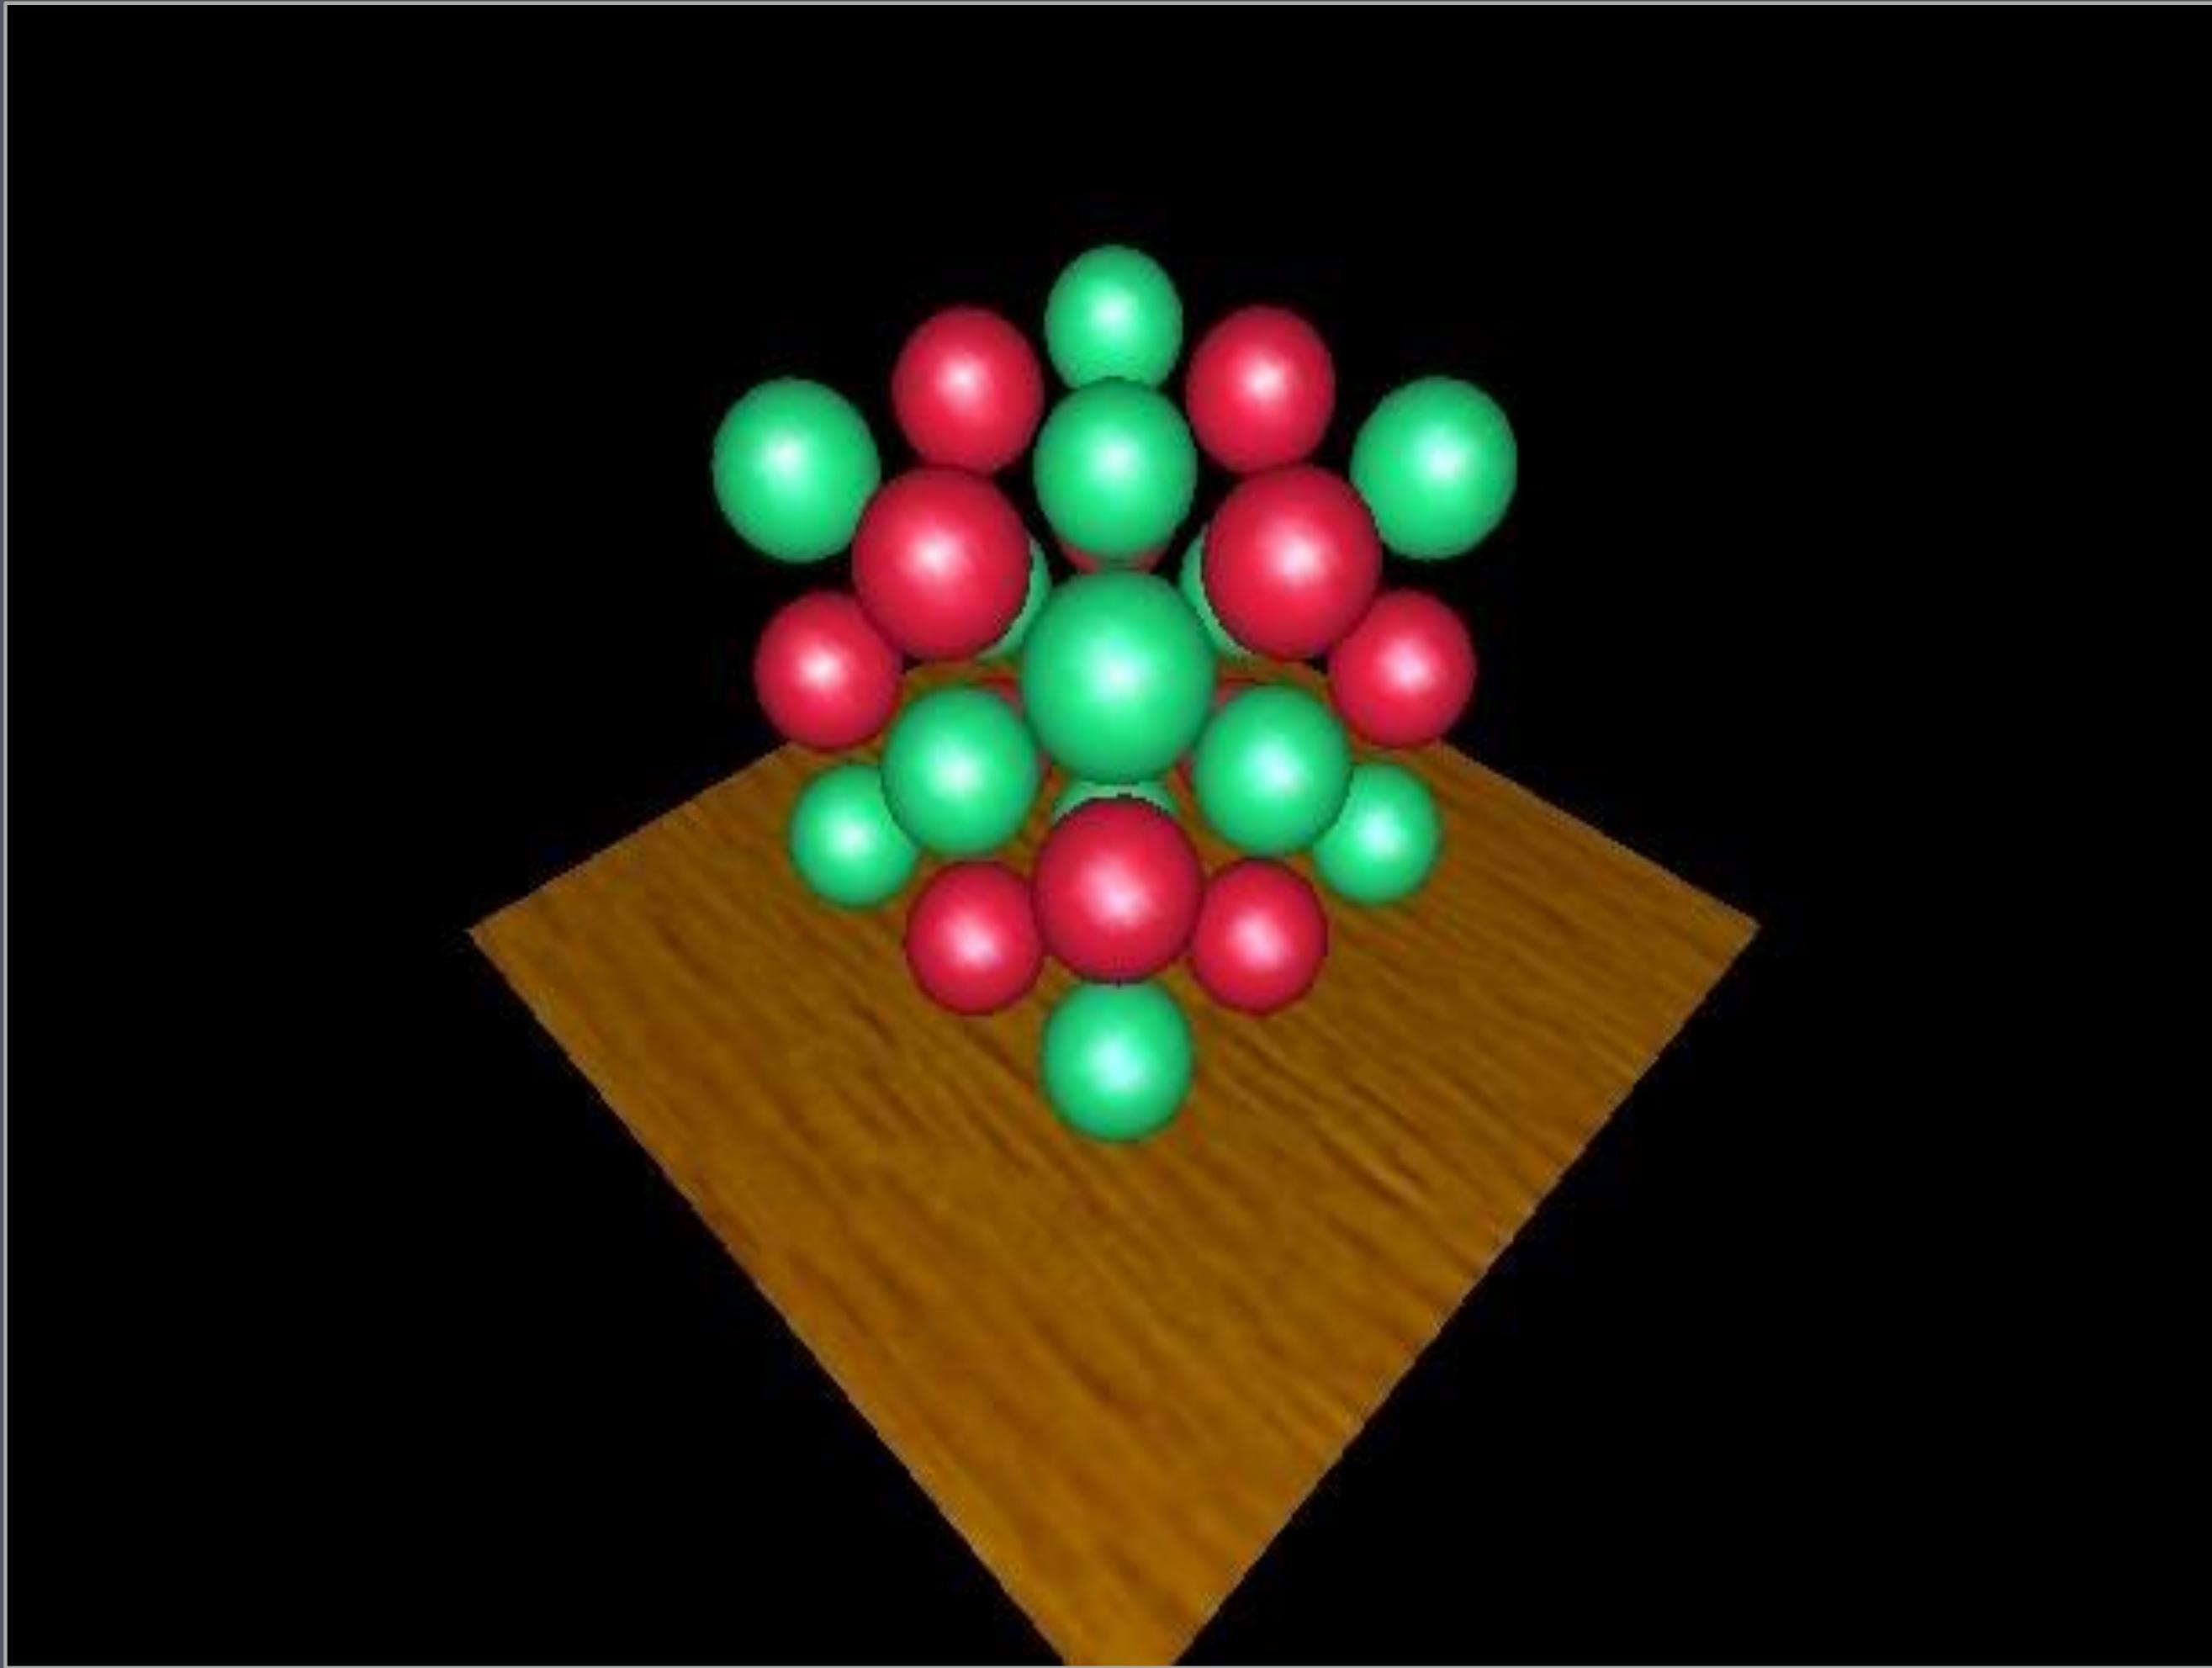




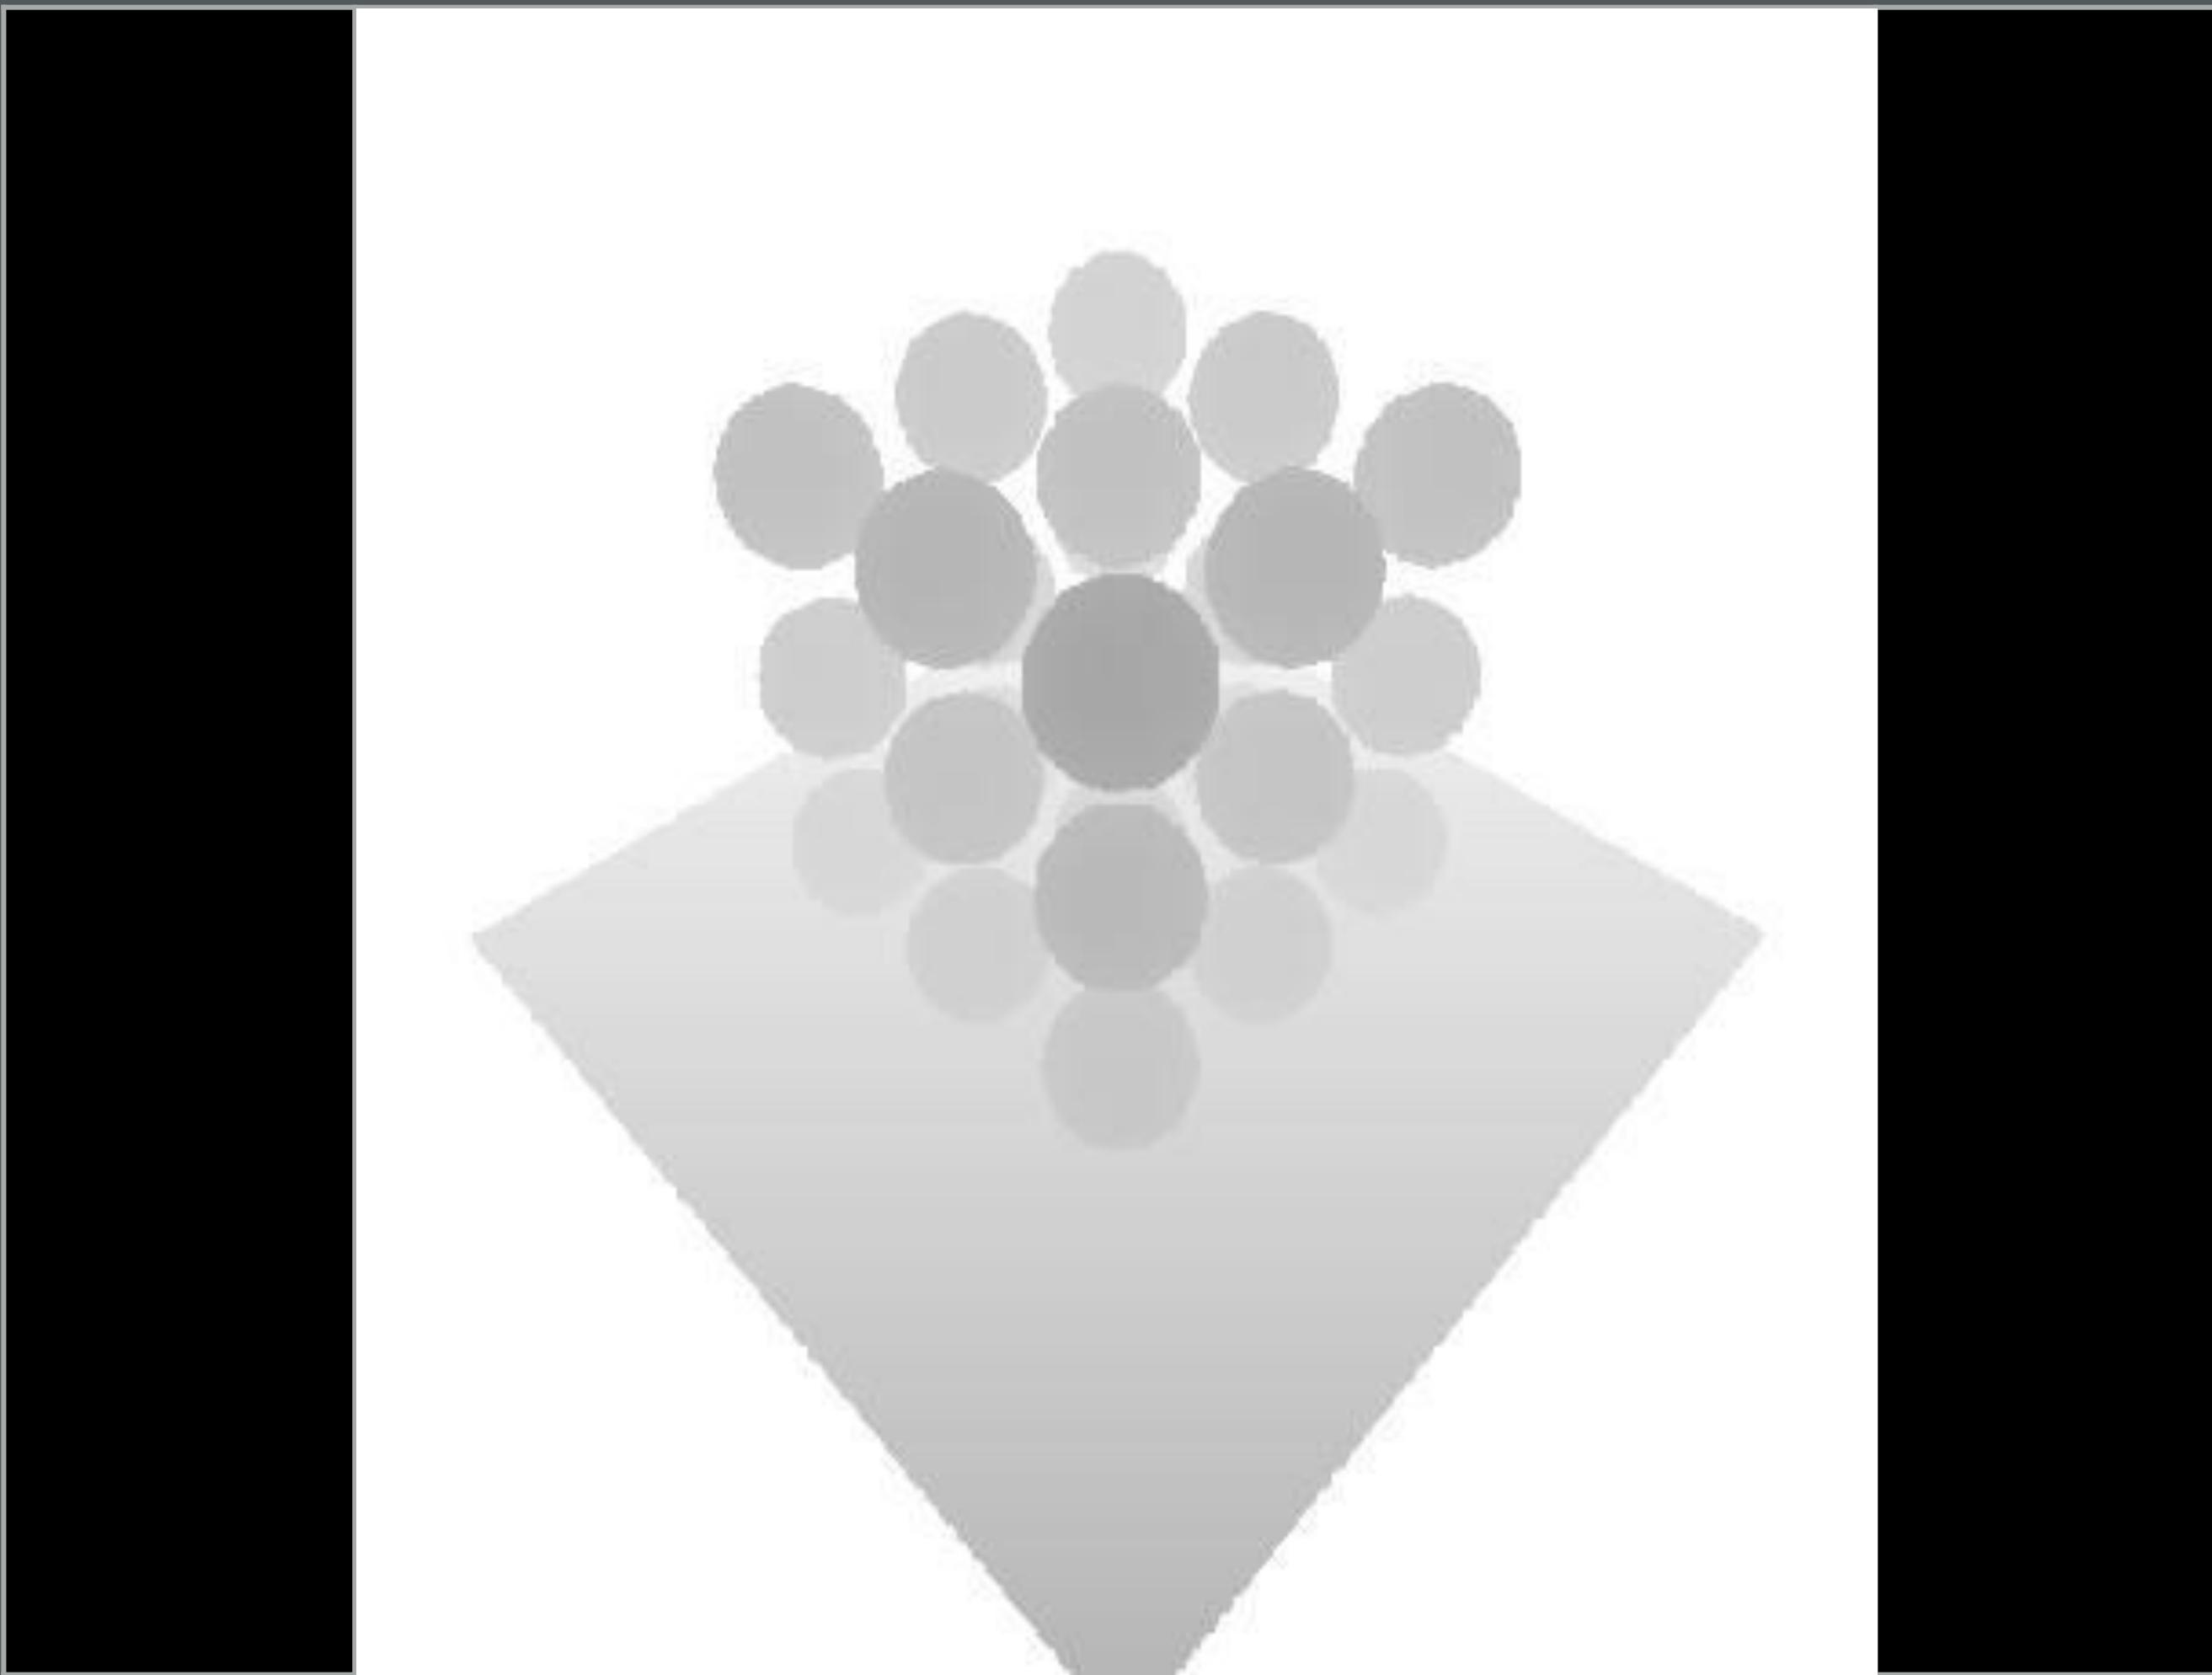
Mark Kilgard



Mark Kilgard



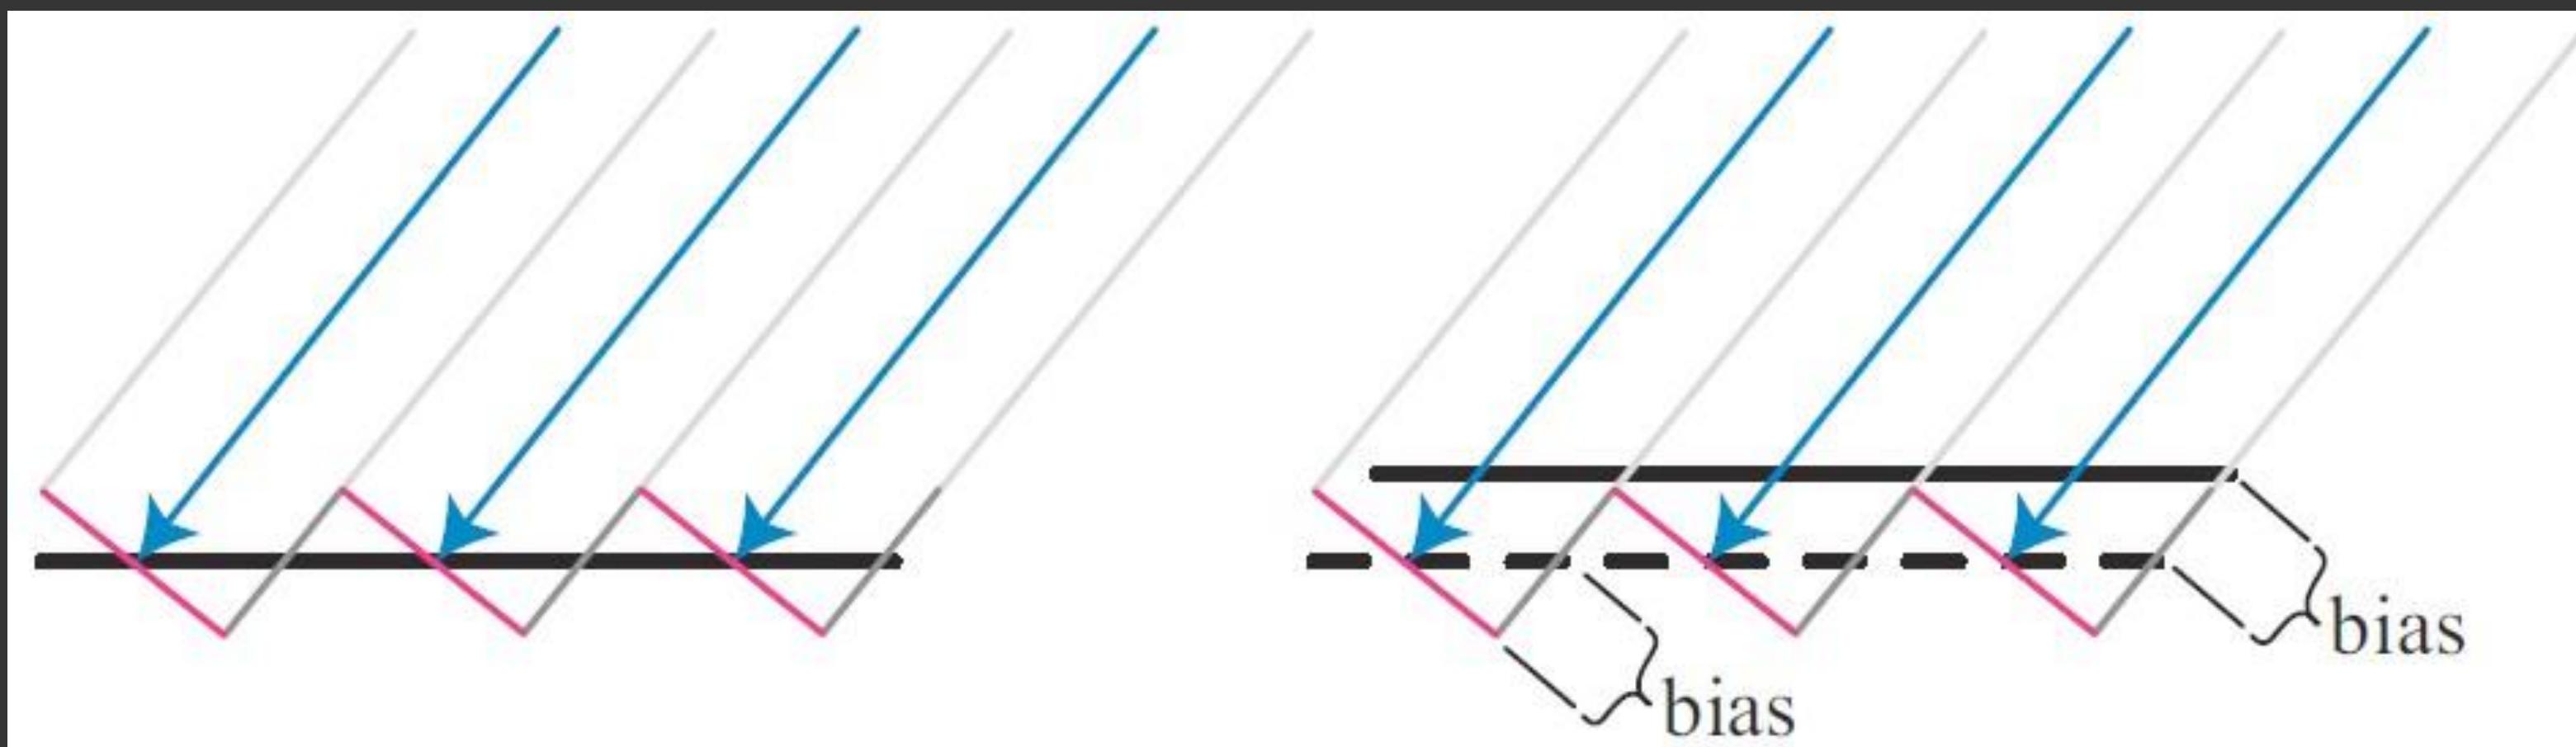
Mark Kilgard

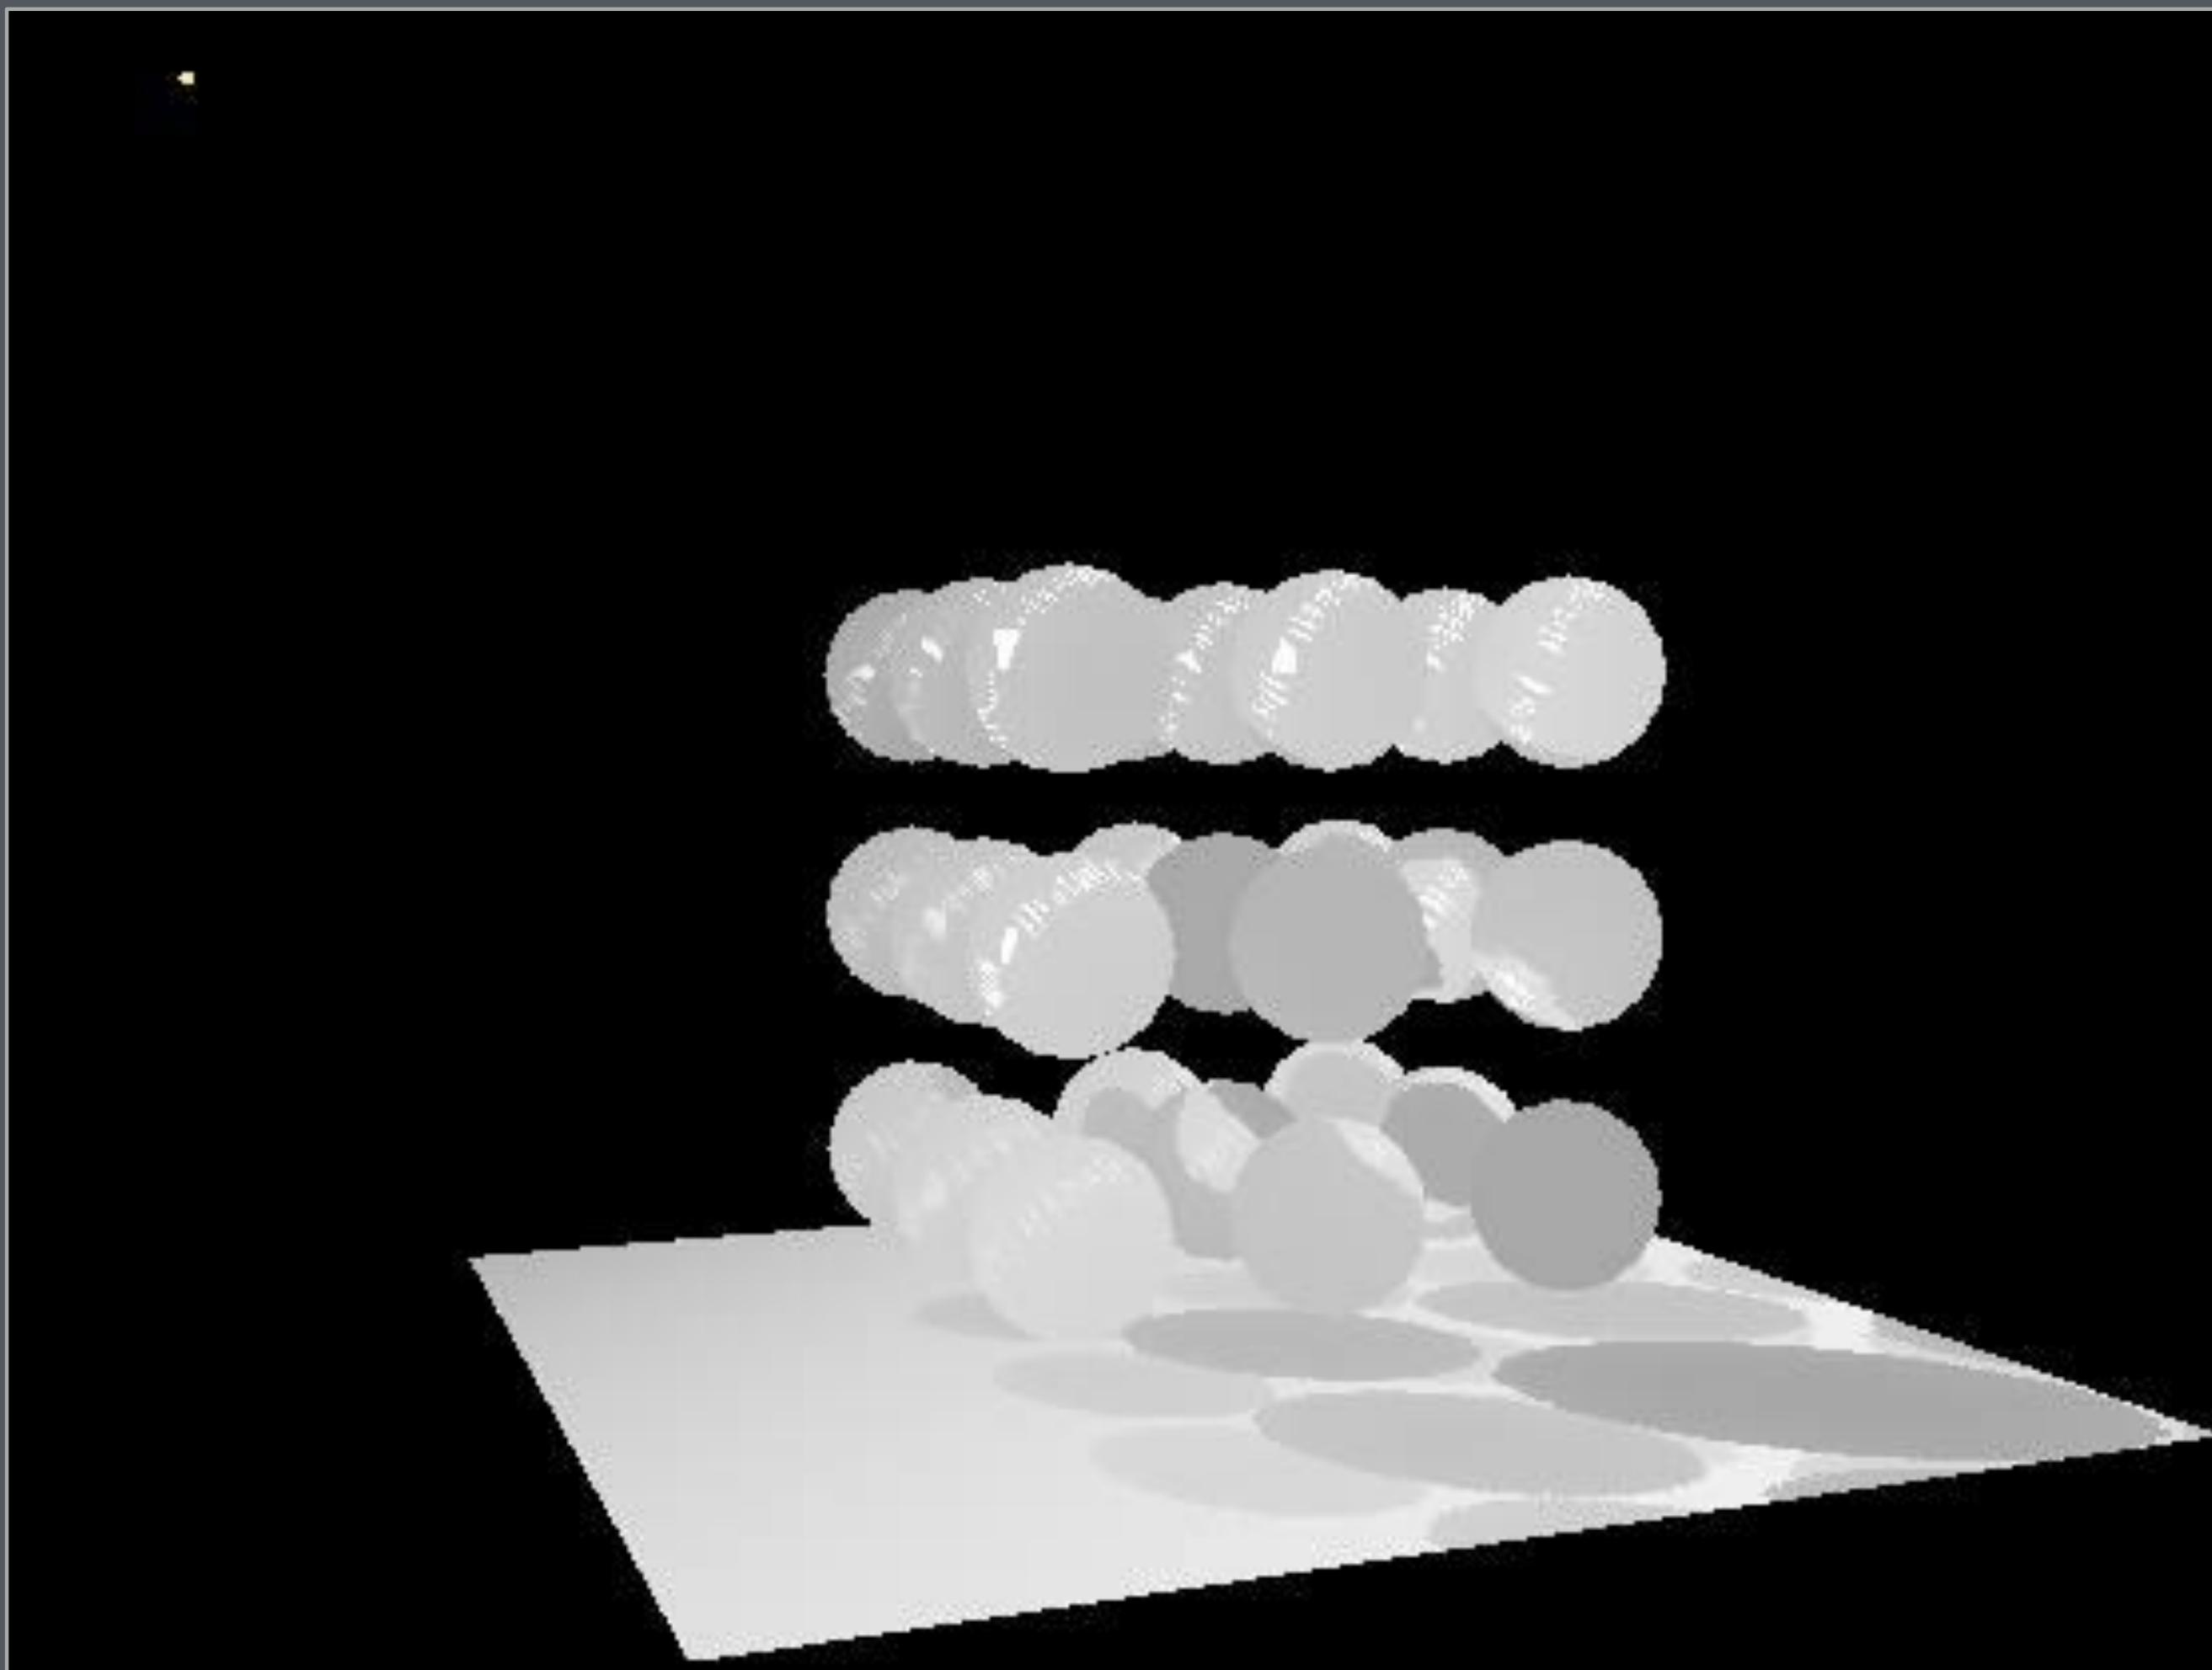


Mark Kilgard

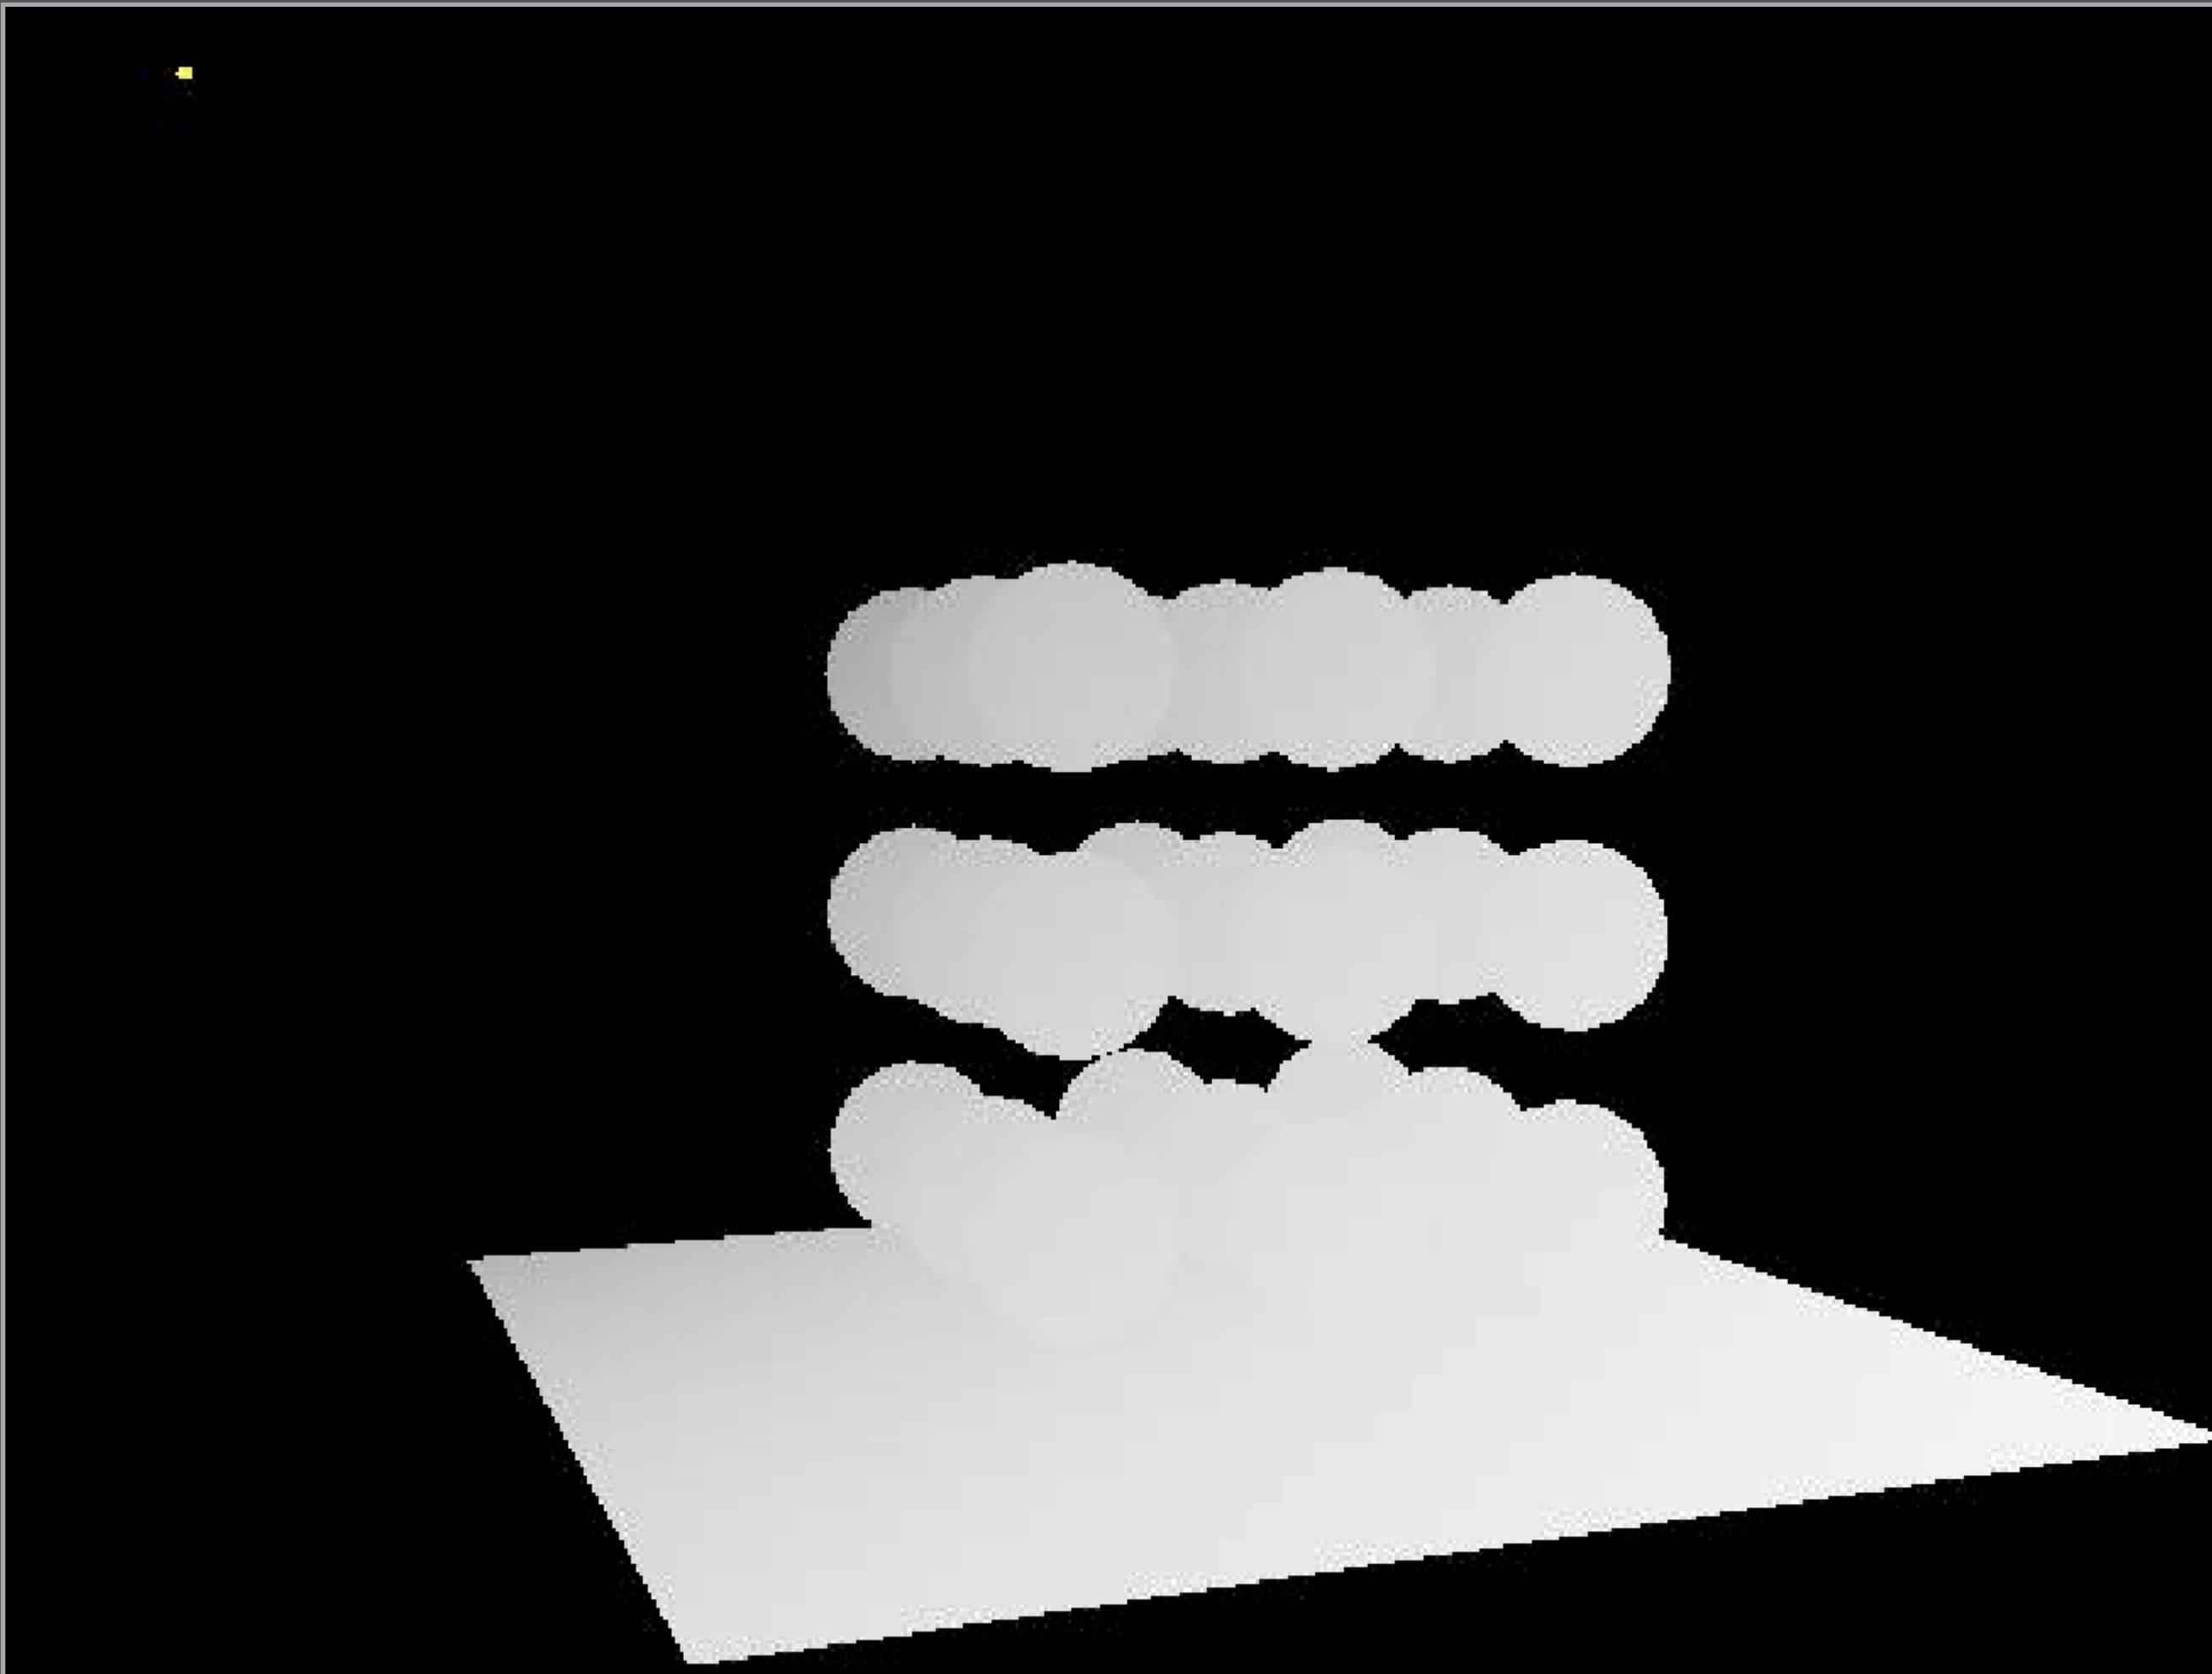
# Shadow Map Issues

- if A and B are approximately equal?
- Speckling

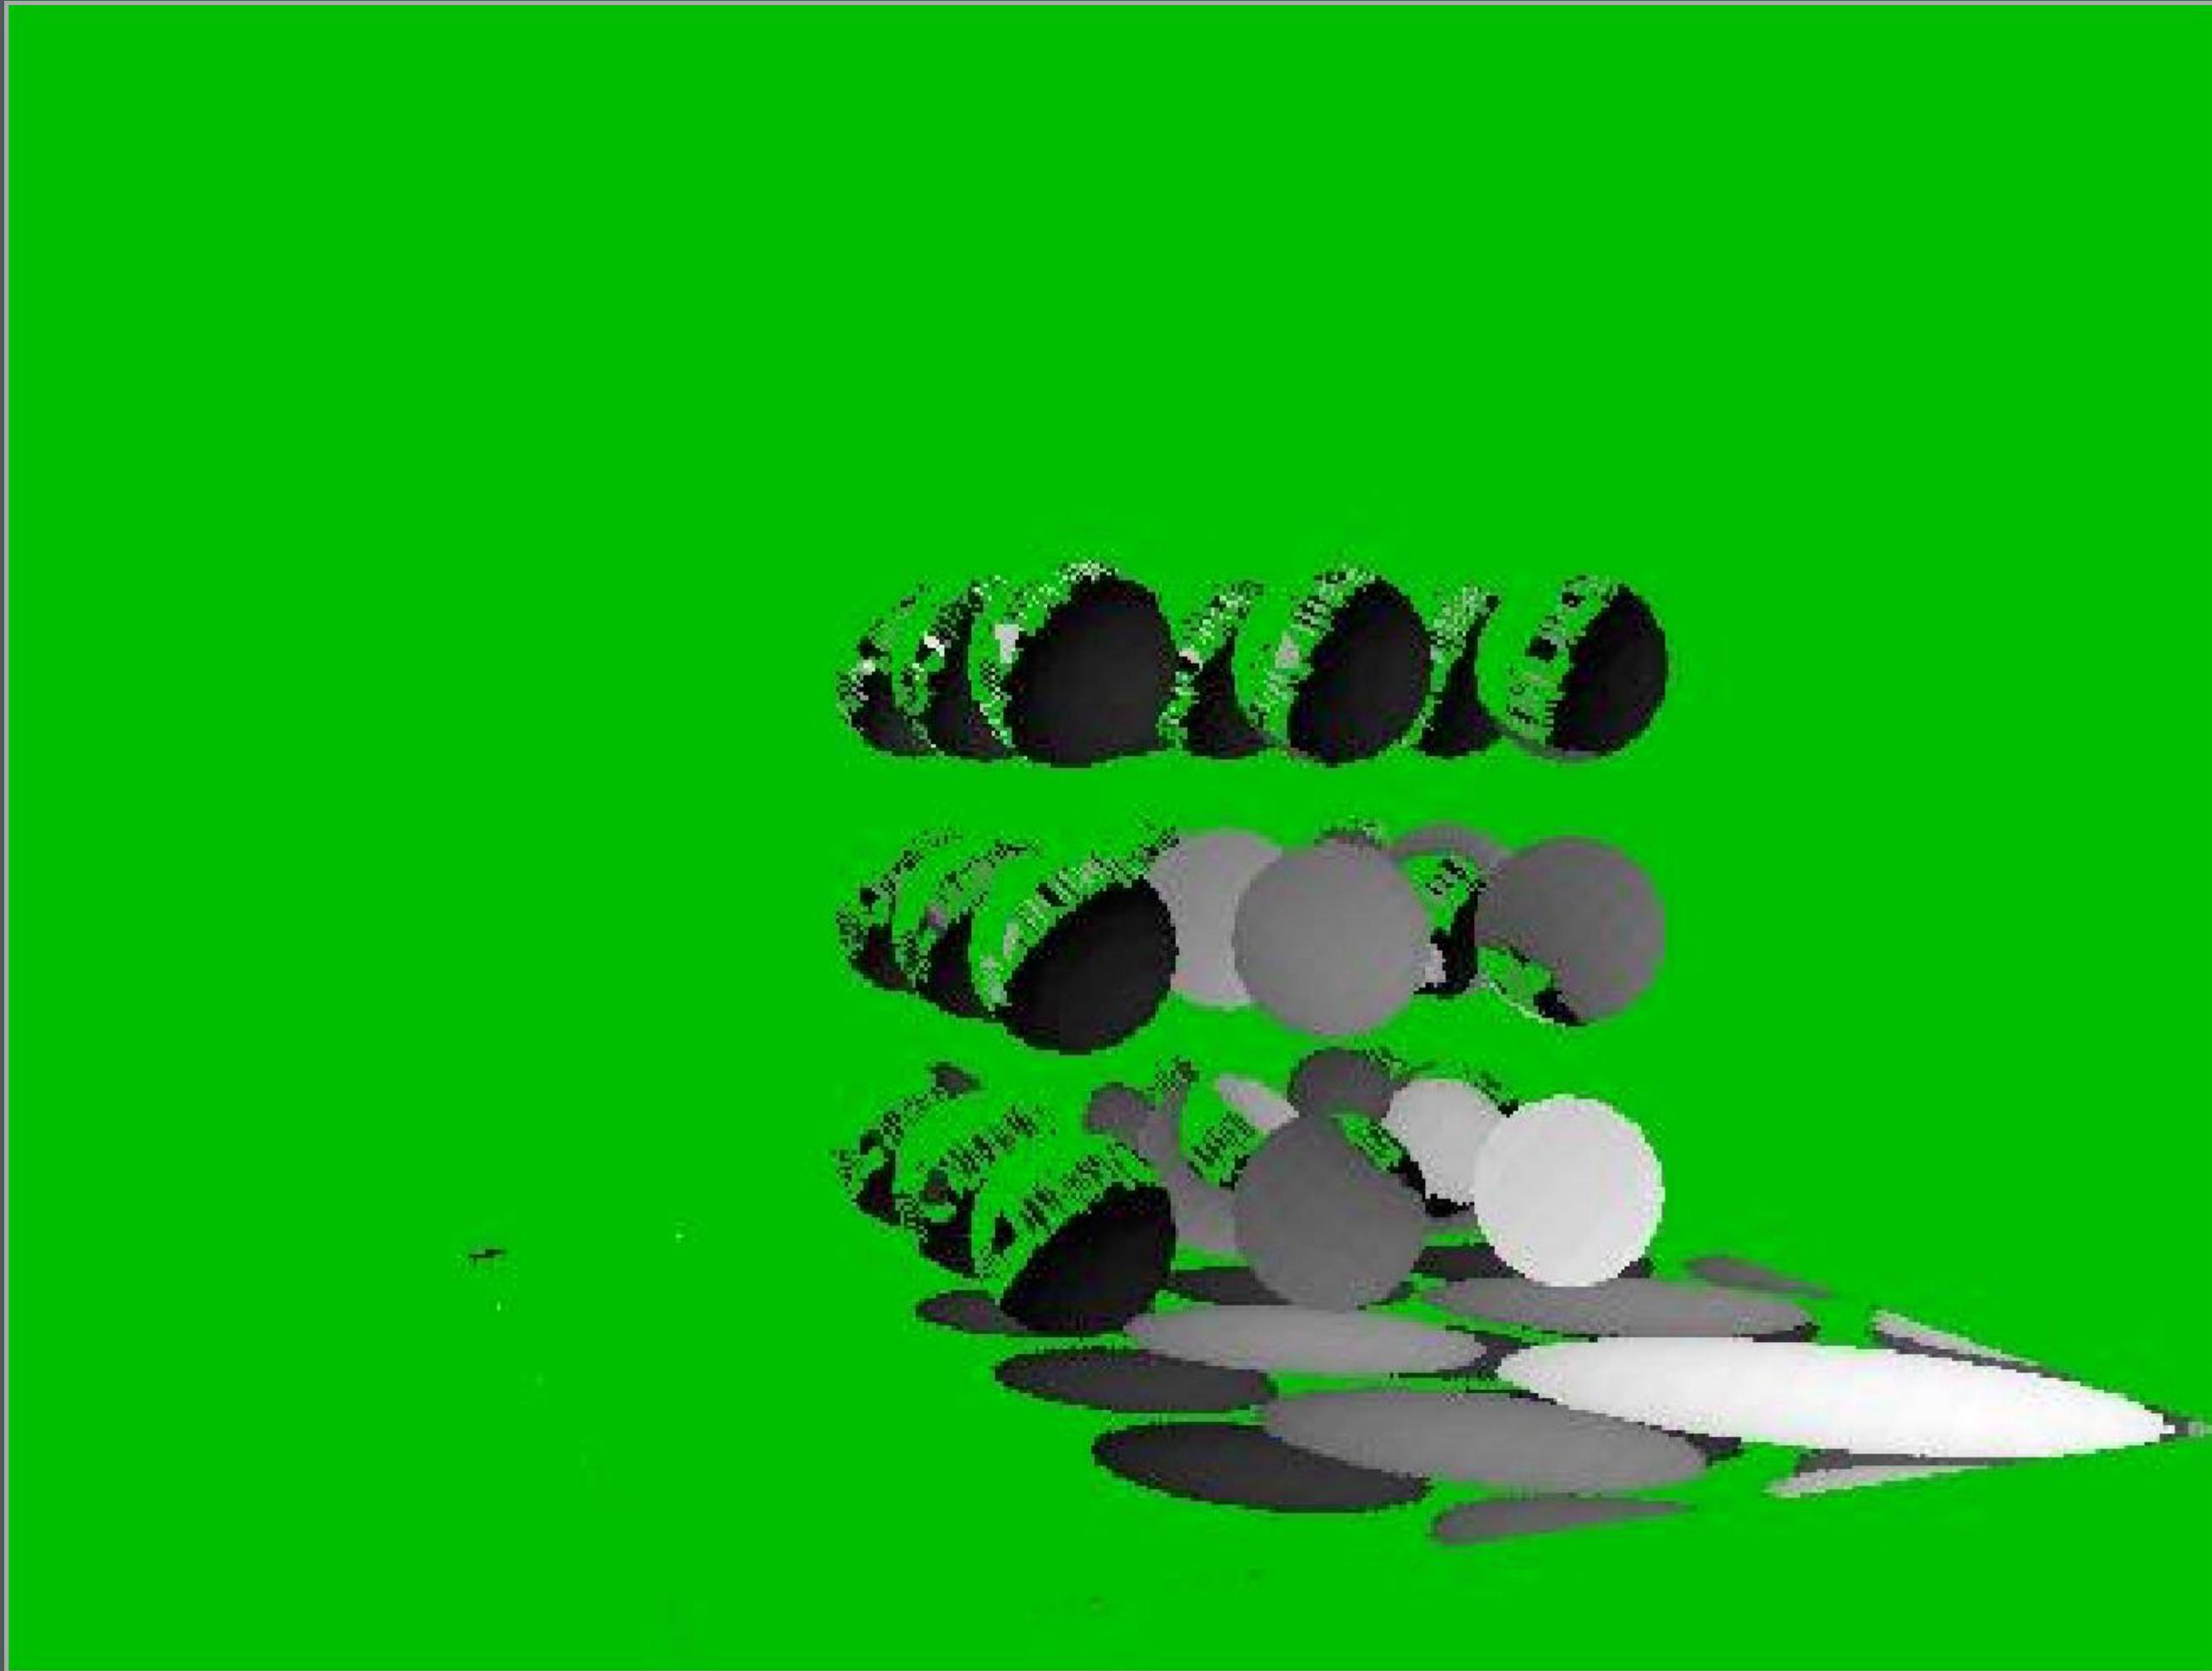




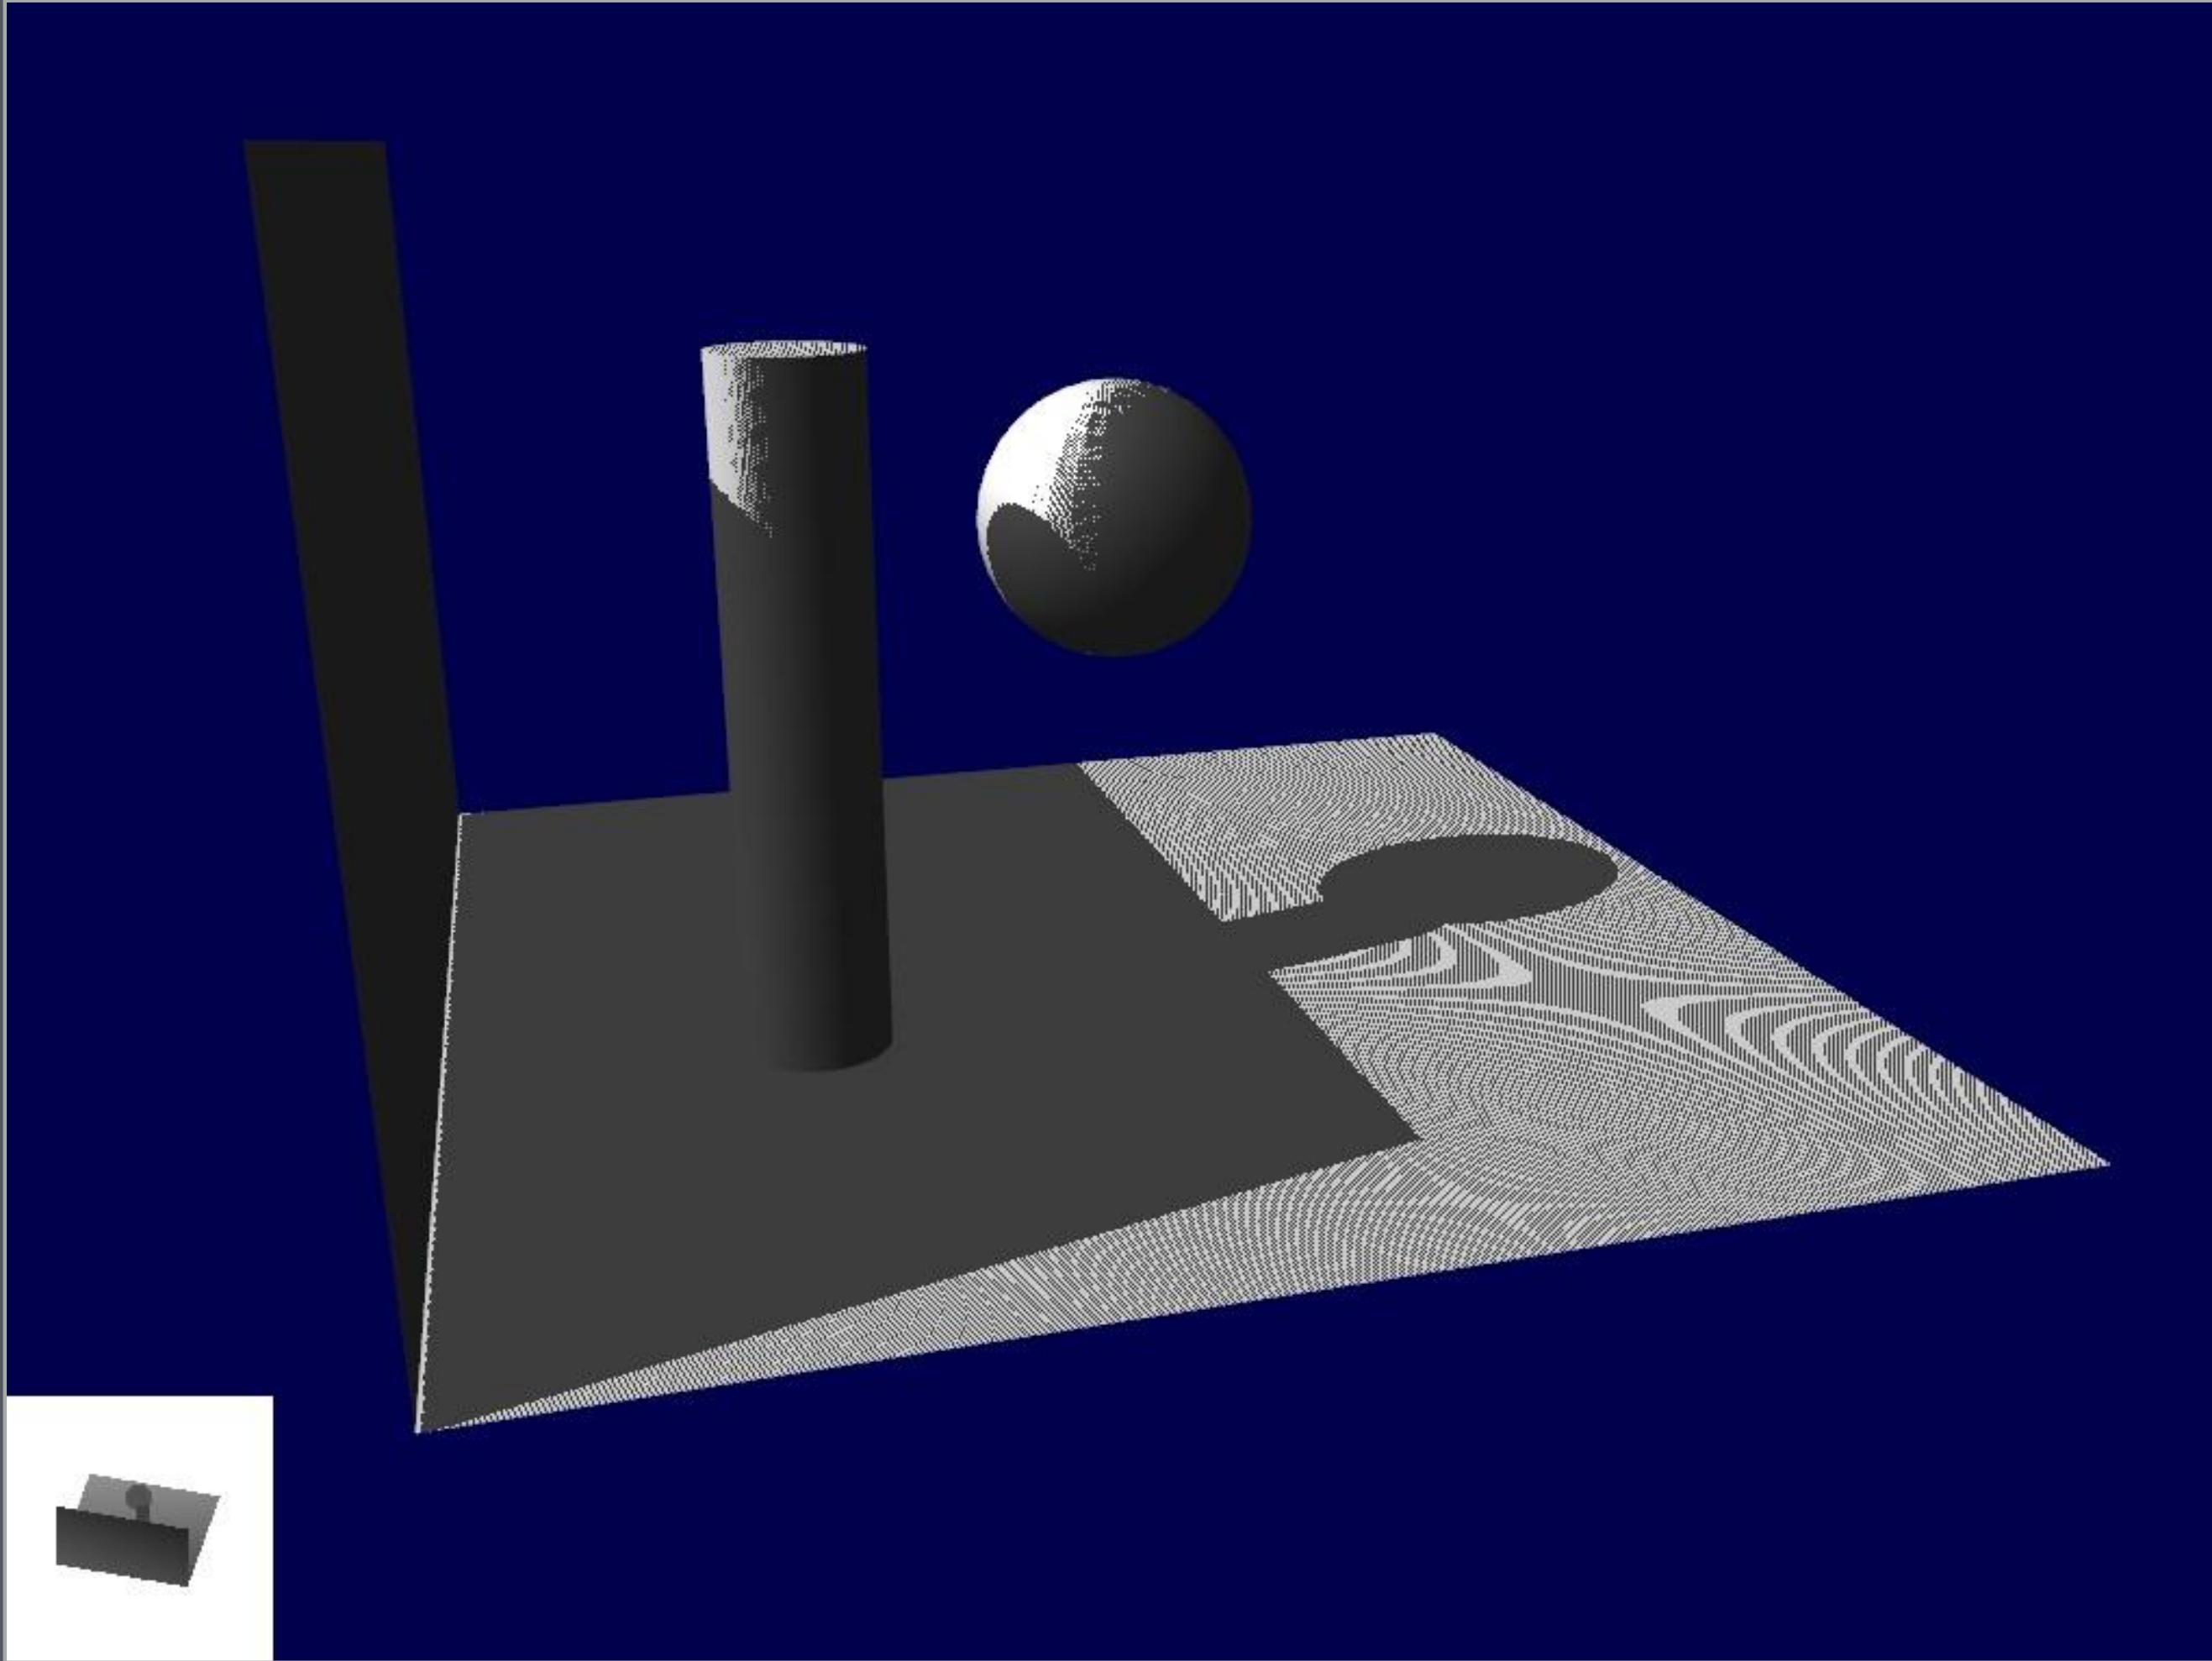
Mark Kilgard



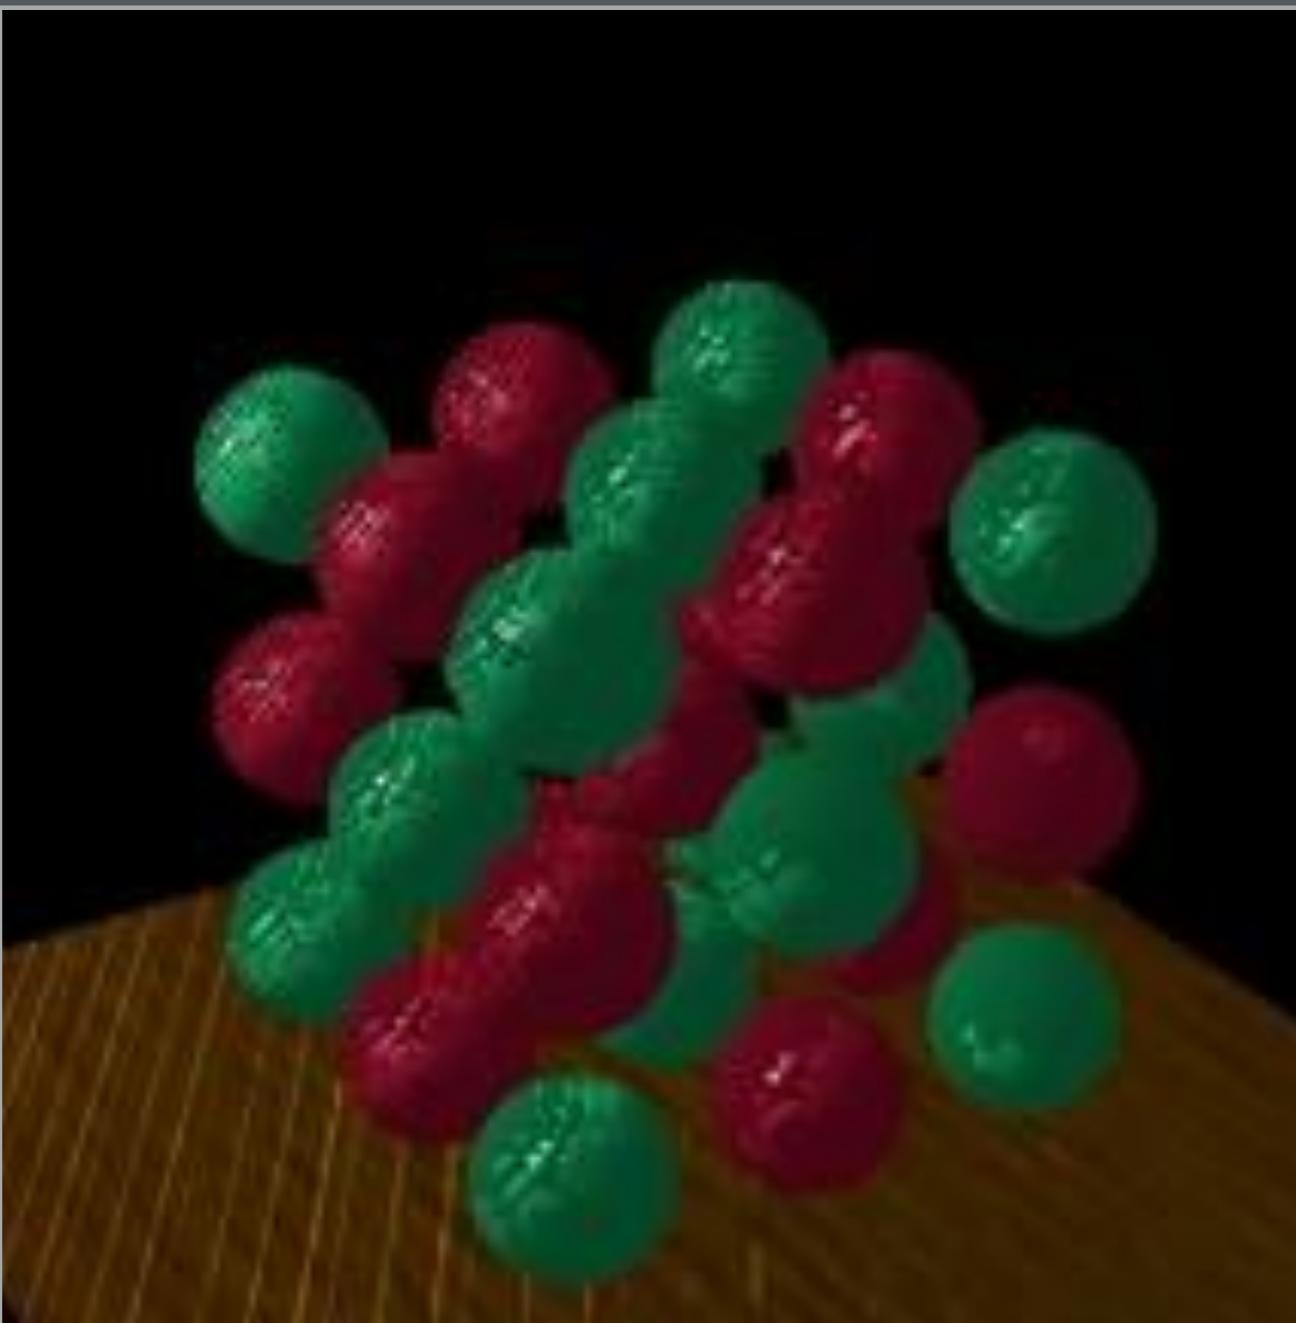
Mark Kilgard



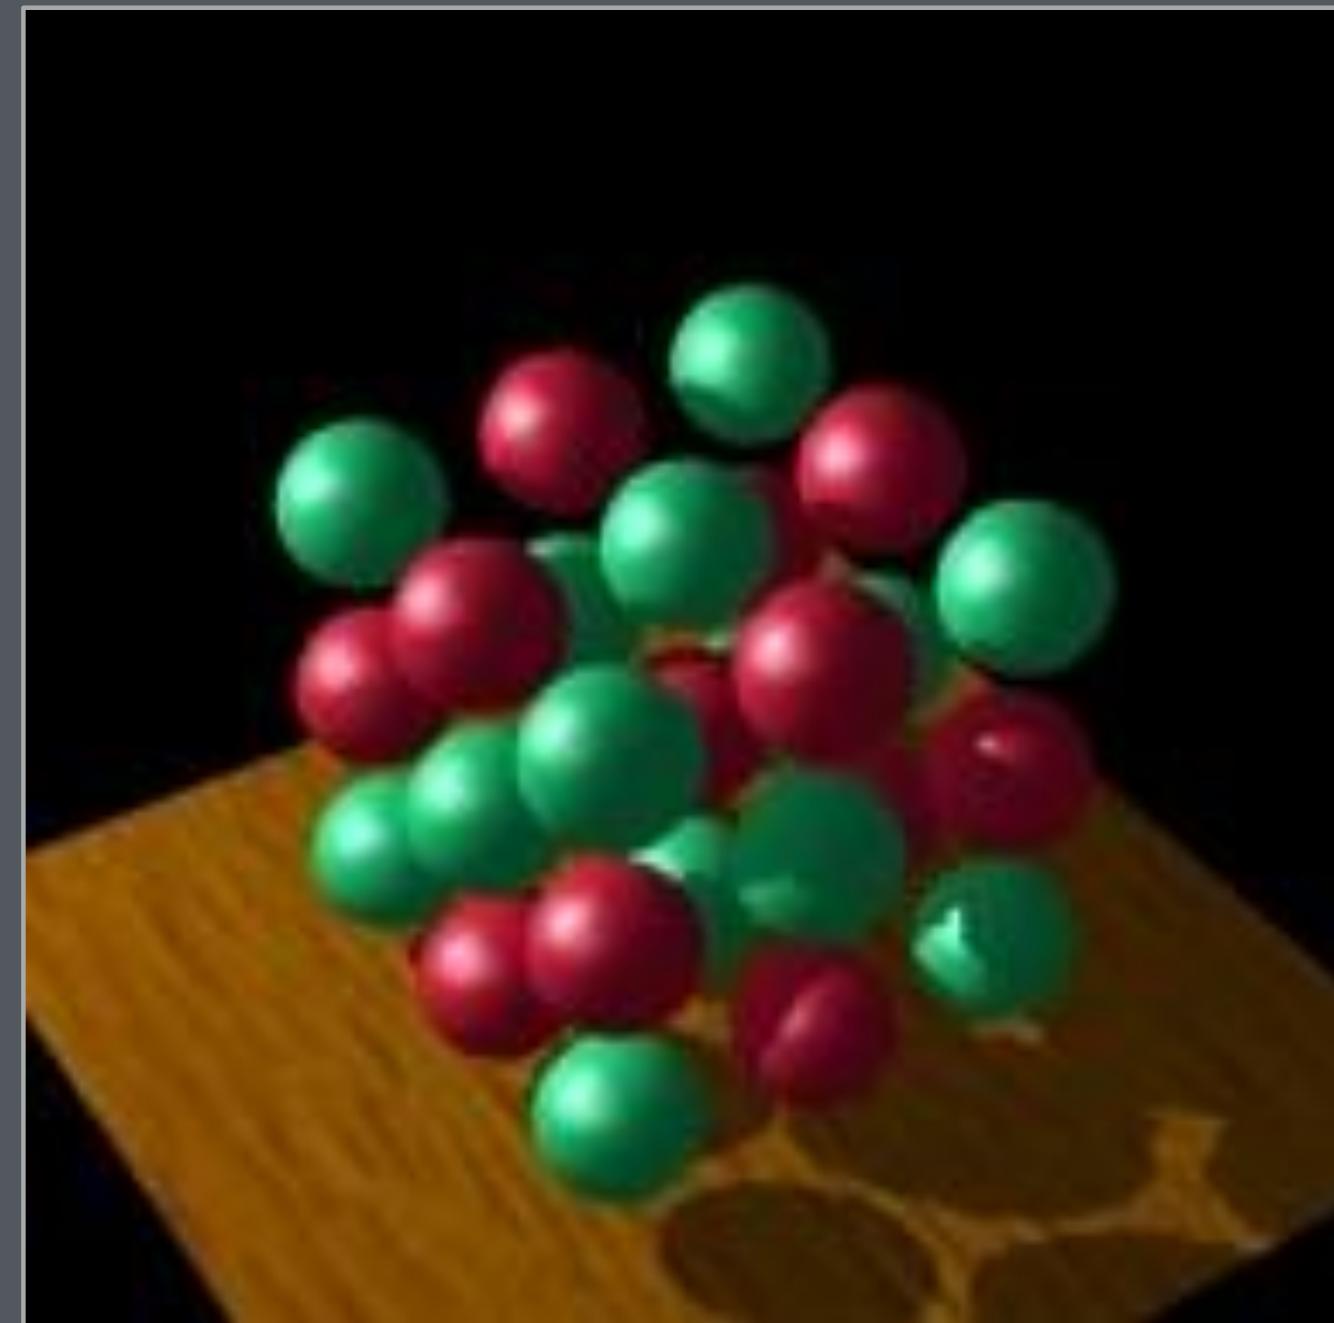
Mark Kilgard



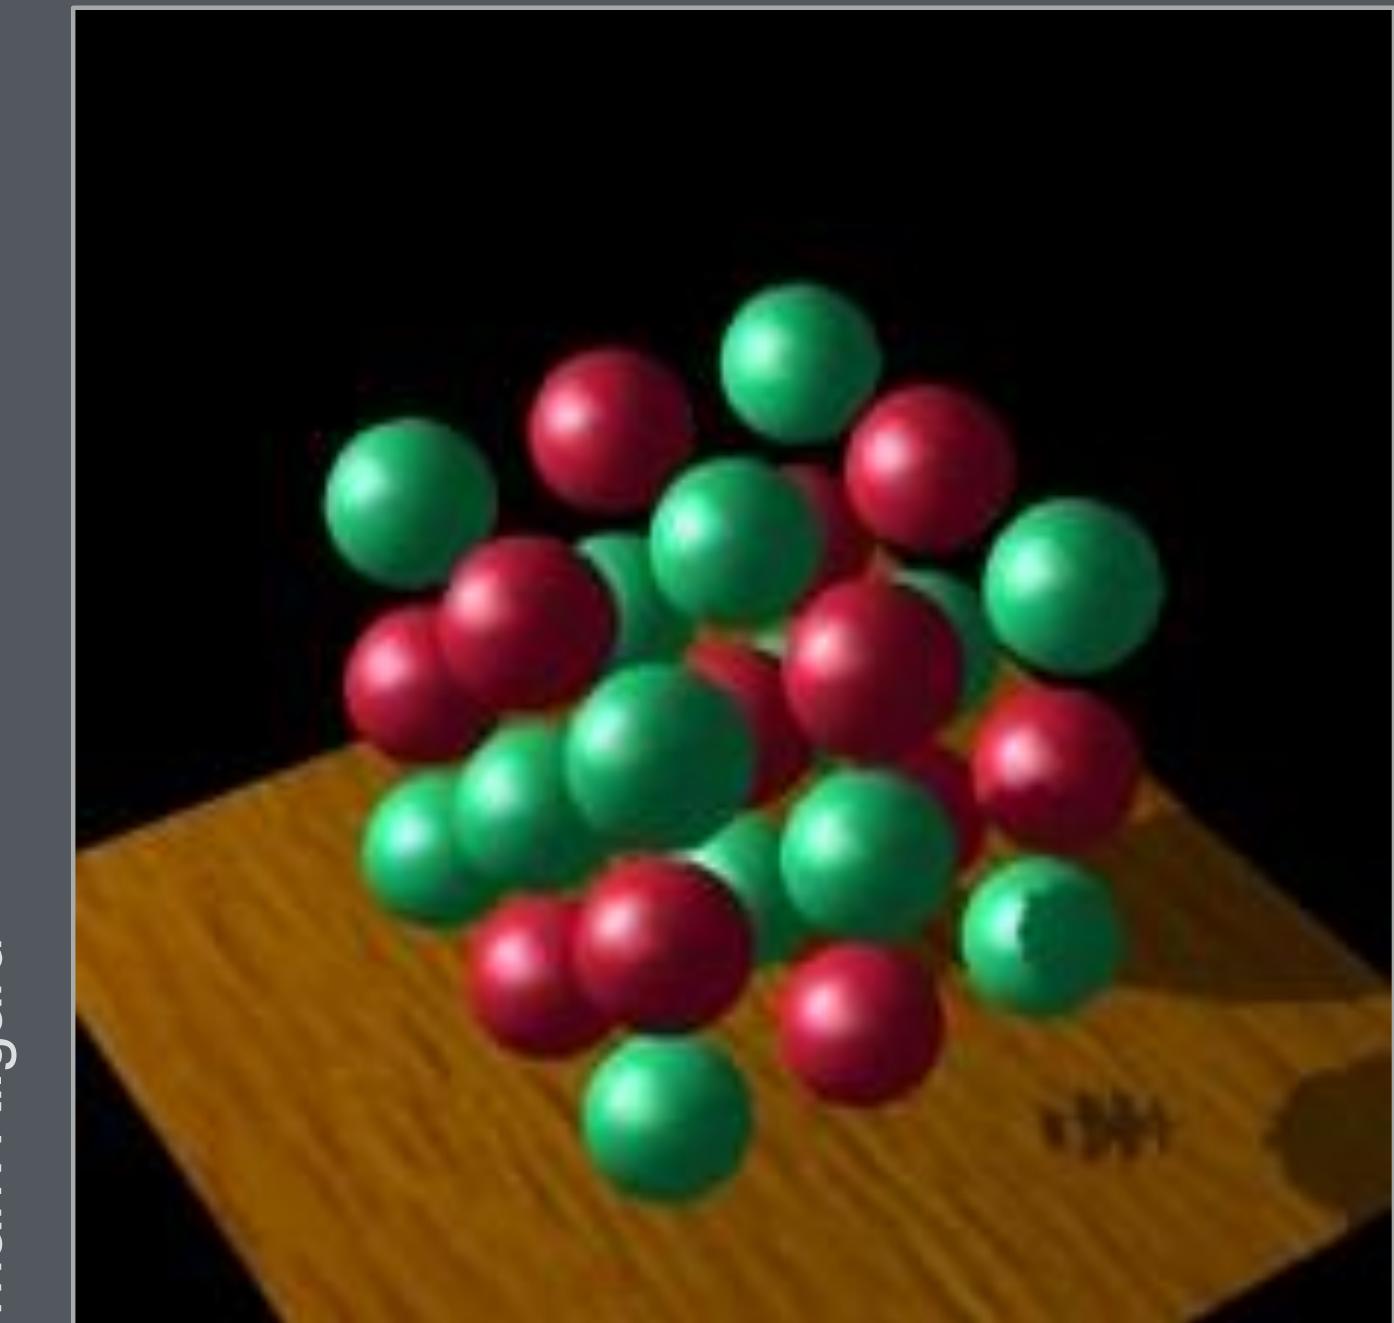
first try at shadow mapping



not enough shadow bias

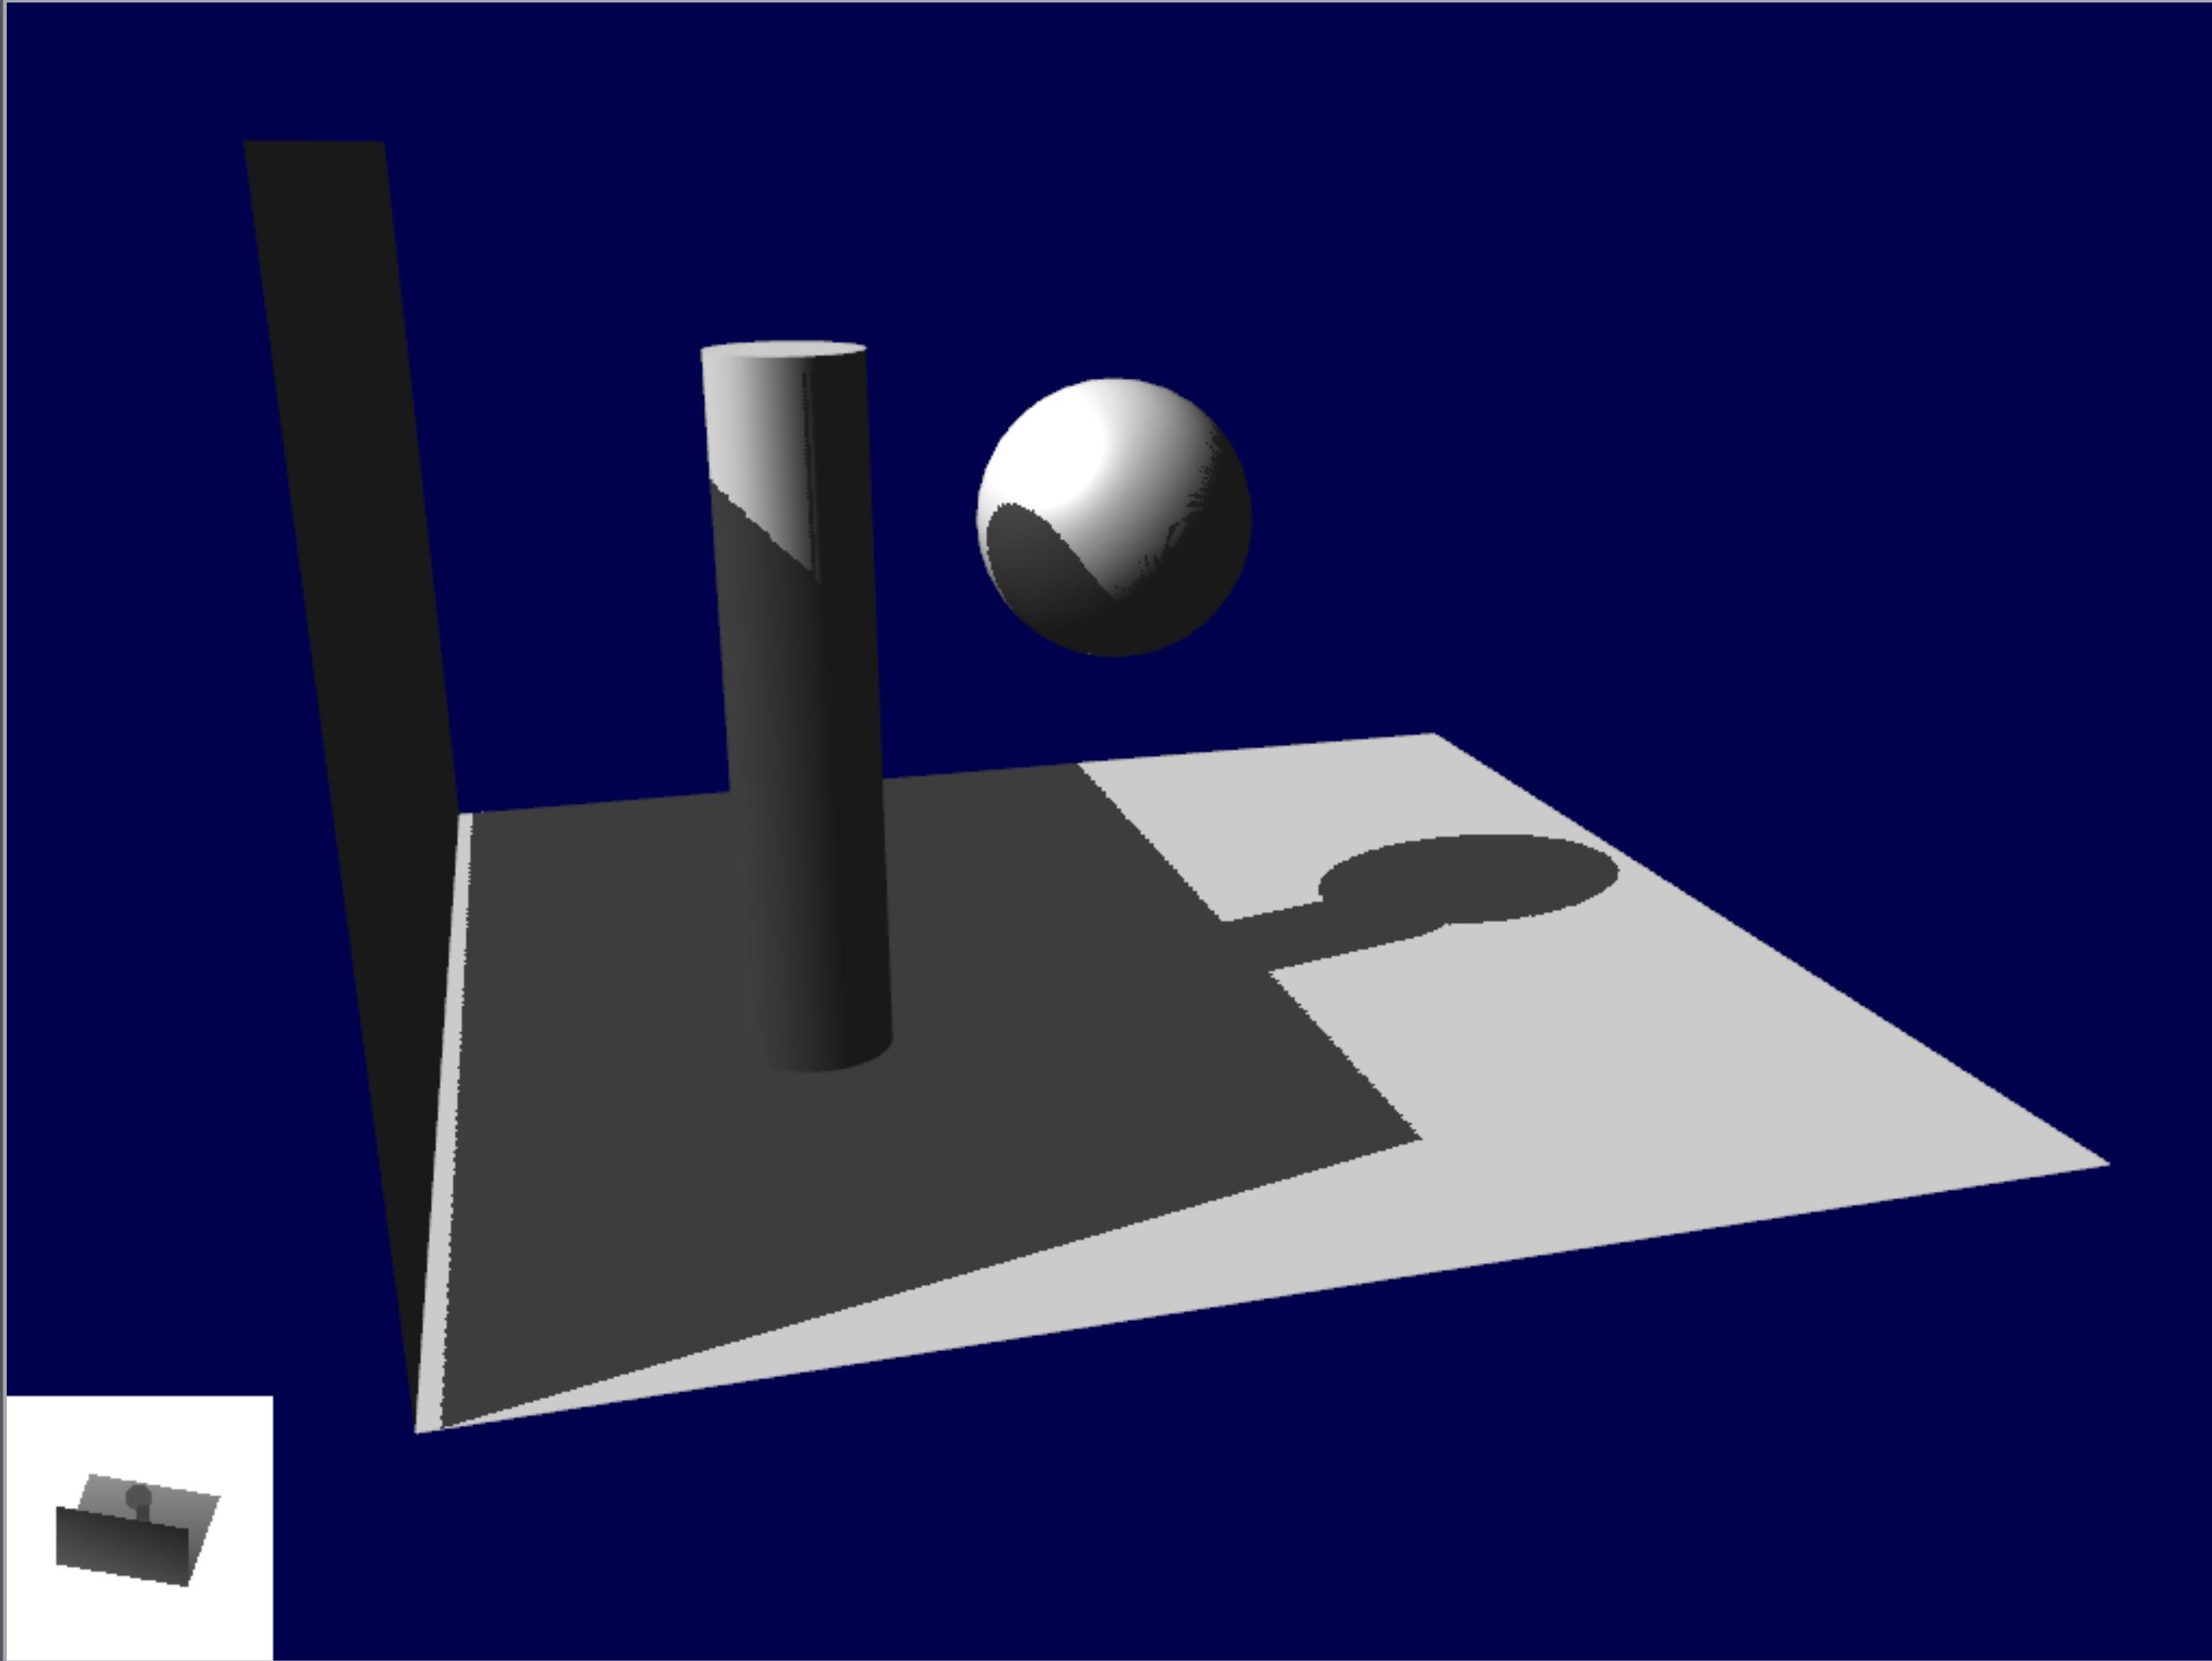


good shadow bias

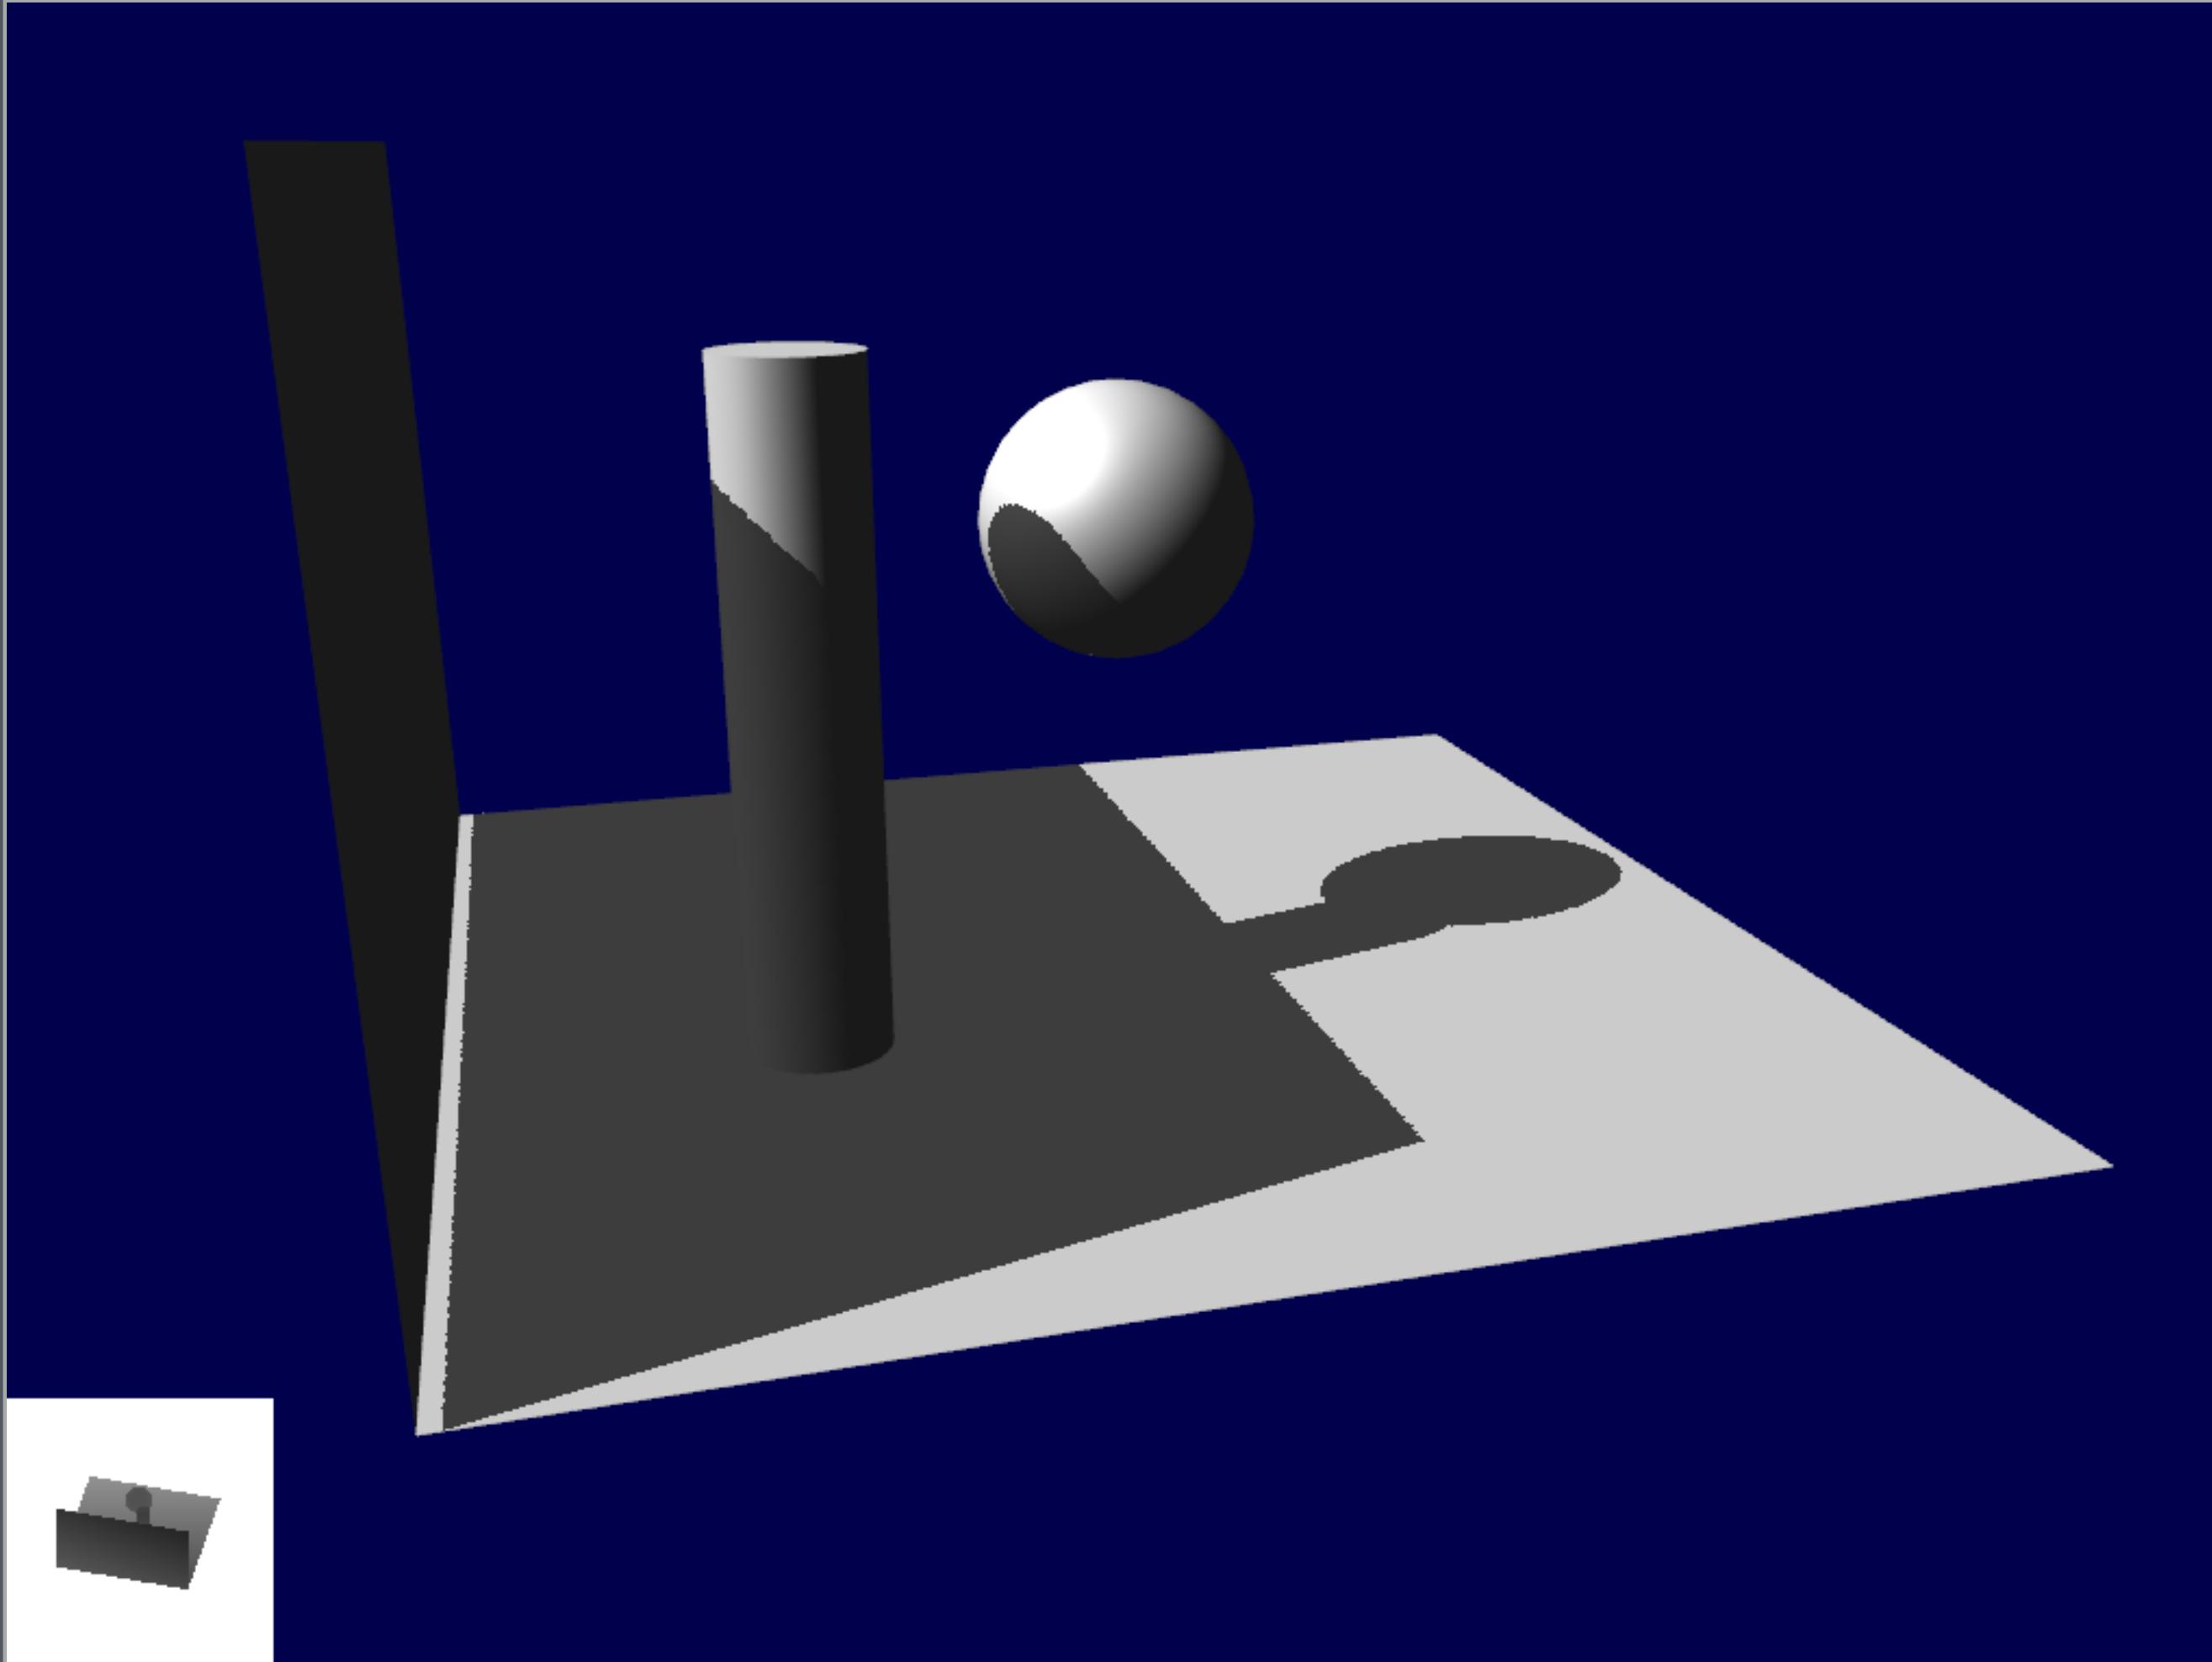


too much shadow bias

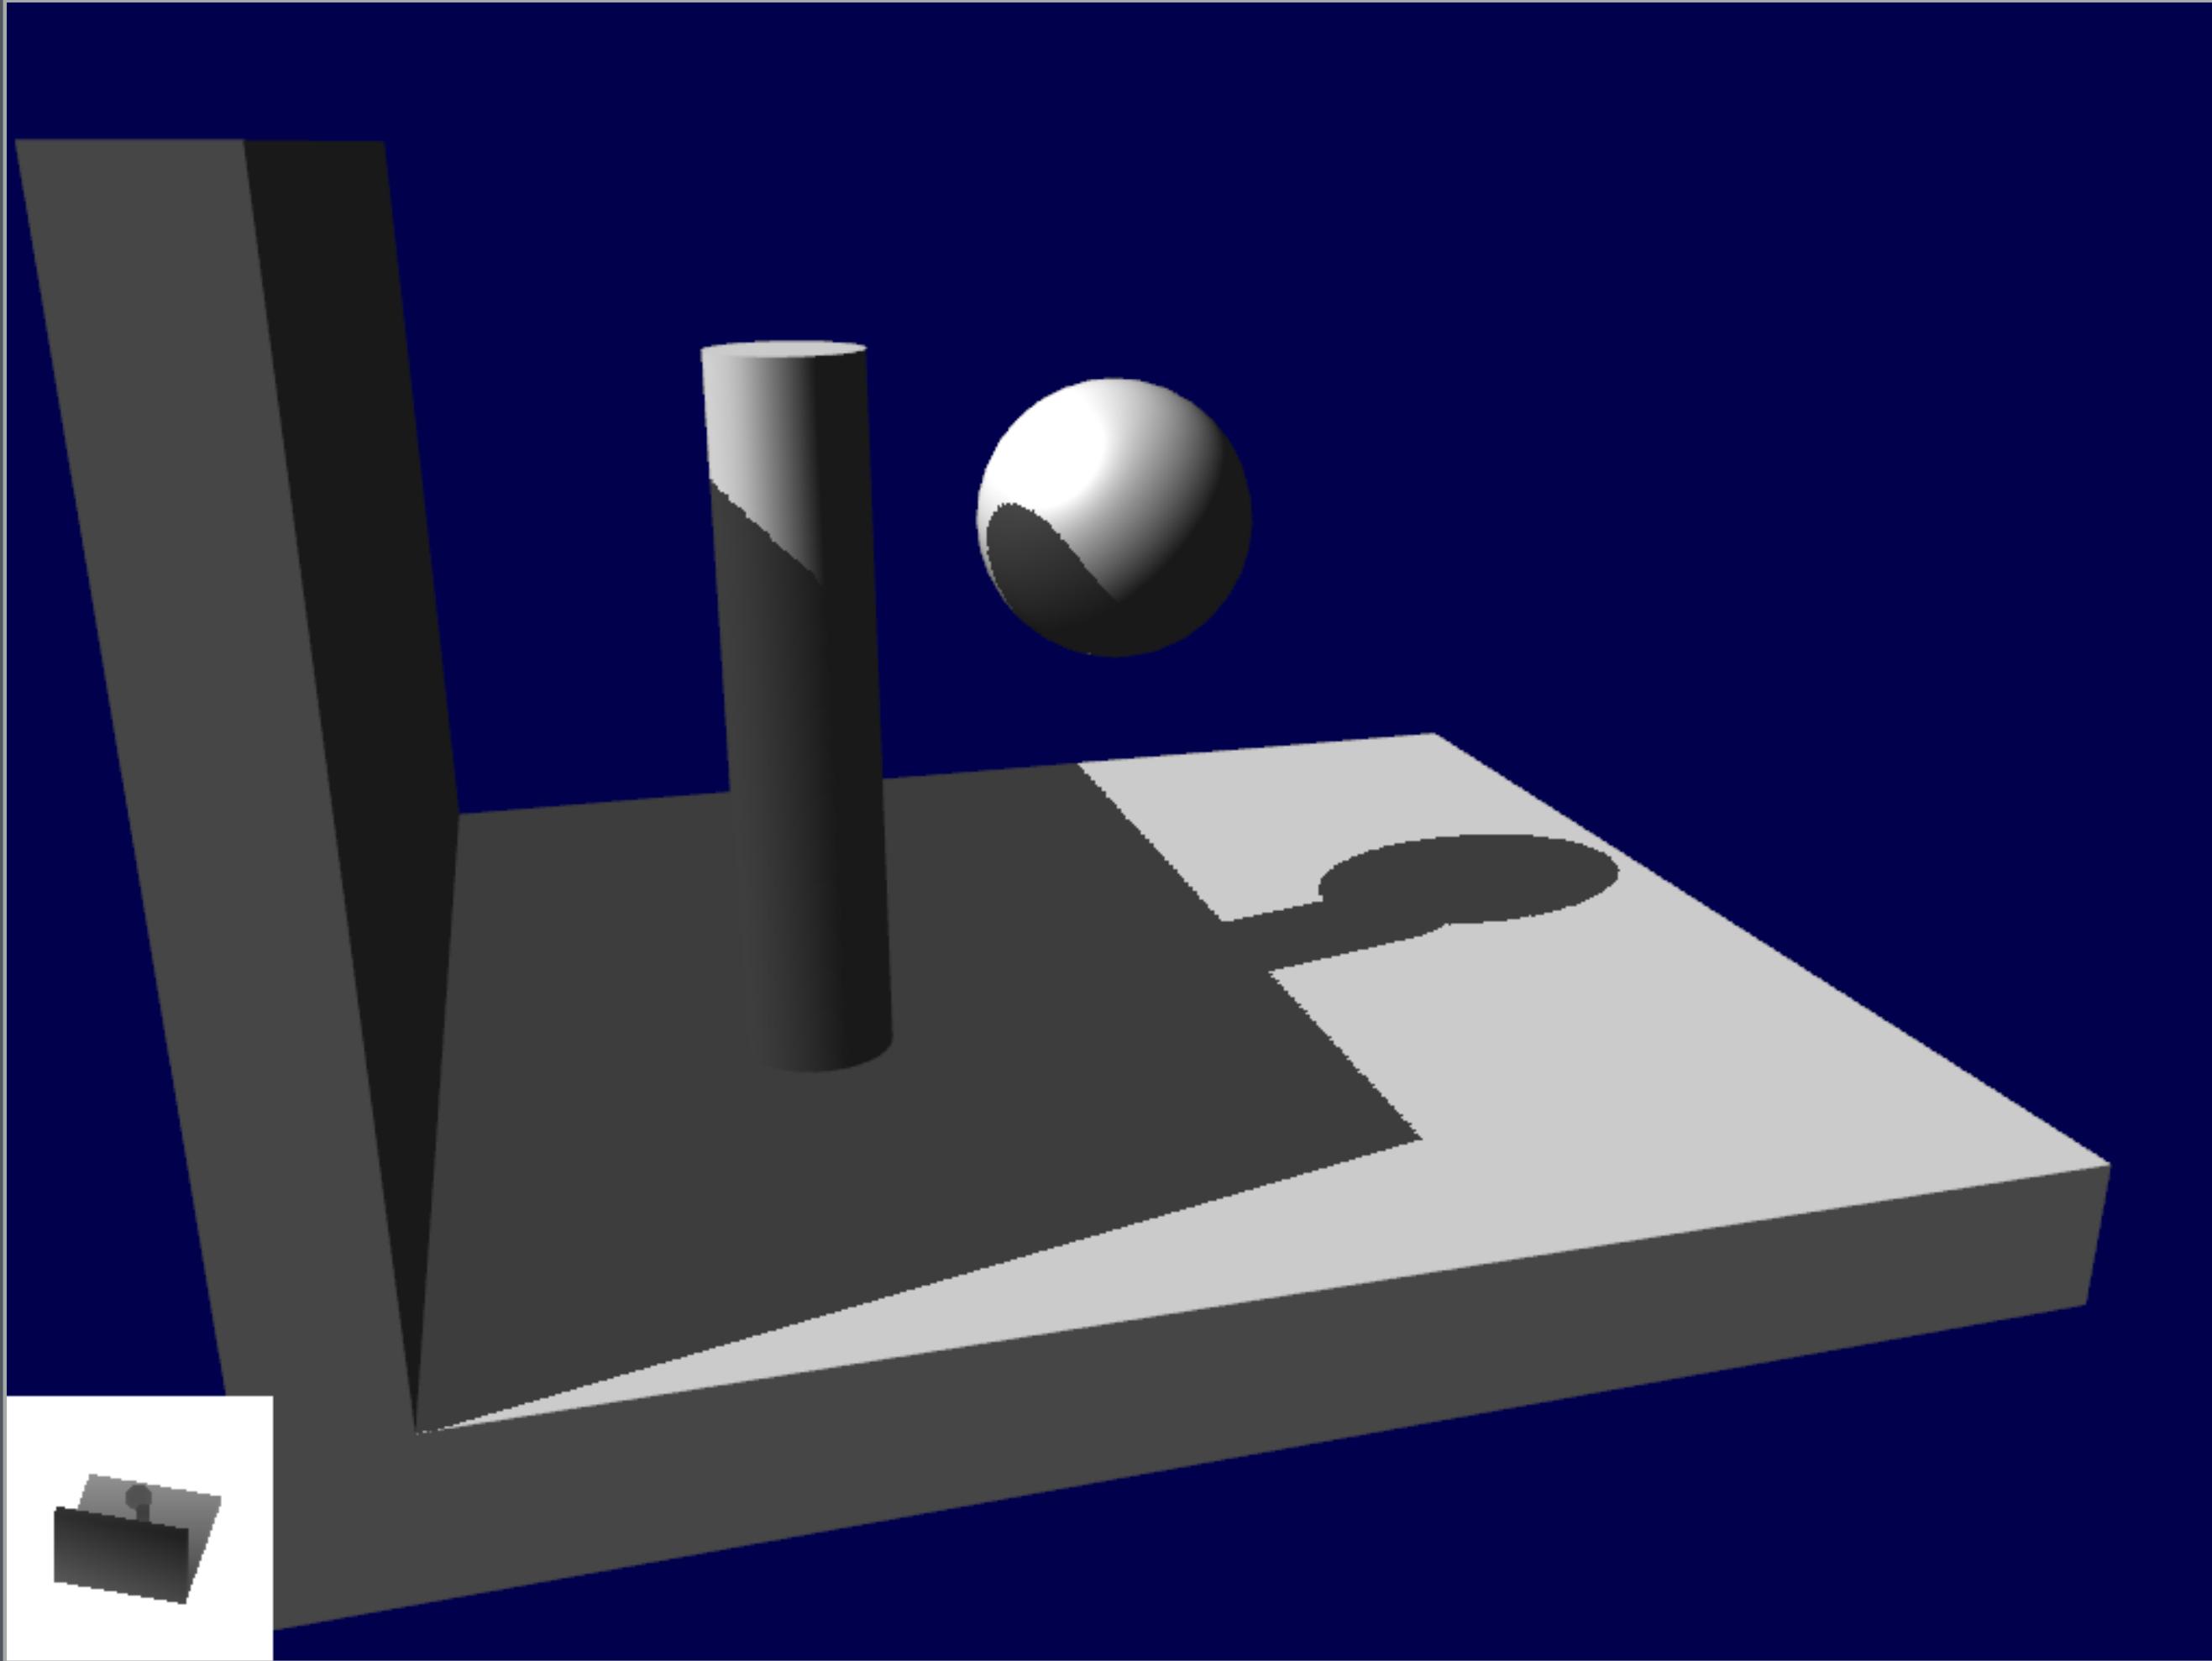
Mark Kilgard



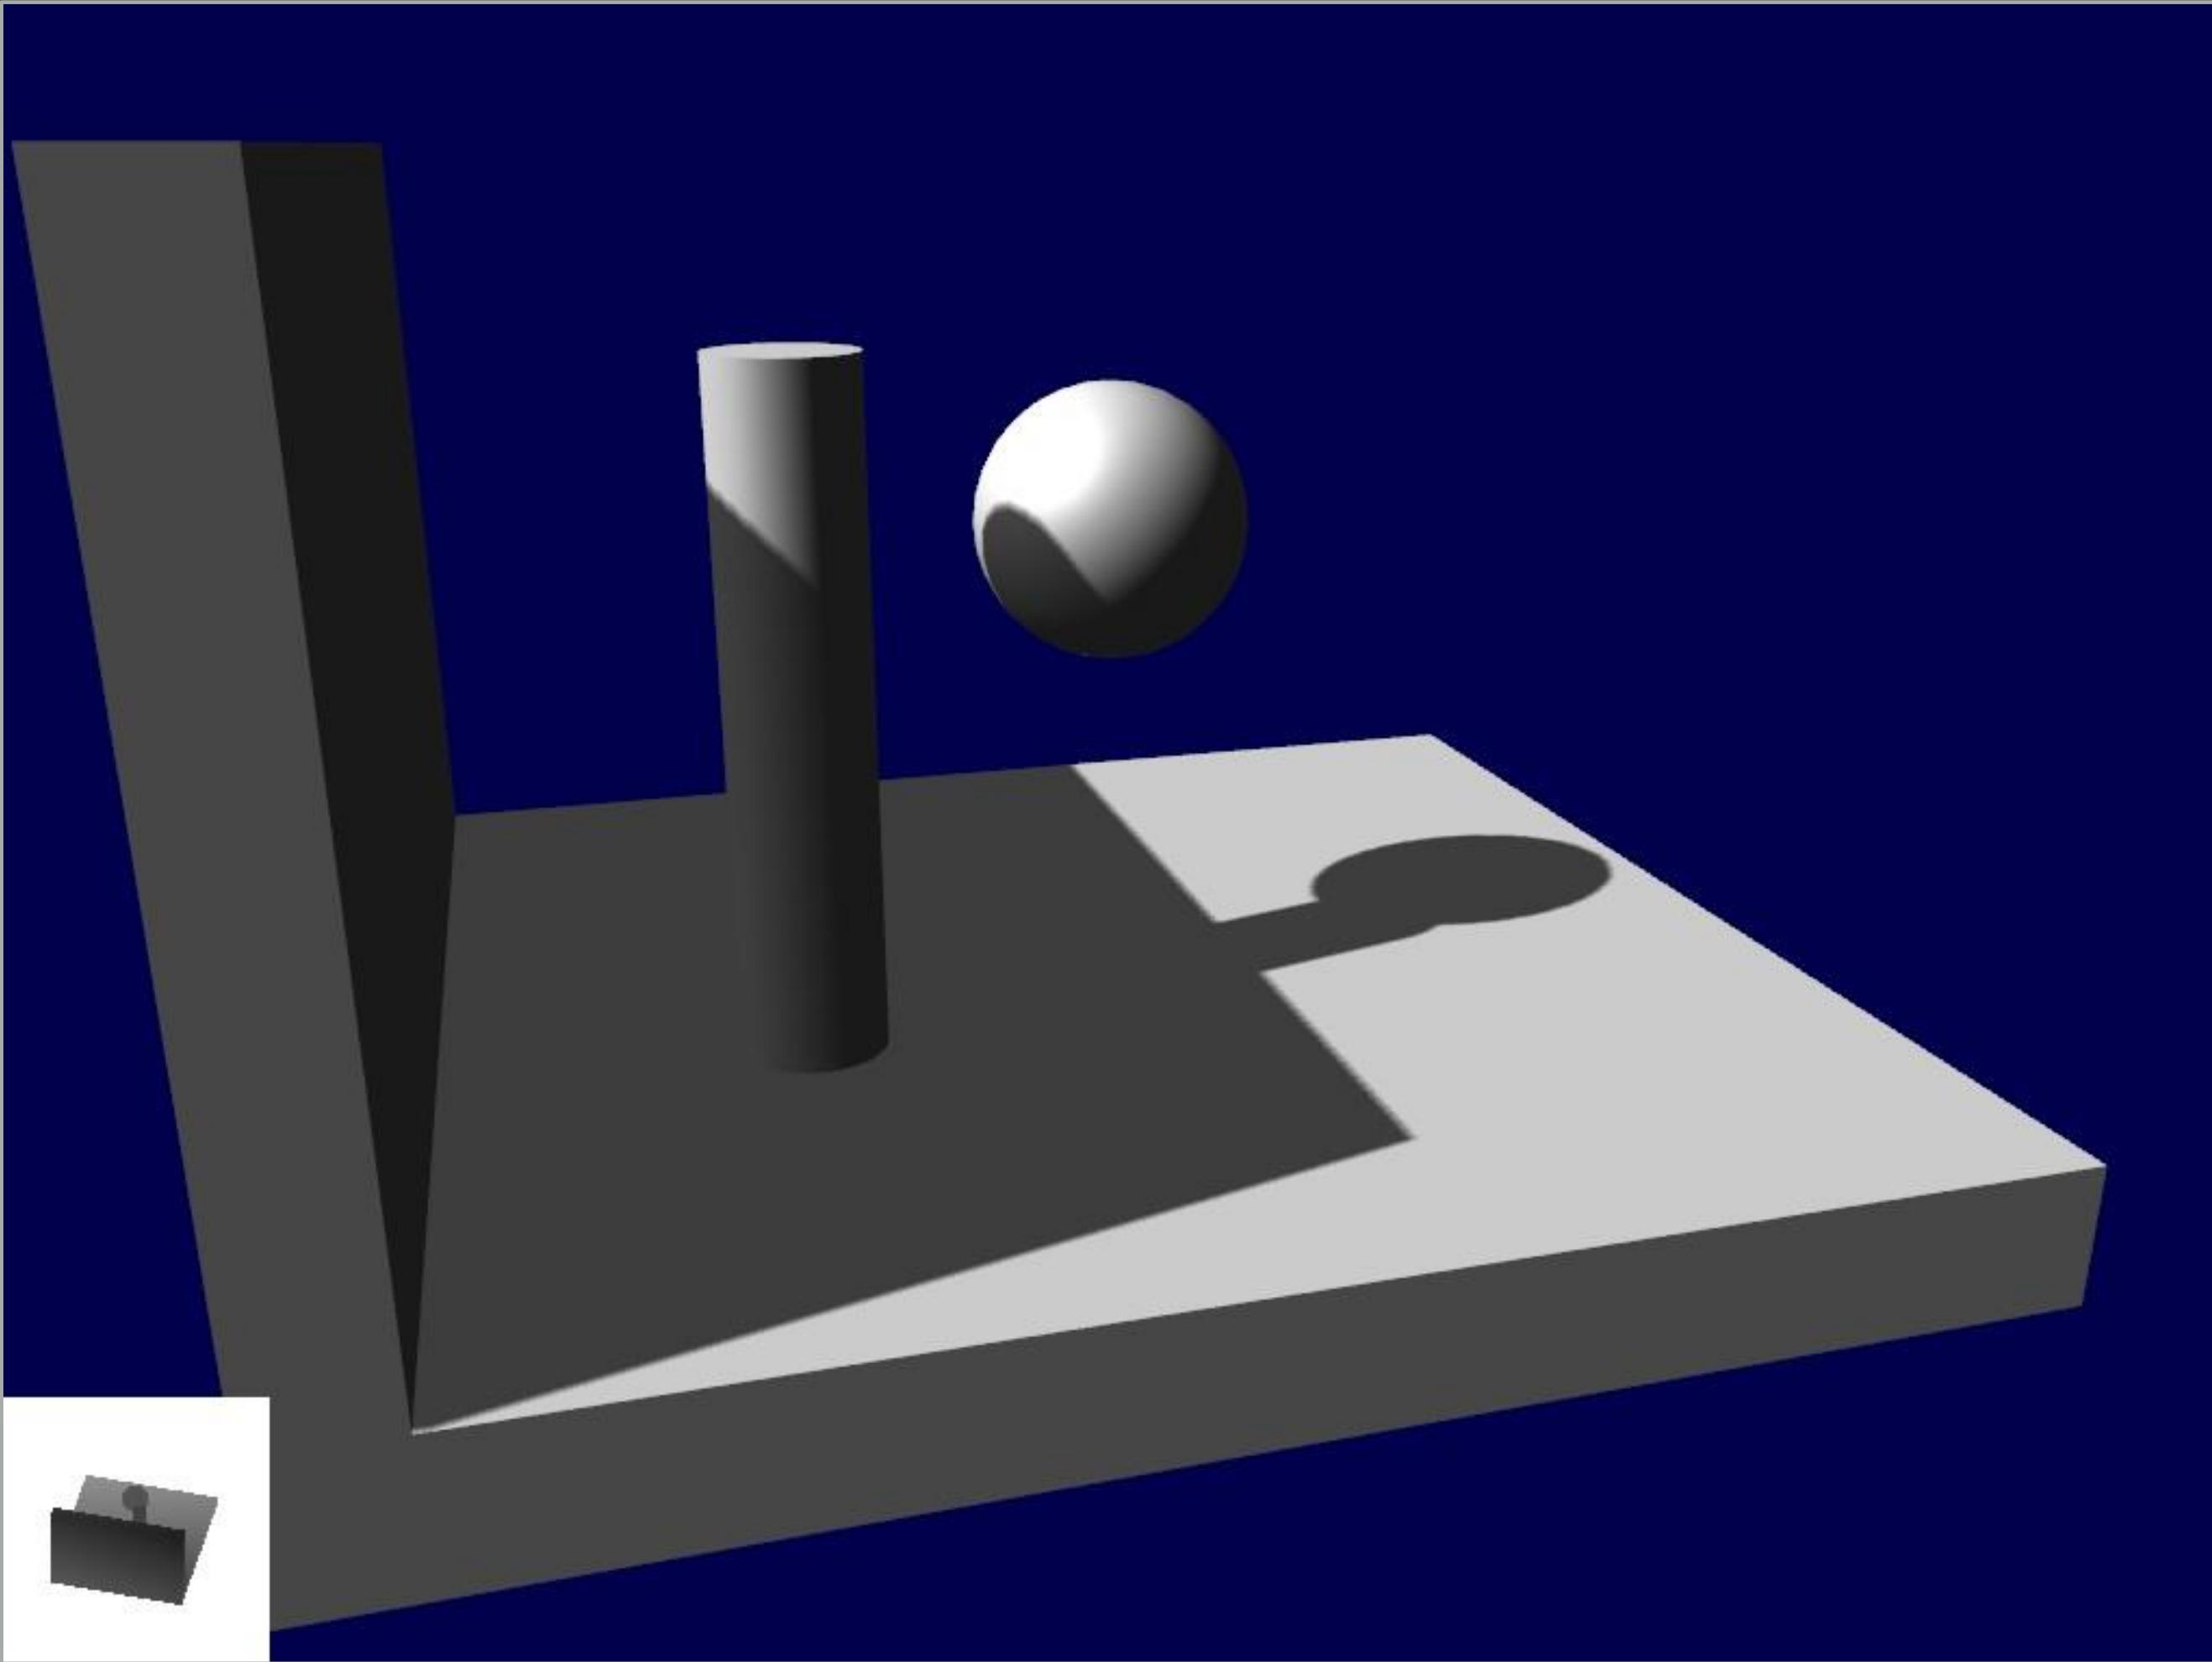
shadow mapping with constant bias



shadow mapping with slope-dependent bias

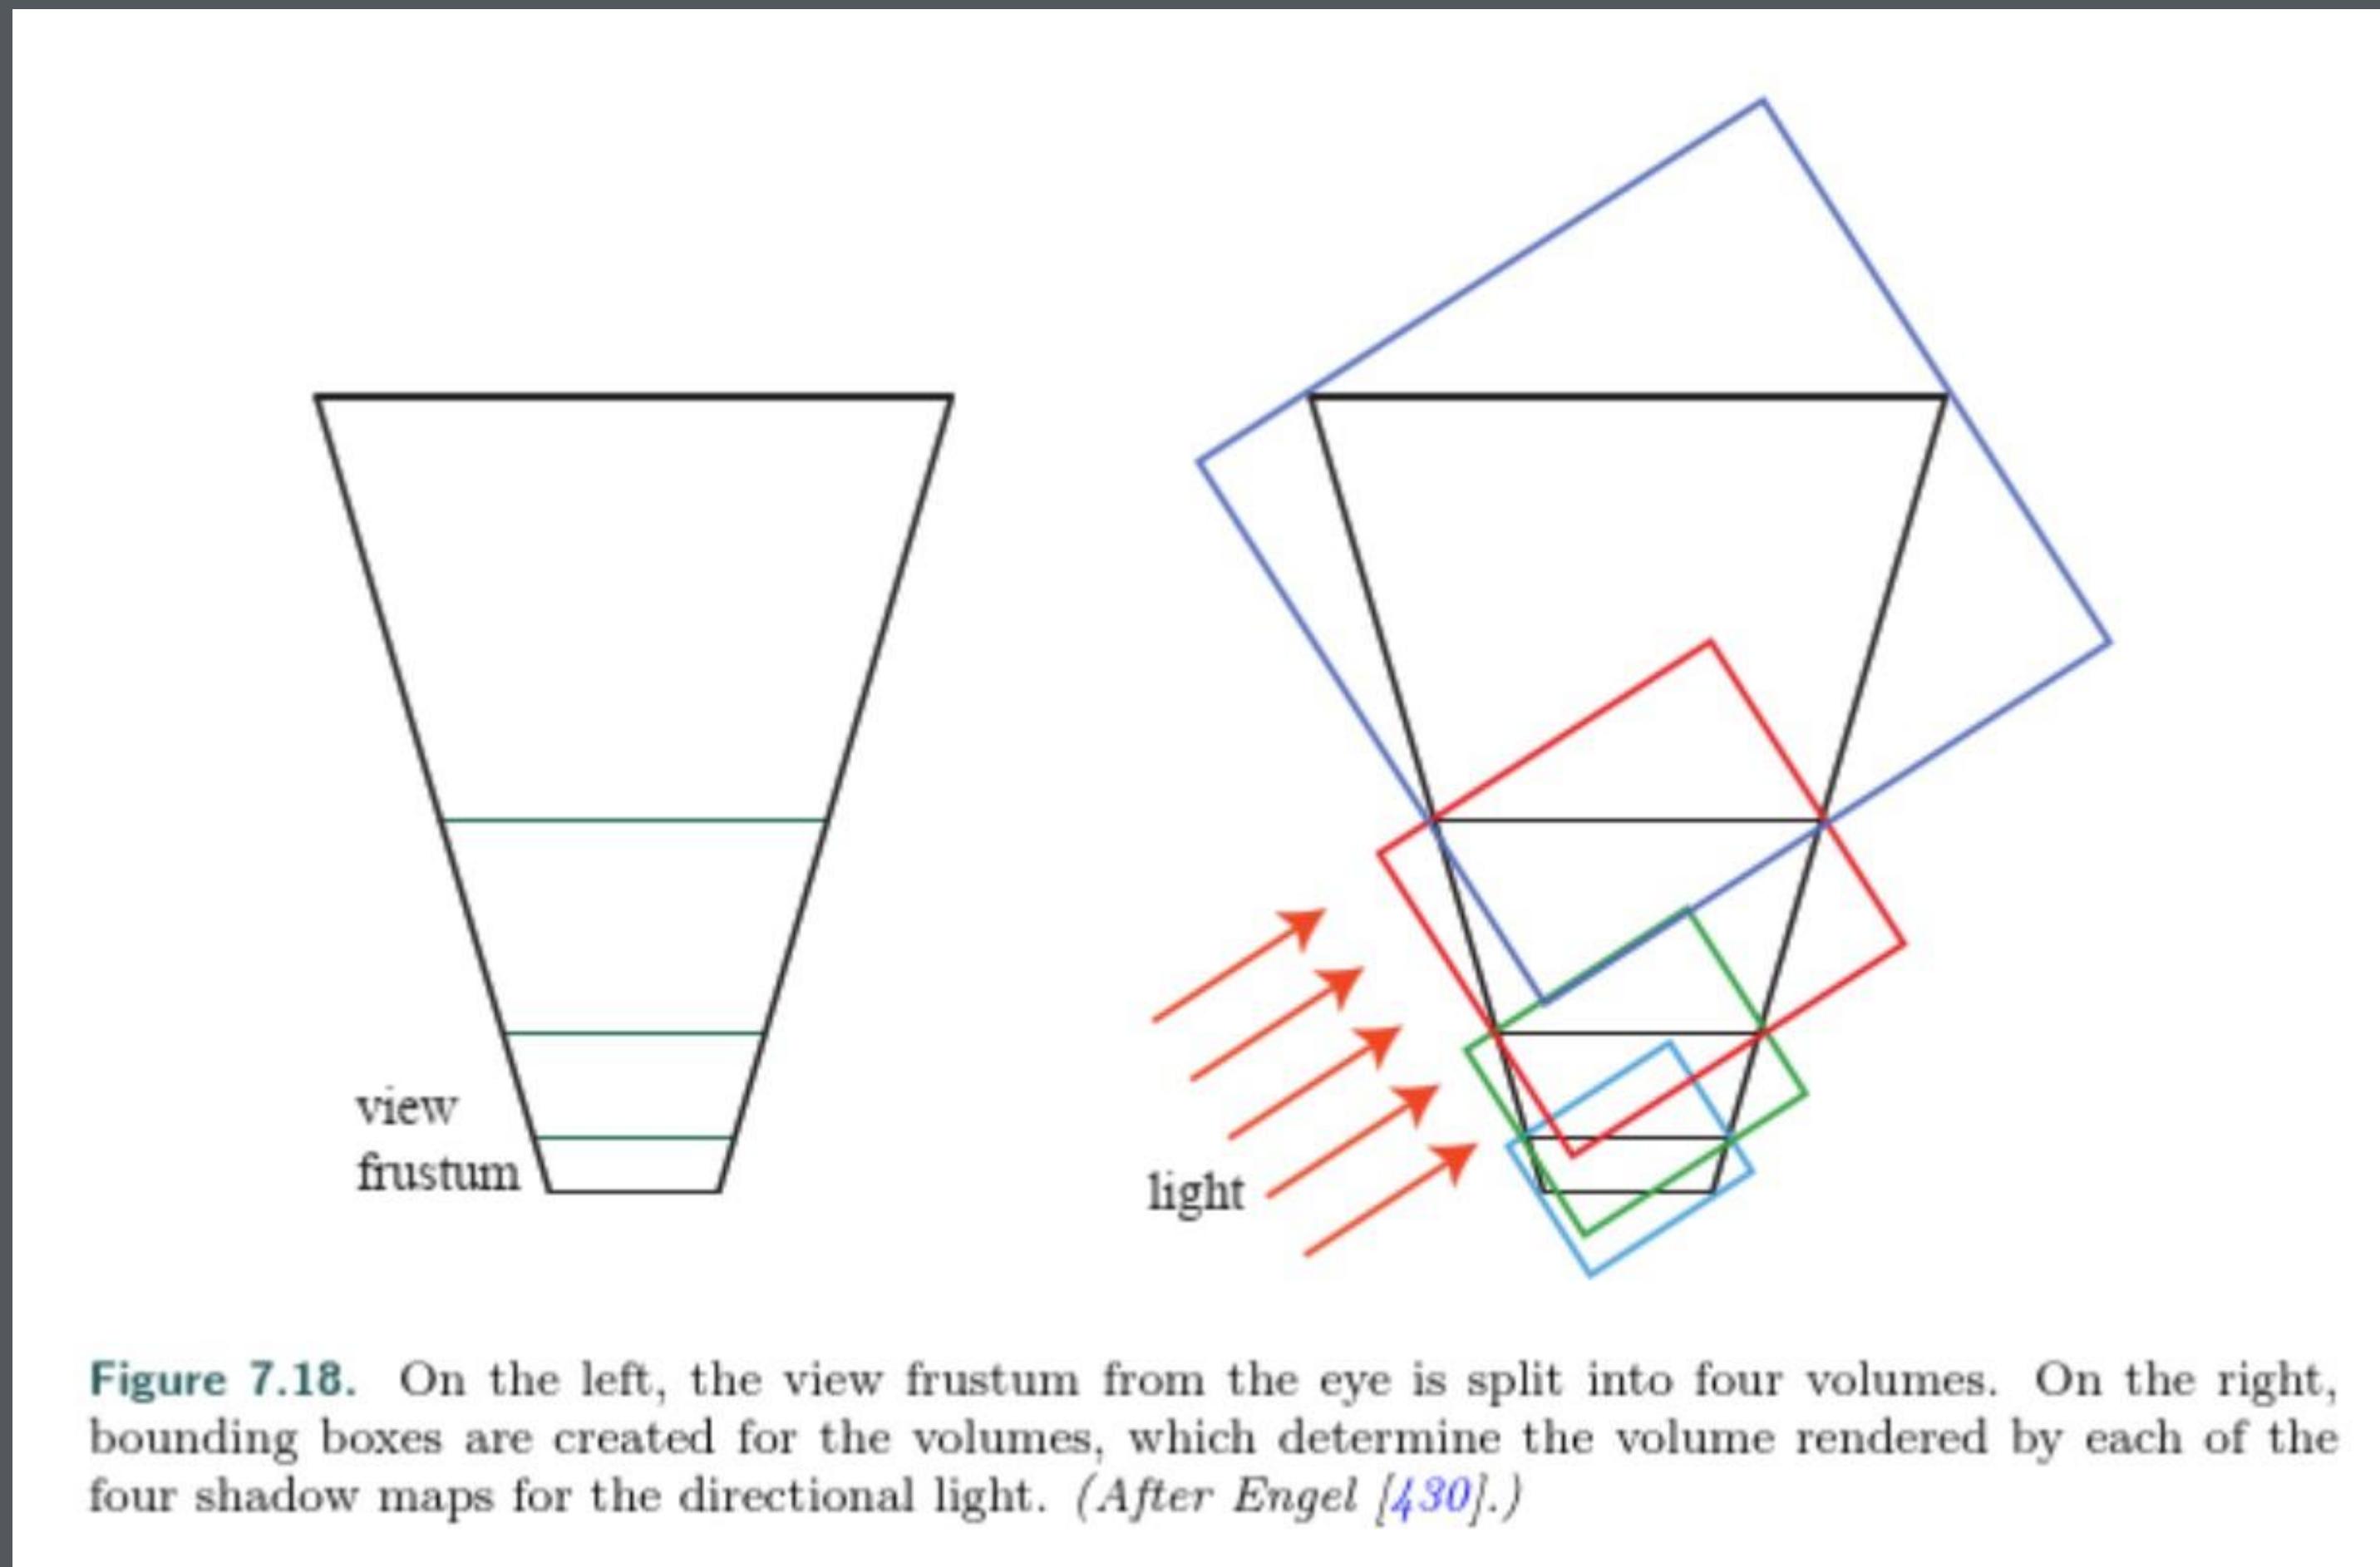


**closed surfaces and slope-dependent bias**



adding percentage-closer filtering

# Cascaded shadow maps (aka. parallel-split SM)





Single shadow map, 2048x2048



Four 1024x1024 shadow maps (equal memory)

Fan Zhang, Chinese U. Hong Kong

# Filtering shadow maps

**Shadow map lookups cause aliasing, need filtering**

**As with normal maps, pixel is a nonlinear function of the shadow depth**

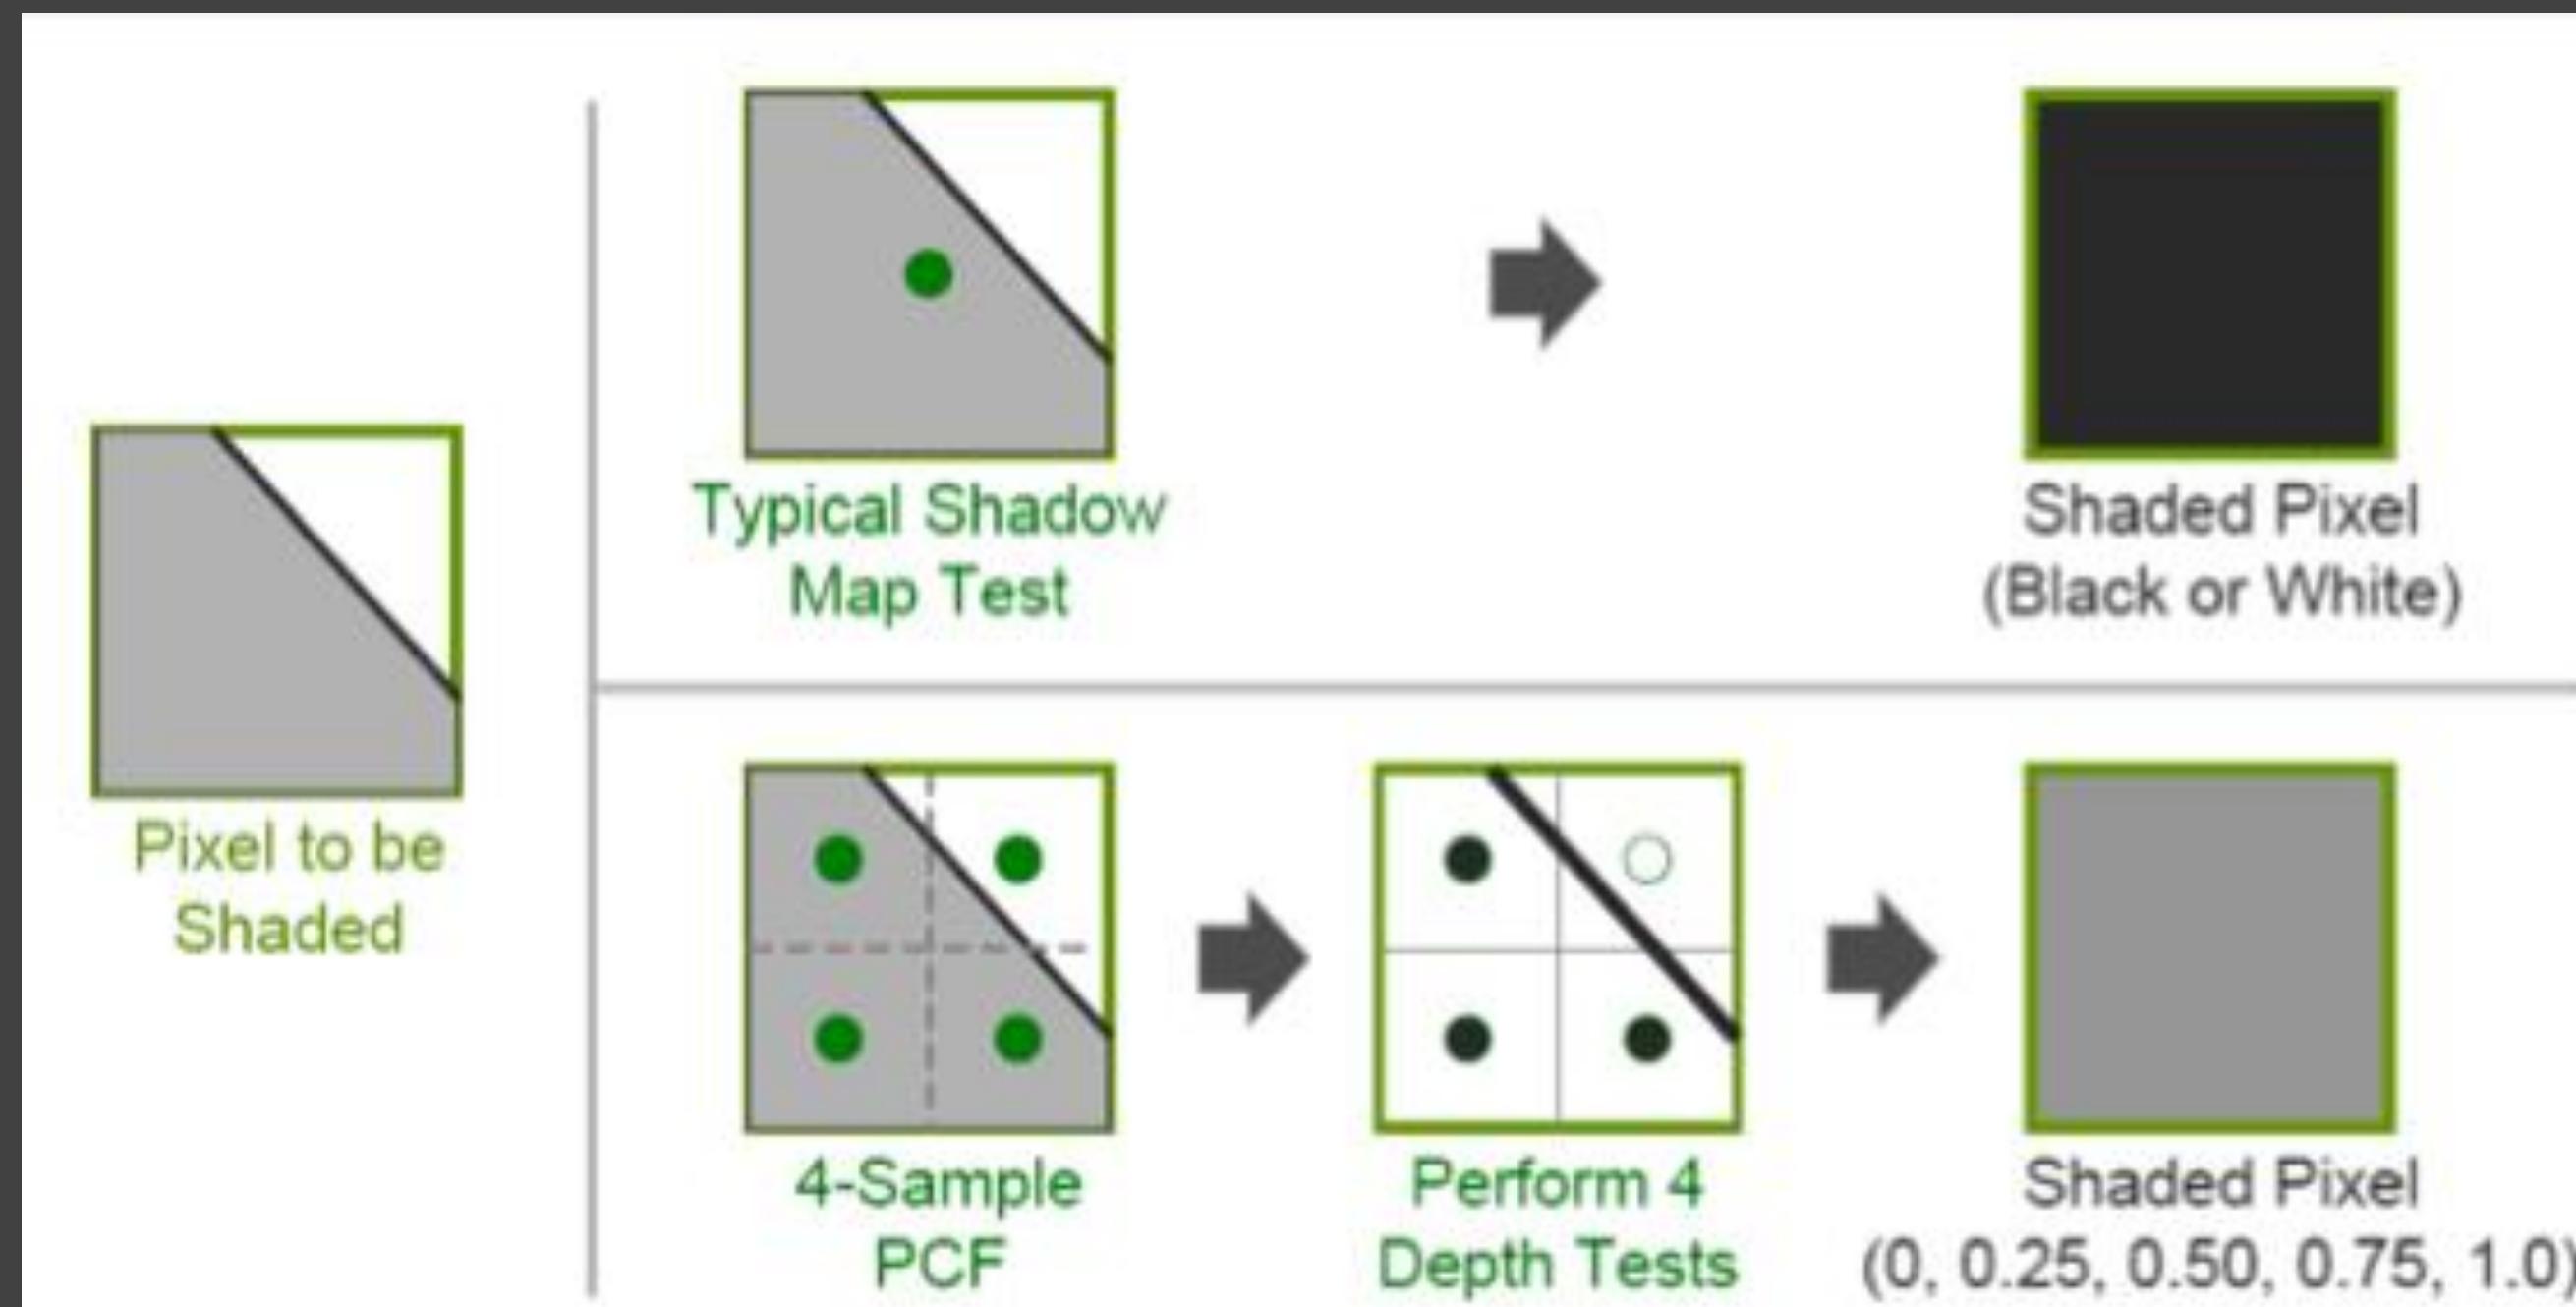
- this means applying a linear filter to the depth is wrong

**We want to filter the output, not the input, of the shadow test**

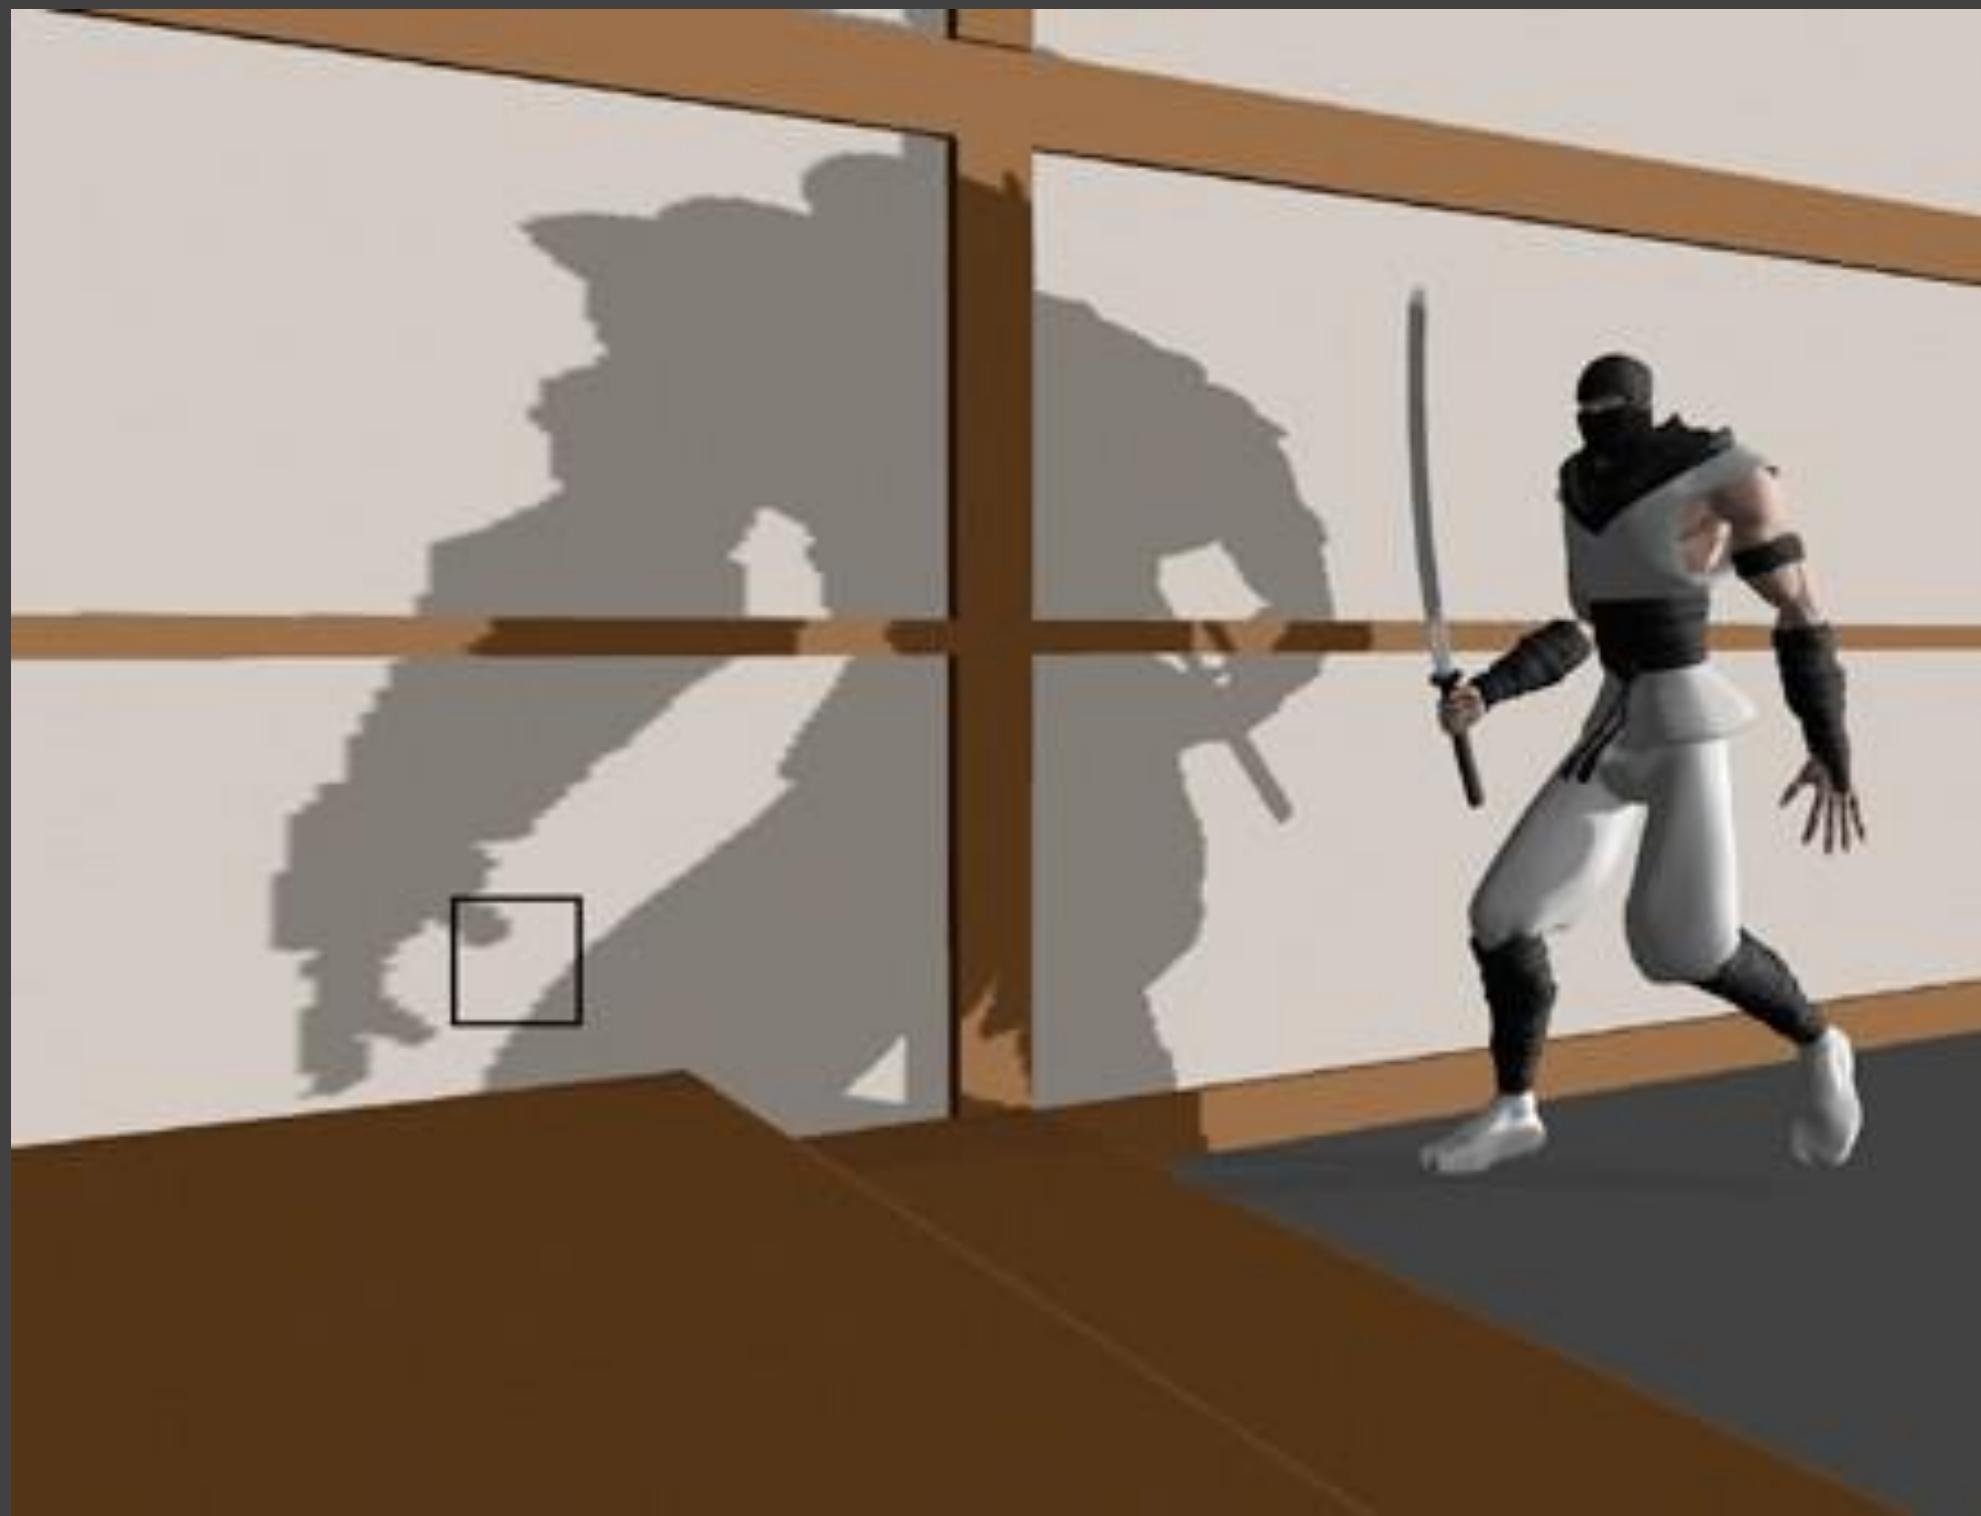
- what fraction of samples pass the test
- samples pass the test if they are closer than the shadow map depth
- therefore “percentage closer filtering” or PCF

# Percentage Closer Filtering

- Soften the shadow to decrease aliasing
  - Reeves, Salesin, Cook 87
  - GPU Gems, Chapter 11



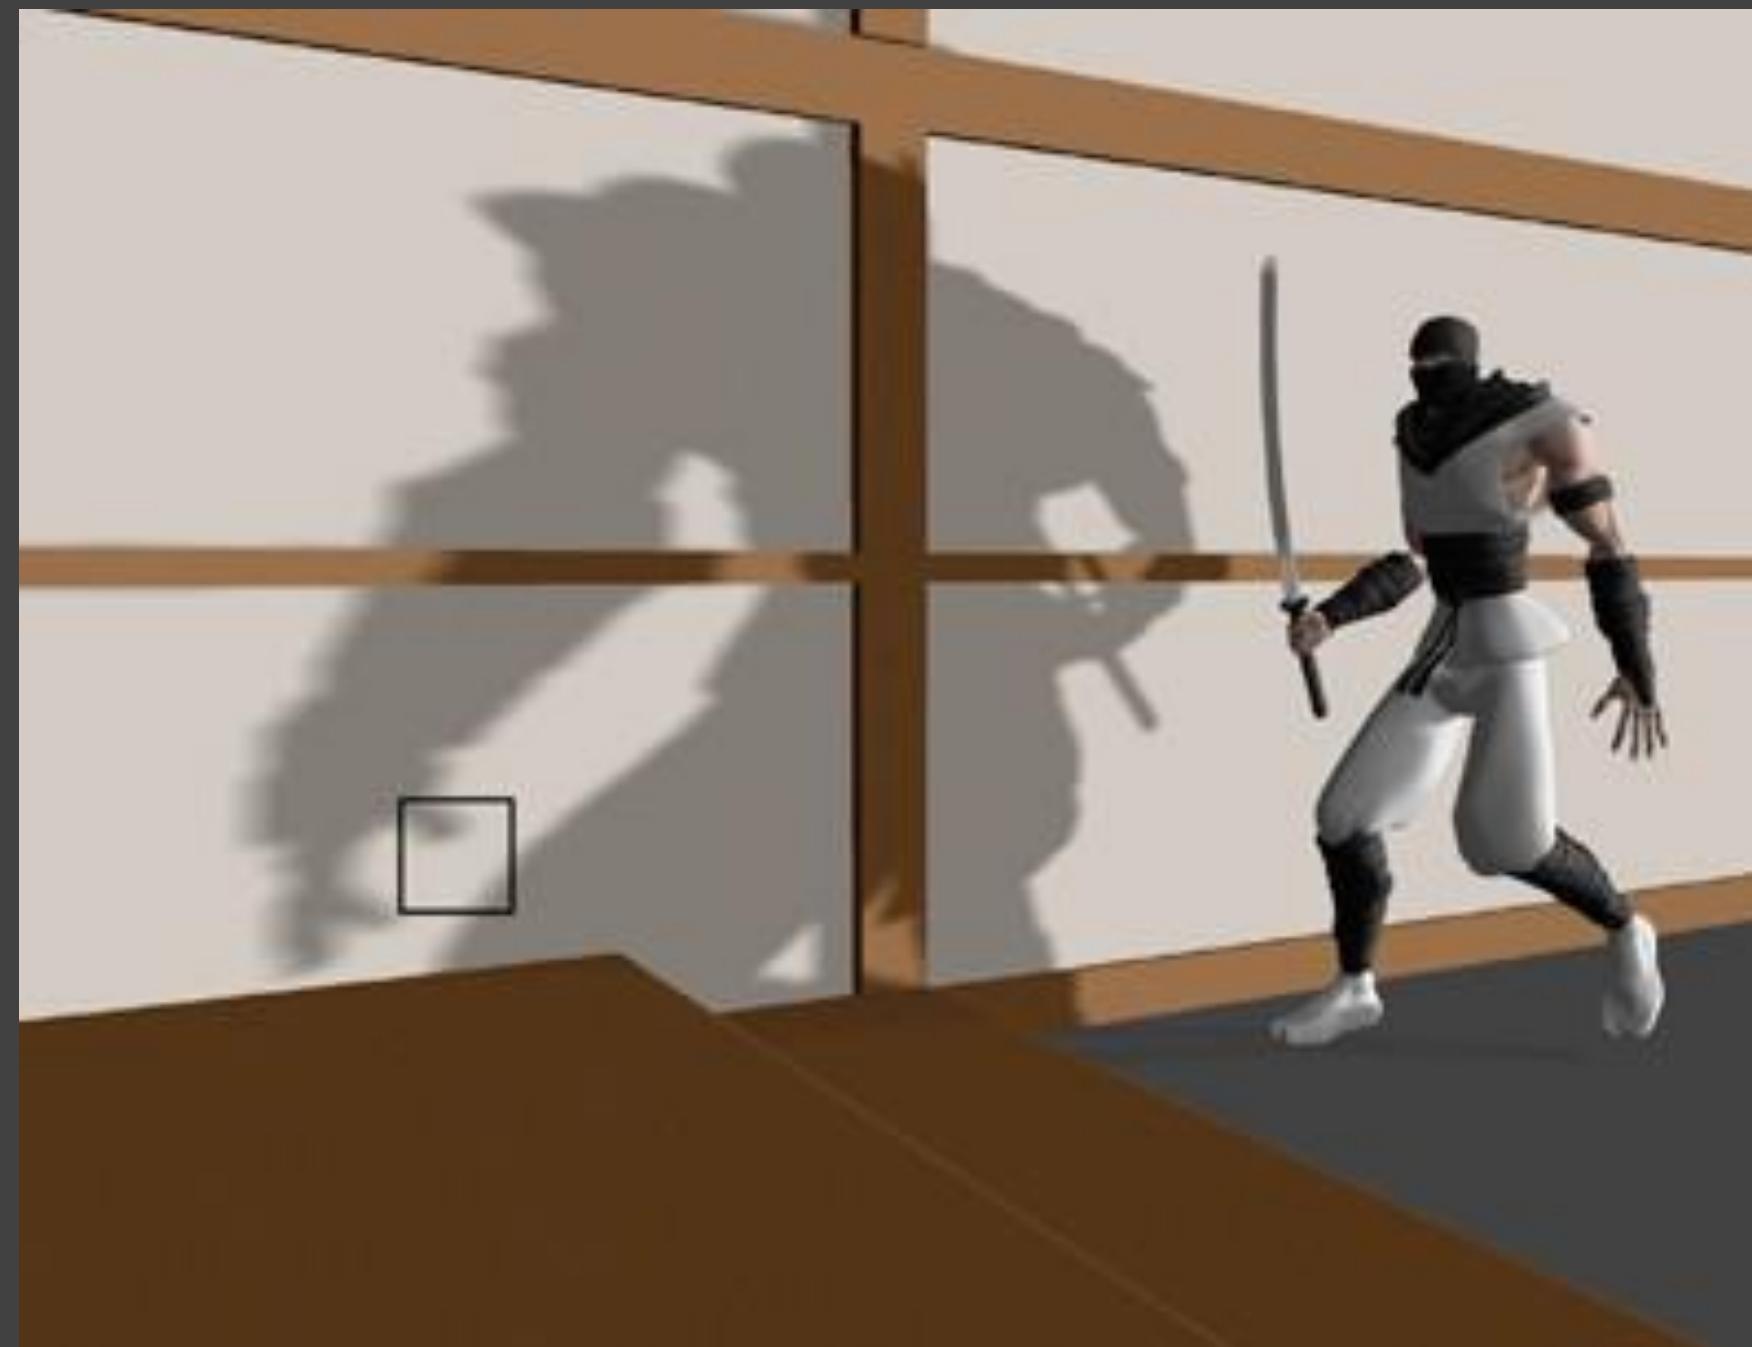
# 1 sample SM

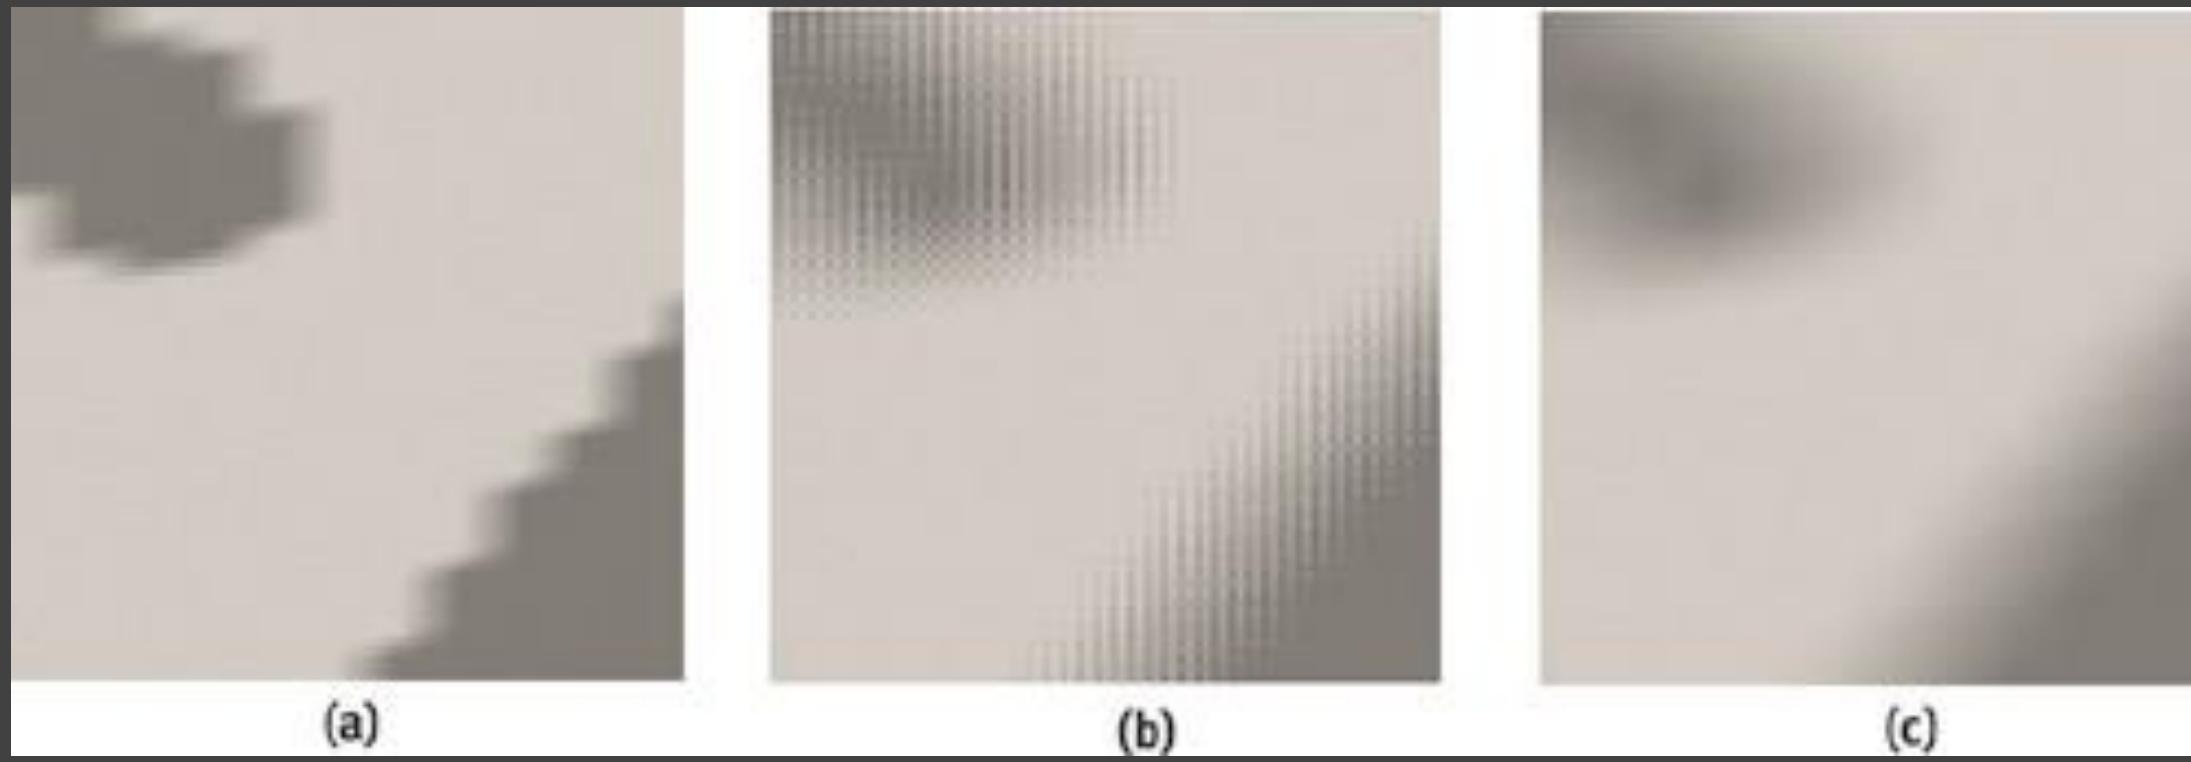


# 4 sample PCF



# 16 sample PCF





# Soft shadows from small sources

**Main effect is to blur shadow boundaries**

- PCF can do this
- ...but how wide to make the filter?

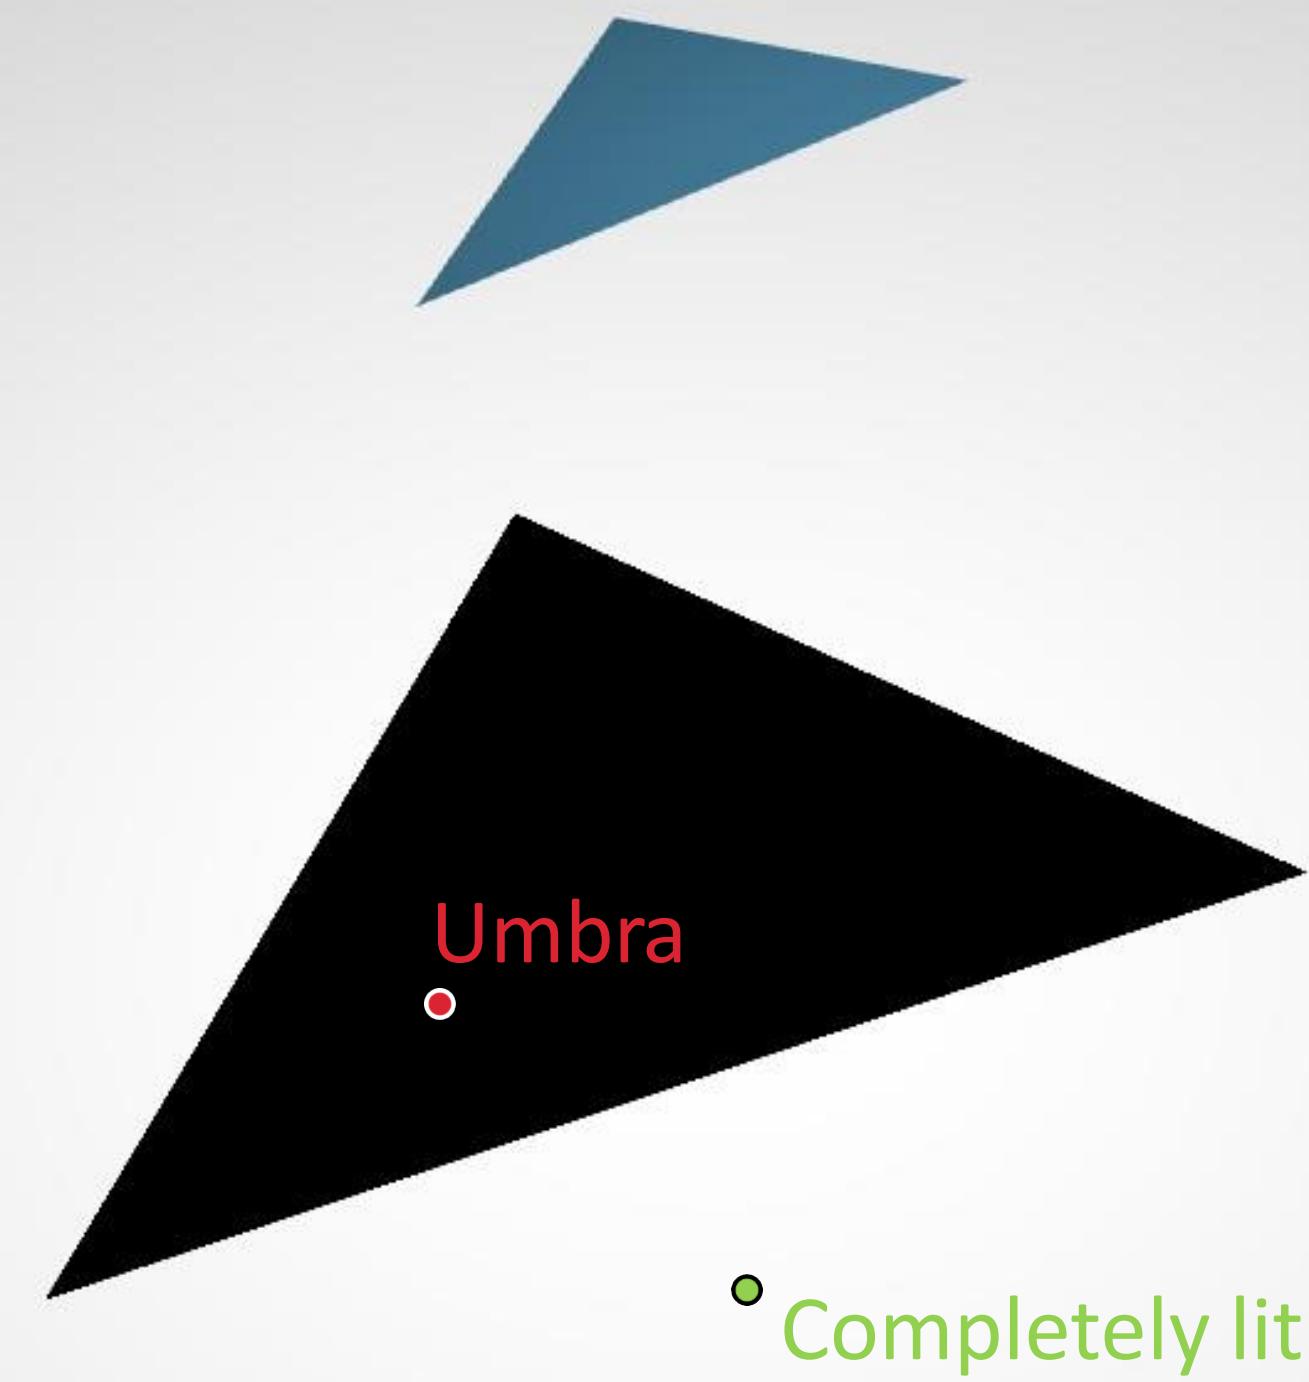
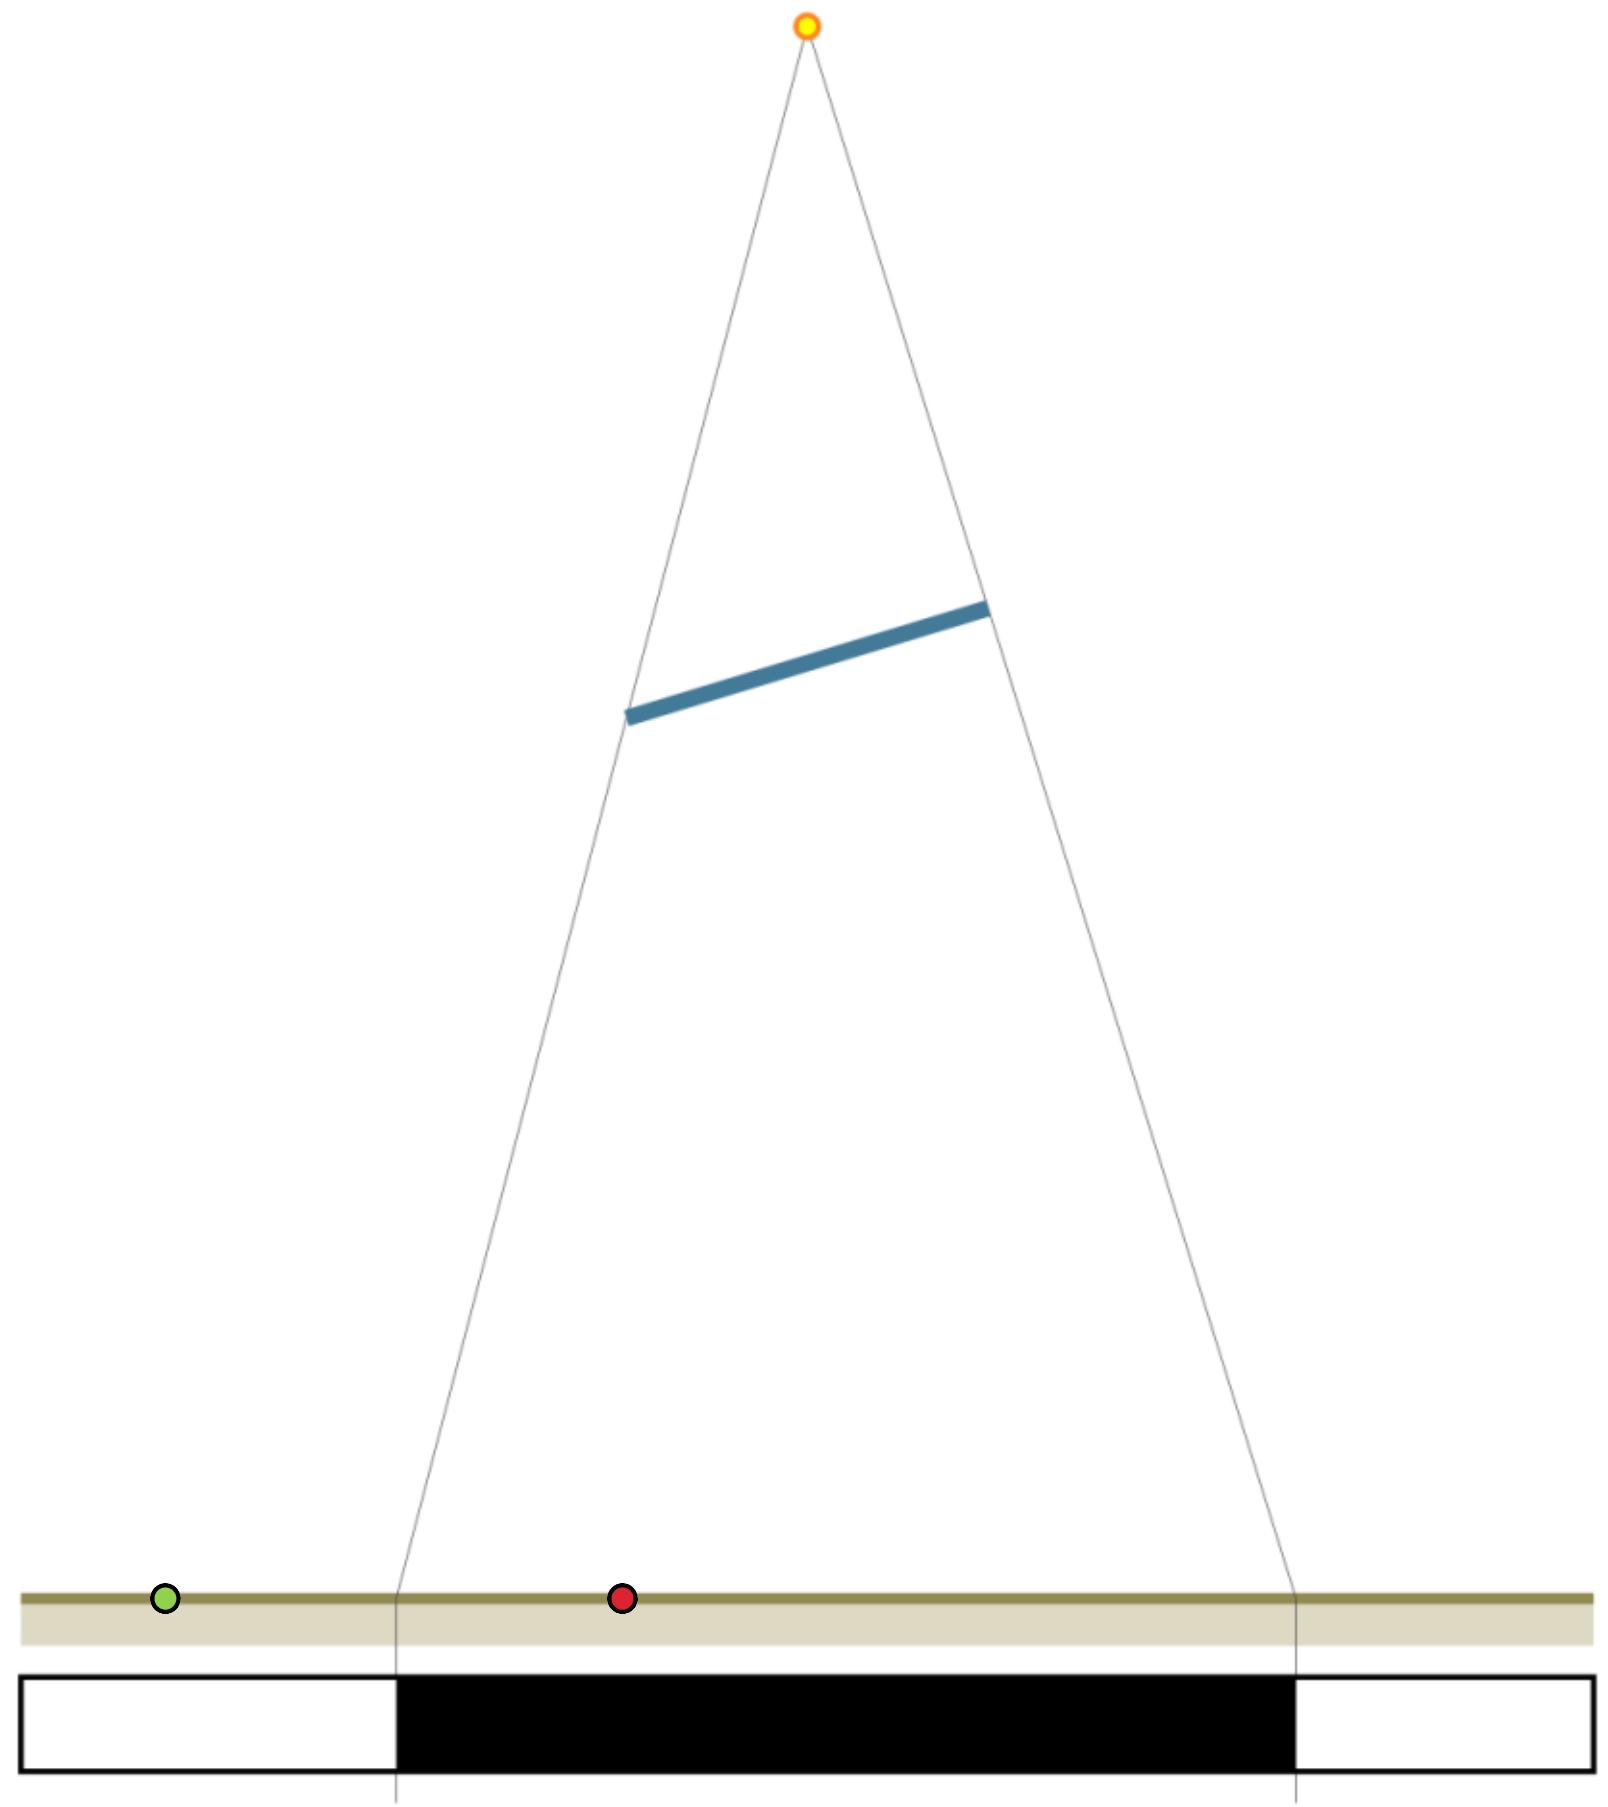
**Real shadows depend on area of light visible from surface**

- this can vary in complex ways
- example: sun viewed through leafy trees

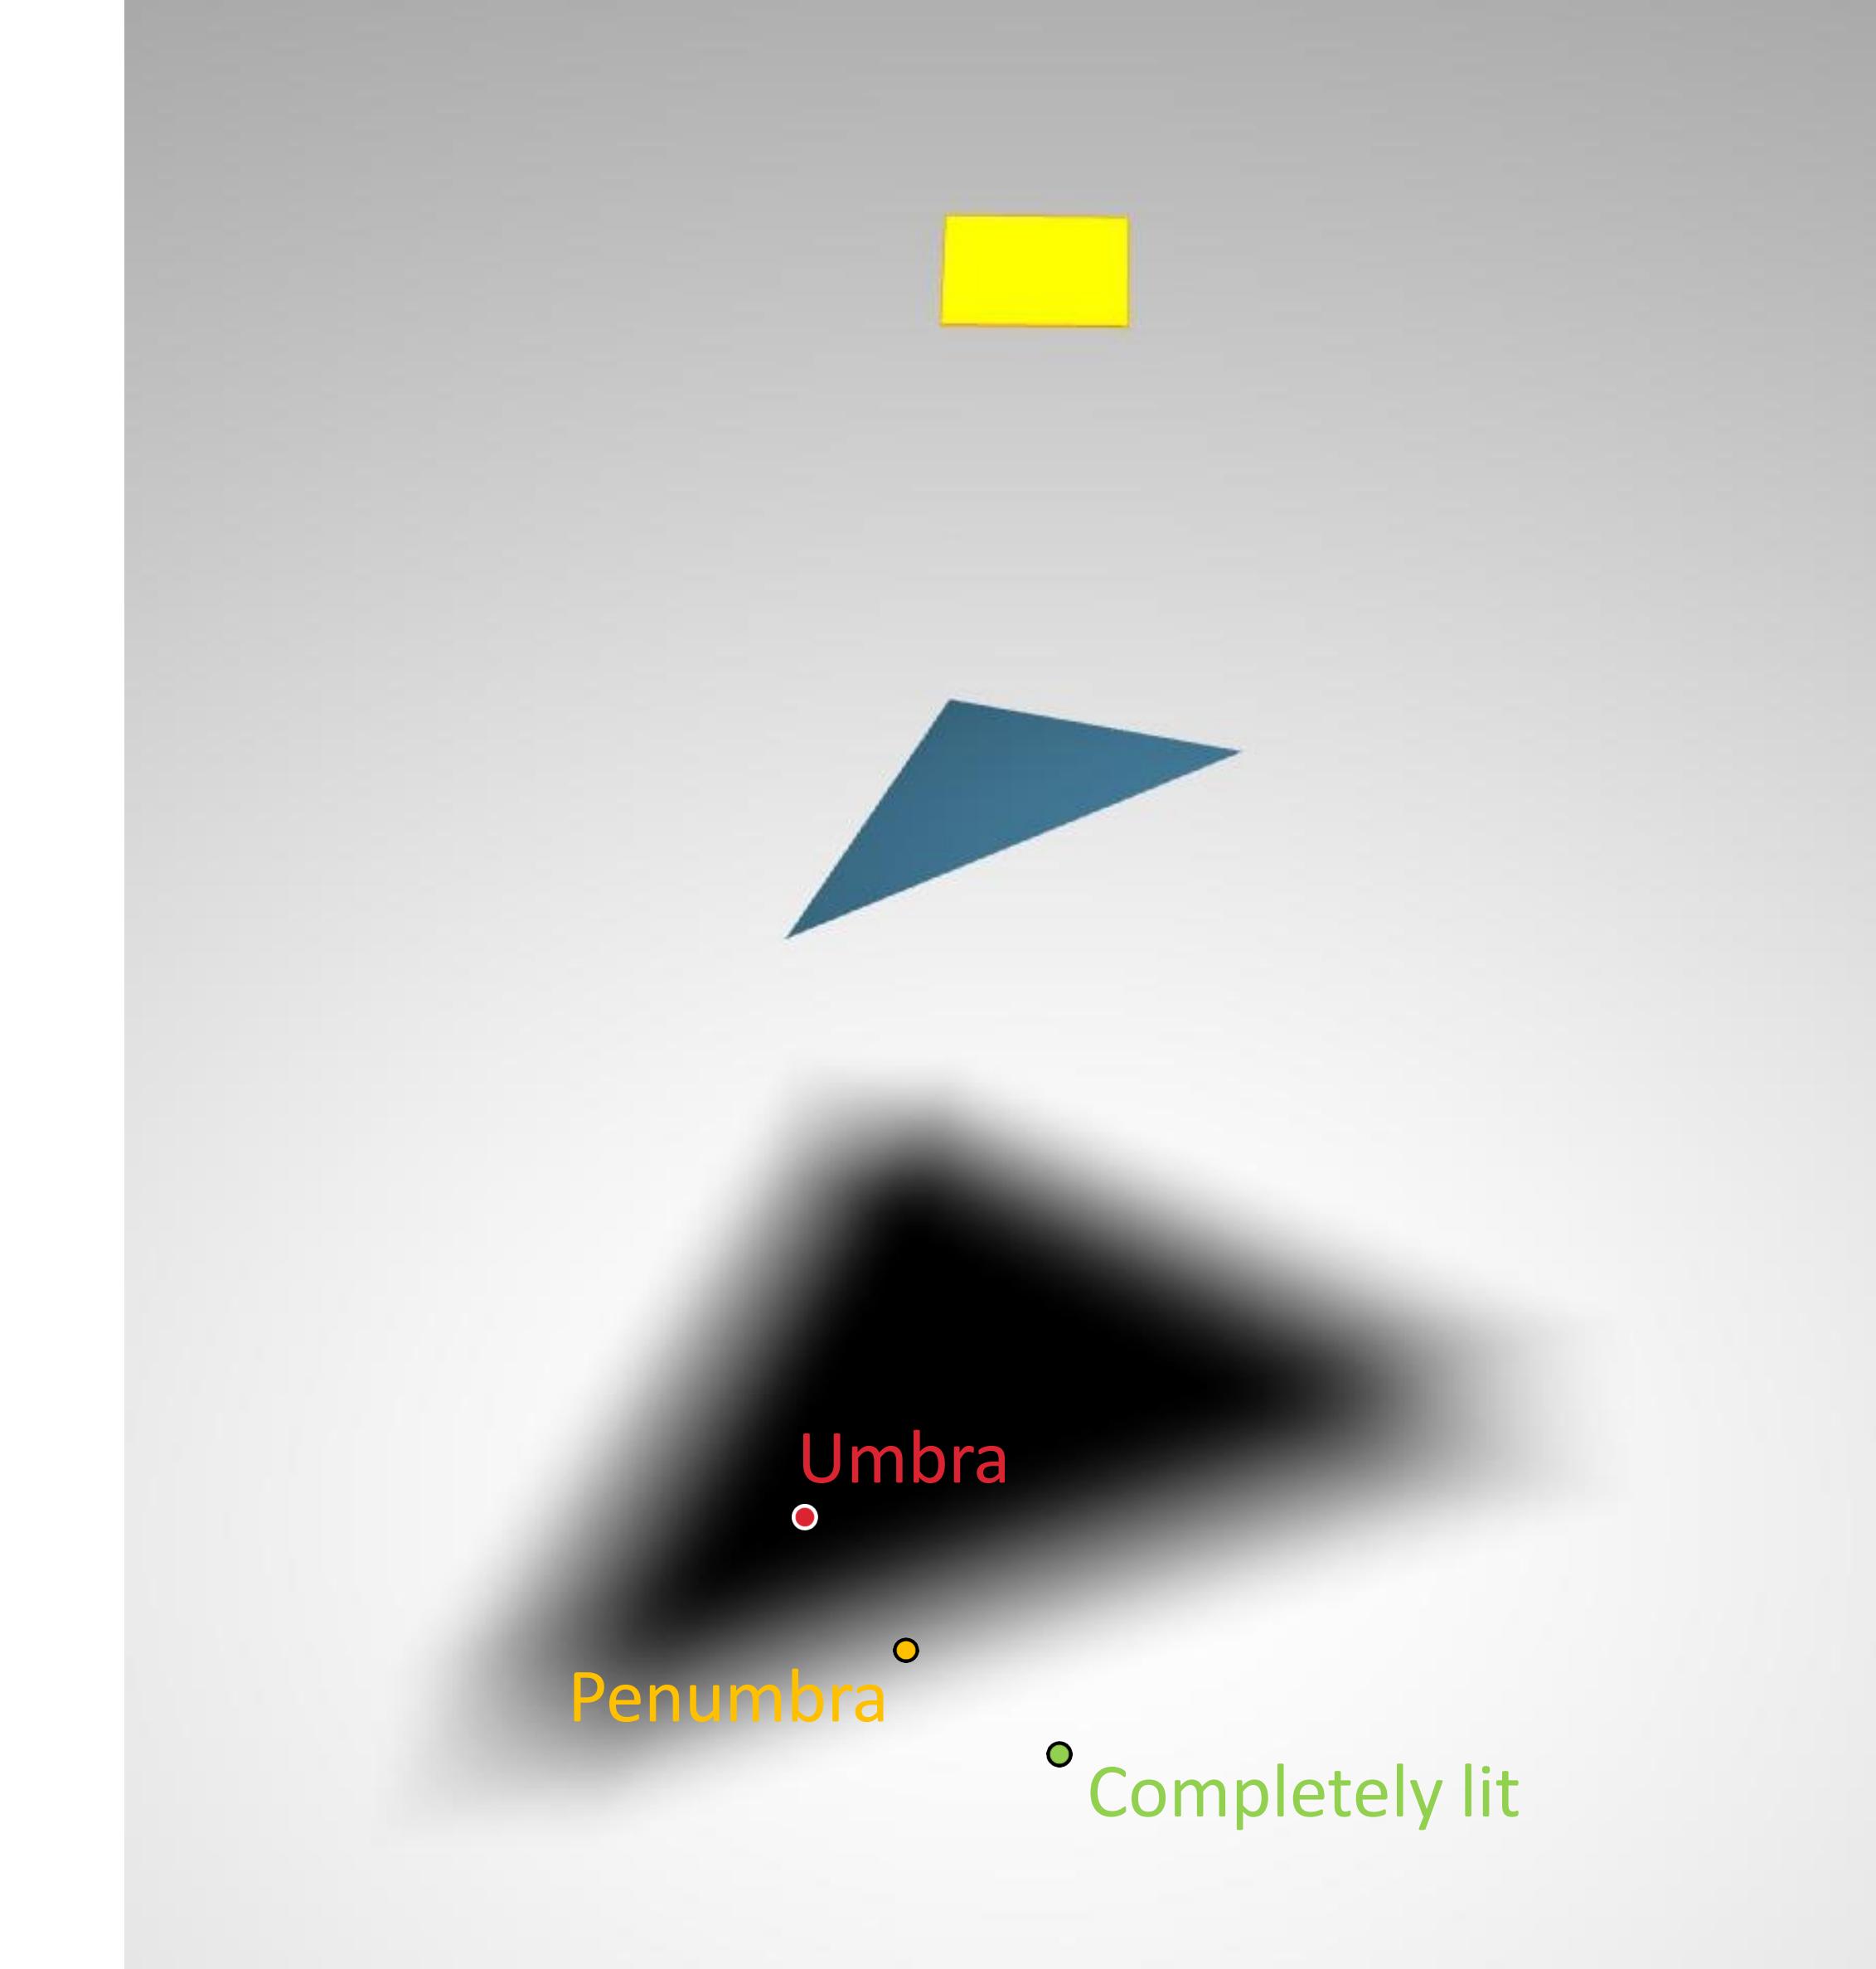
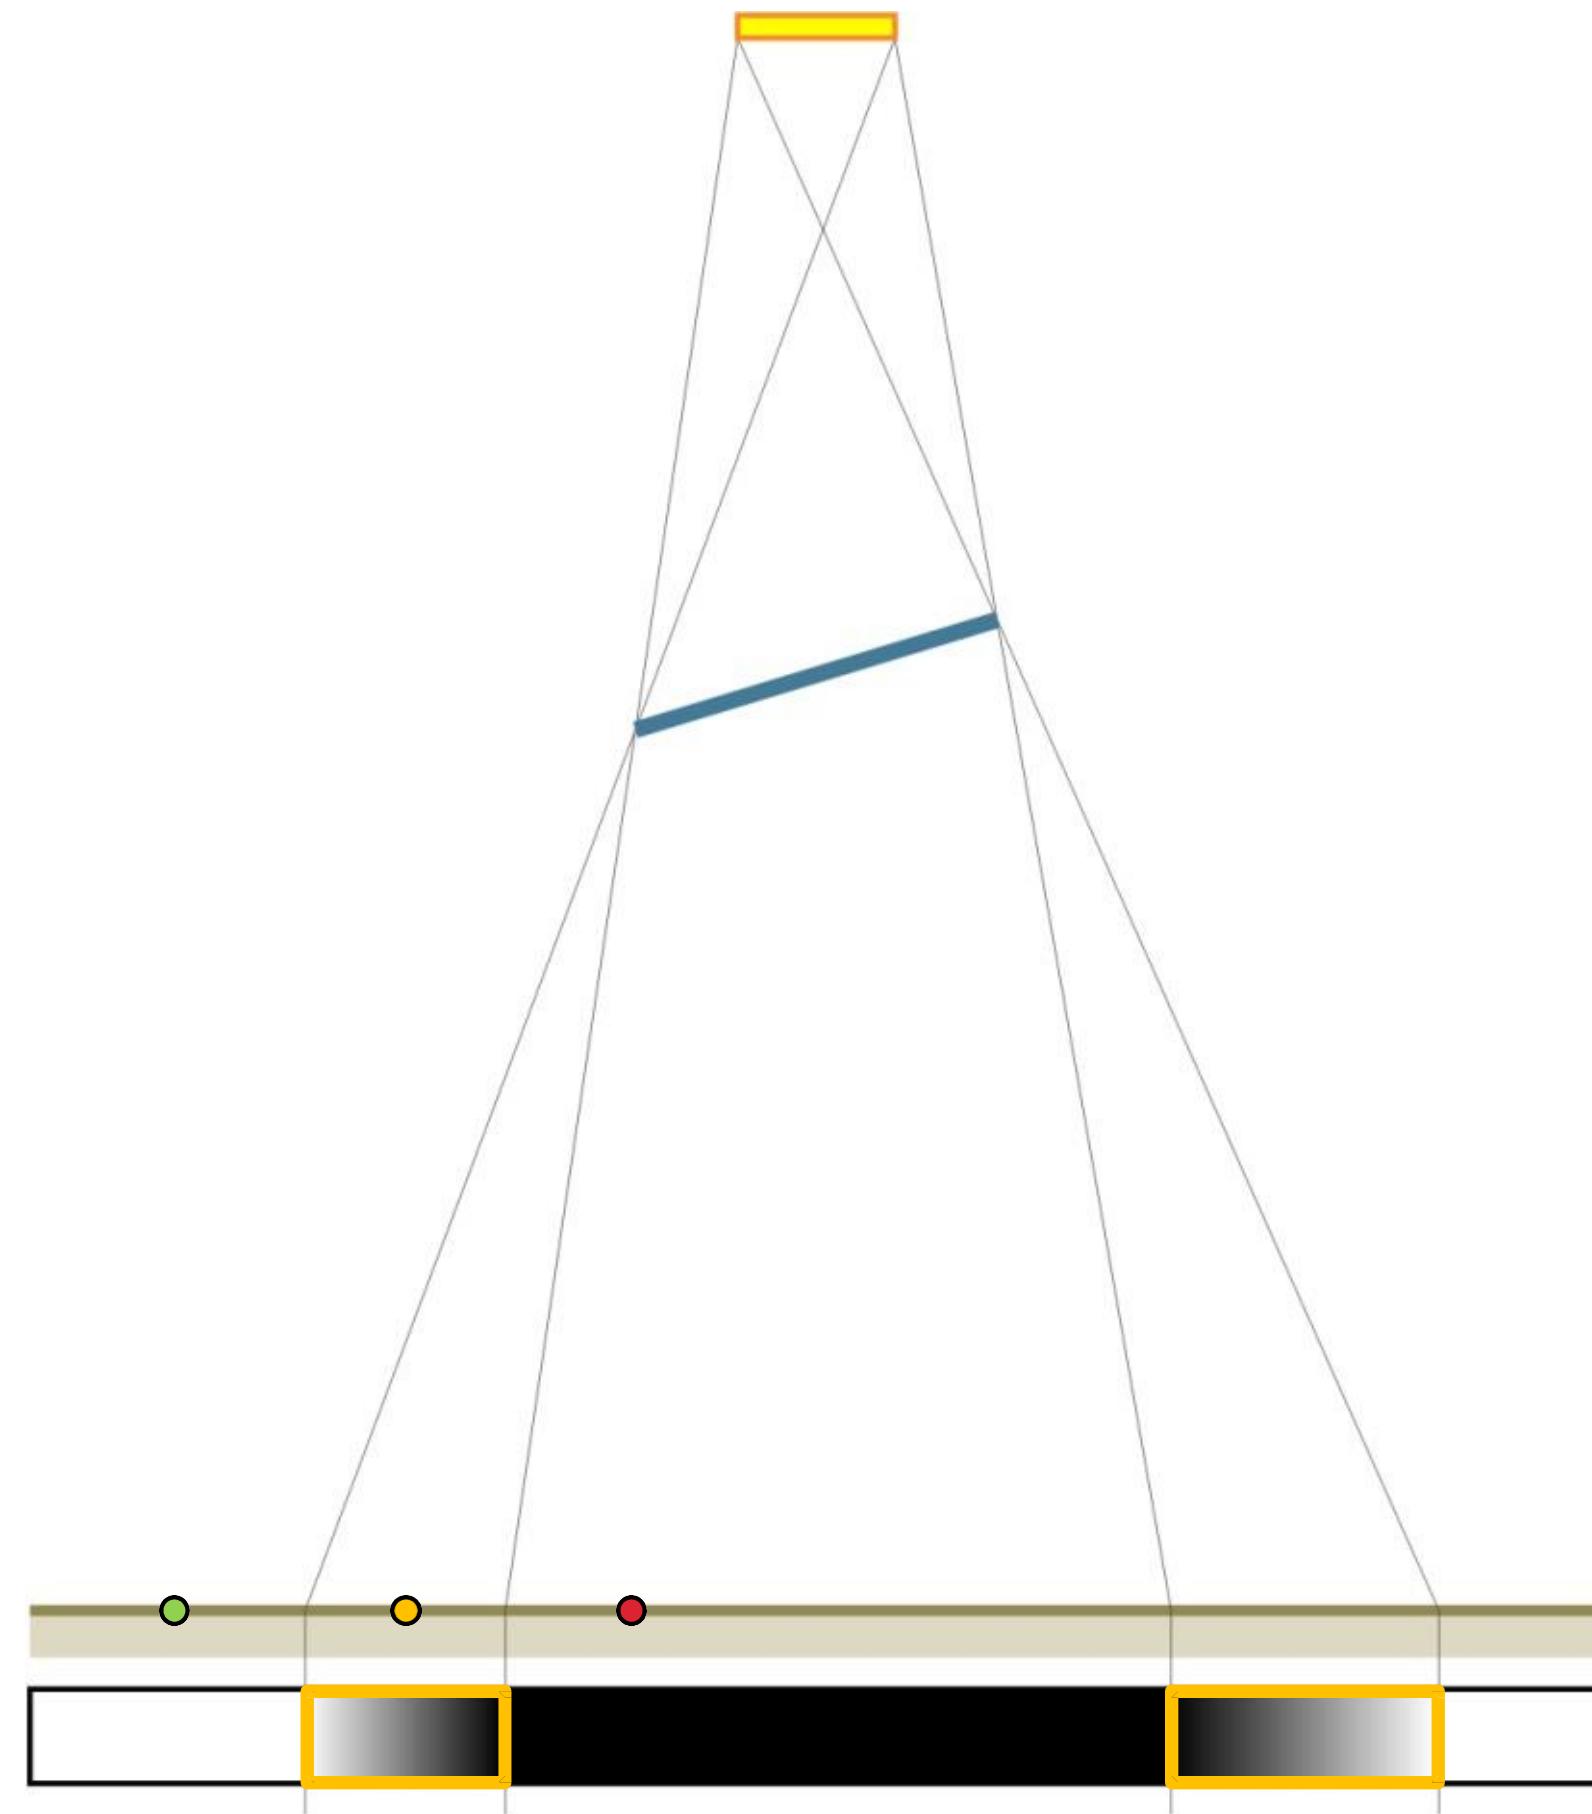
**Useful approximation: convolution**

- shadows are convolutions when the blocker and source are parallel and planar
- occluder fusion: approximating some occluding geometry as a planar blocker

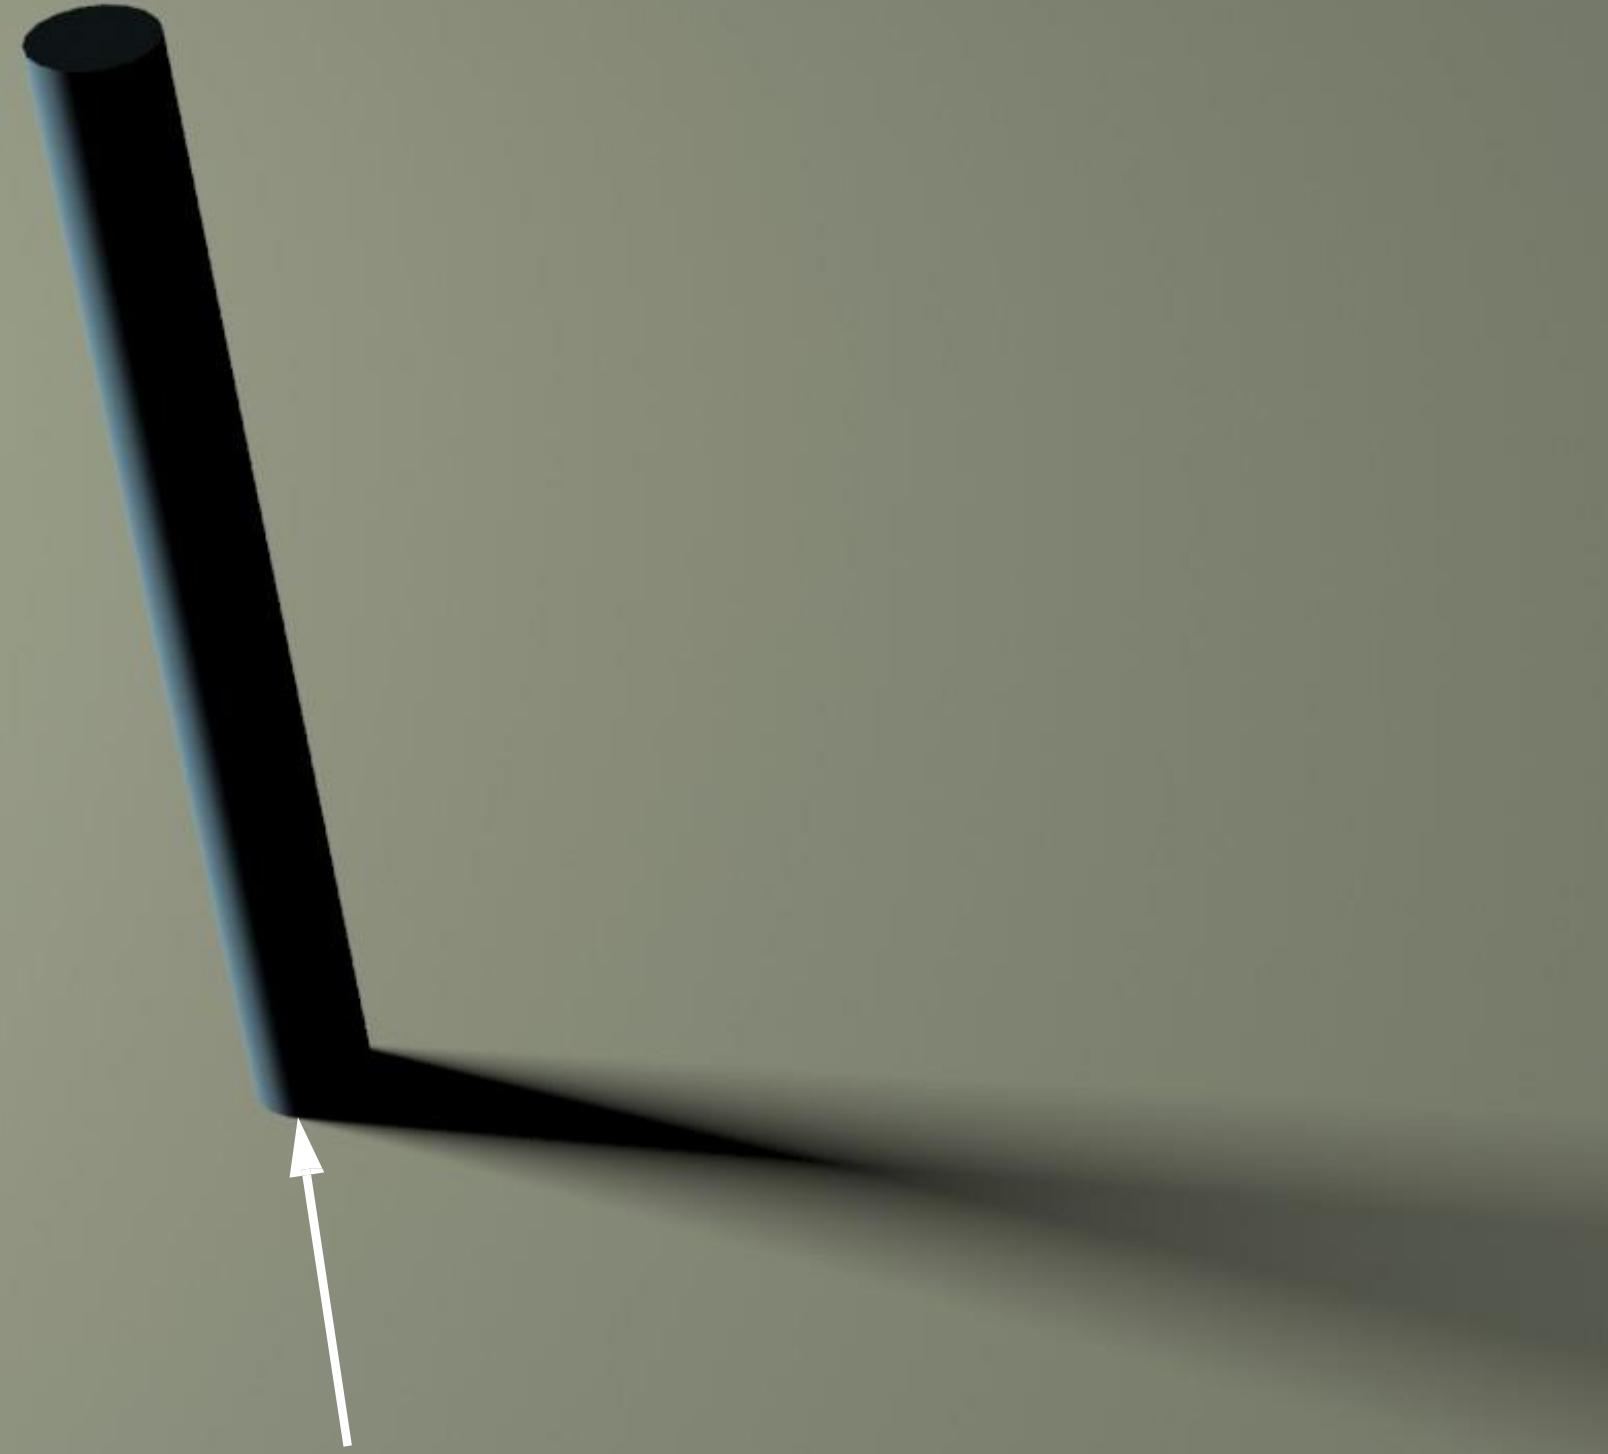
# Hard Shadows



# Soft Shadows

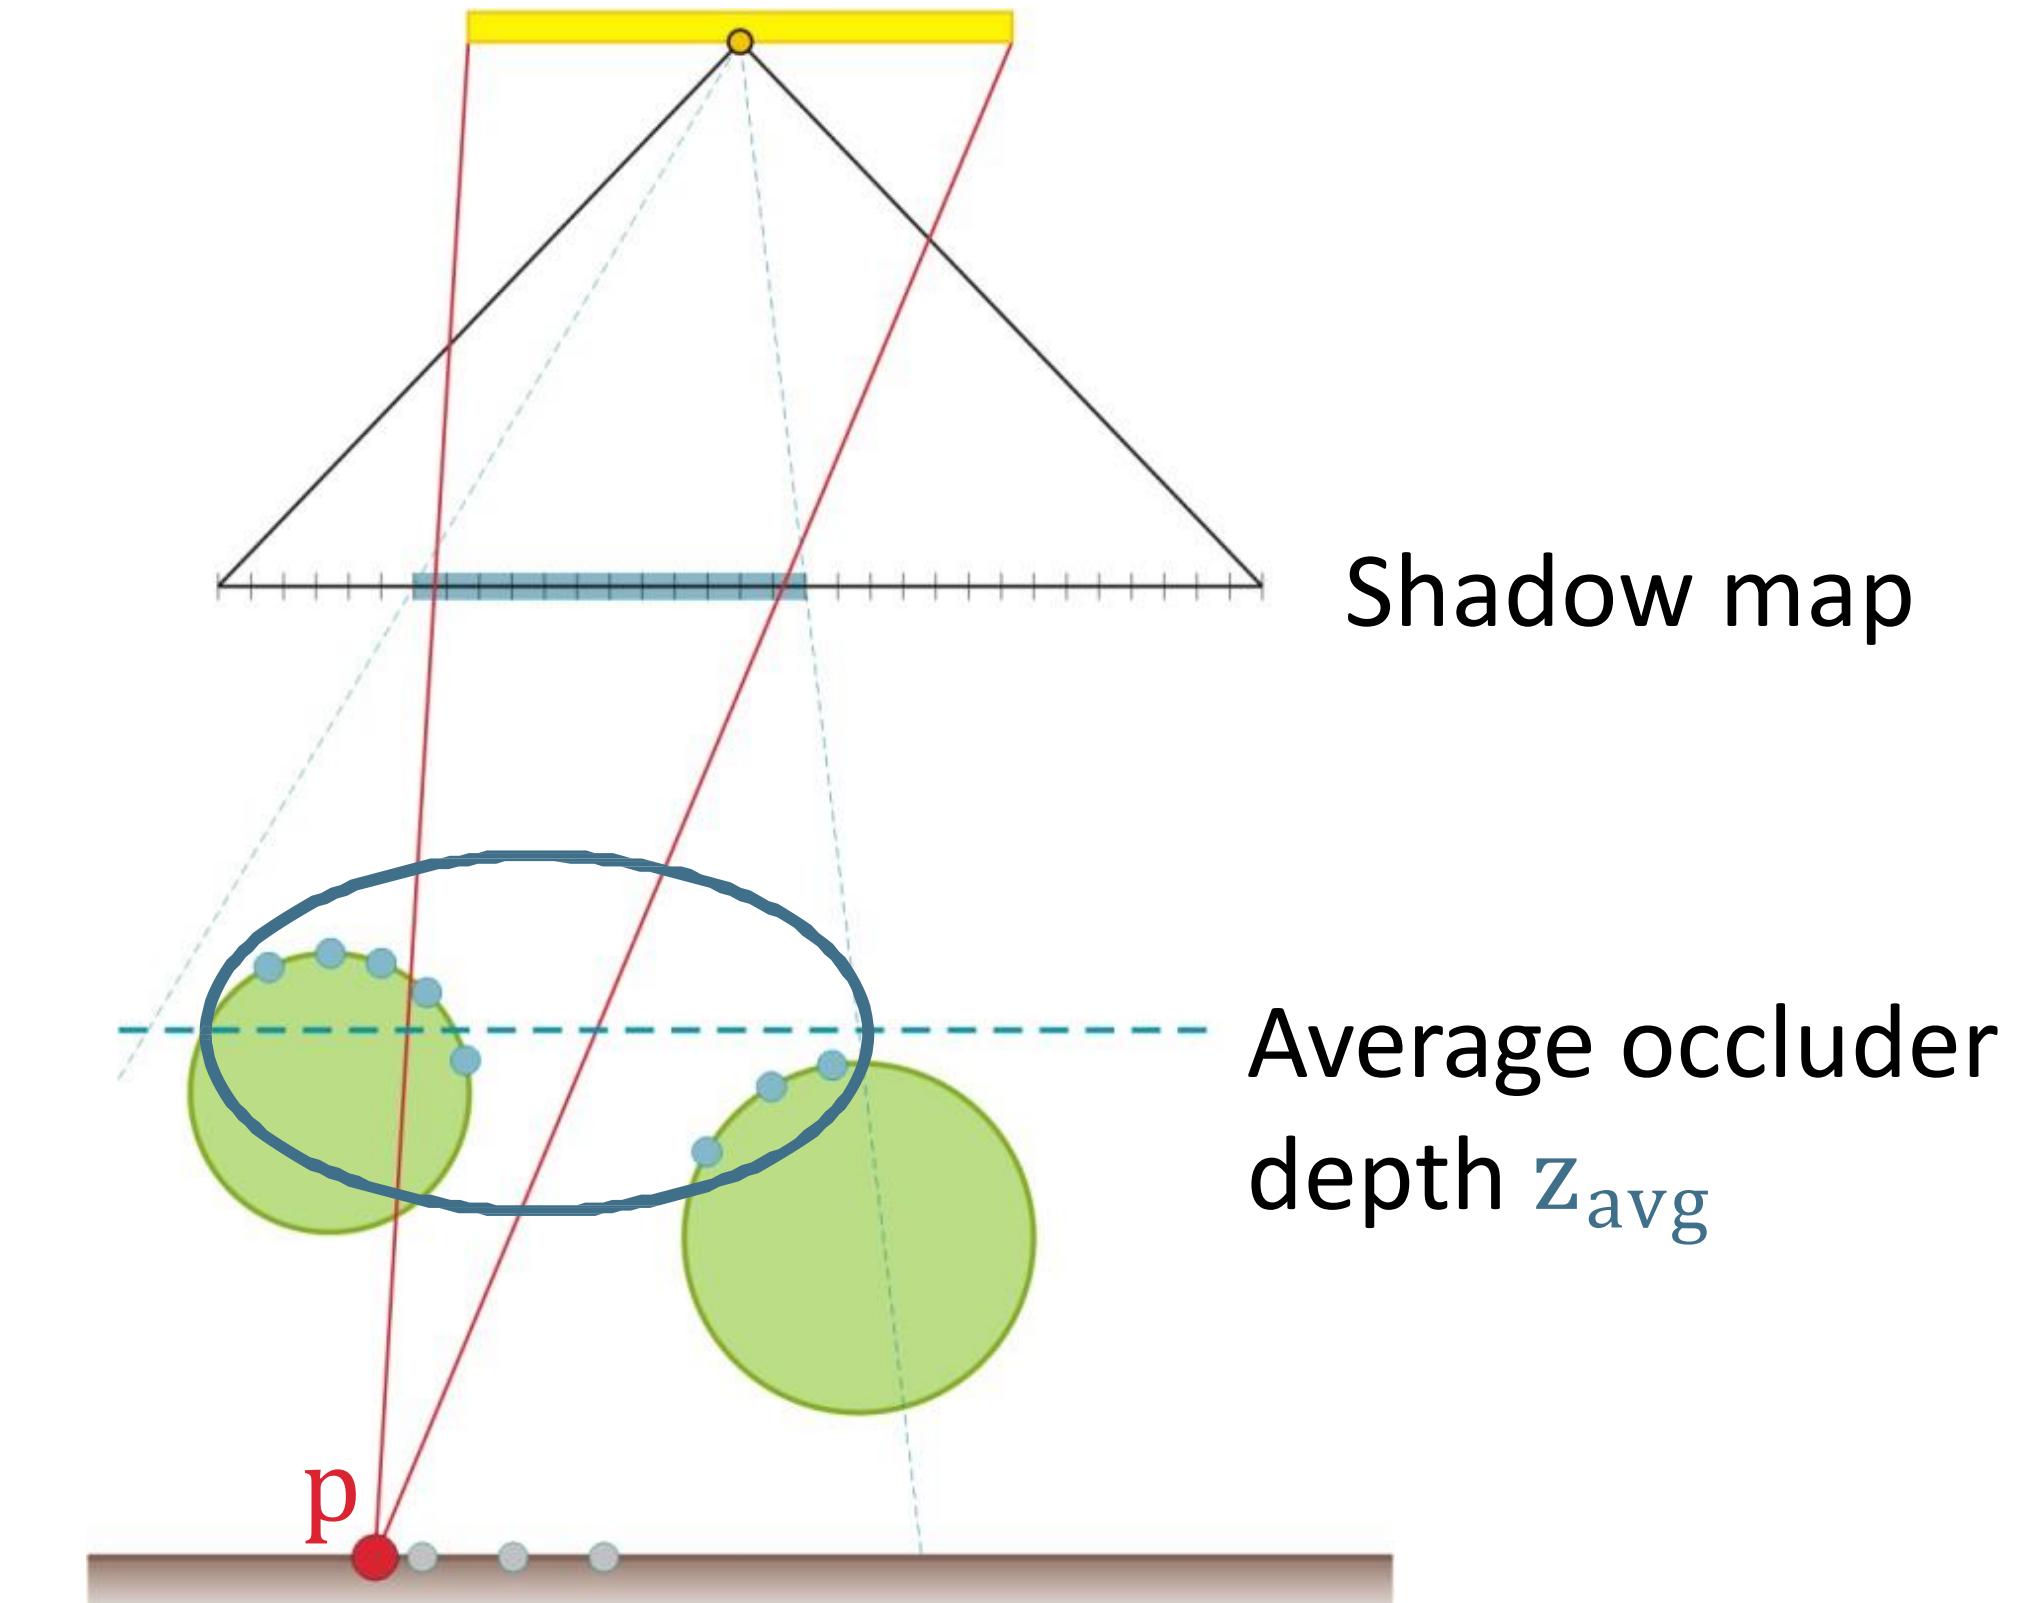


# Shadow Hardening on Contact



# Percentage-Closer Soft Shadows

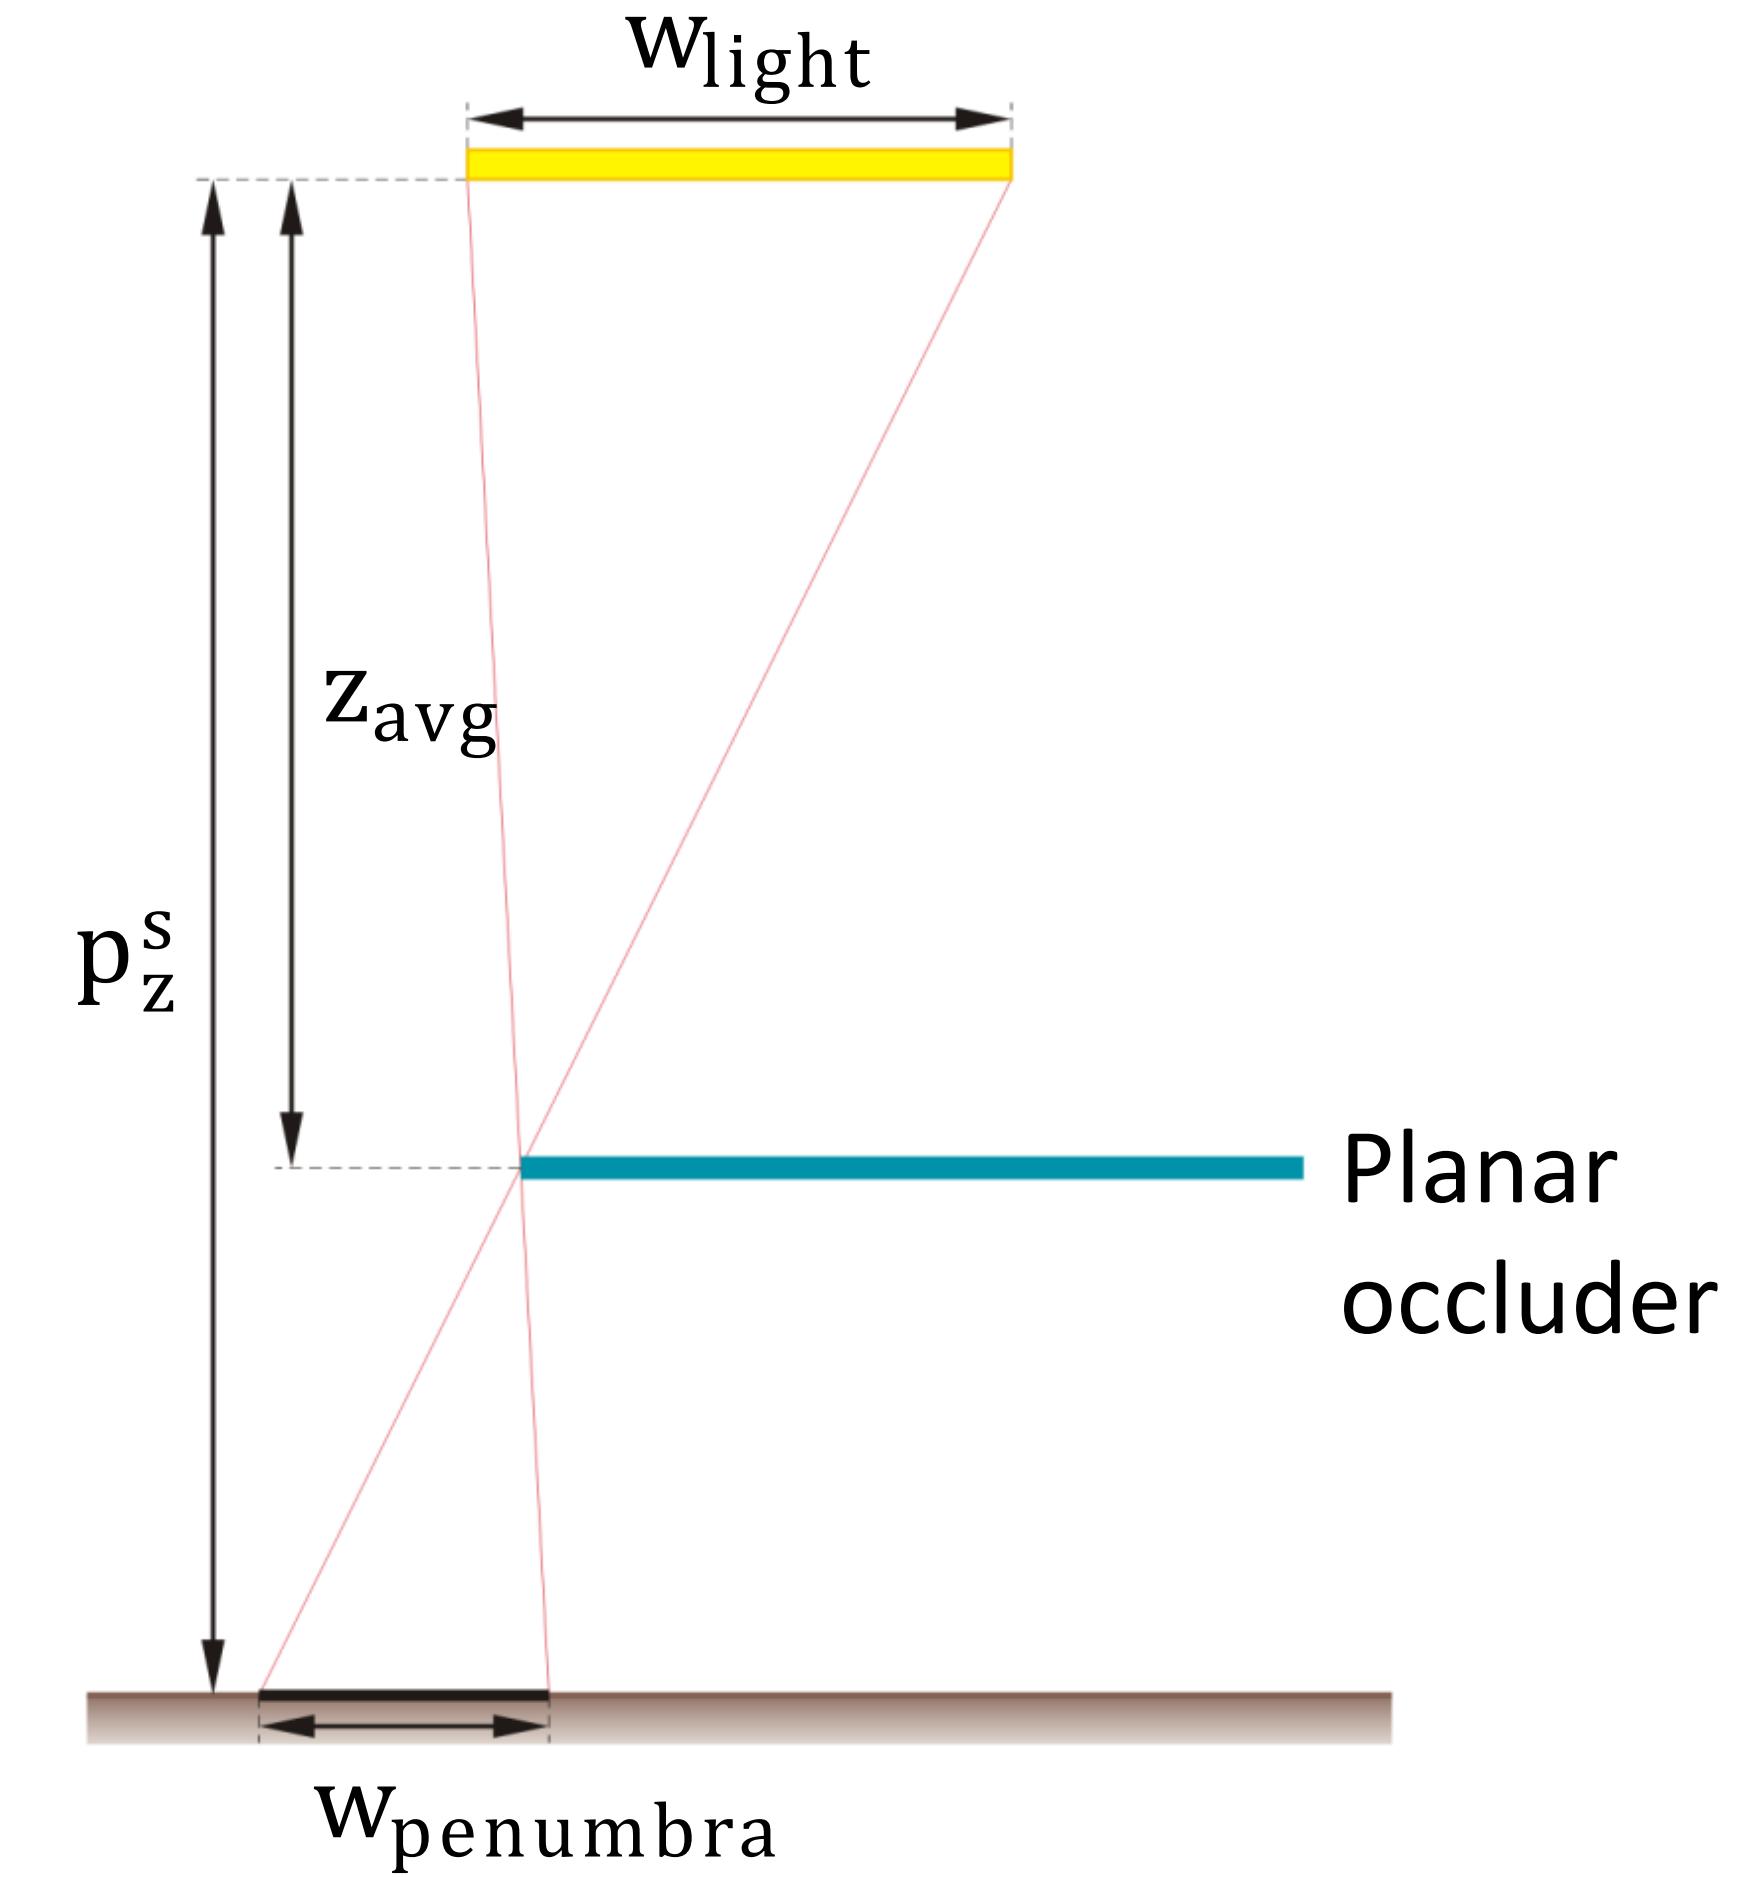
## 1. Blocker search



# Percentage-Closer Soft Shadows

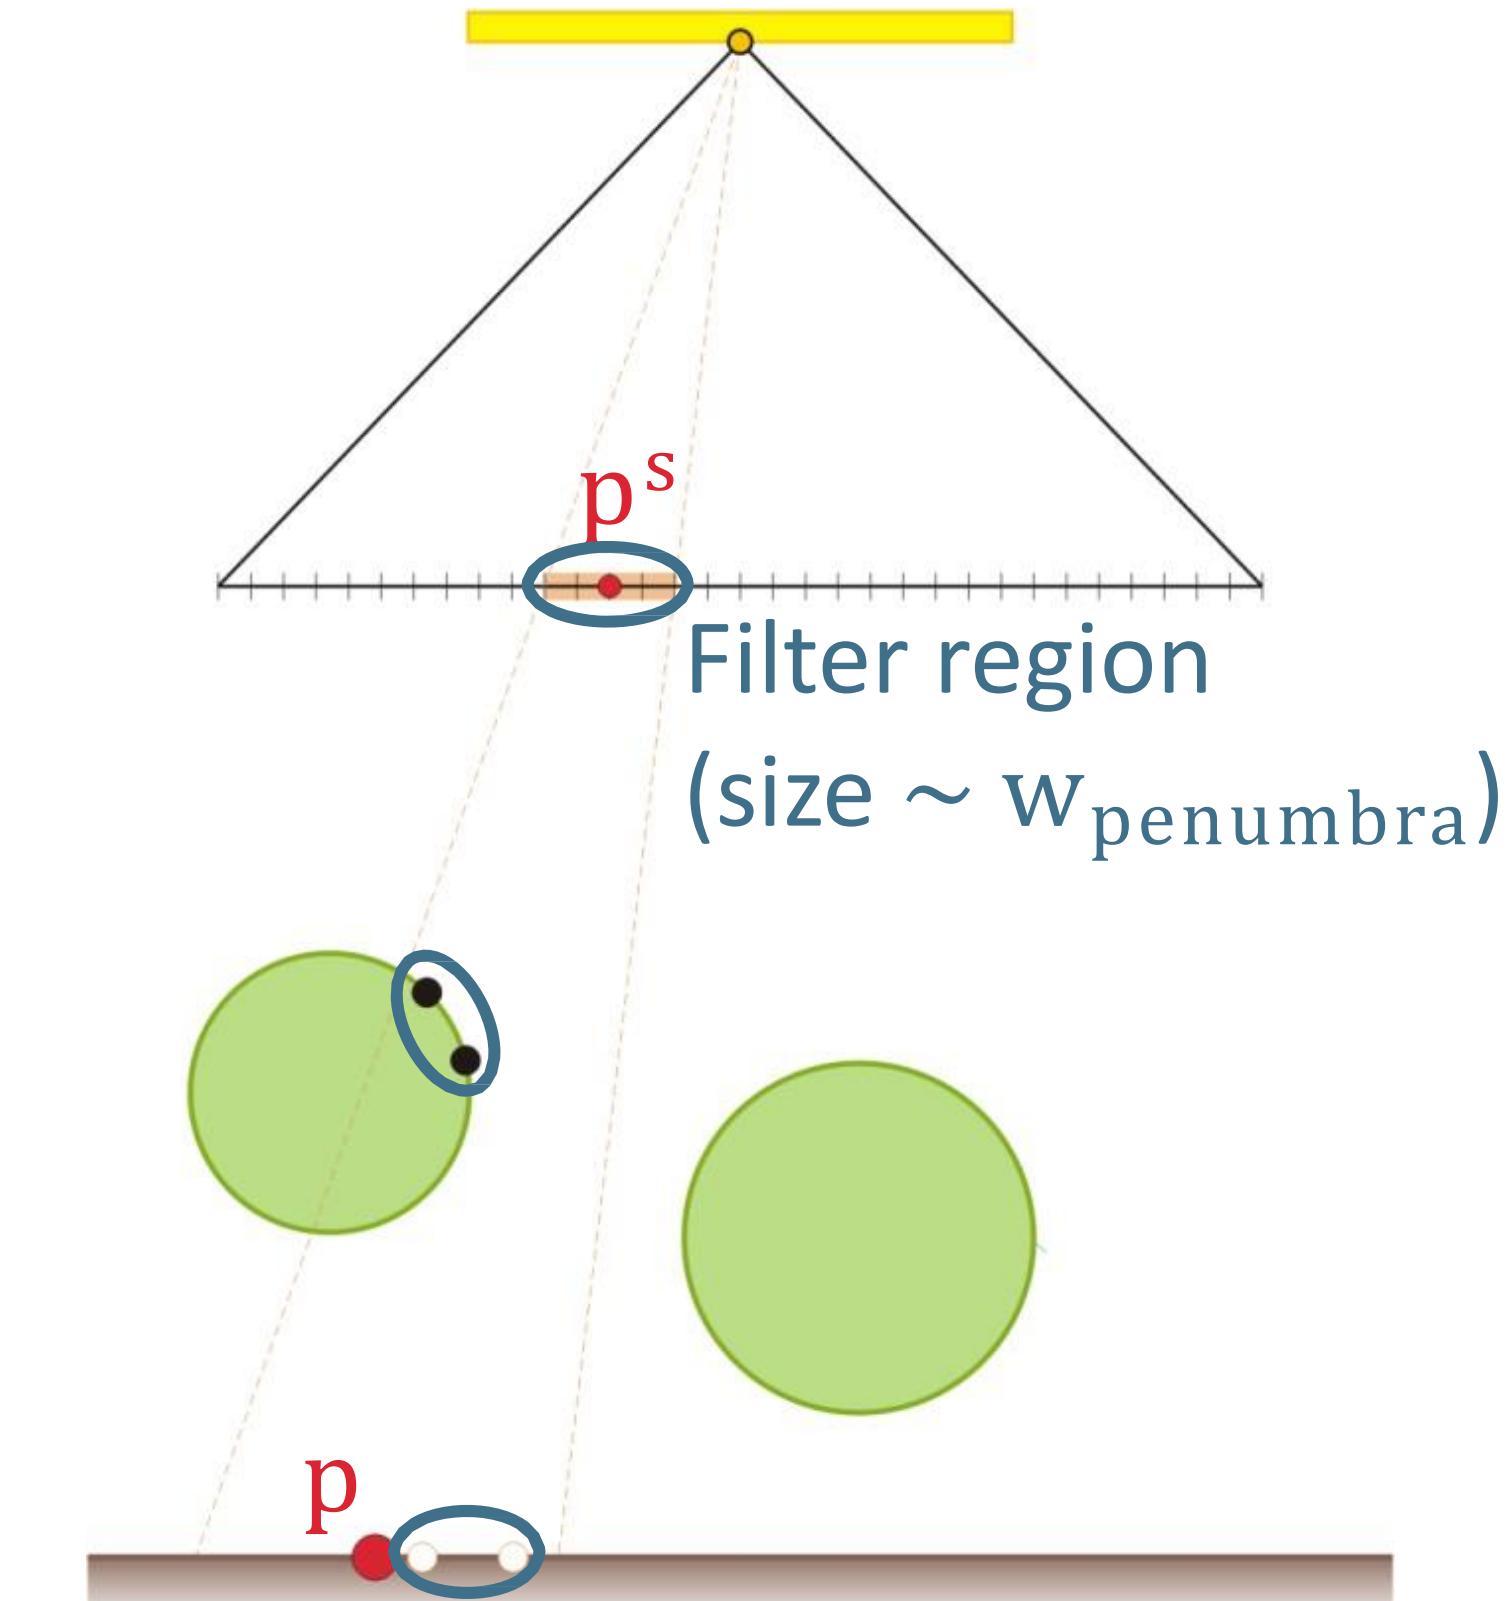
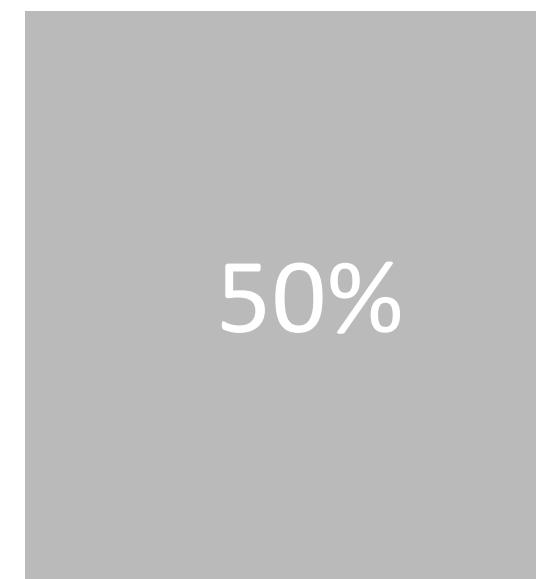
1. Blocker search
2. Penumbra width estimation

$$w_{\text{penumbra}} = \frac{p_z^s - z_{\text{avg}}}{z_{\text{avg}}} w_{\text{light}}$$



# Percentage-Closer Soft Shadows

1. Blocker search
2. Penumbra width estimation
3. Filtering



# Percentage-closer soft shadows



Fernando, NVidia whitepaper ~2005

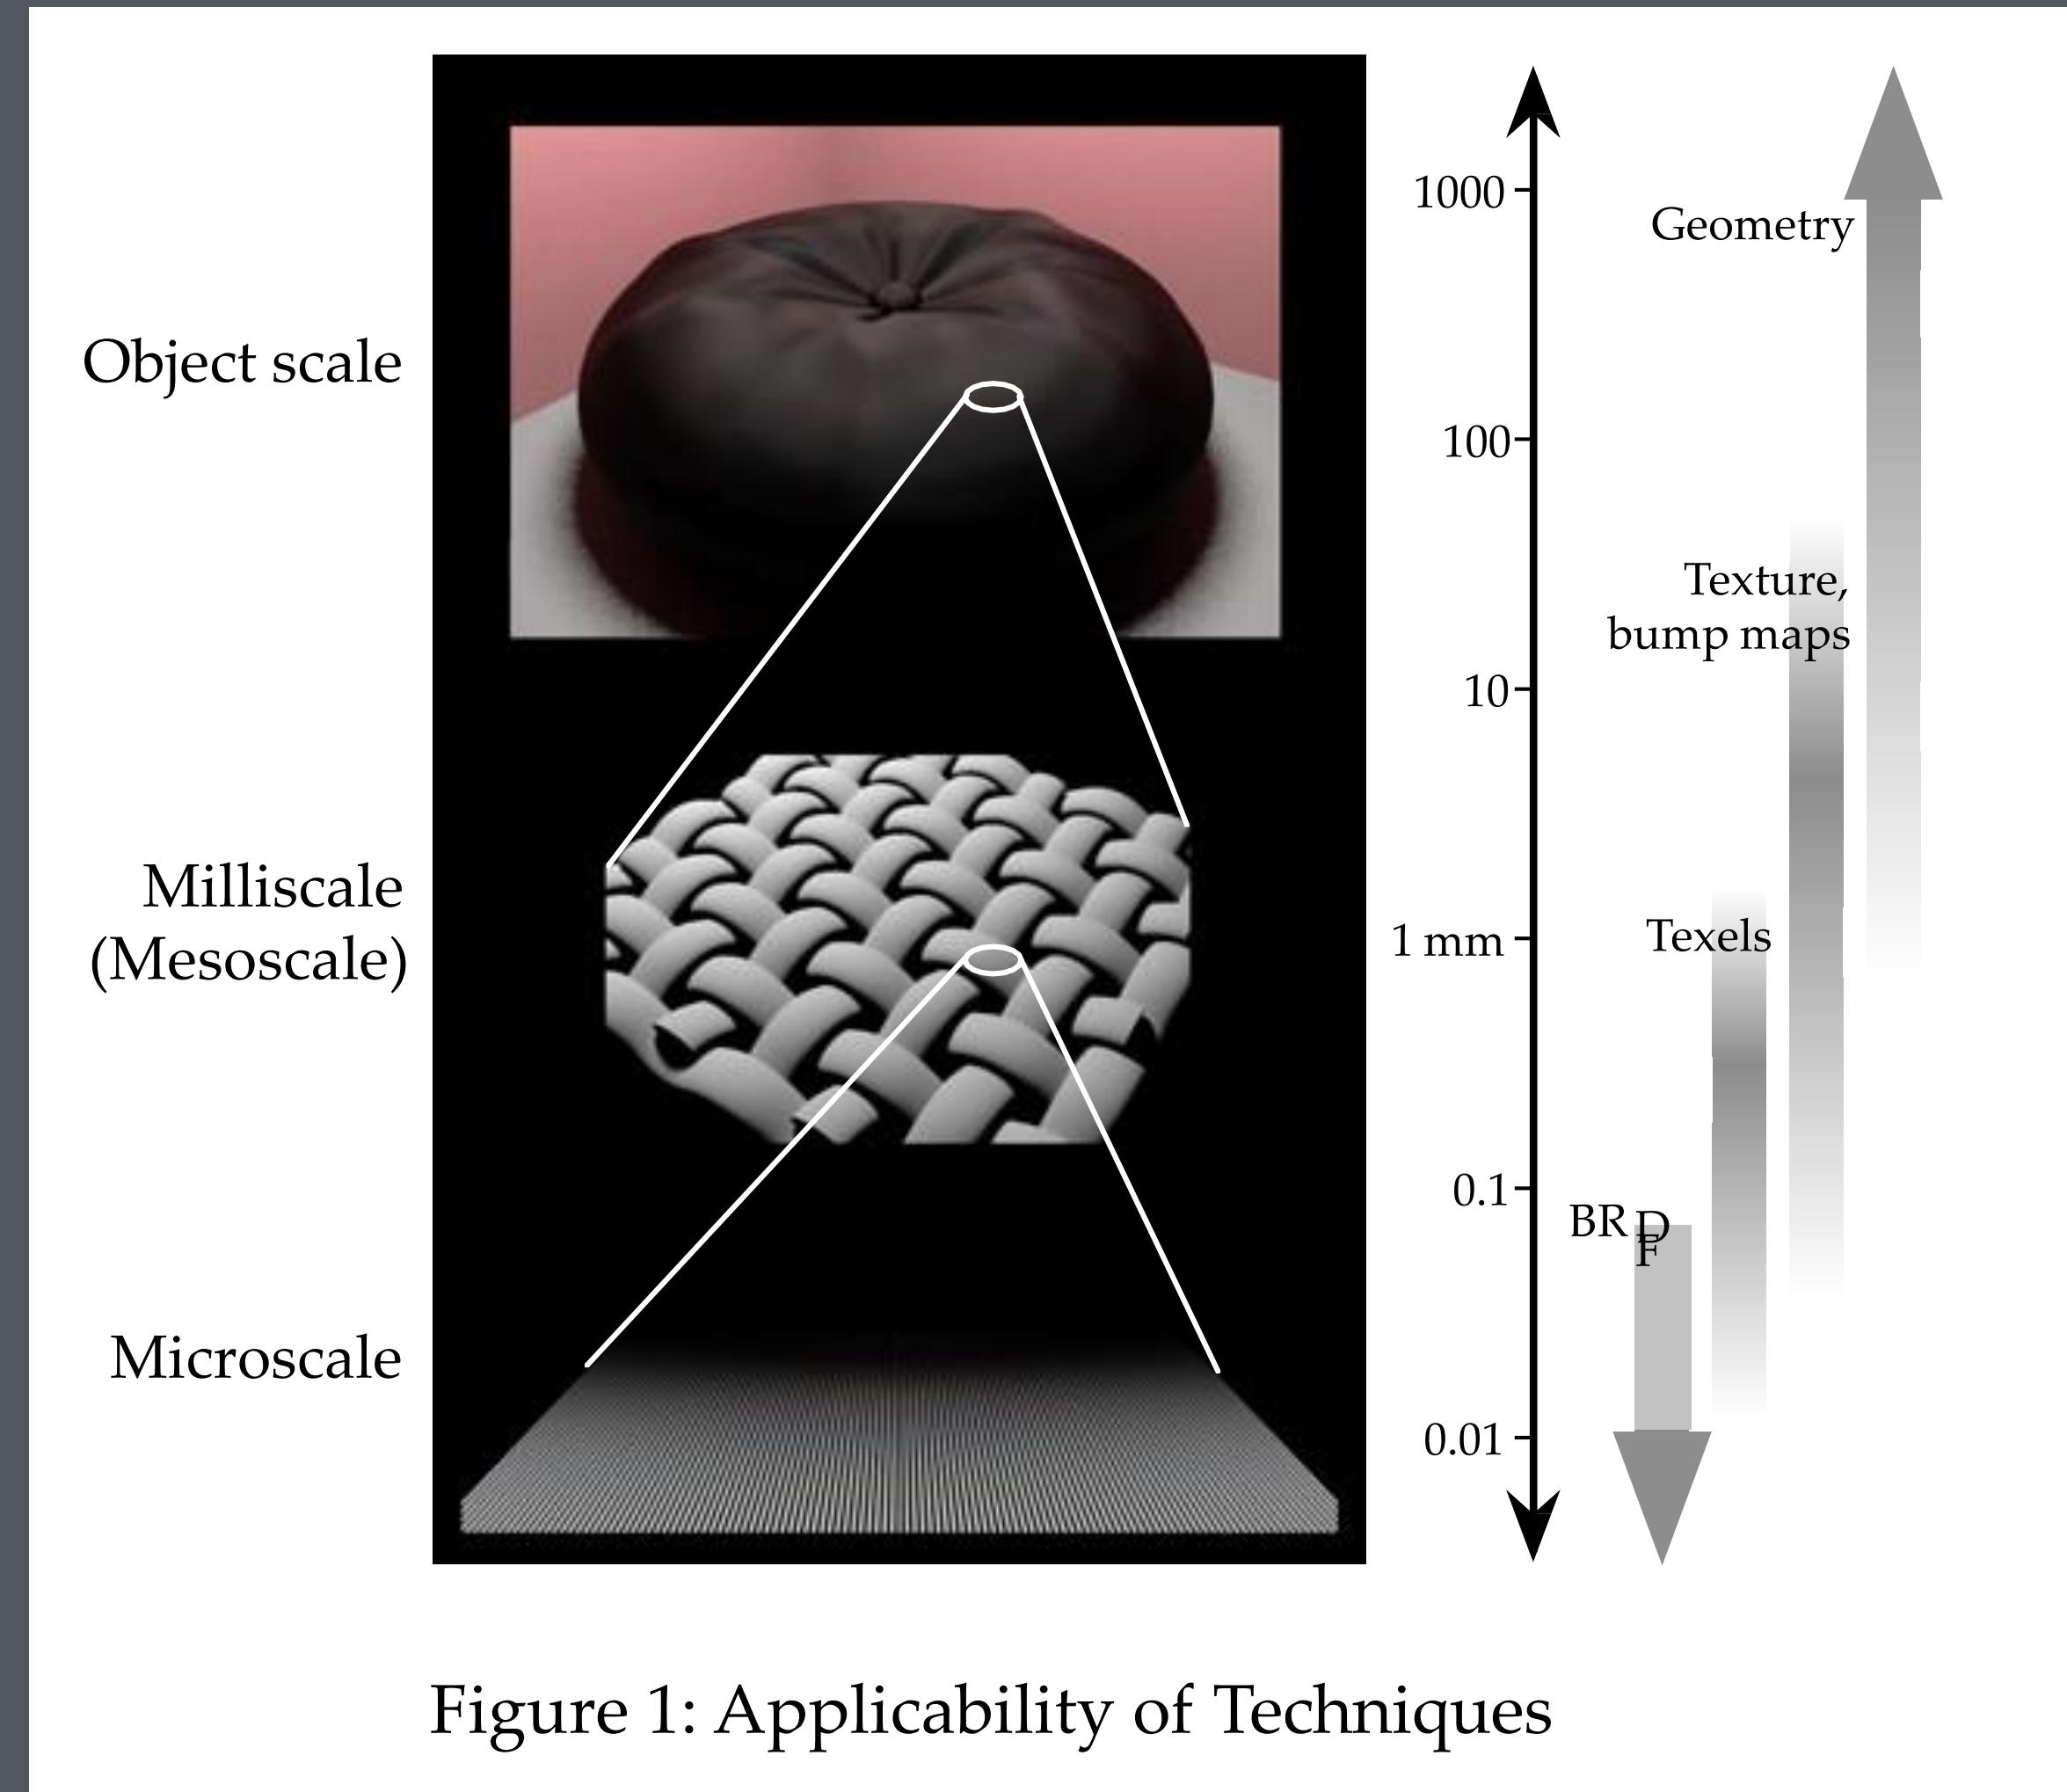
# Detail map

# Hierarchy of scales

macroscopic

mesoscopic

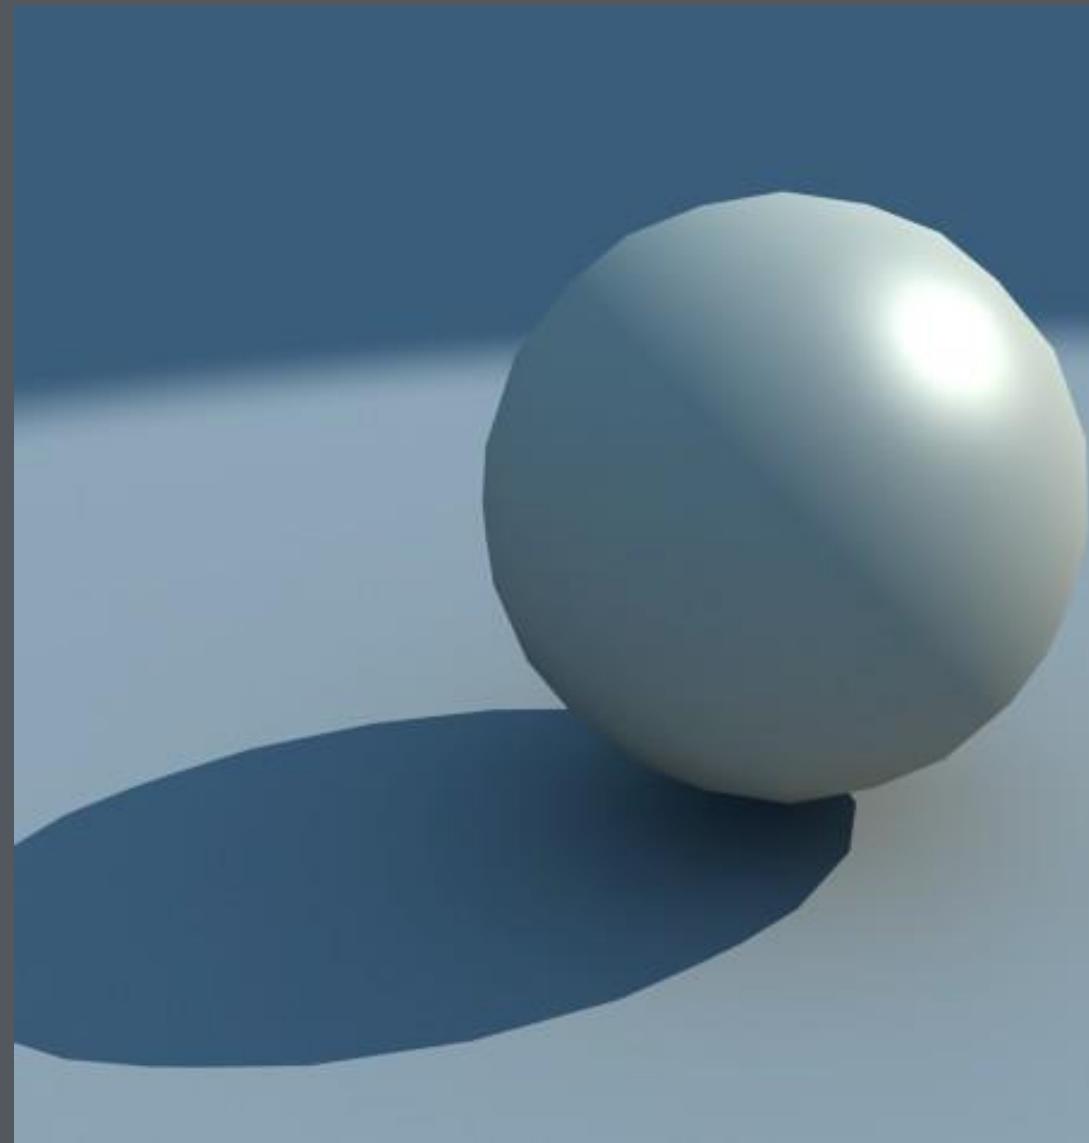
microscopic



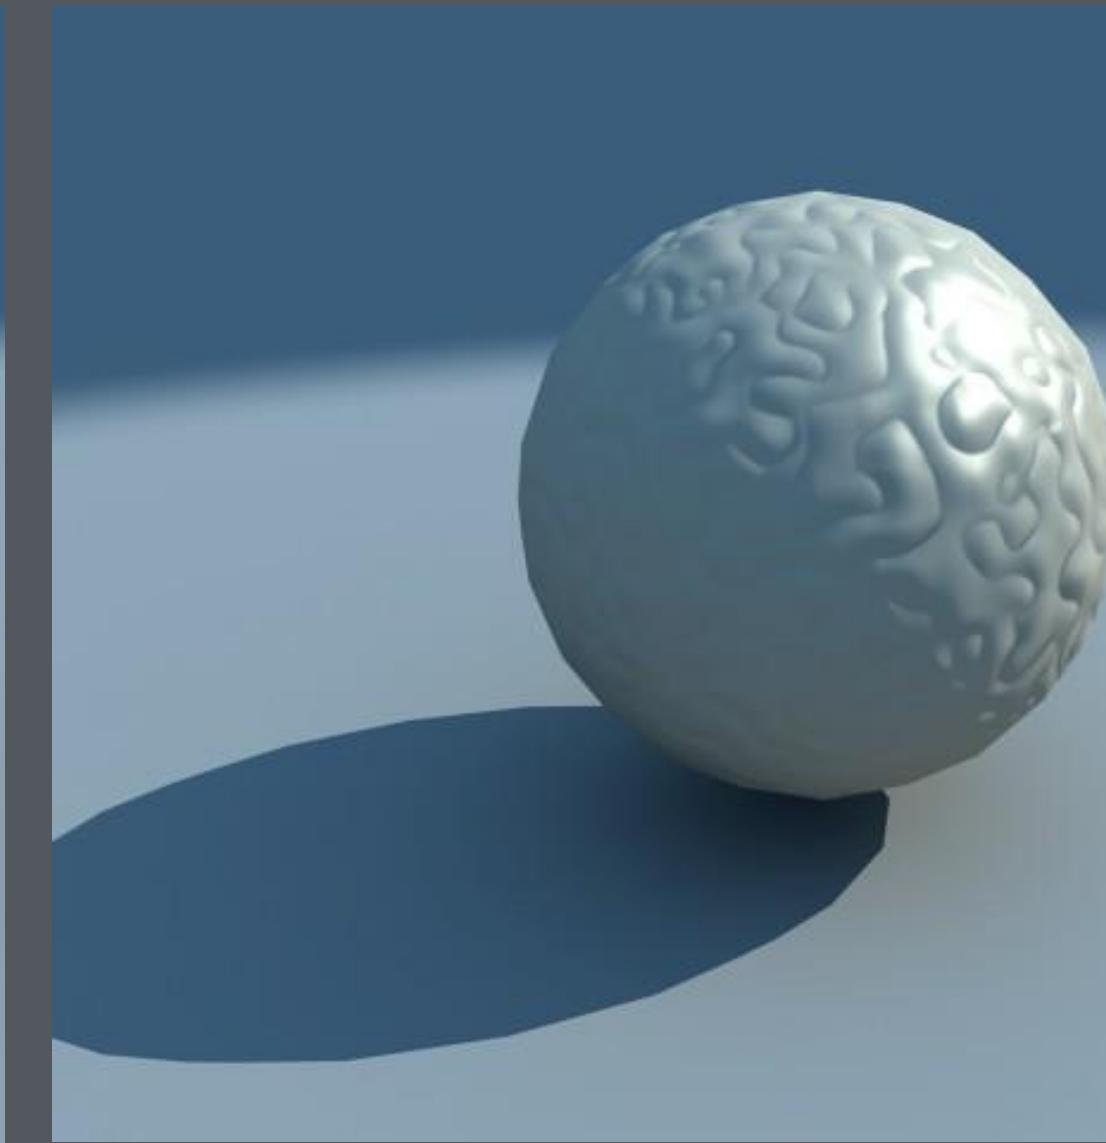
# Displacement and Bump/Normal Mapping

**Mimic effect of geometric detail/meso geometry**

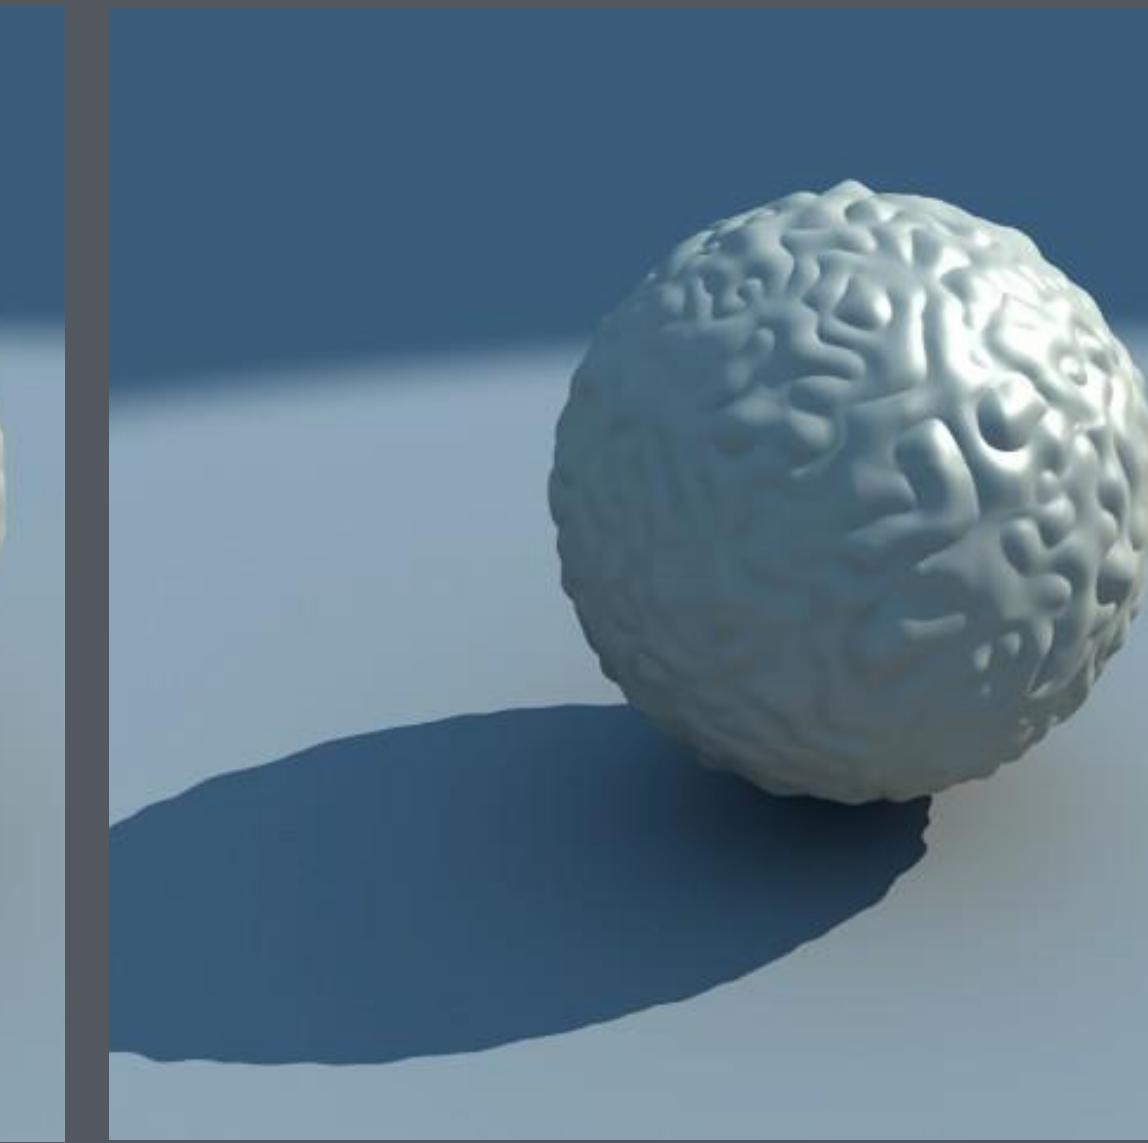
- Also detail mapping



Geometry



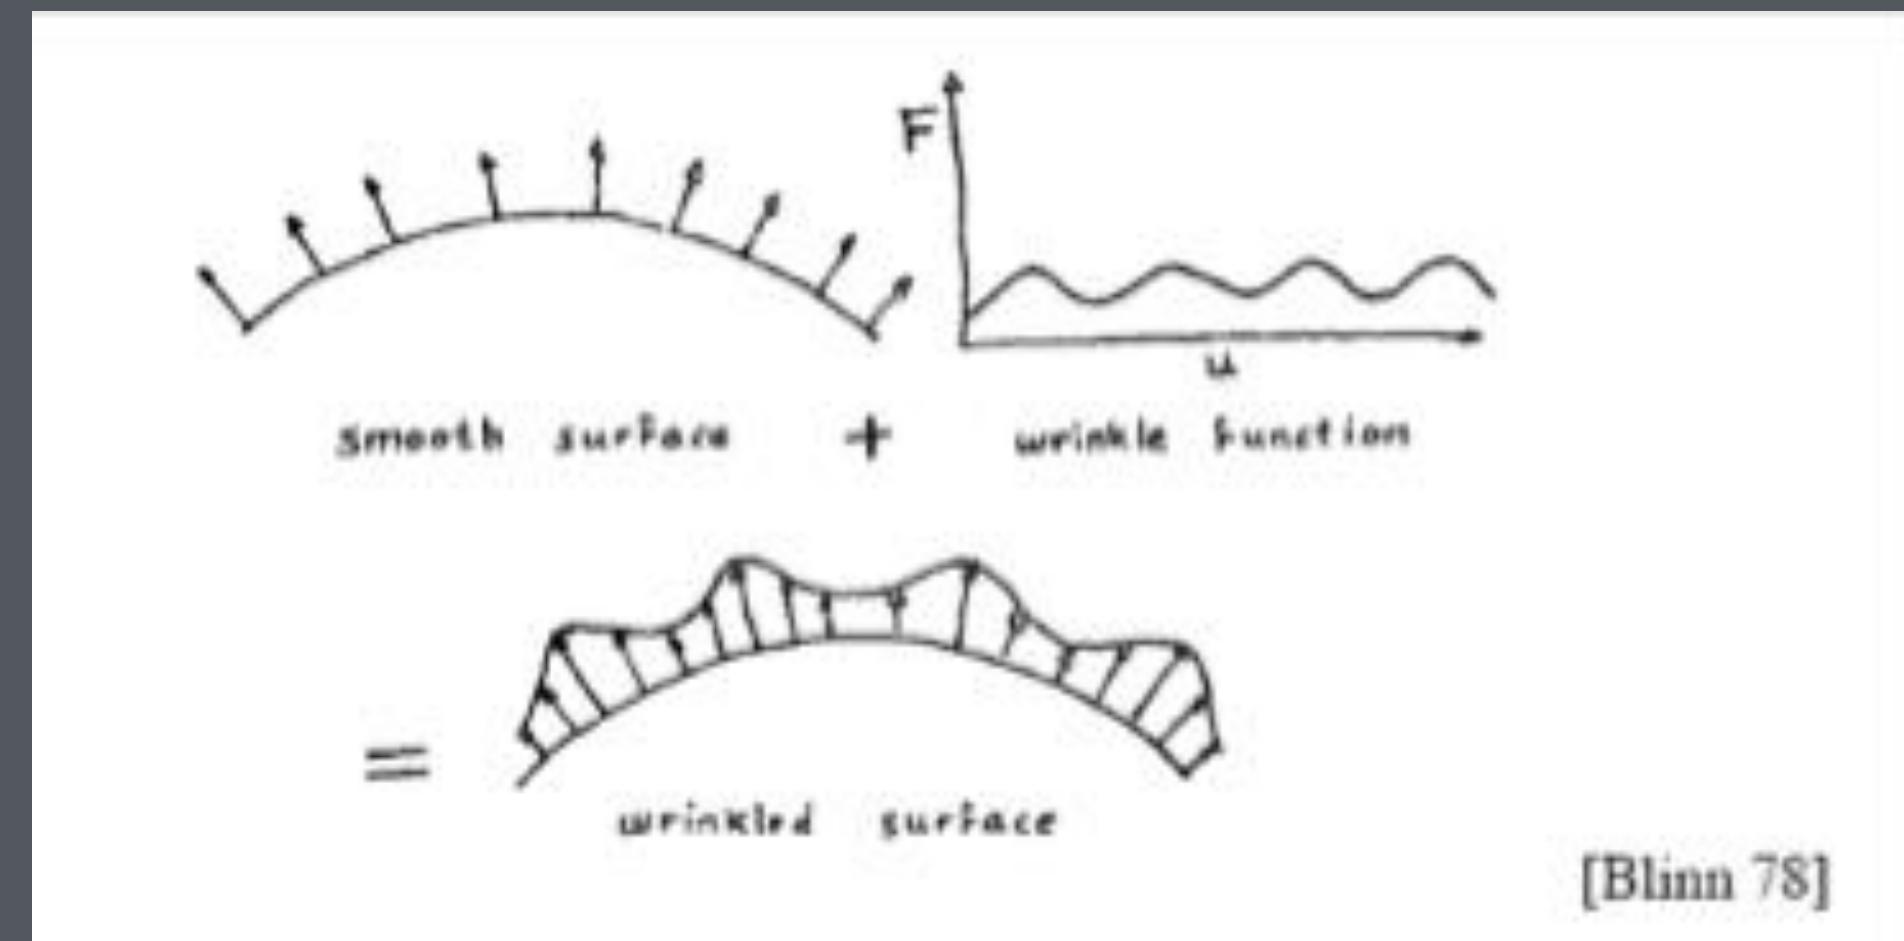
Bump  
mapping



Displacement  
mapping

# Displacement Mapping

- $P_{\text{new}} = P_{\text{old}} + DM(u) * N$



$$\mathbf{p}^0(u, v) = \mathbf{p}(u, v) + h(u, v)\mathbf{n}(u, v)$$

# Displacement in vertex shader



**Without Vertex Textures**



**With Vertex Textures**

Images used with permission from *Pacific Fighters*. © 2004 Developed by 1C:Maddox Games.  
All rights reserved. © 2004 Ubi Soft Entertainment.

# Displacement Maps

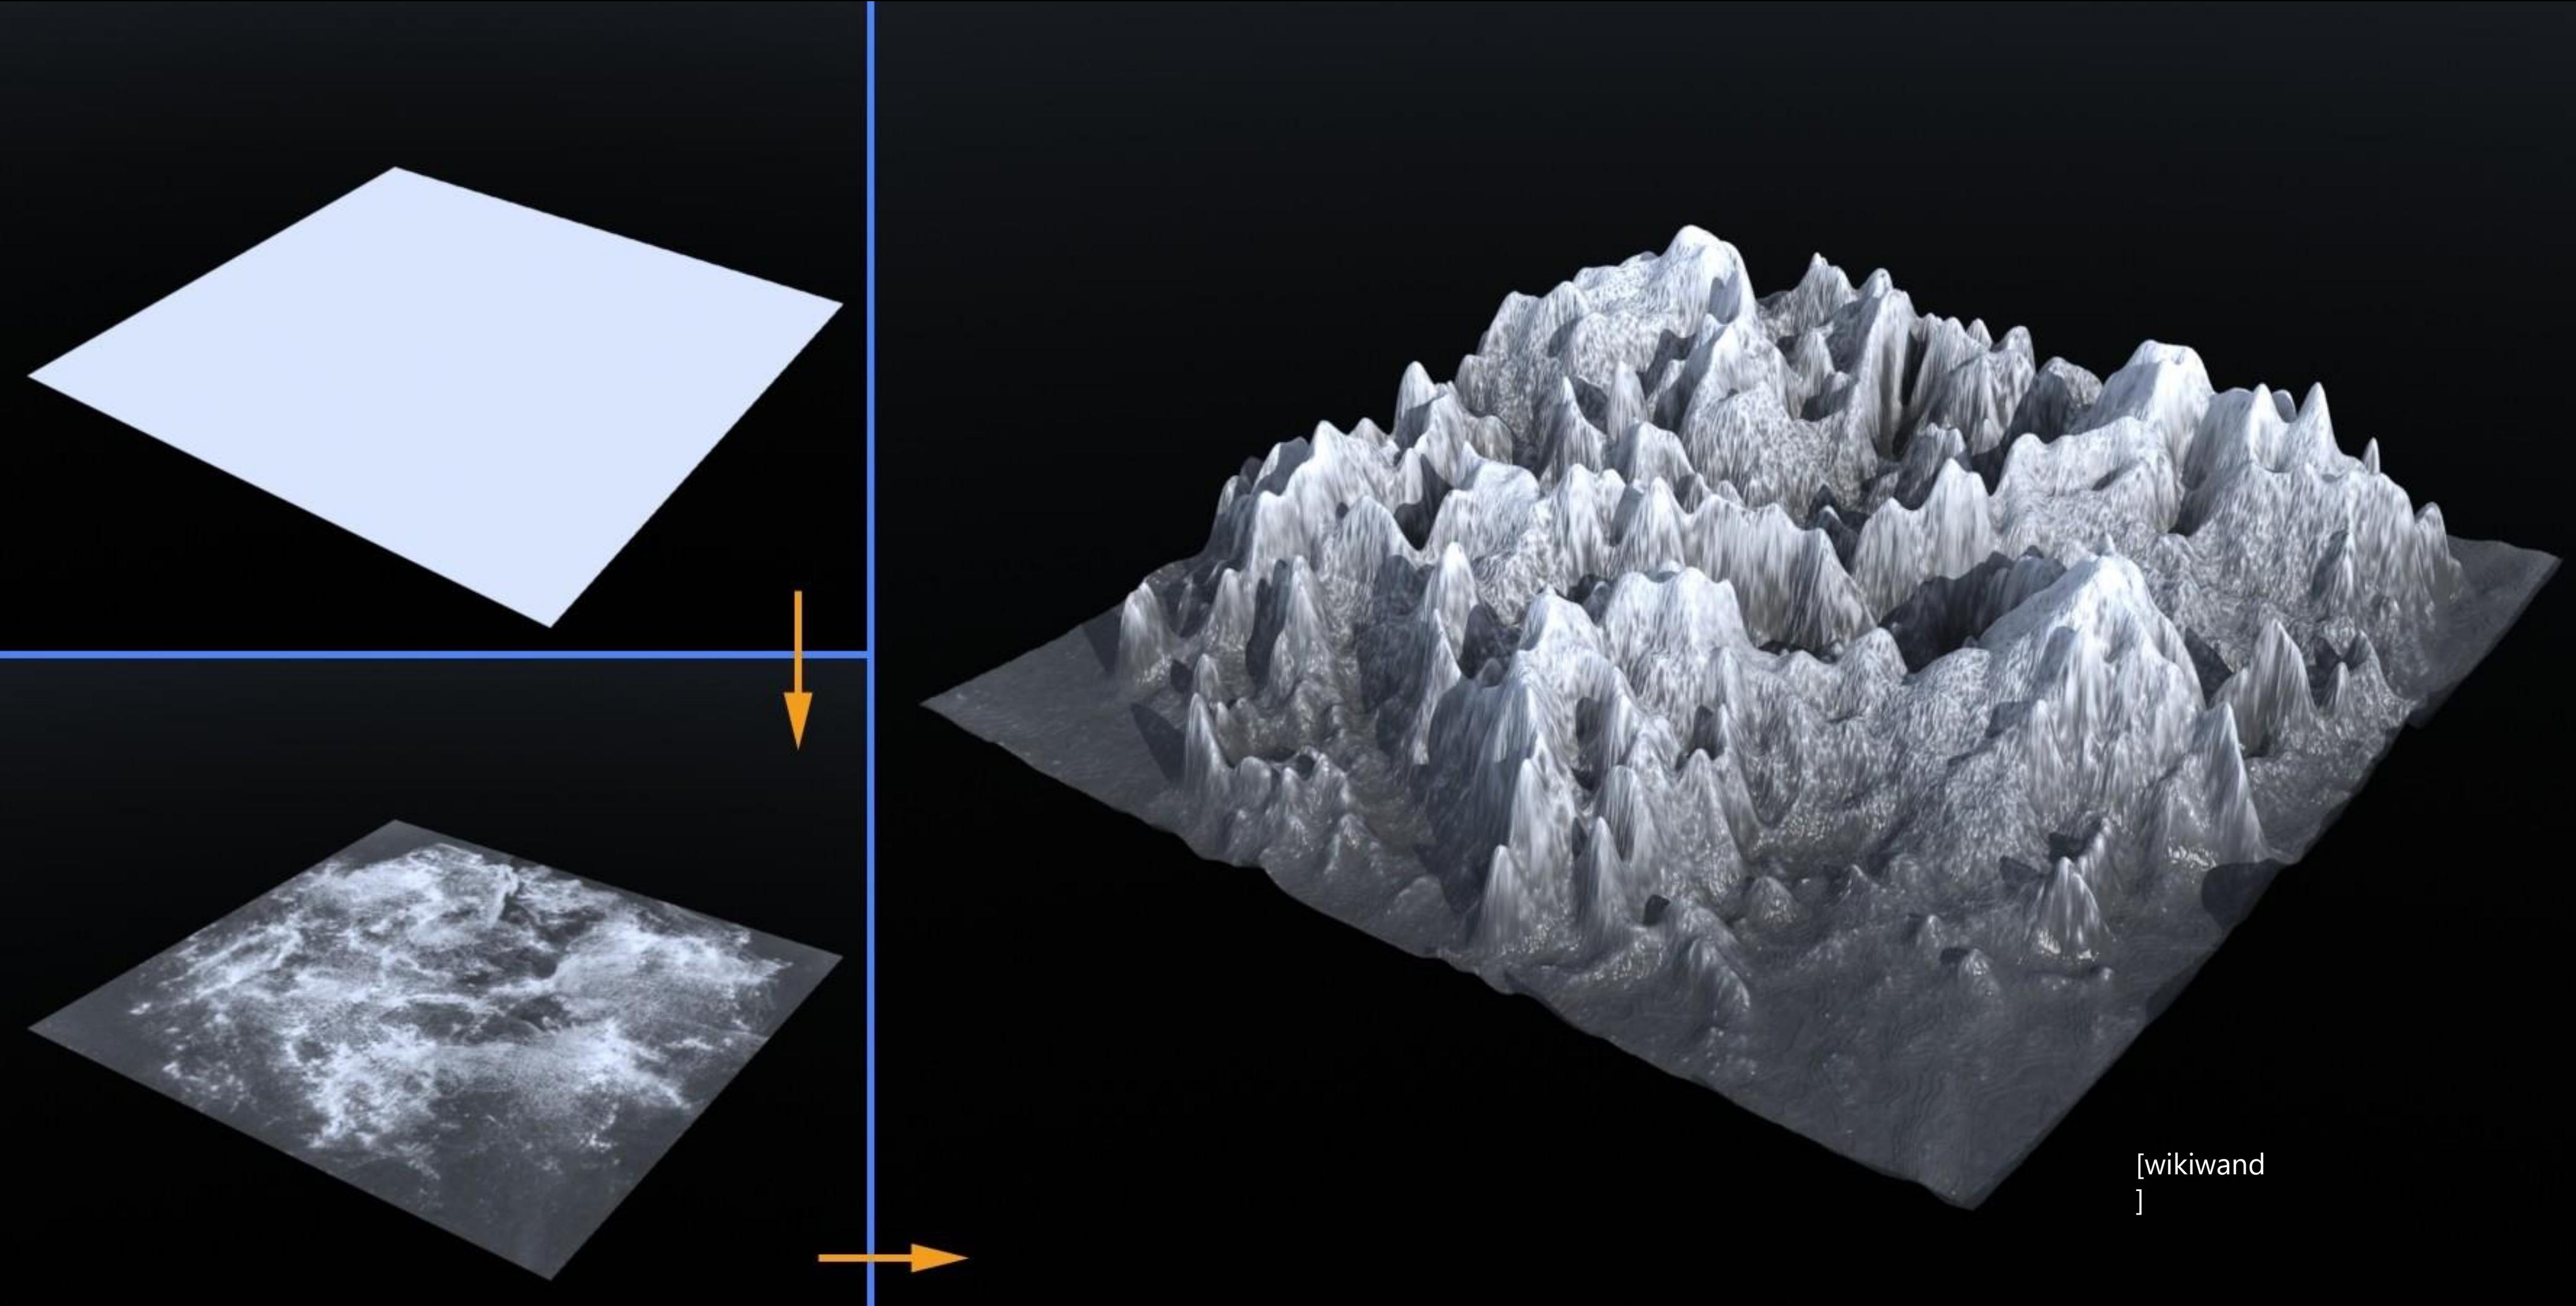
## Pros

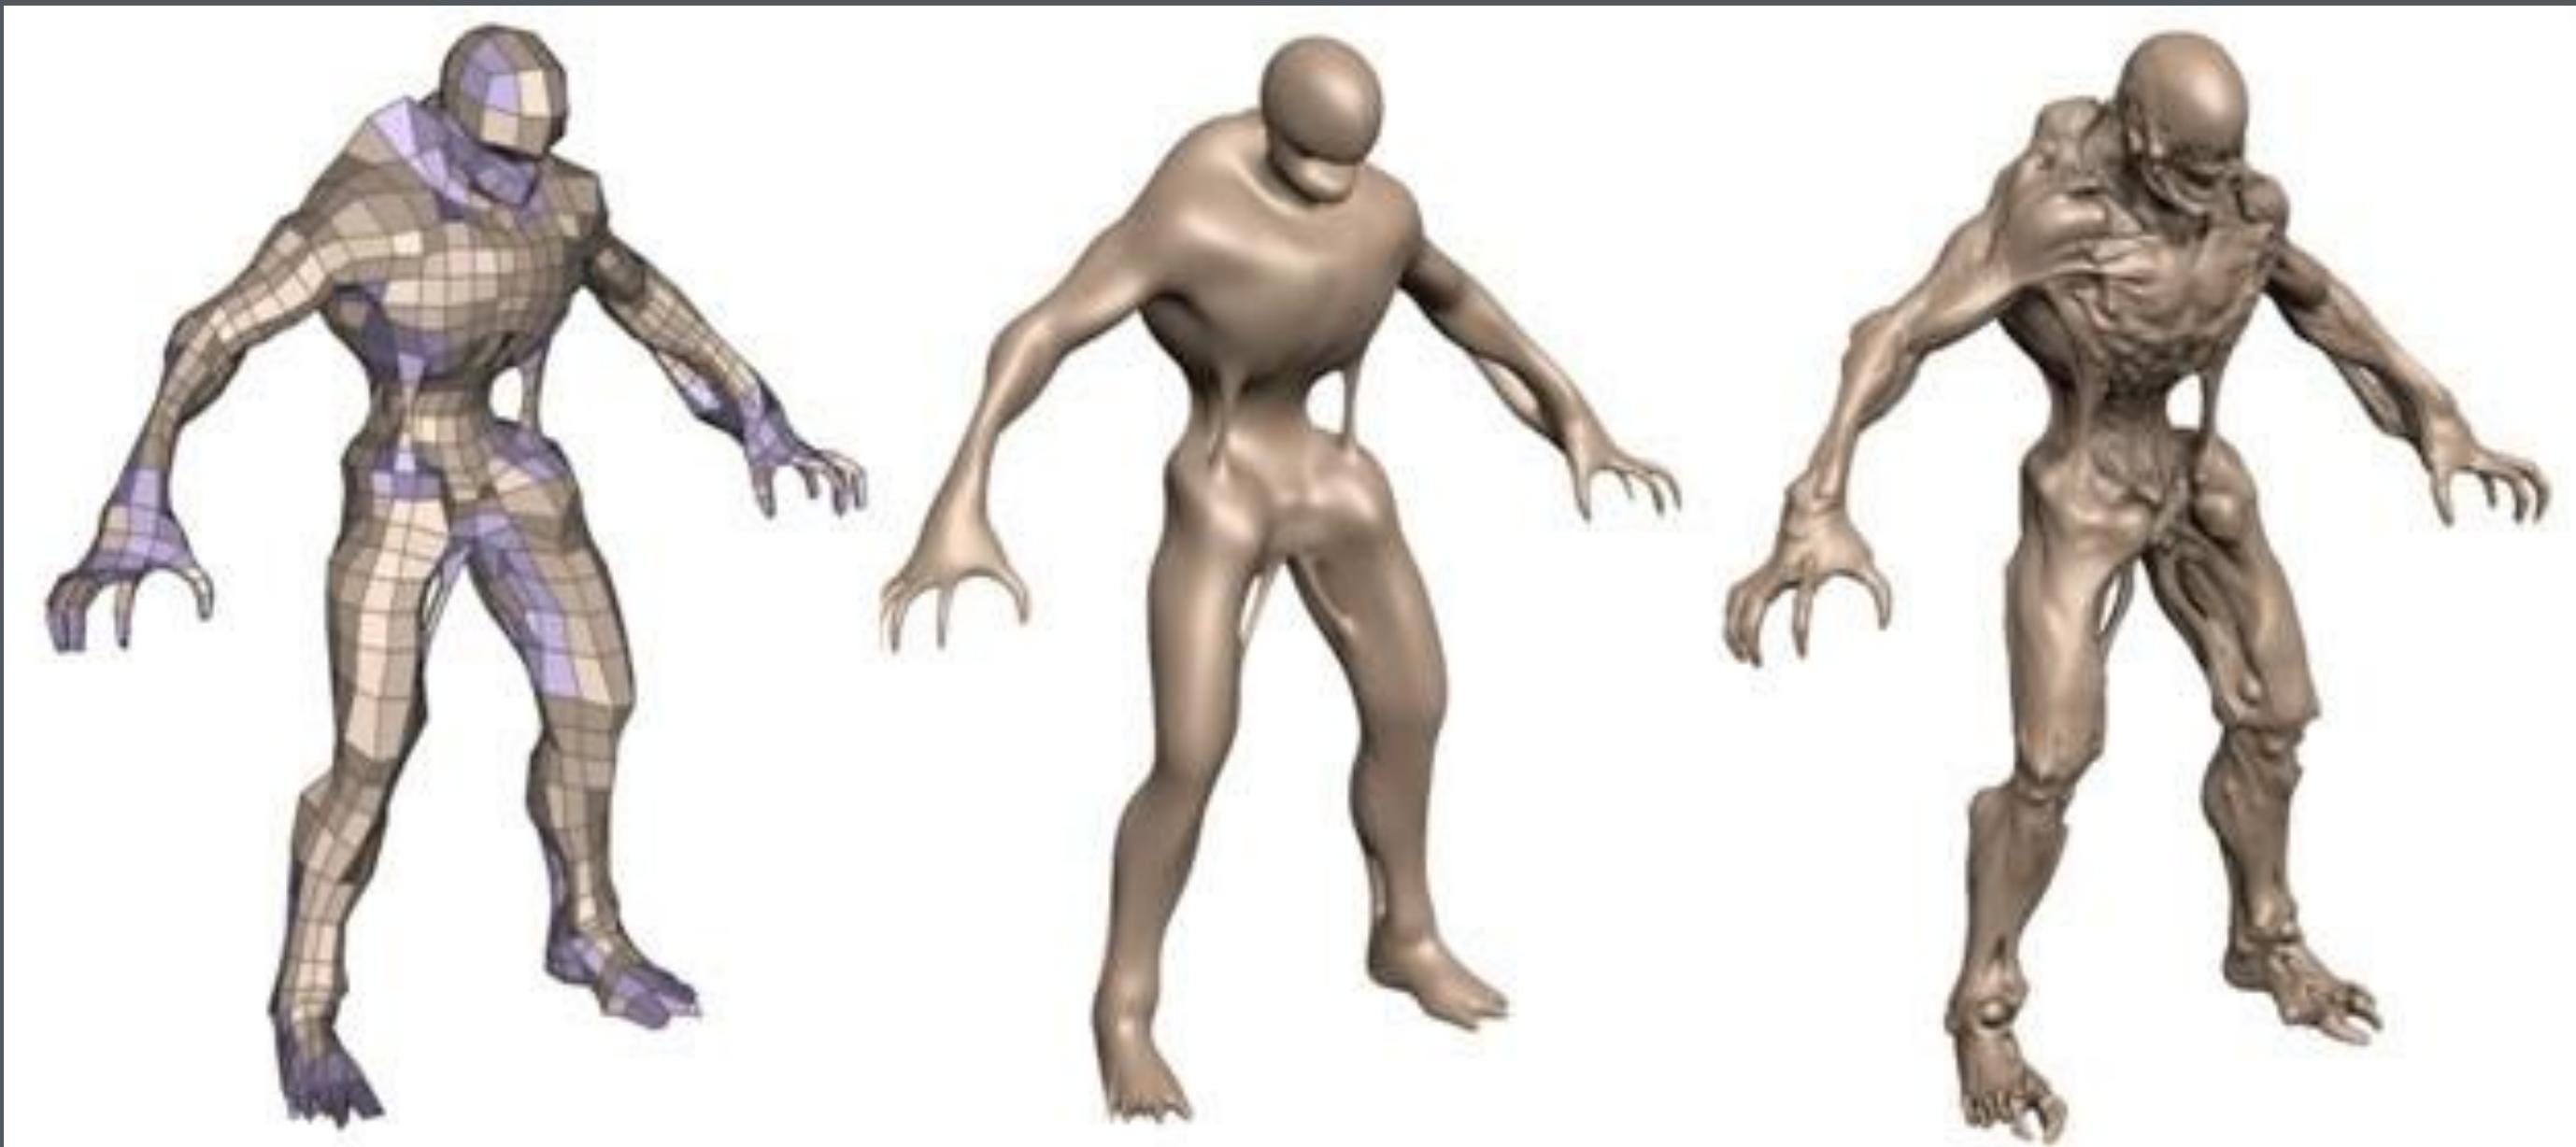
- Gives you very complex surfaces

## Cons

- Gives you very complex surfaces
- Or boring with small numbers of vertices

## Relationship with tessellation shaders





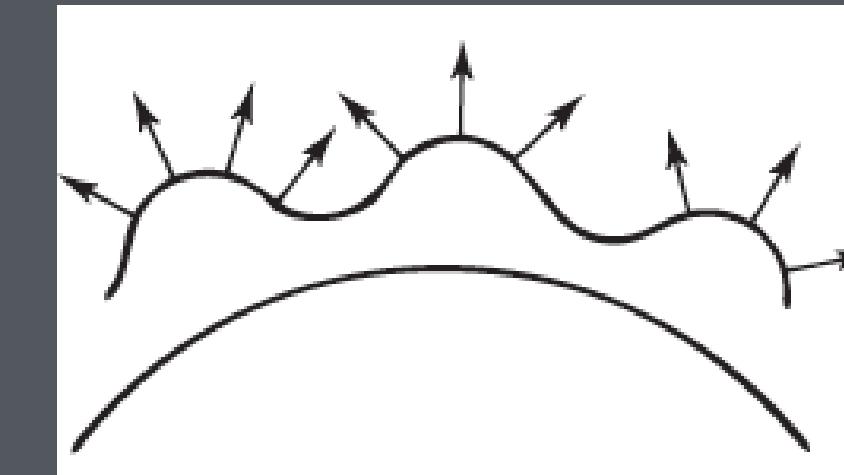
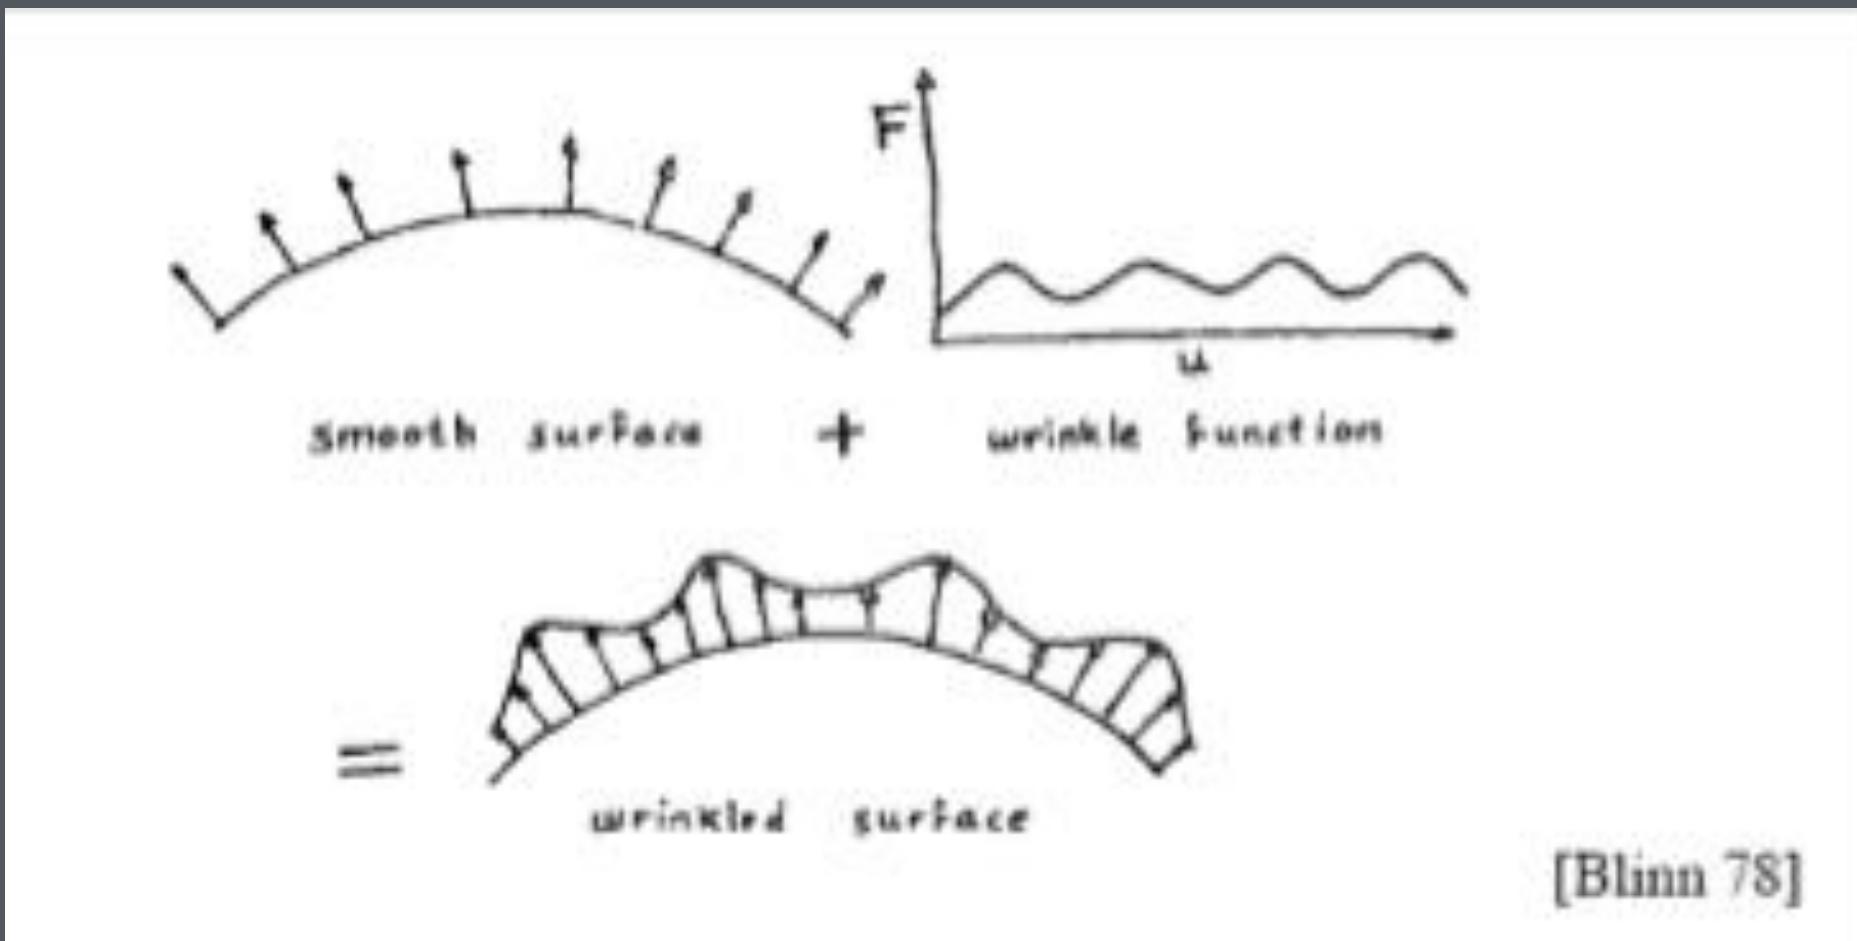
Original

Tesselated

Displac  
ement

# Bump mapping

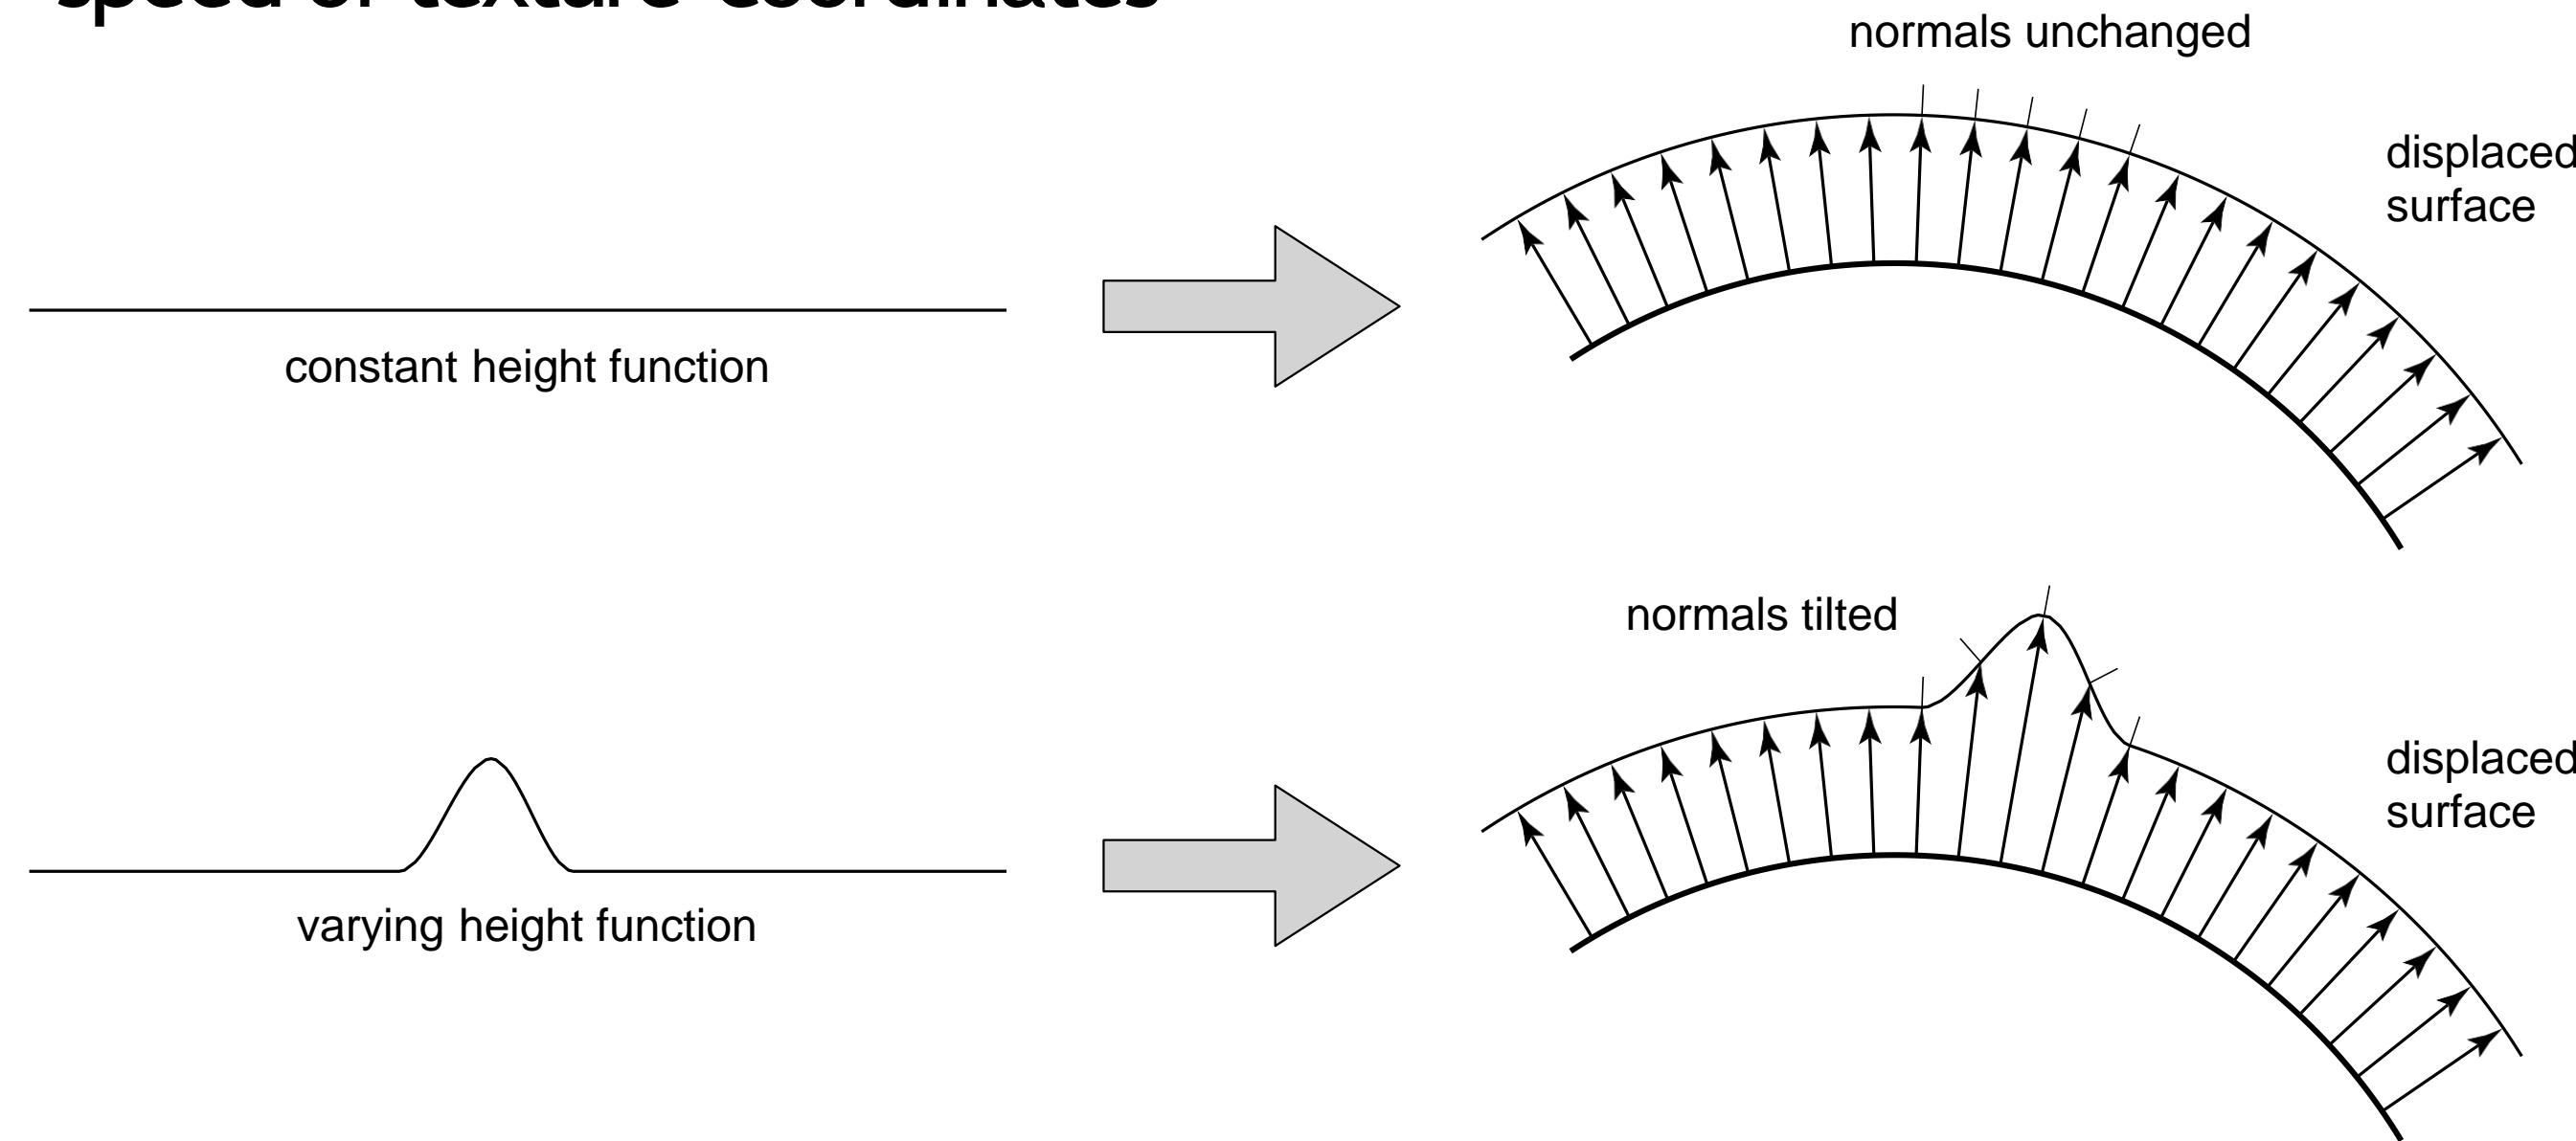
“Simulation of Wrinkled Surfaces” Blinn 78



Blinn: keep surface, use new normals

# Normals in displacement mapping

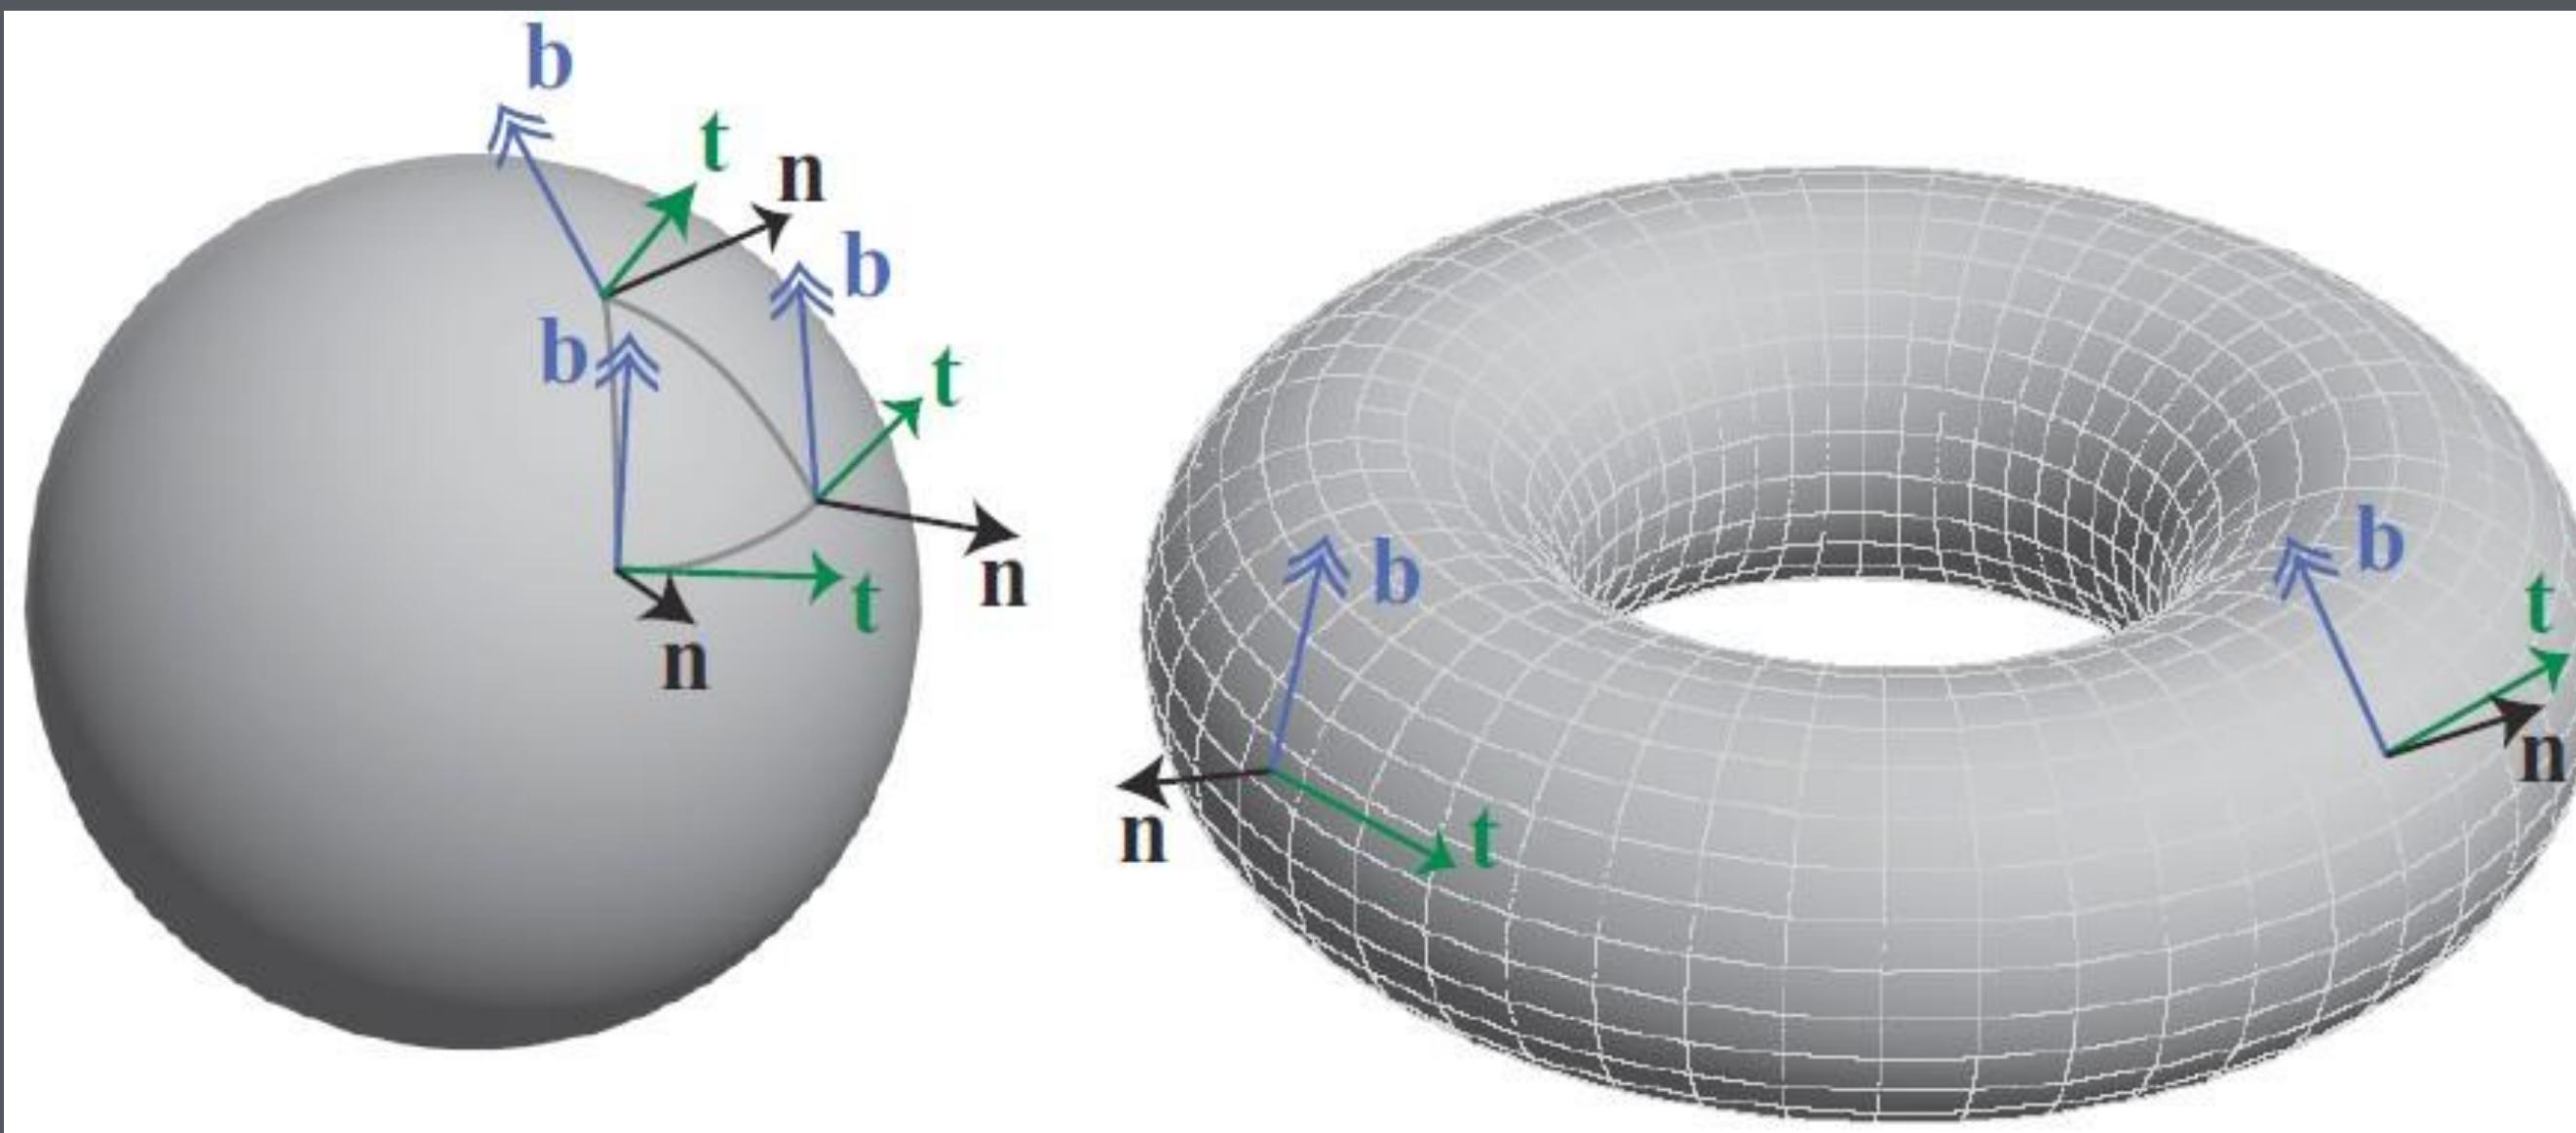
- Displacement changes the surface normal, depending on:
  - derivative of height function
  - orientation of texture coordinates
  - speed of texture coordinates



# How to change the normal?

**First, need some frame of reference**

- Normal is modified with respect to that
- Have tangent space basis:  $t$  and  $b$
- Normal, tangent and bitangent vectors



# Geometry for displacement

- **geometric inputs**
  - $utangent$  (unnormalized) as vertex attribute
  - $vtangent$  (unnormalized) as vertex attribute
  - height field as a texture
- **vertex stage**
  - compute displaced vertex position
  - look up displacement value from texture
  - compute normal to displaced surface
    - compute derivatives of height by finite differences
    - add offset to the base surface tangents
    - normalized cross product is the shading normal
- **fragment stage: just compute shading**

(or compute them  
ahead of time  
and store height and  
derivatives in a  
3-channel texture)

# Bump Mapping

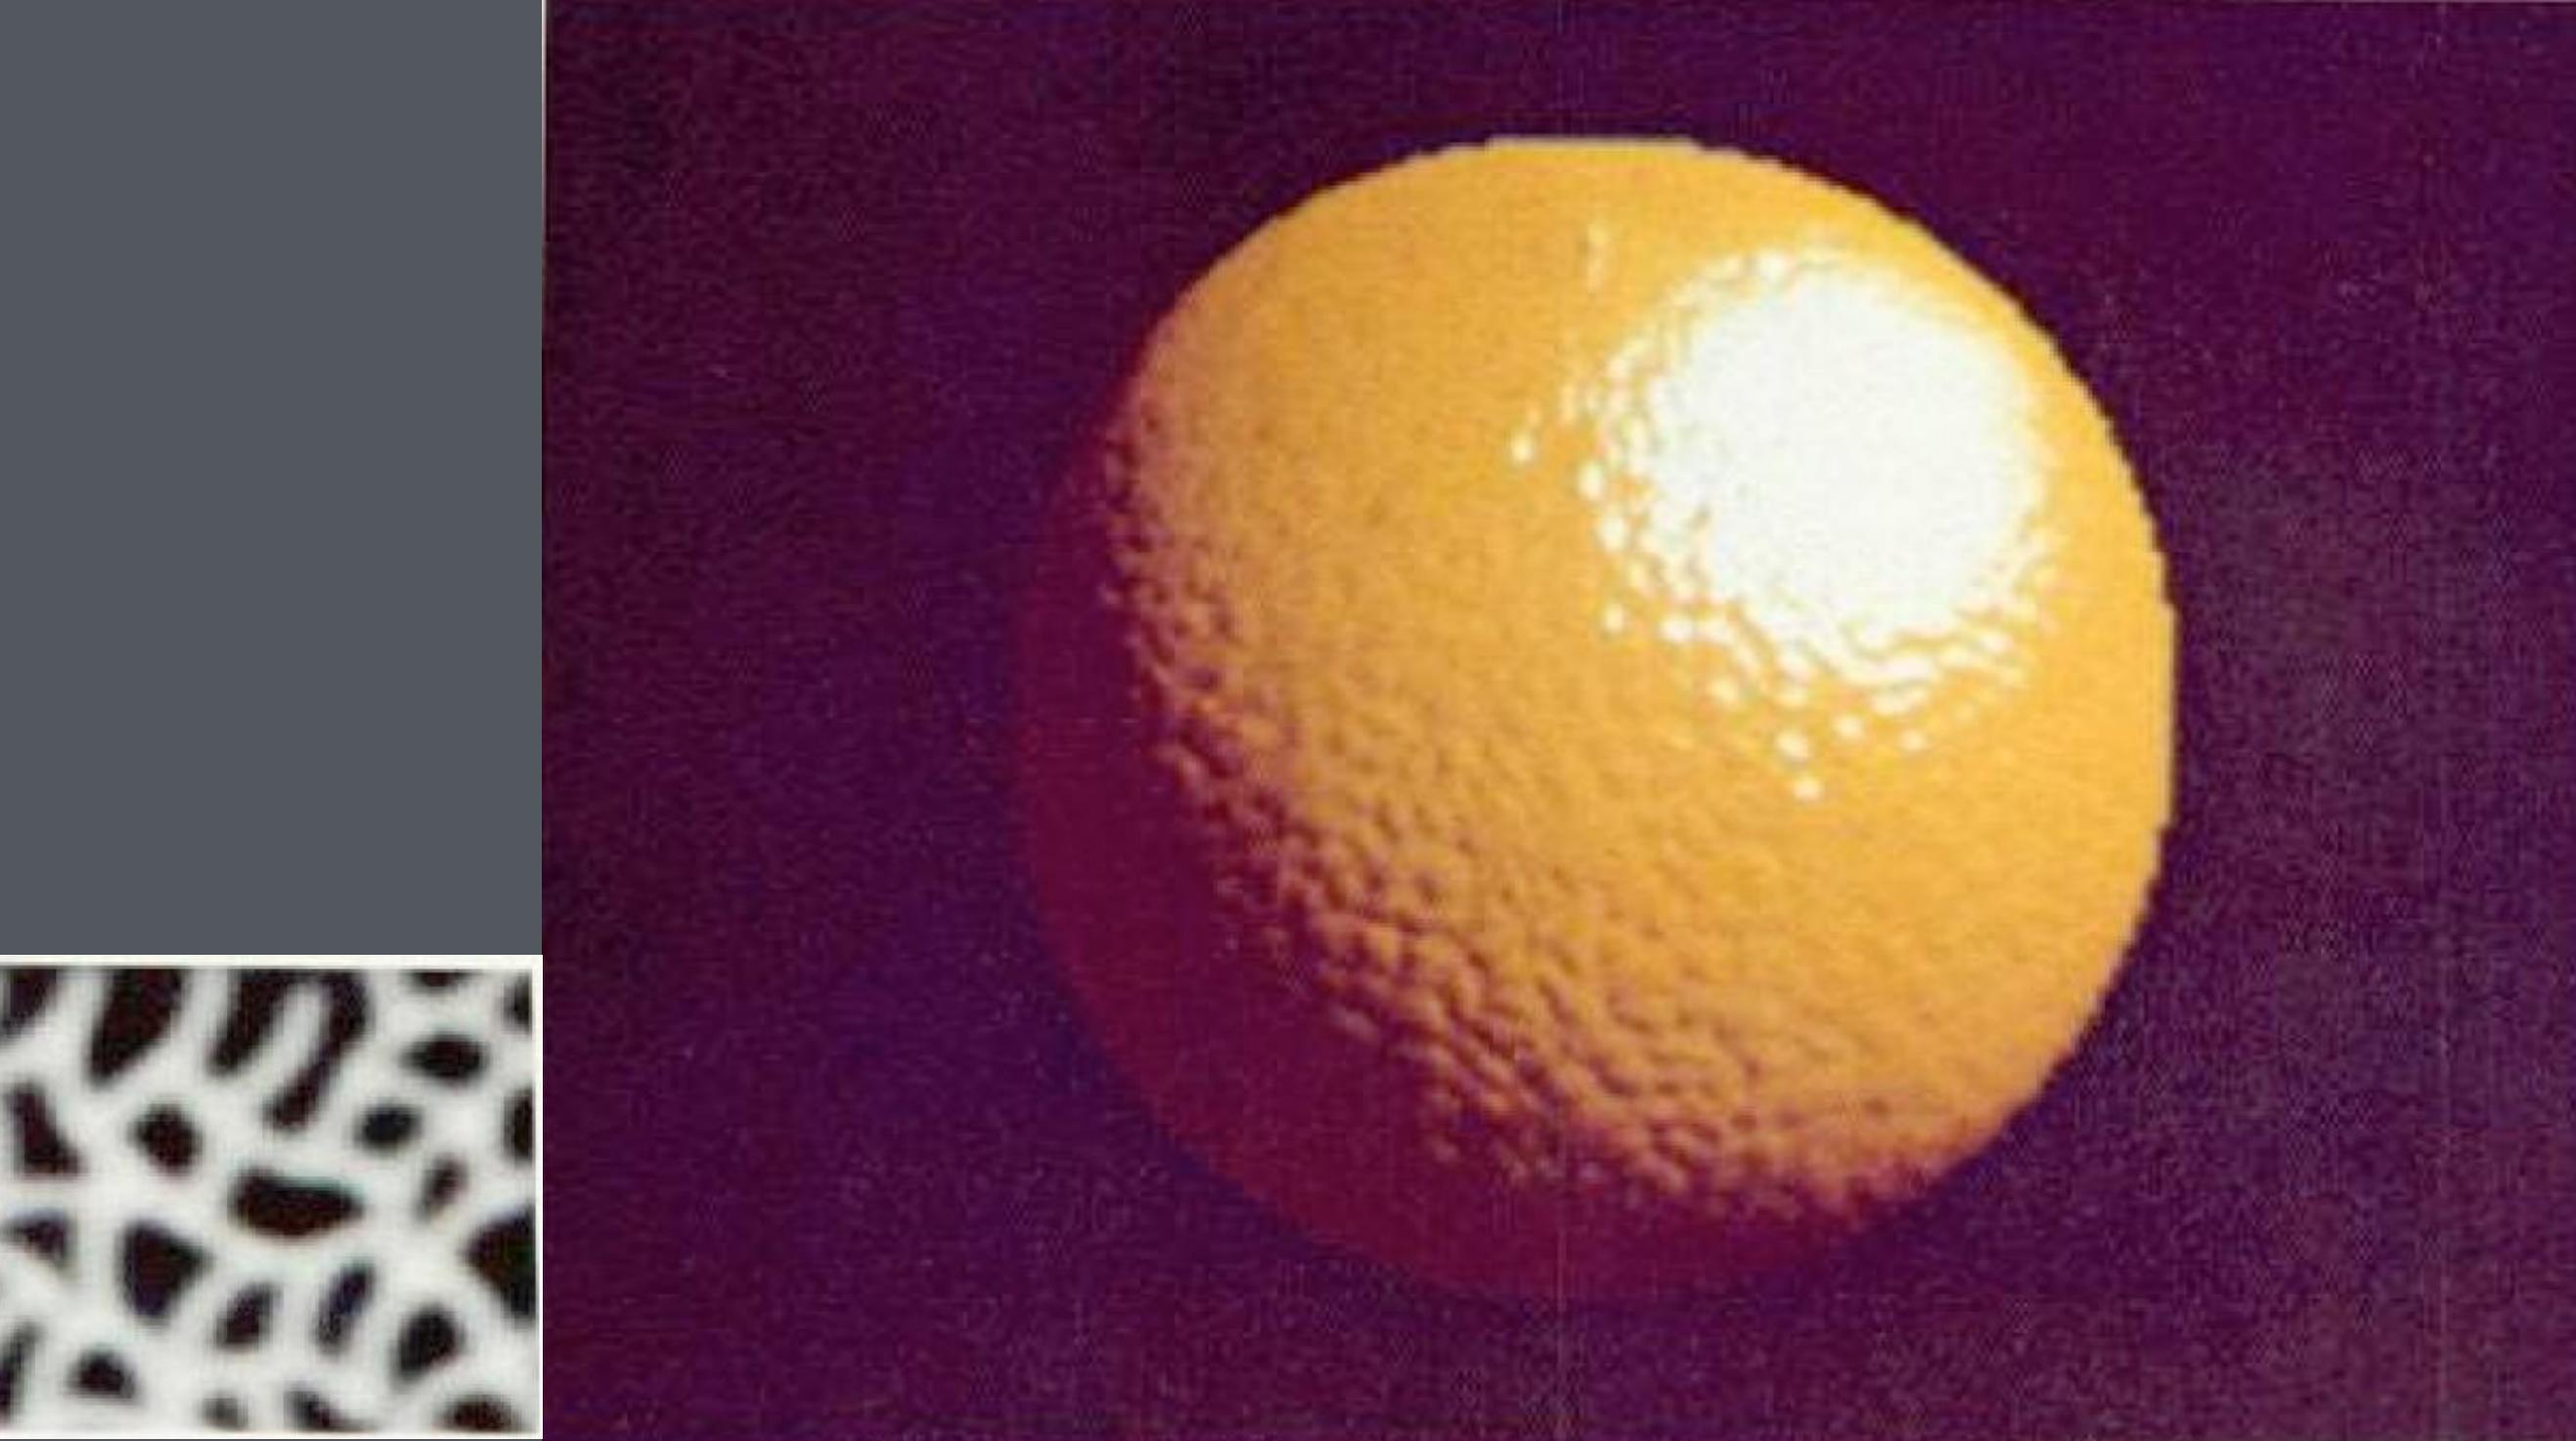
## **Alter normals of surface**

- Only affects shading normals

## **Also, mimics effect of small scale geometry**

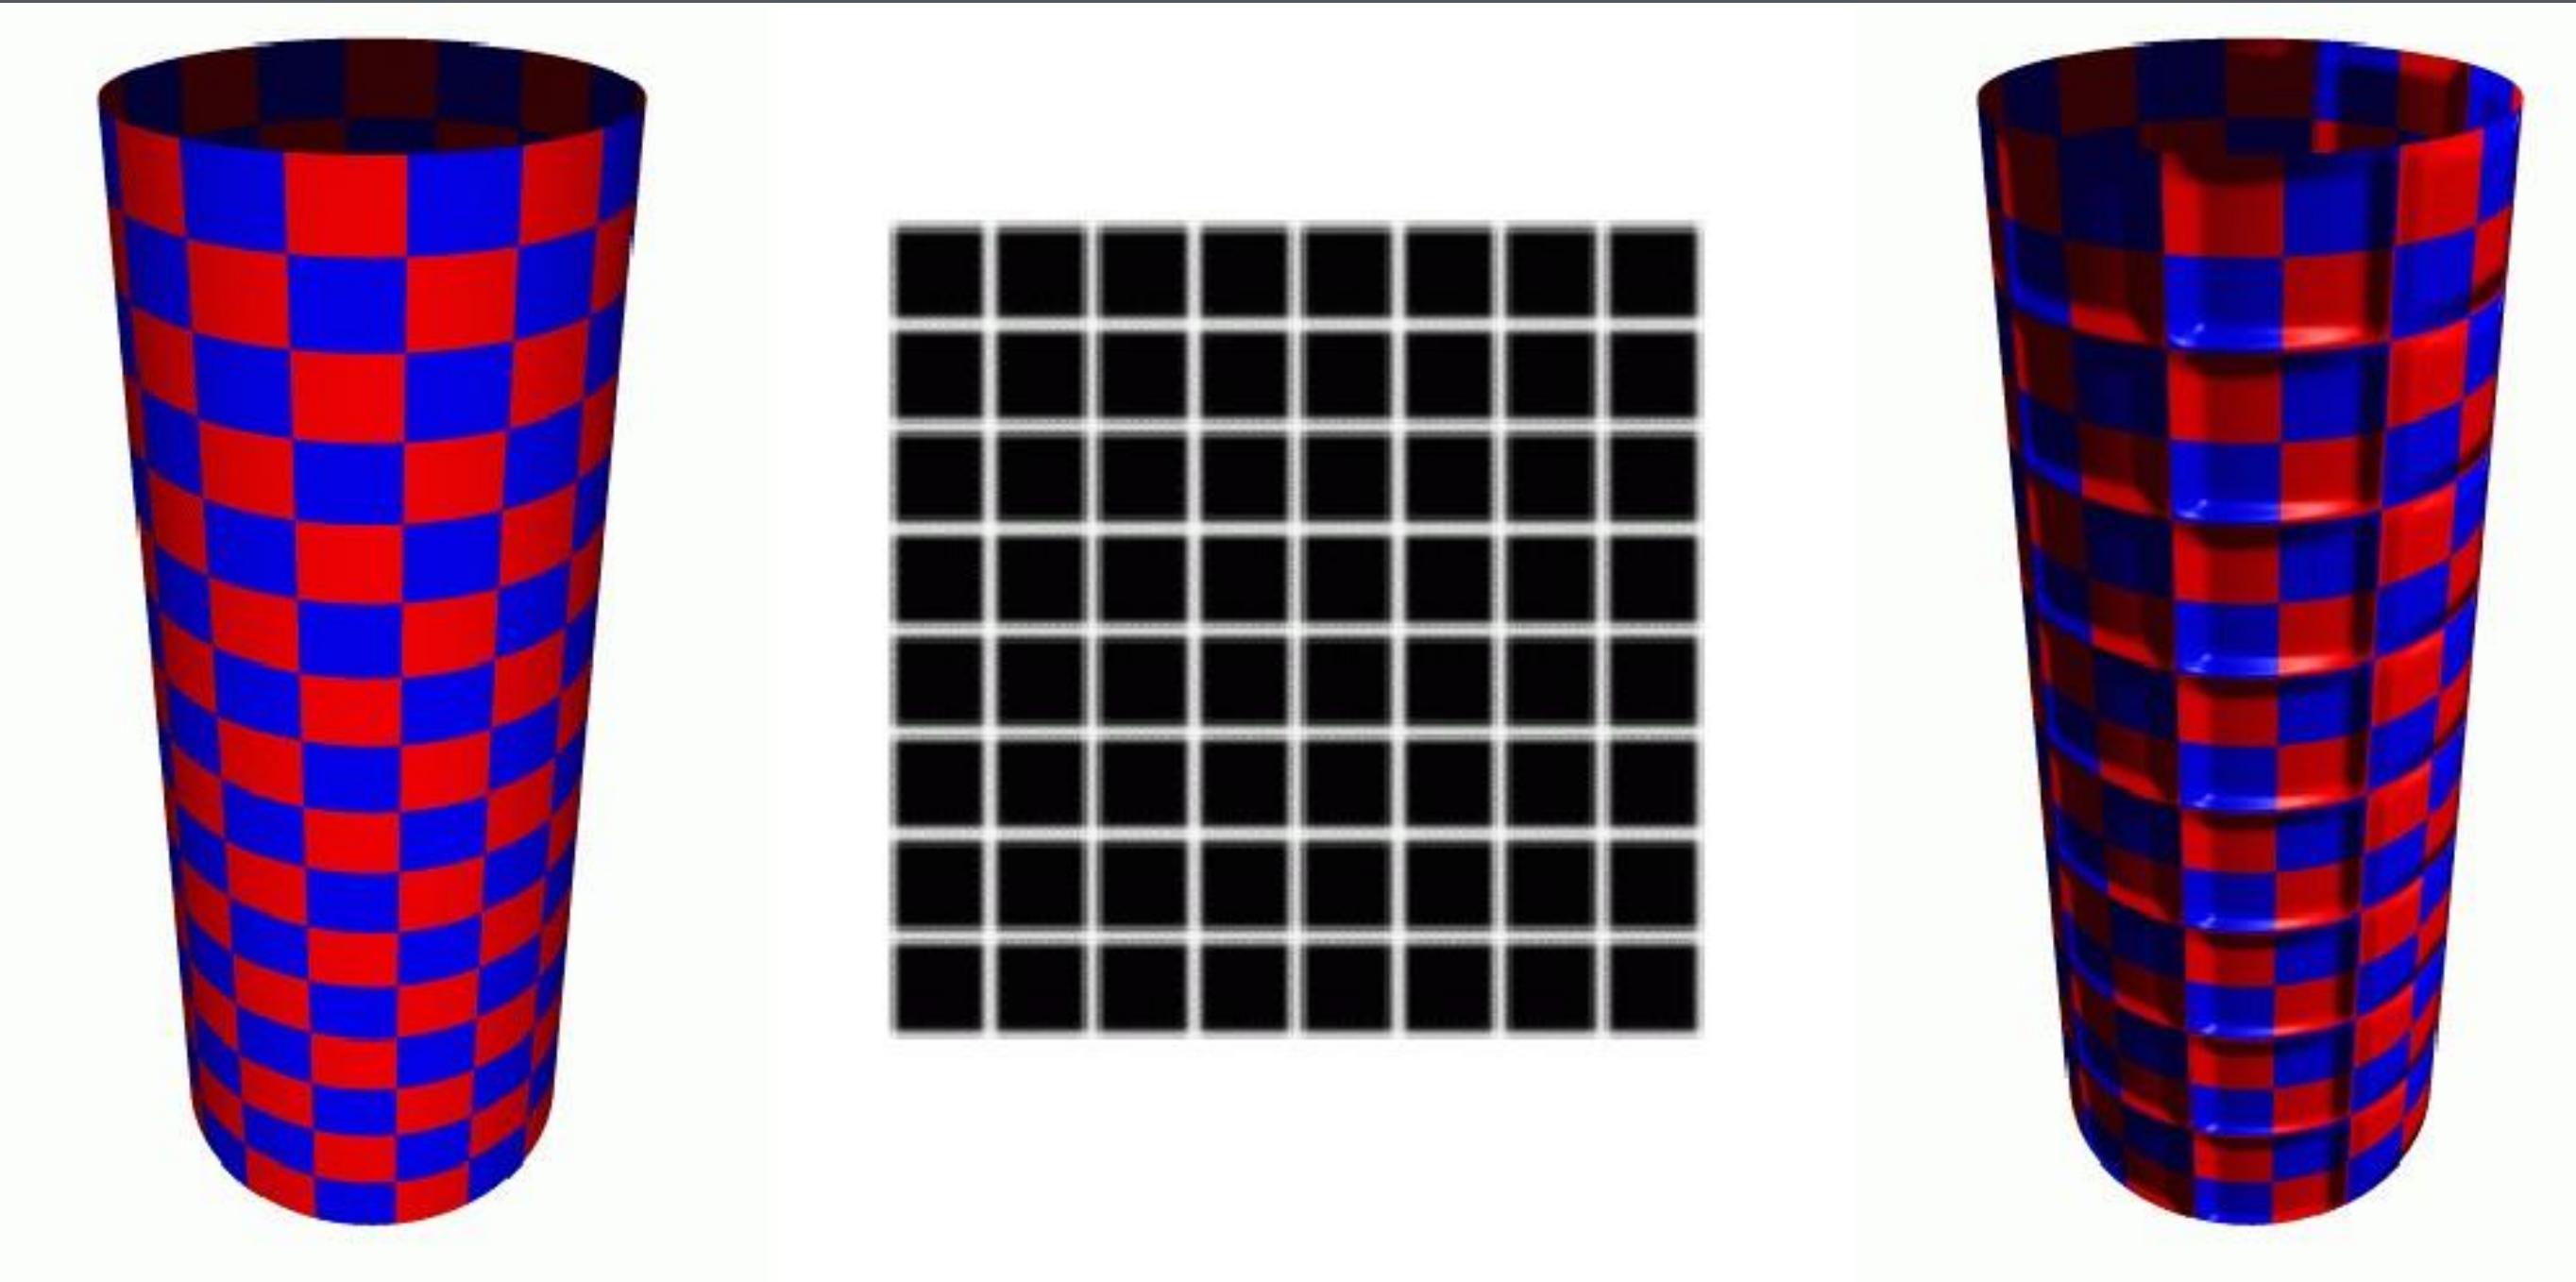
- Detail map
- Except at silhouette
- Adds perceived bumps, wrinkles

# Bump mapping



[Blinn  
1978]

# Bump Mapping



# Bump mapping

## **Displacement mapping is expensive**

- requires densely tessellated geometry
- many triangles to rasterize

## **For small displacements, the most important effect is on the normal**

- so just do that part; don't displace the surface

## **Bump mapping is then a fragment operation**

- doesn't require dense tessellation
- doesn't actually displace the surface
- gives shading that looks just like displaced surface

# Bump Mapping



# Bump mapping

## Geometric inputs

- tangent vectors (unnormalized) as vertex attributes
- height field as a texture
- no dense triangulation needed

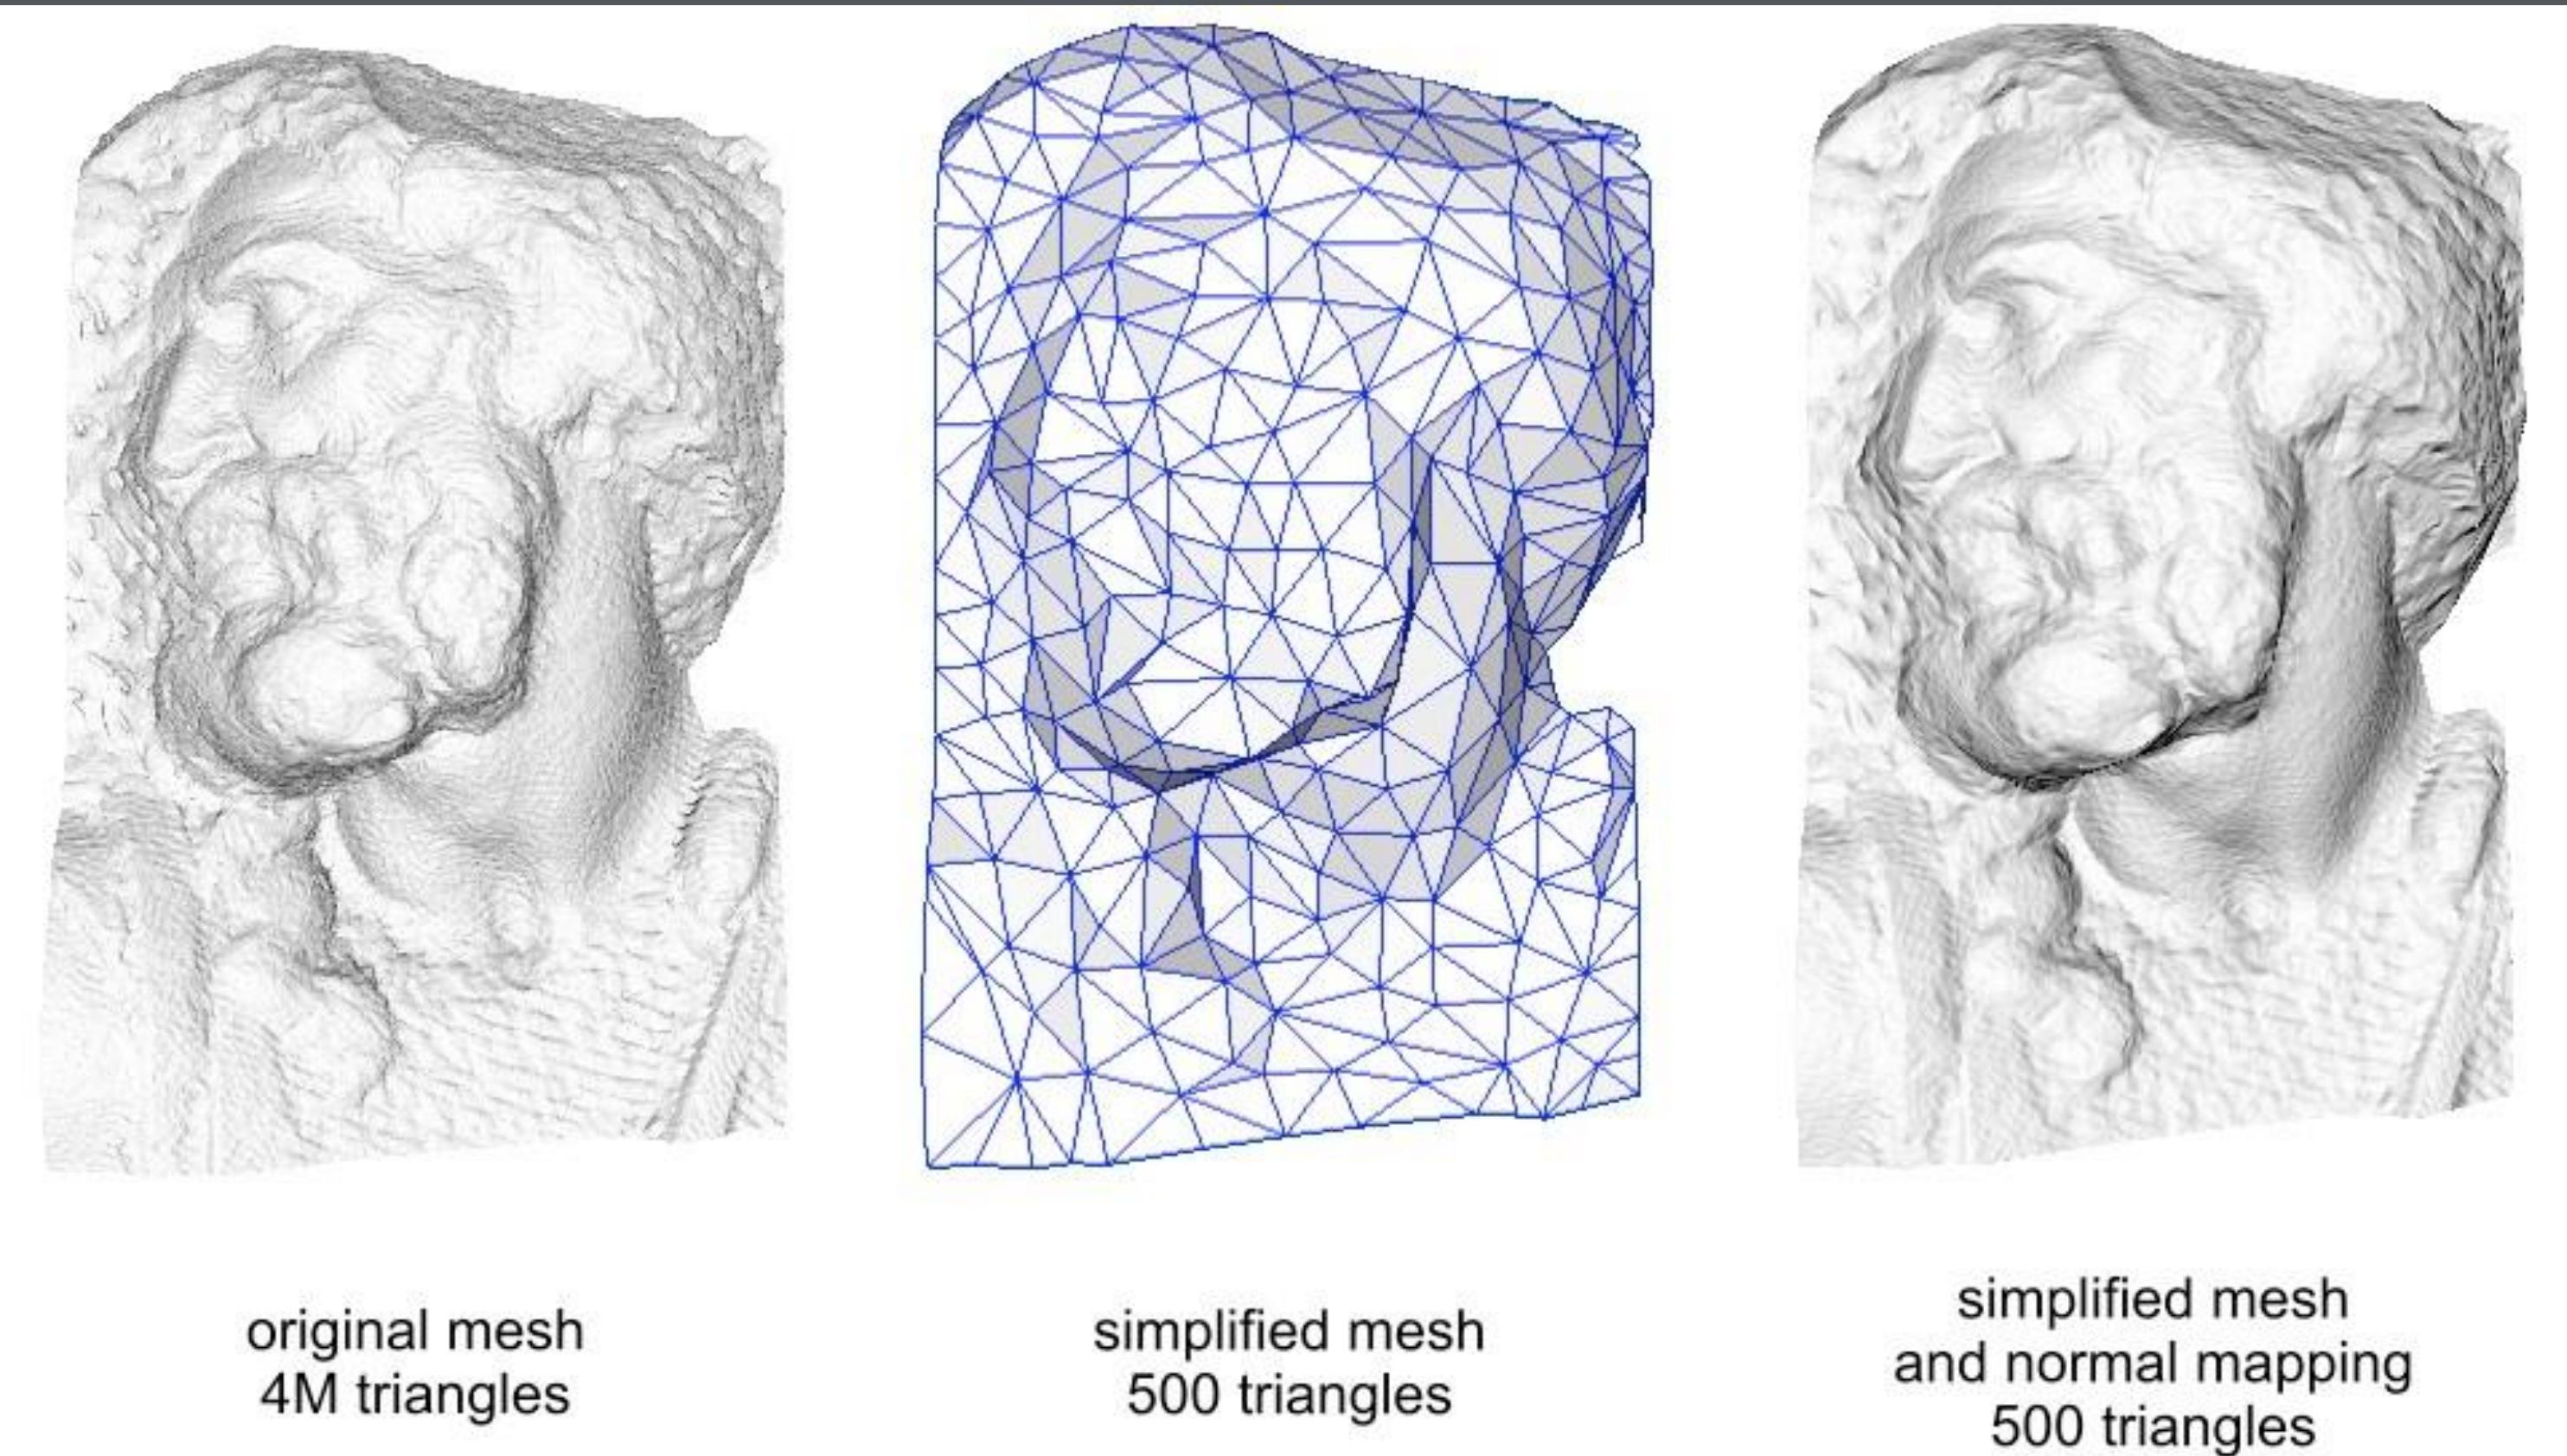
## Vertex phase

- simply transform and pass through the position and tangents

## Fragment phase

- compute normal to displaced surface
  - compute derivatives of height by finite differences
  - add offset to the base surface tangents; cross product is the shading normal
- compute shading using displaced normal

# Normal mapping



[Paolo  
Cignoni]

# Normal mapping

## Geometric prerequisites

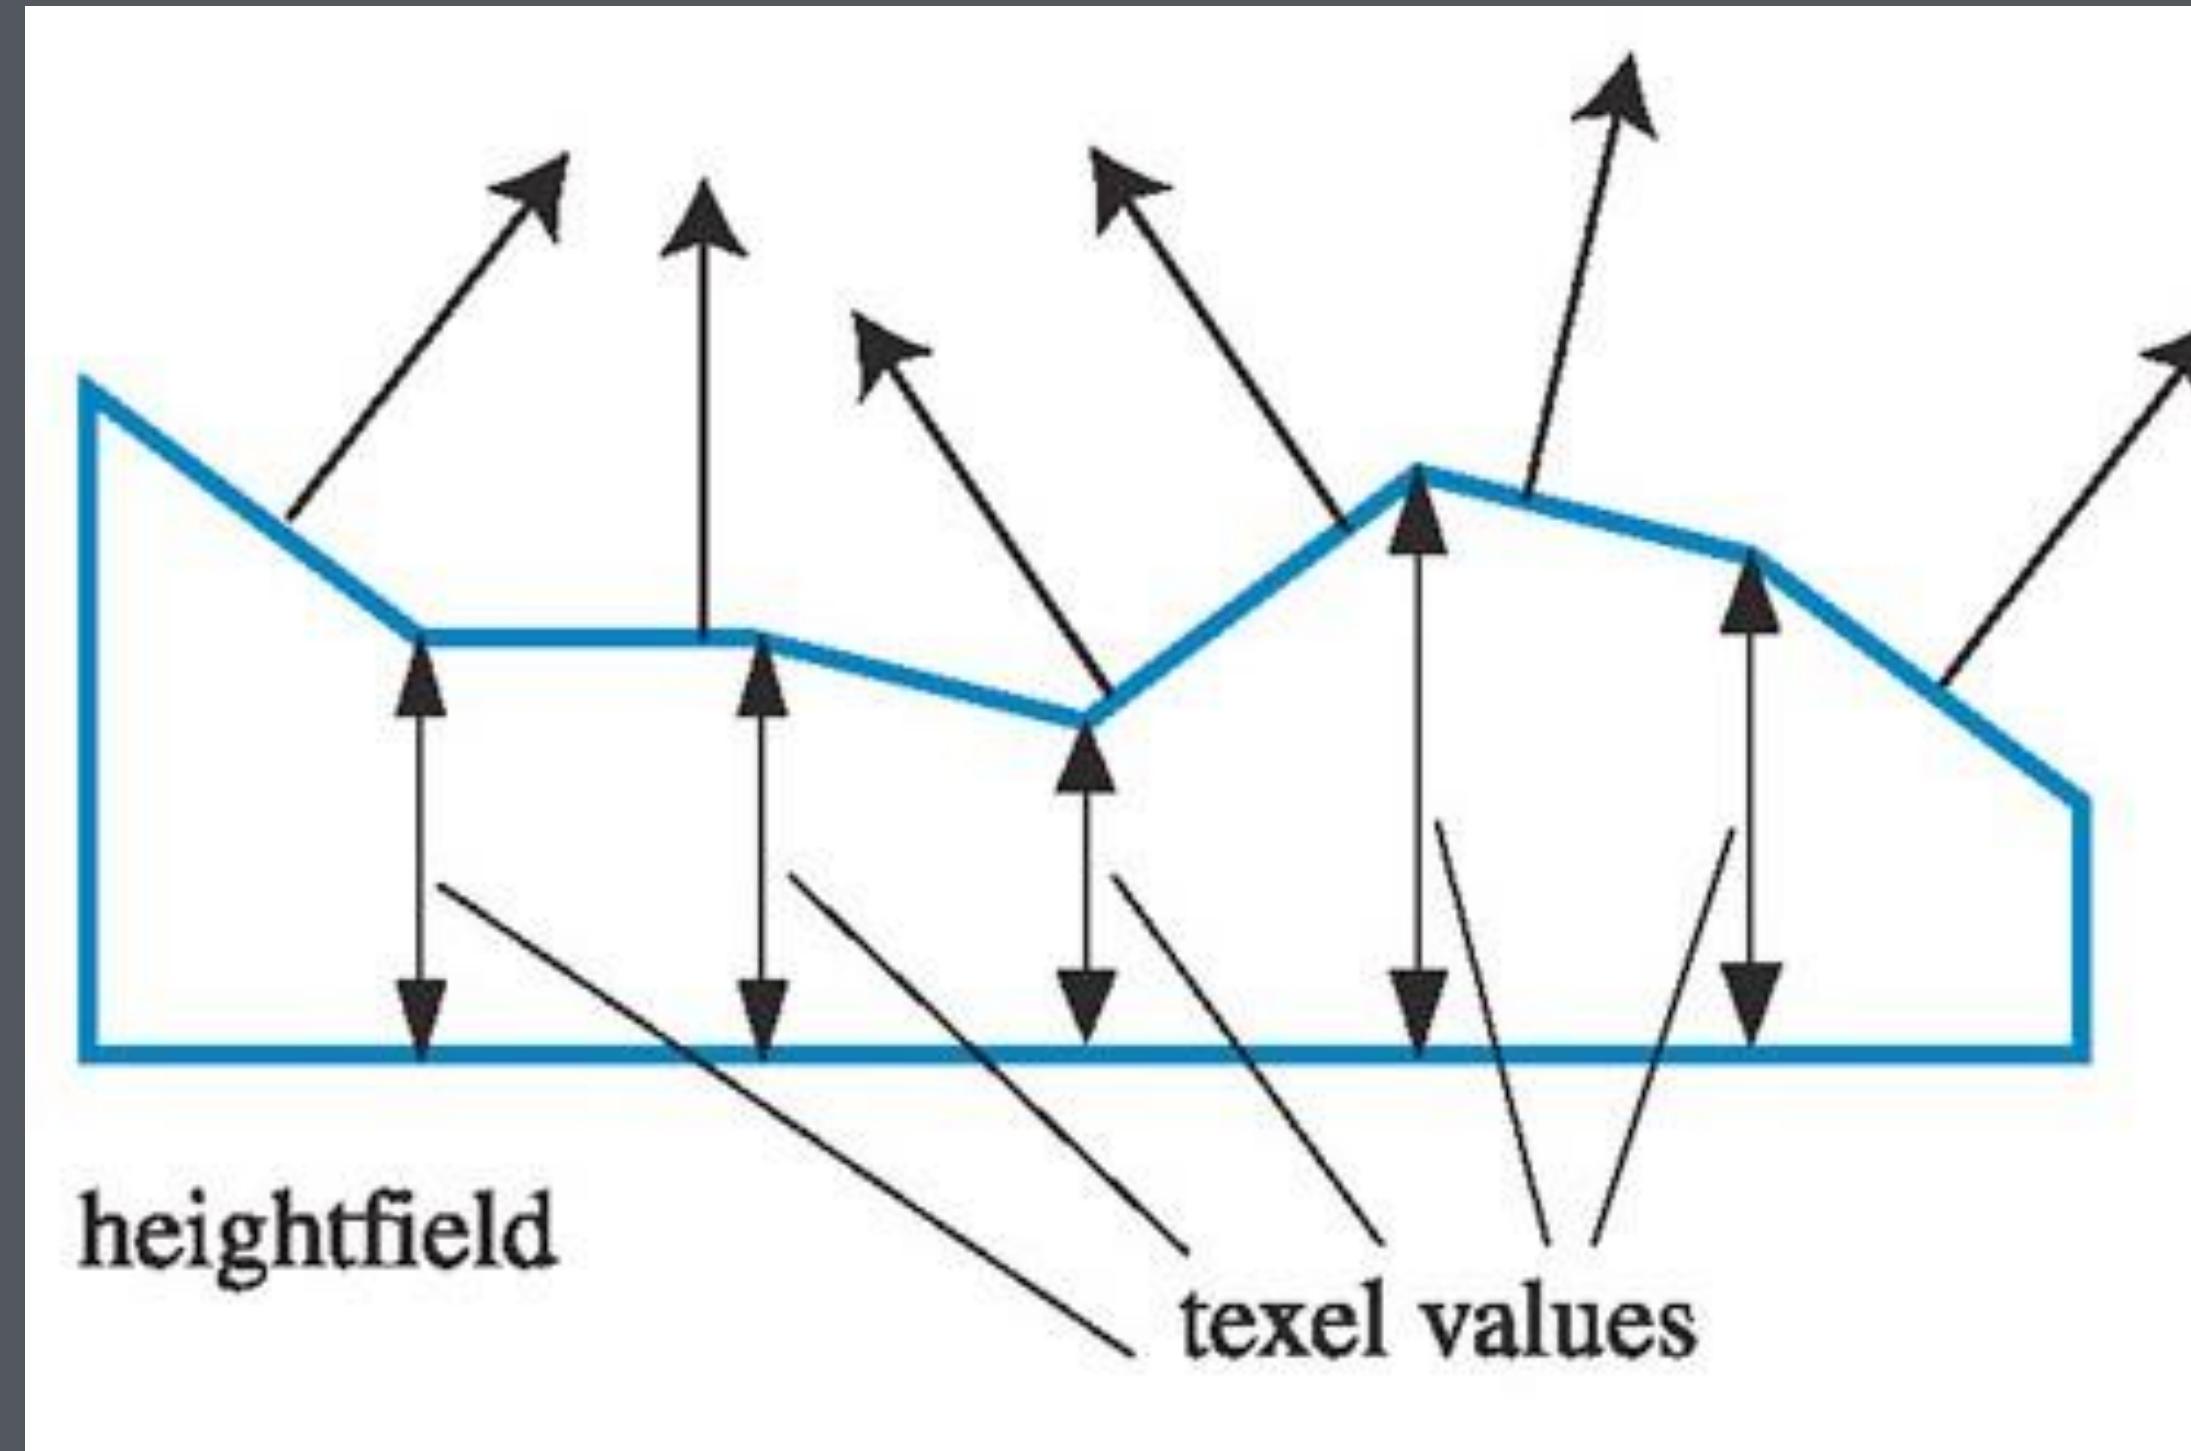
- Texture map (3 channels) representing normal field
  - single lookups into normal map required
- Smooth normals
- Unit tangent vectors
  - if you want to store normals in tangent space (and you do)
- No dense triangulation needed
- No finite differencing needed

## Geometric logic

- look up normal from map
- transform into (tangent-u, tangent-v, normal) space

# Heightfield: Blinn's original idea

Single scalar, more computation to infer  $\mathbf{N}'$



# Perturbed normal given height map

## Normal is determined by partial derivatives of height

- in the local frame of the displacement map:

$$\mathbf{n}_{\text{disp}} = (h_u, h_v, 1)$$

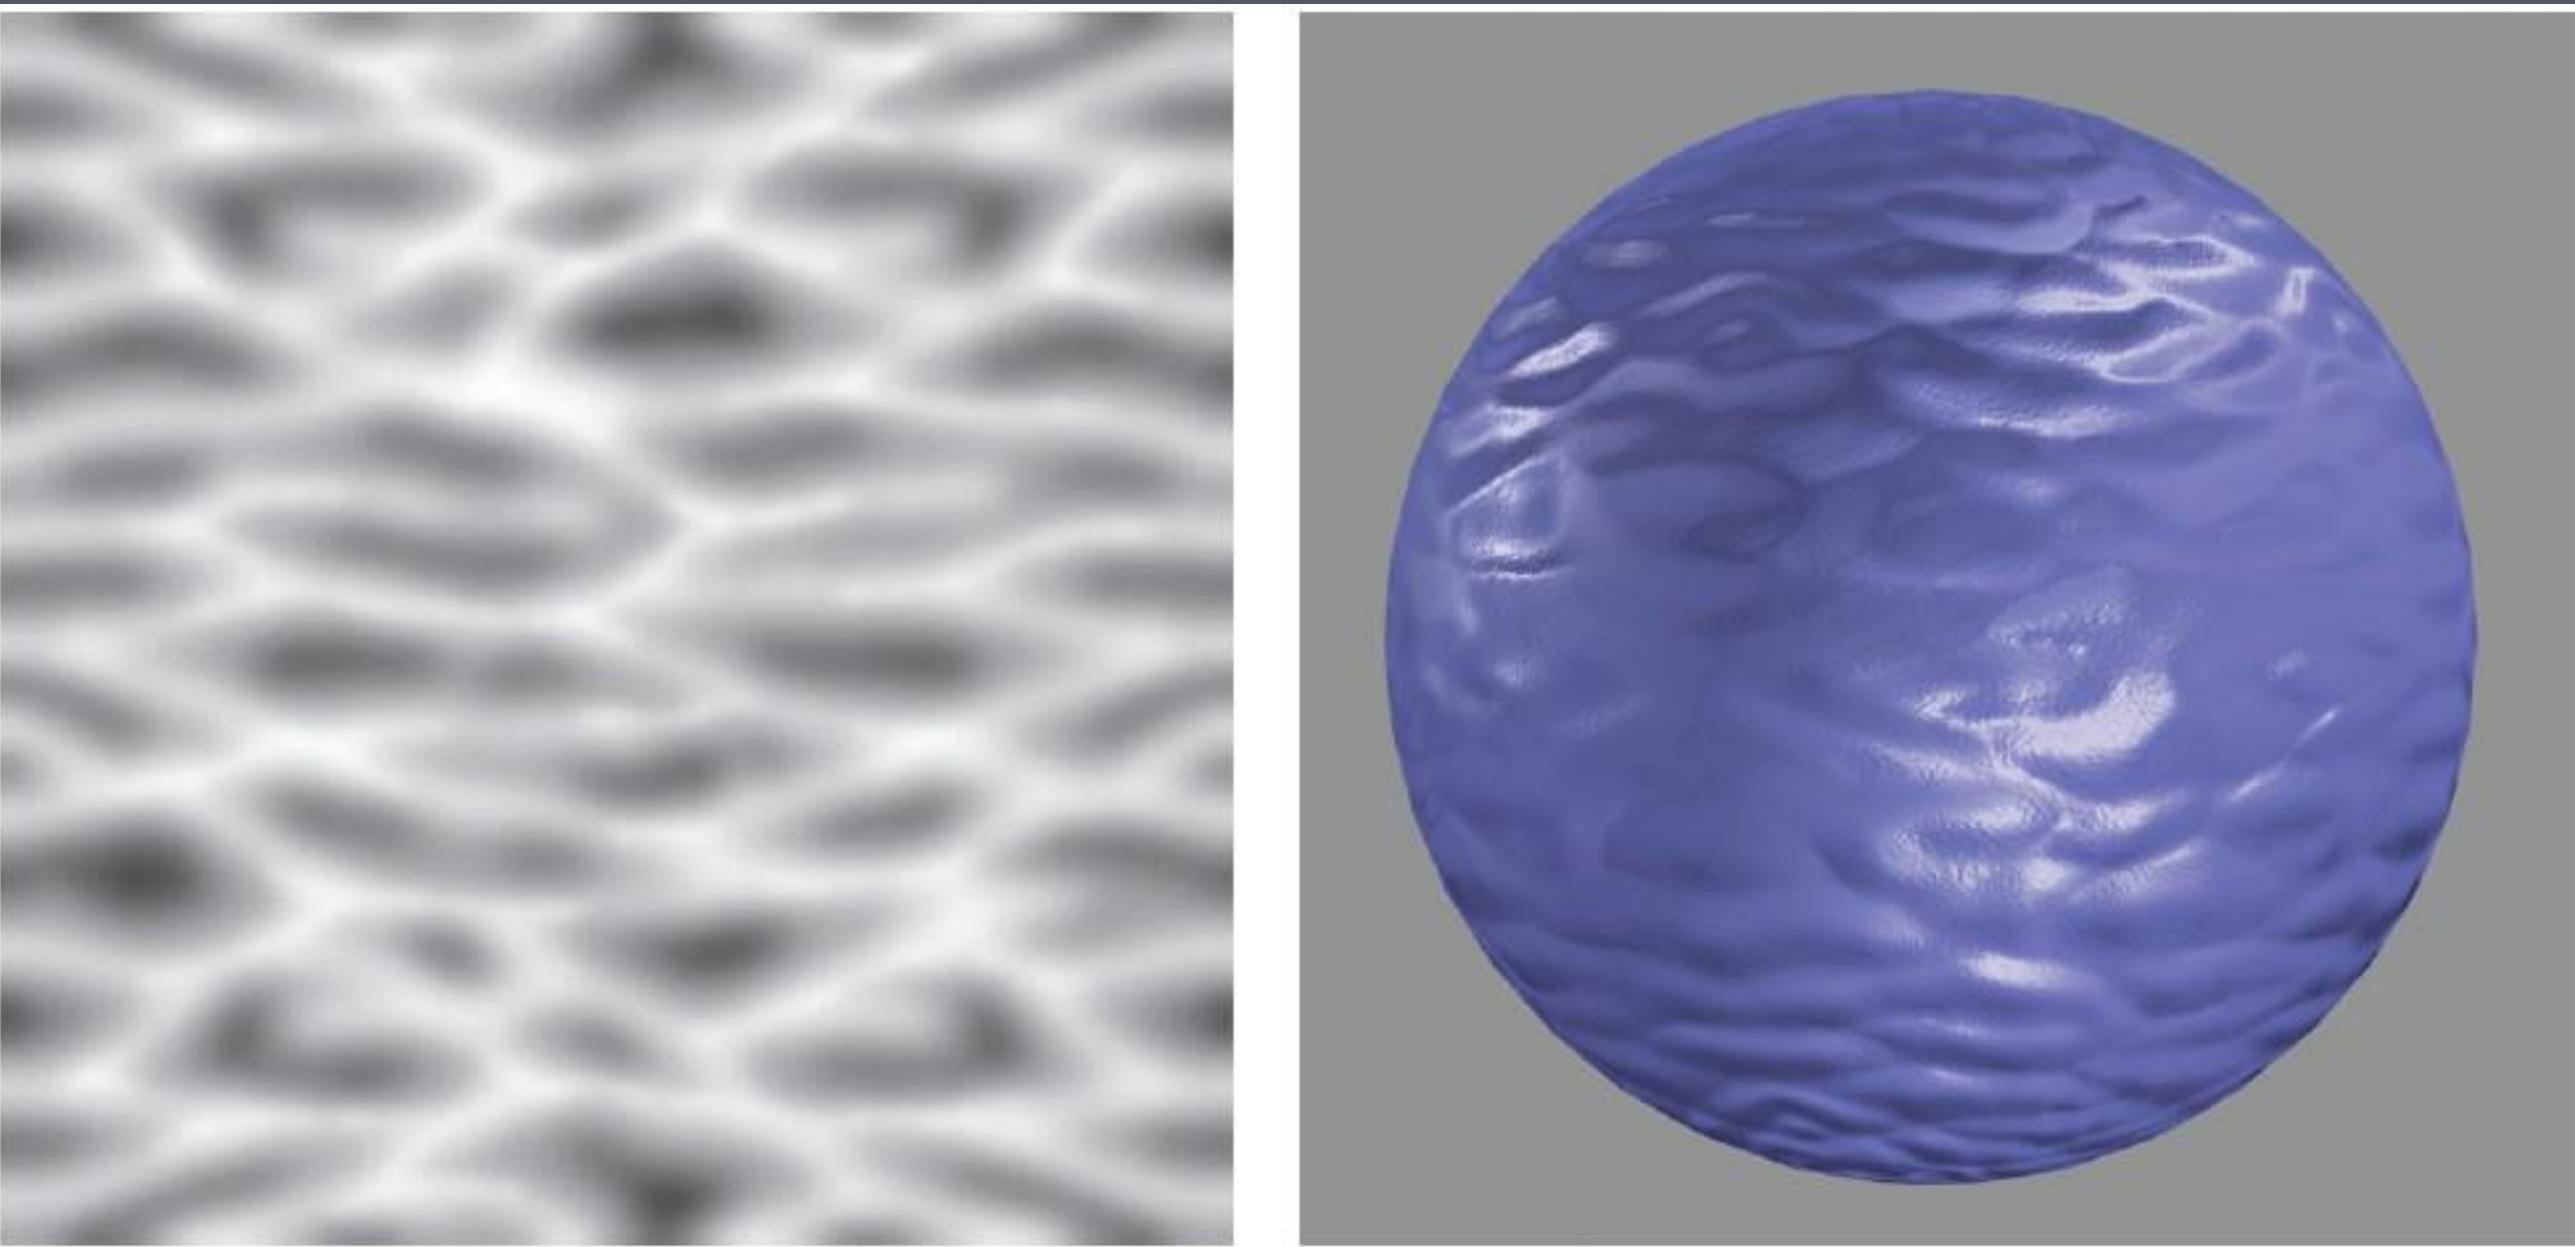
- approx: heights are small compared to radius of curvature (constant normal)
- then the displaced surface is locally a linear transformation of the height field
- normal transforms by the adjoint matrix (as normals always to)
- perform 4 lookups to get 4 neighboring height values
- subtract to obtain finite difference derivatives

# Height Field Bump Maps

Older technique, less memory

Texture map value is a height

Gray scale value: light is +, dark is -

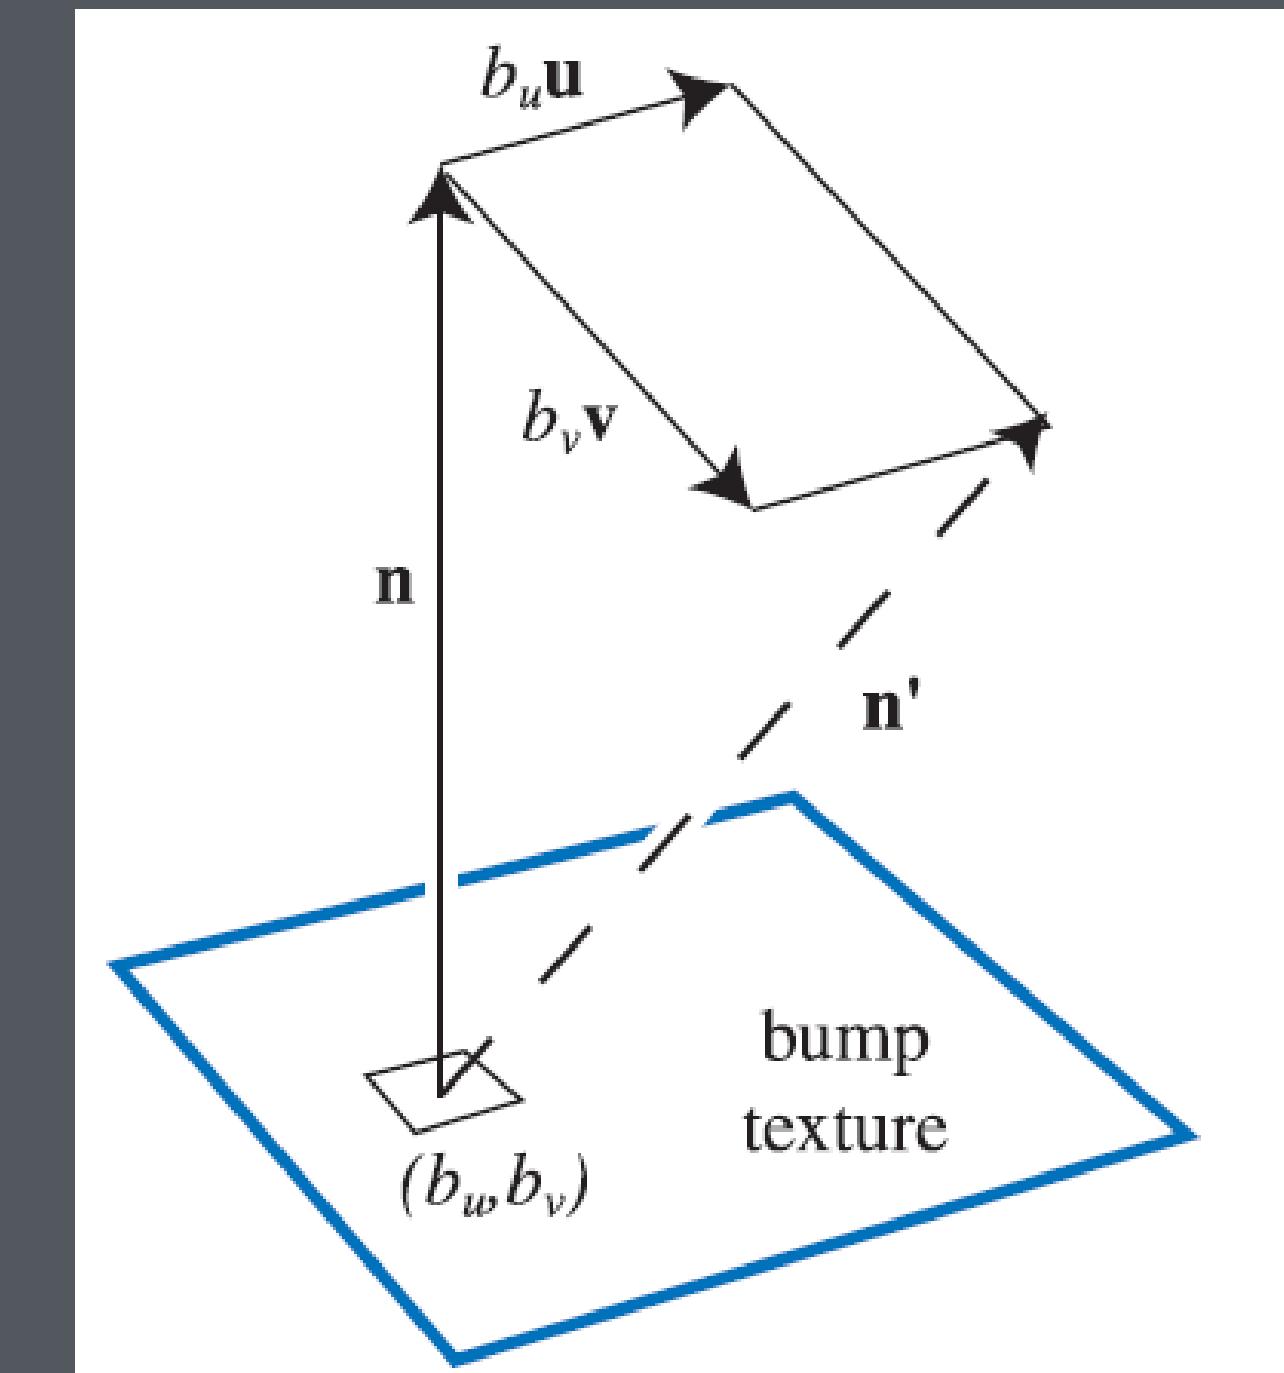


# Bump Mapping

Look up  $b_u$  and  $b_v$

$N'$  is not normalized

$$N' = N + b_u T + b_v B$$



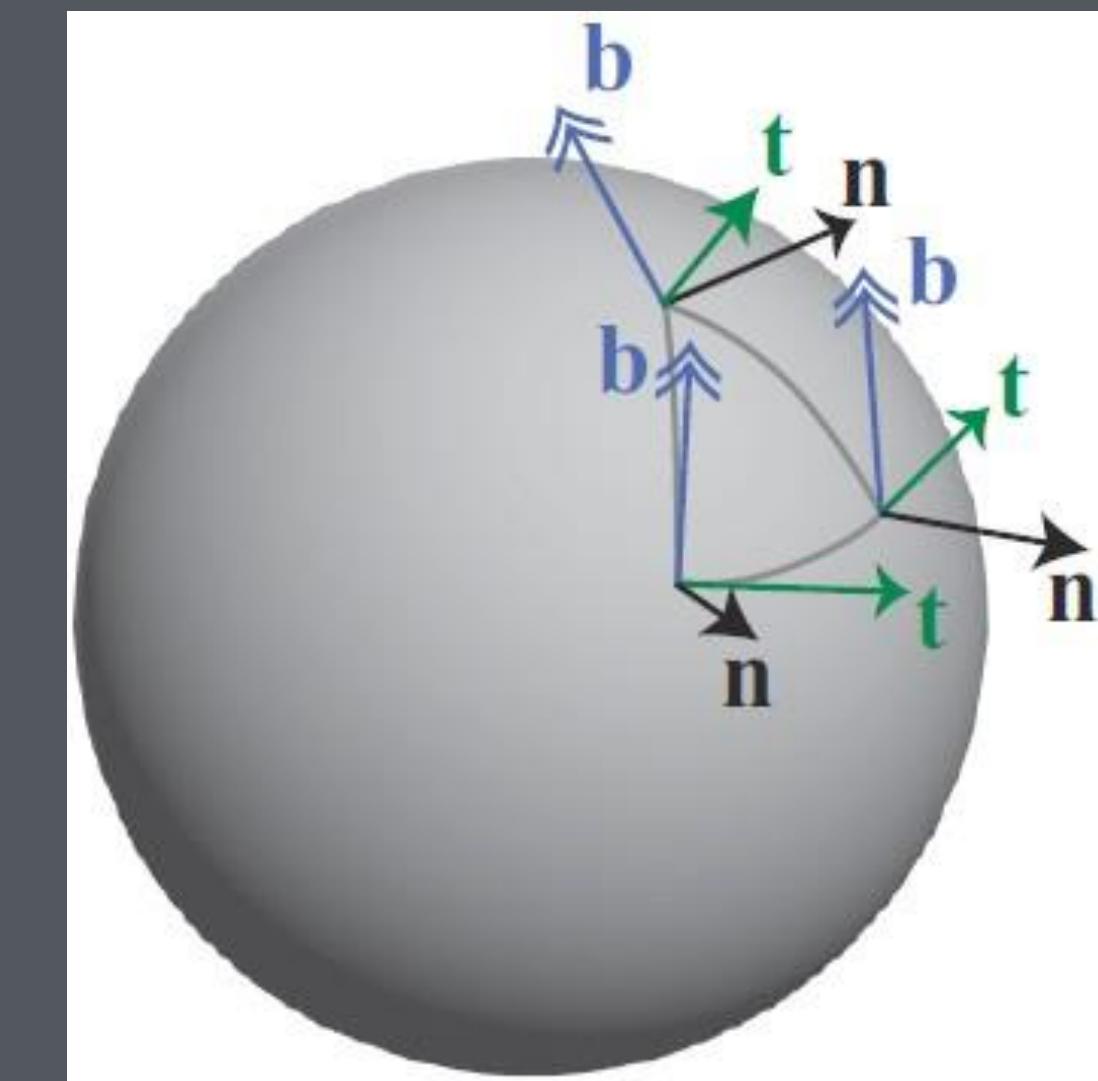
# Rendering with Bump Maps

$N' \cdot L$

Perturb  $N$  to get  $N'$  using bump map

Transform  $L$  to tangent space of surface

- Have  $N$ ,  $T$  (tangent), bitangent  $B = T \times N$



# Normal Maps

**Preferred technique for bump mapping for modern graphics cards**

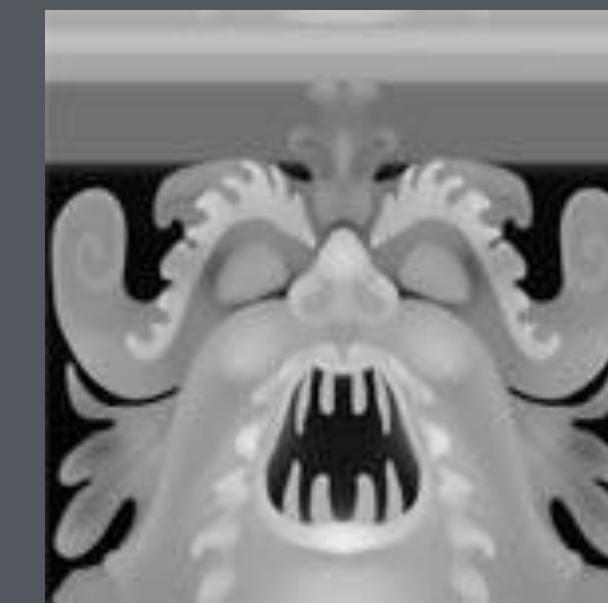
**Store new normals in texture map**

- Encodes  $(x, y, z)$  mapped to  $[-1, 1]$

**More memory but lower computation**



Normal Map



Height Map

- Store

```
colorComponent = 0.5 * normalComponent + 0.5
```

- 

- Use

```
normalComponent = 2* colorComponent -1
```

# Normal Map

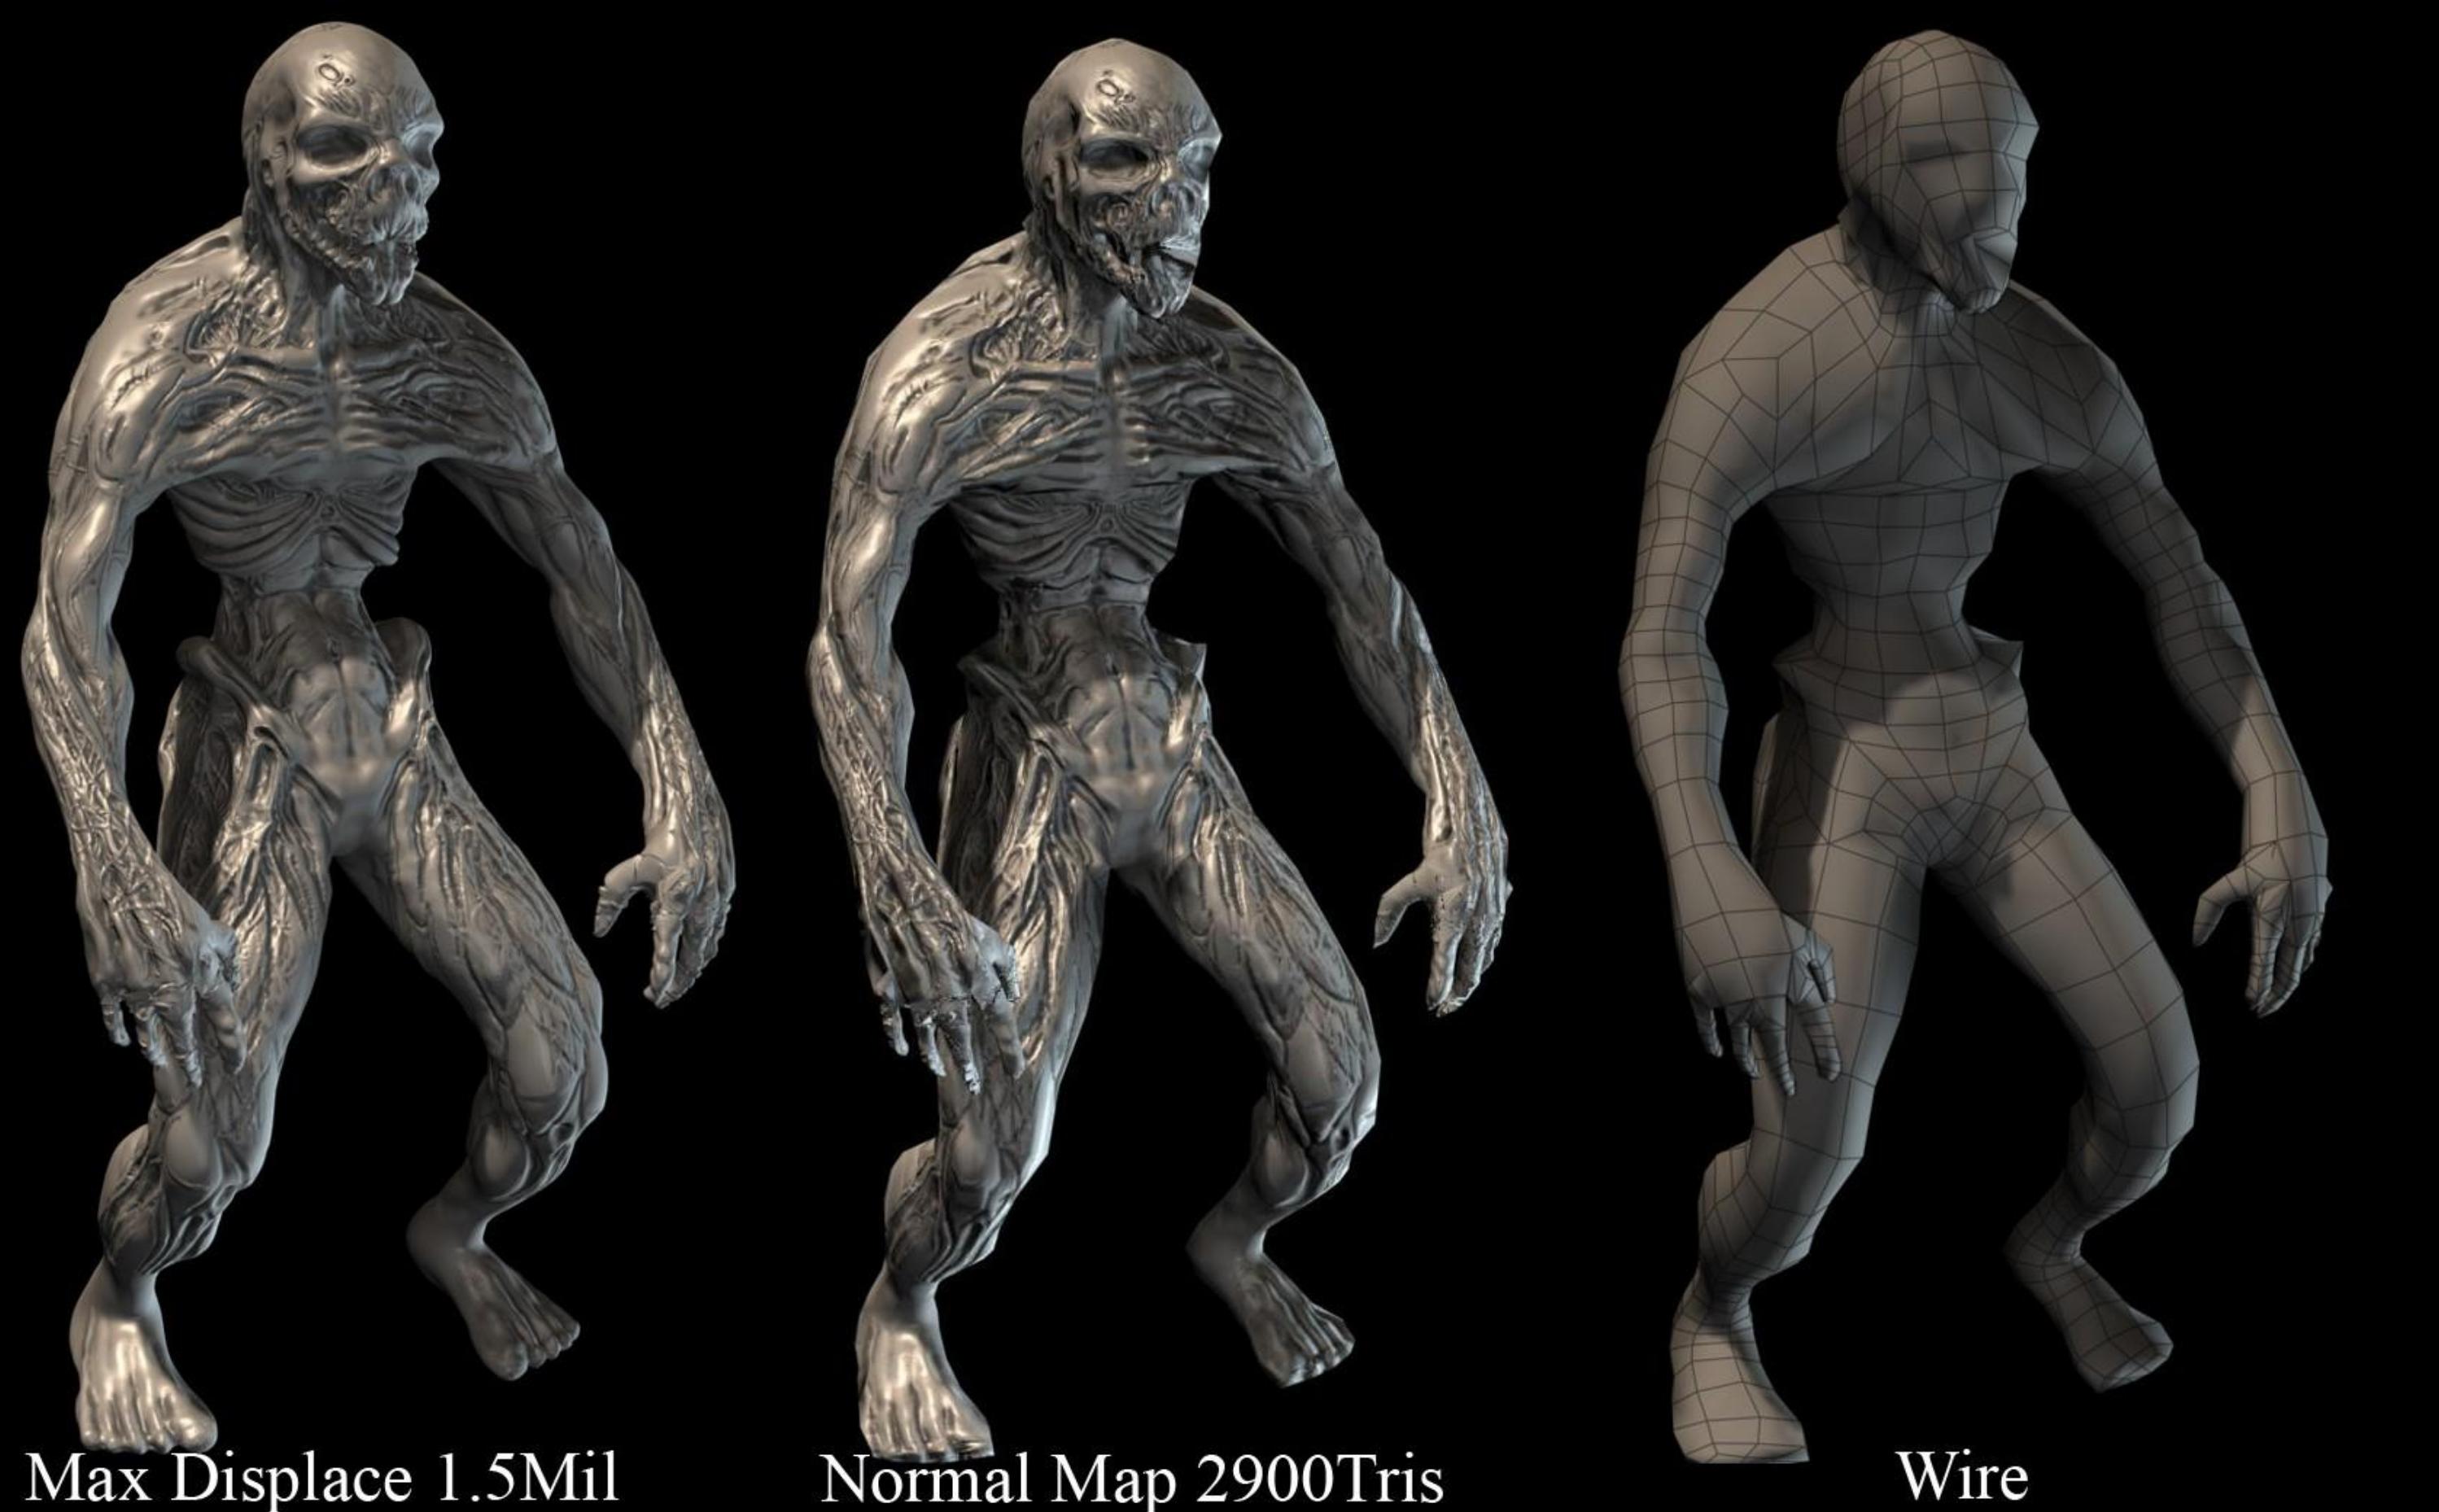


# Creating Normal Maps

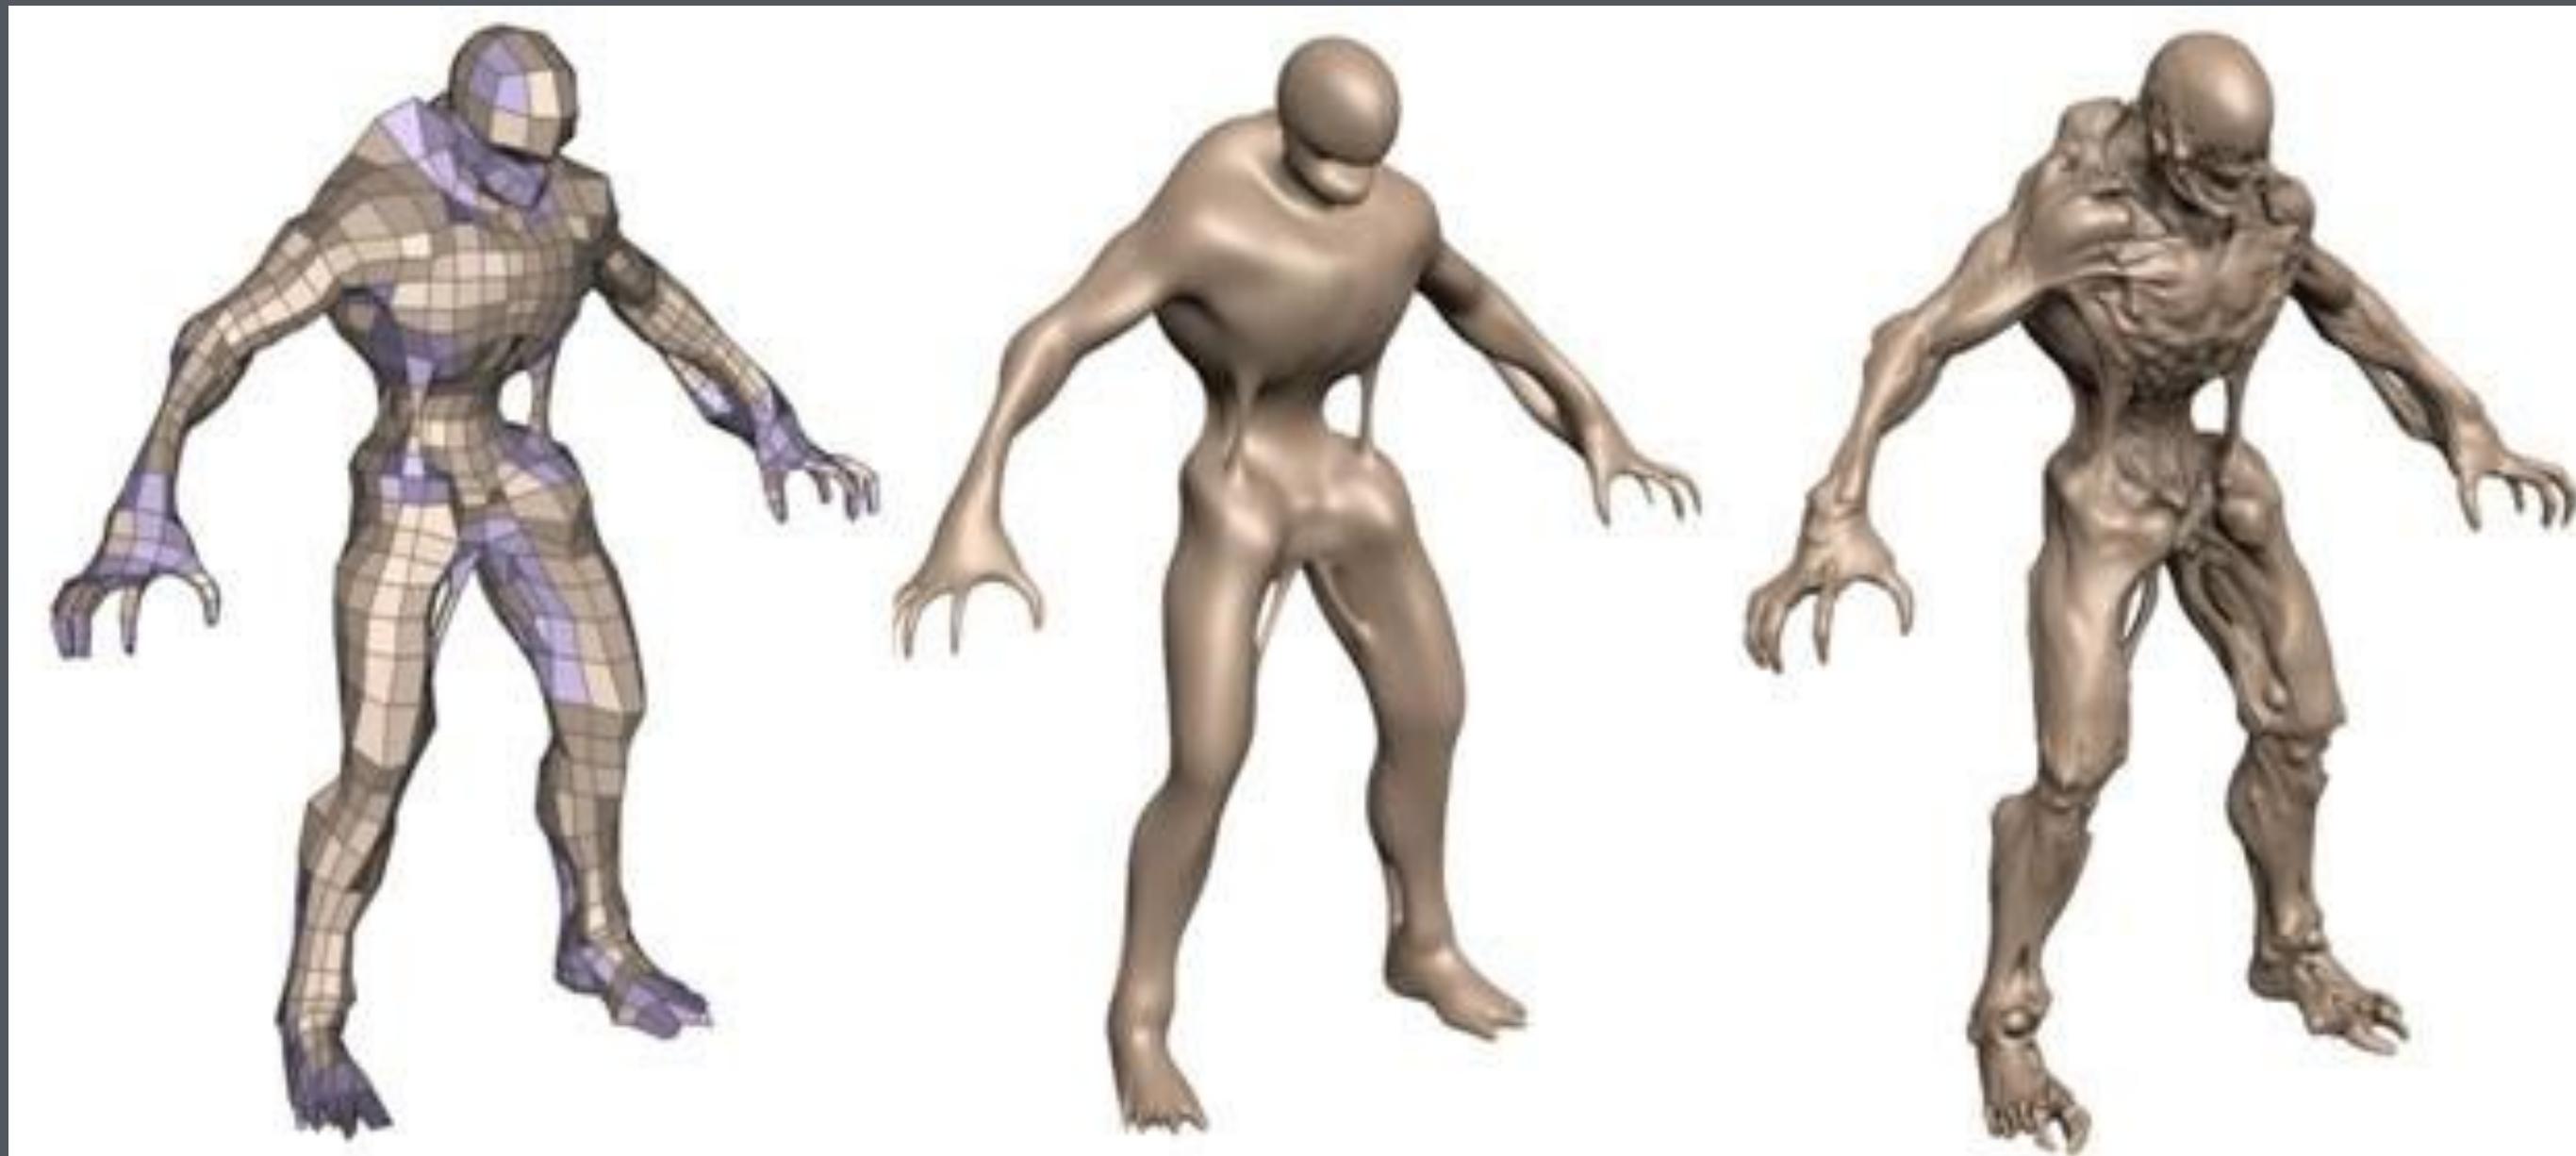
**First create complex geometry**

**Simplify (in modeling time) to simple mesh with normal map**

# Displacement Maps vs. Normal Maps



# Compare with the opposite view

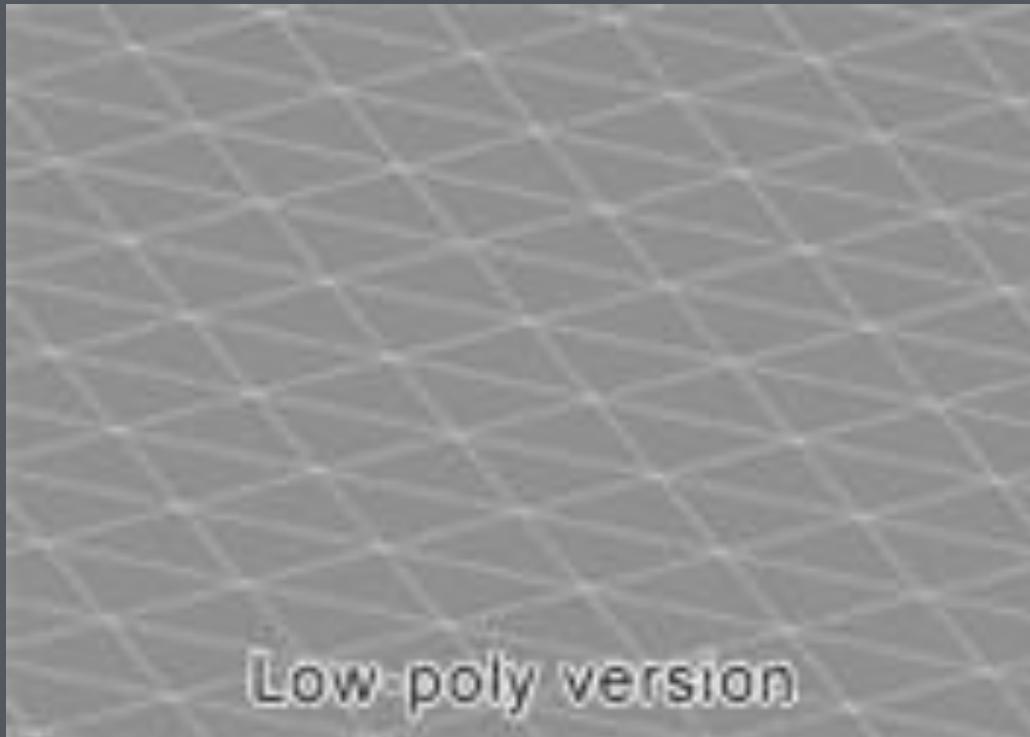


Original

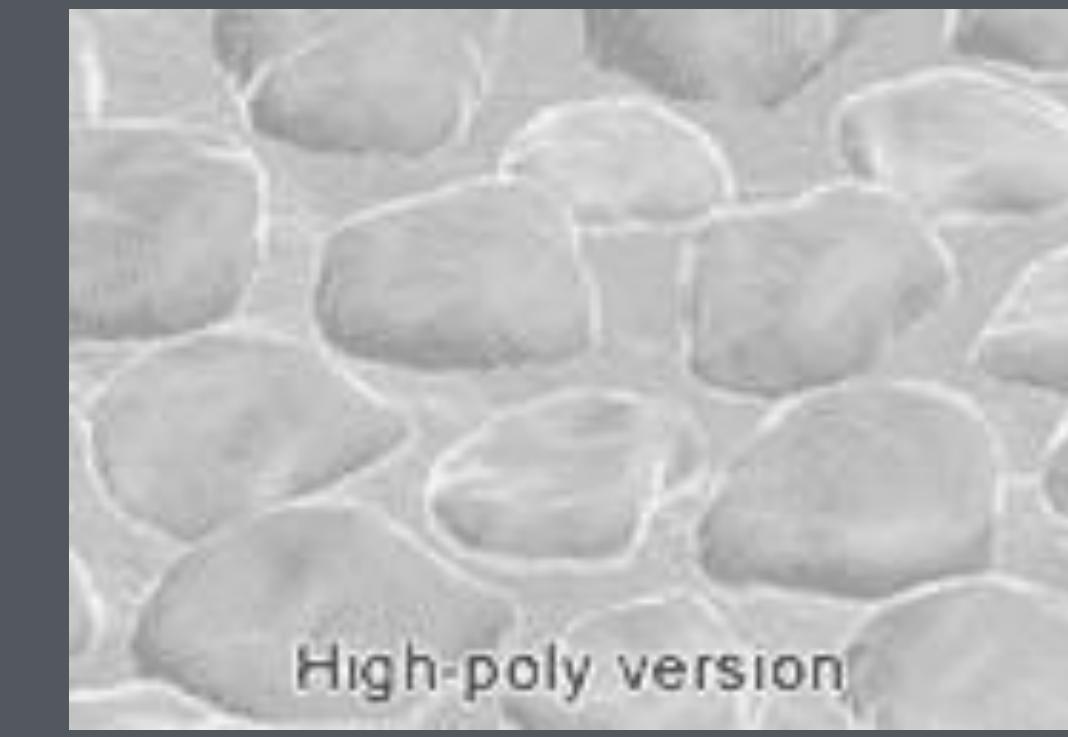
Tesselated

Displacement  
Mapped

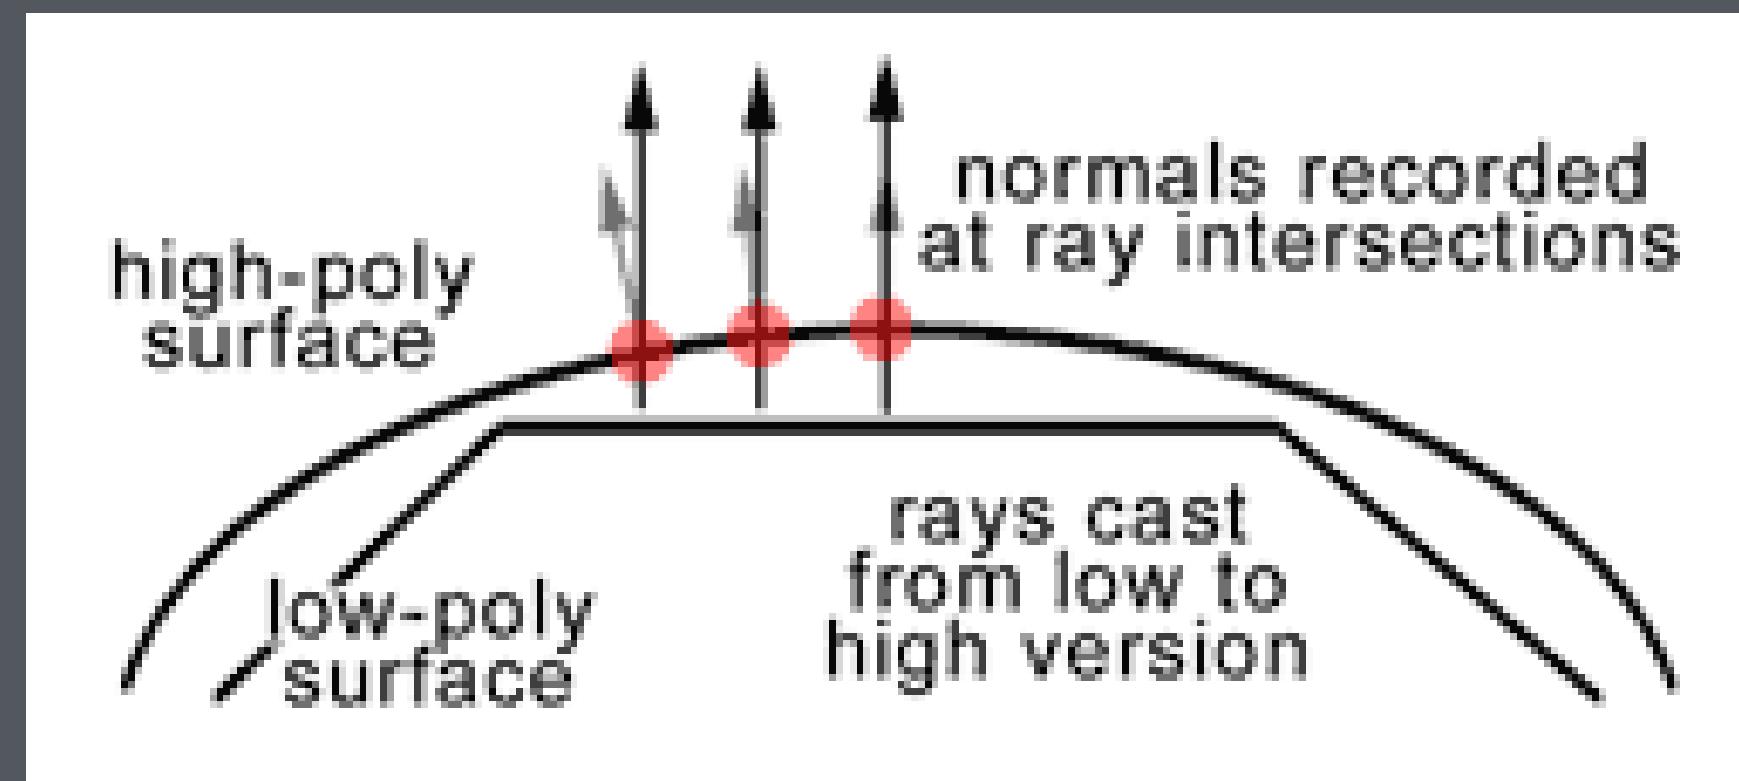
# Creating Normal Maps



Low-poly version



High-poly version



# Which space is normal map in?

## World space

- Easy computation, but can't use the same normal map for...
  - two walls
  - A rotating object

## Object space

- Better, but cannot be reused for...
  - deforming objects
  - different objects with similar material

## Tangent space

- Can reuse for deforming surfaces
- Transform lighting to this space and shade

# Parallax Mapping

## Problem with normal mapping

- No self-occlusion
- Supposed to be a height field but never see this occlusion across different viewing angles

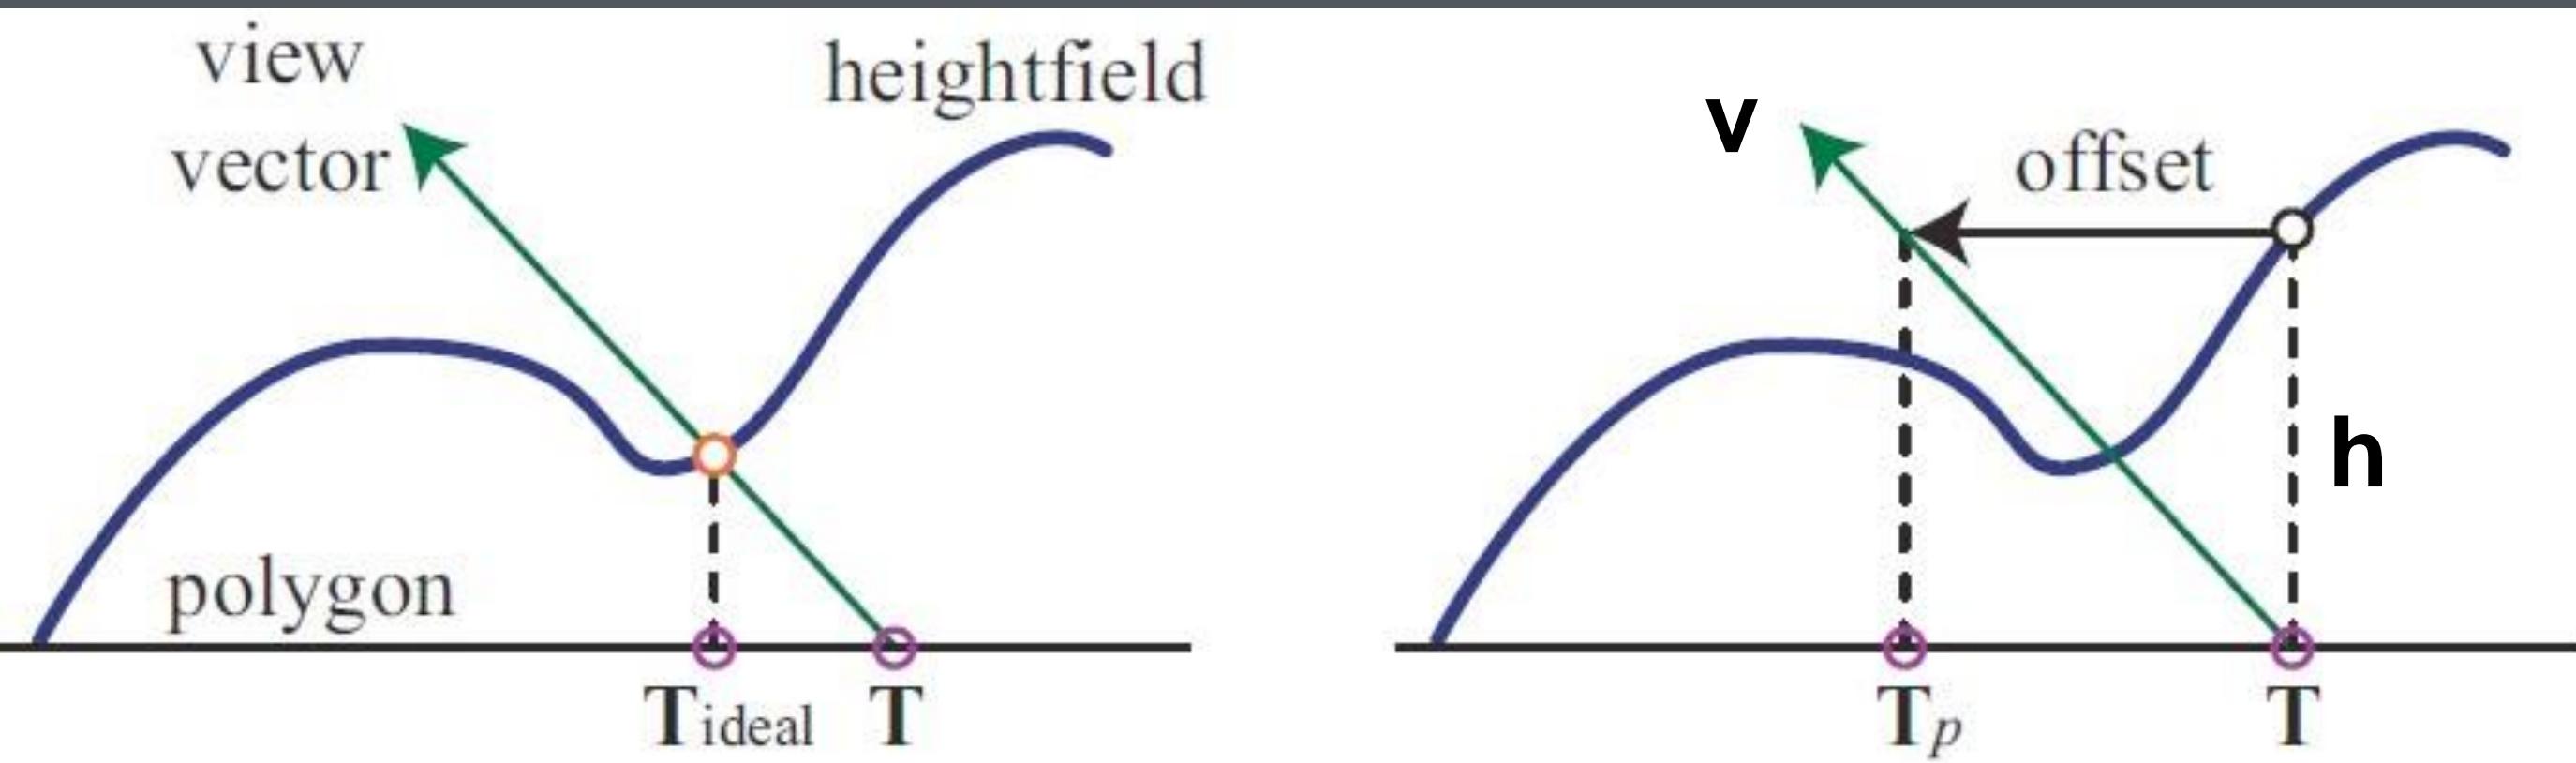
## Parallax mapping

- Positions of objects move relative to one other as viewpoint changes

# Parallax Mapping

Want  $T_{\text{ideal}}$

Use  $T_p$  to approximate it



$$T_p = T + h \frac{\mathbf{v}_{xy}}{\mathbf{v}_z}$$

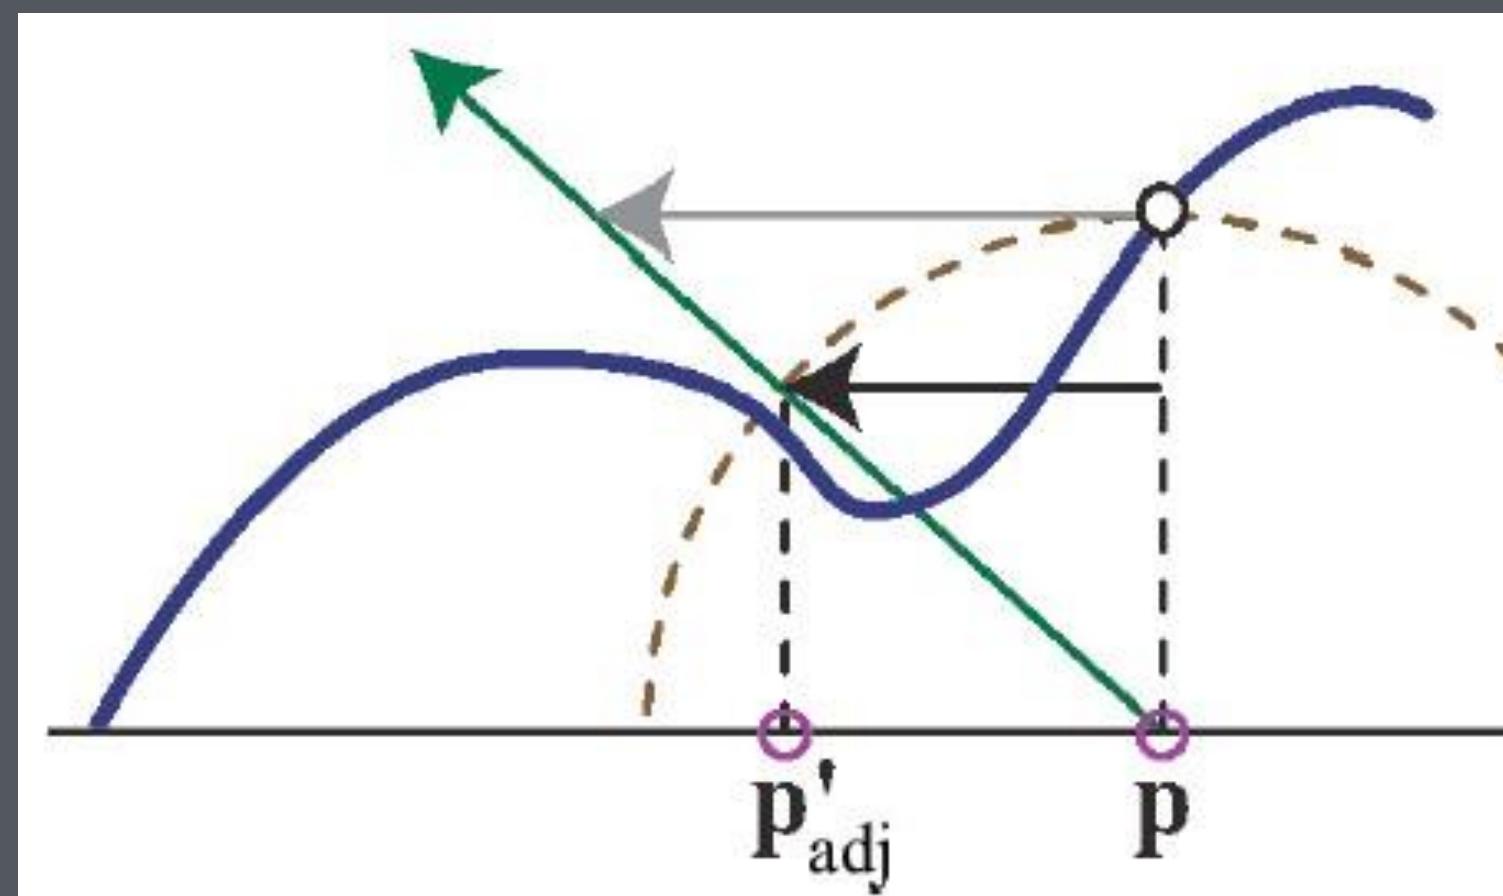
# Parallax Offset Limiting

**Problem:** at steep viewing, can offset too much

**Limit offset**

- results in a simpler formula

$$\mathbf{p}_{\text{adj}} = \mathbf{p} + h\mathbf{v}_{xy}$$



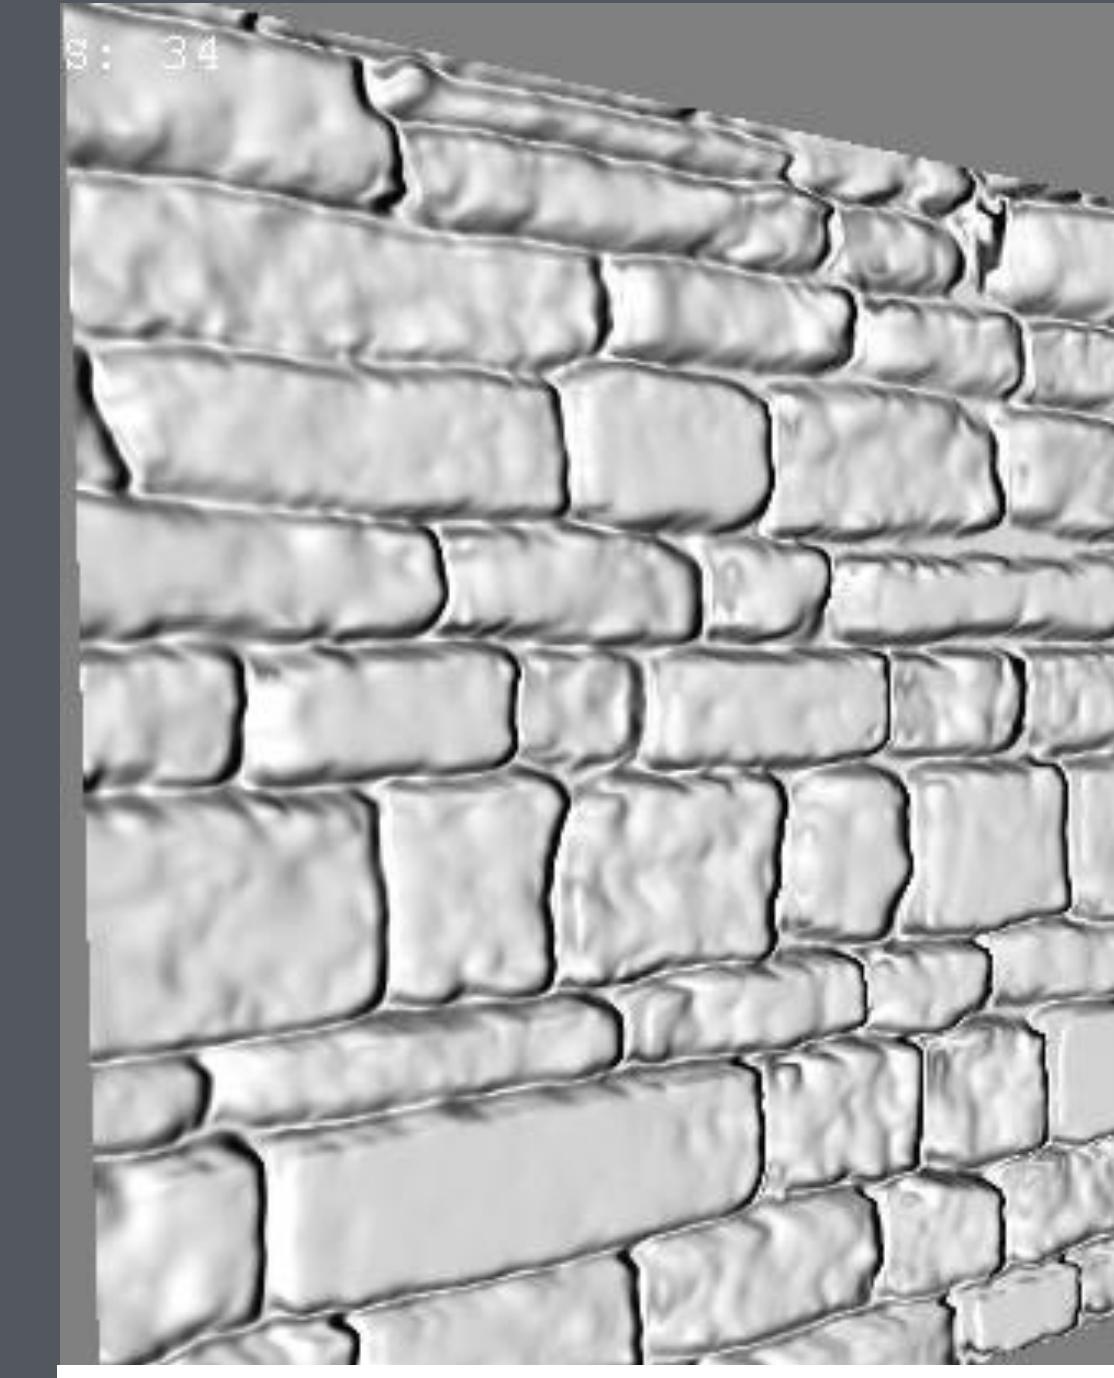
# Parallax Offset Limiting

Widely used in games

- the standard in simple bump mapping



Normal Mapping



Parallax Mapping Offset Limiting



1,100 polygon object w/  
parallax occlusion mapping



1.5 million polygon

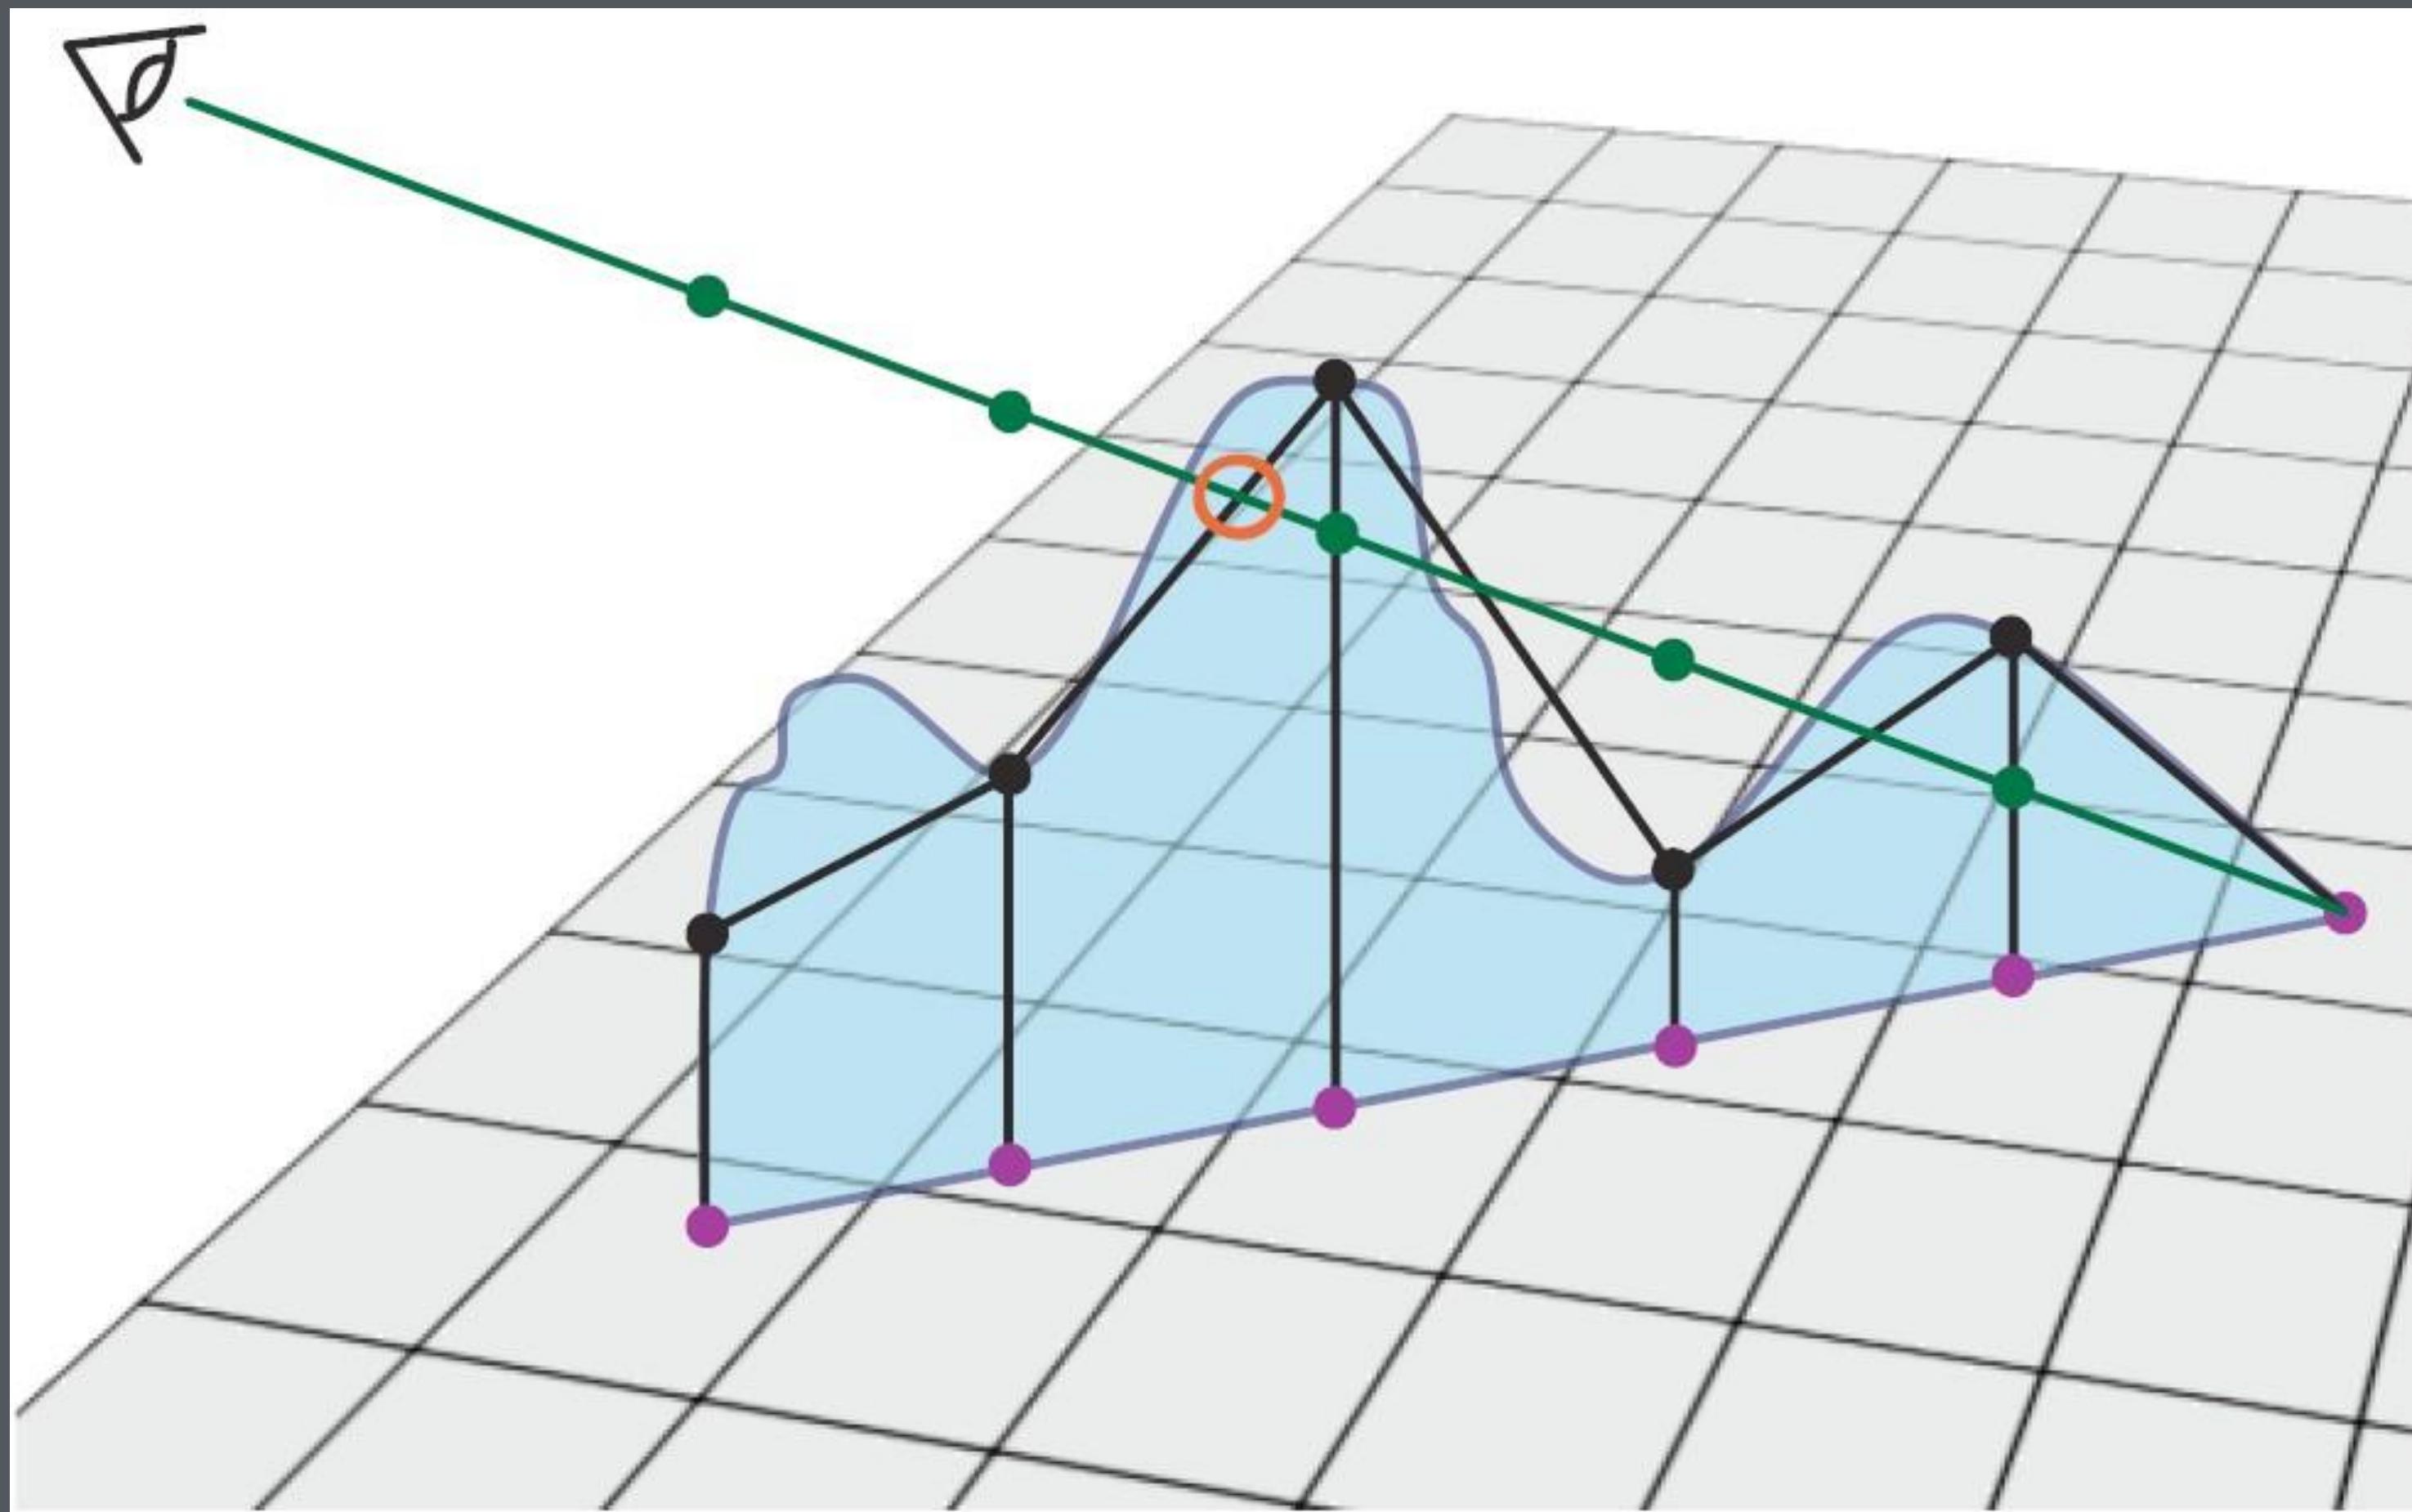
# Relief Mapping

**Aka Parallax occlusion mapping, relief mapping, steep parallax mapping**

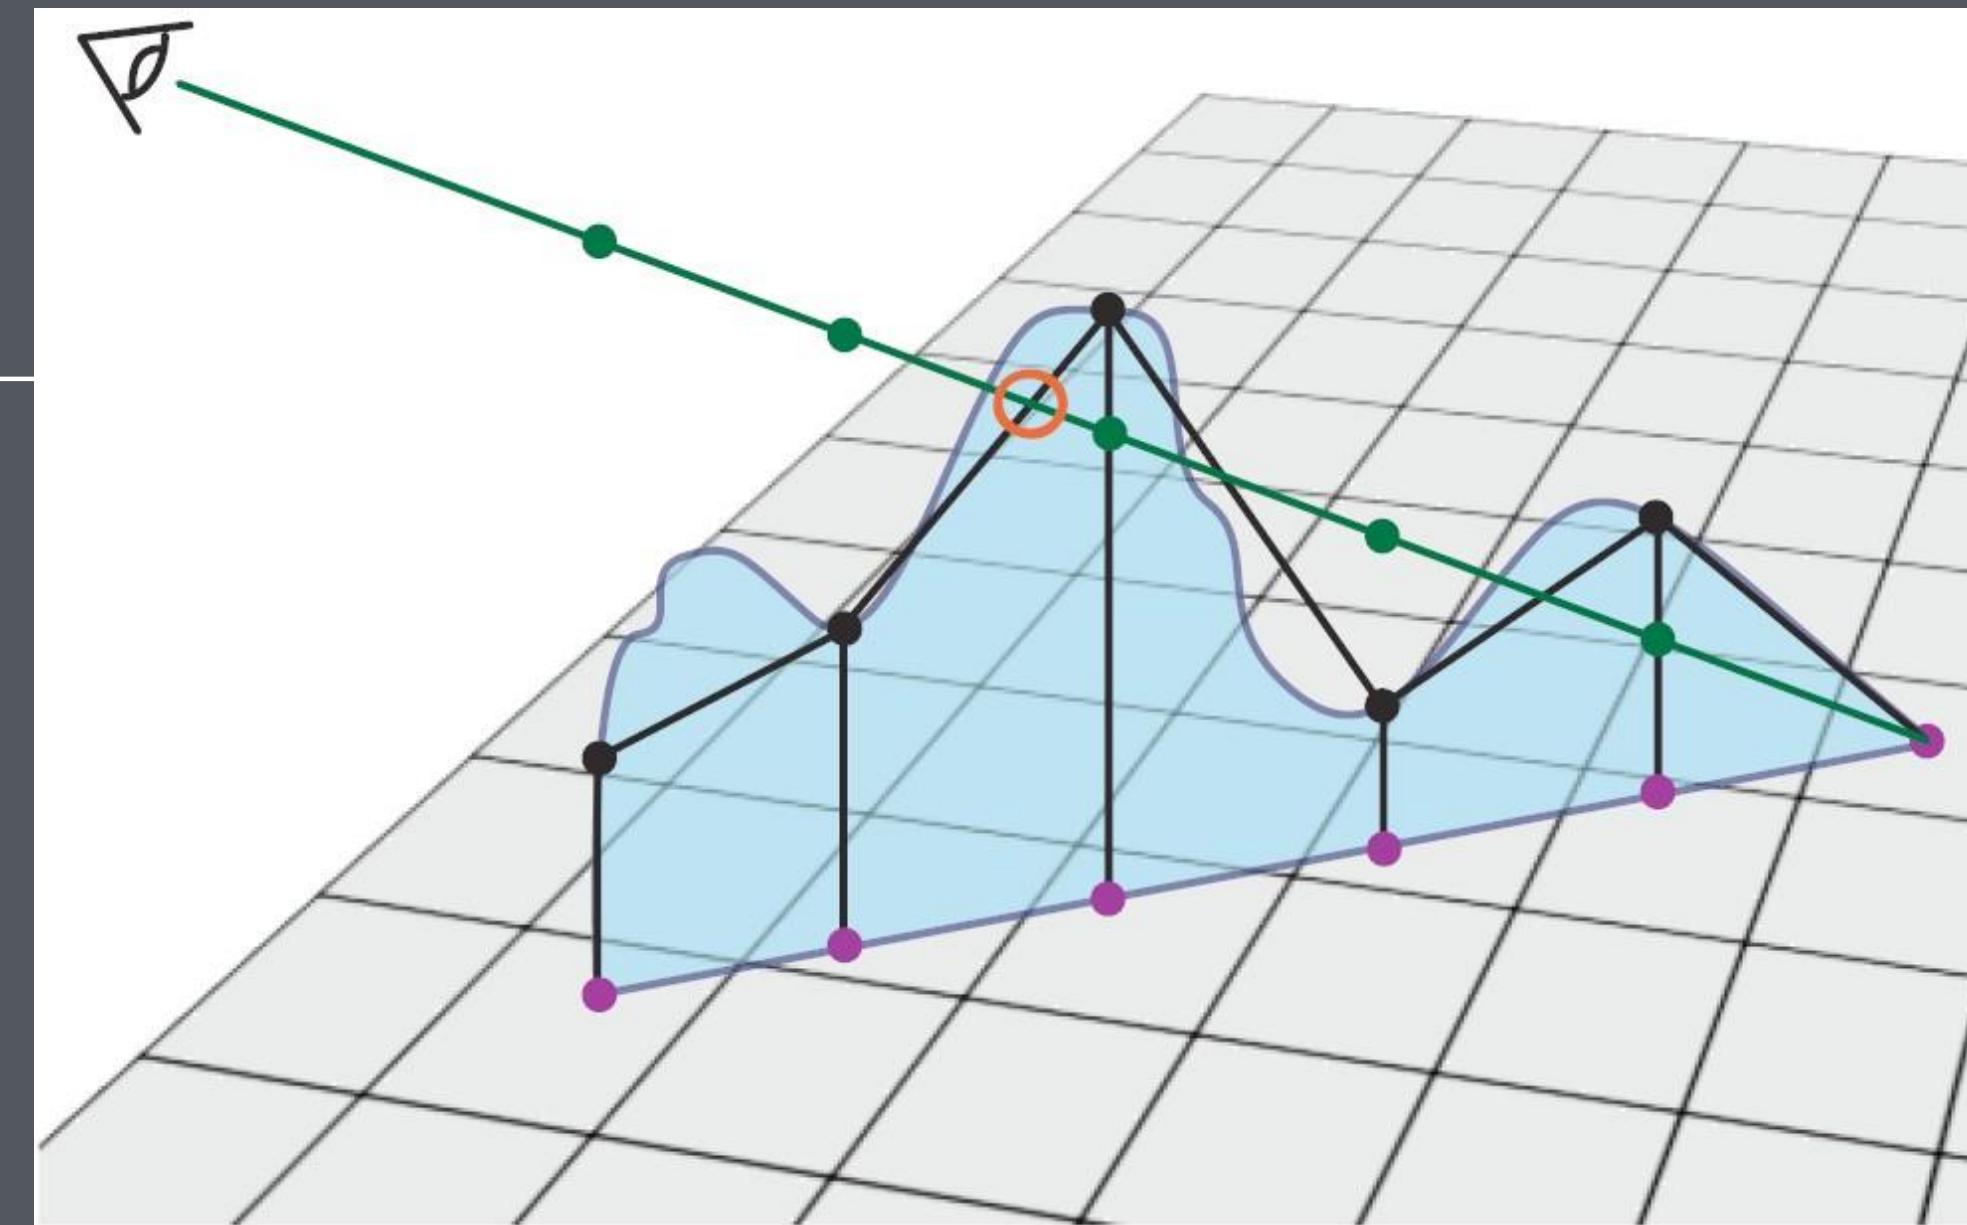
**Tries to find where the view ray intersects the height field**

- Kinda

# Relief Mapping

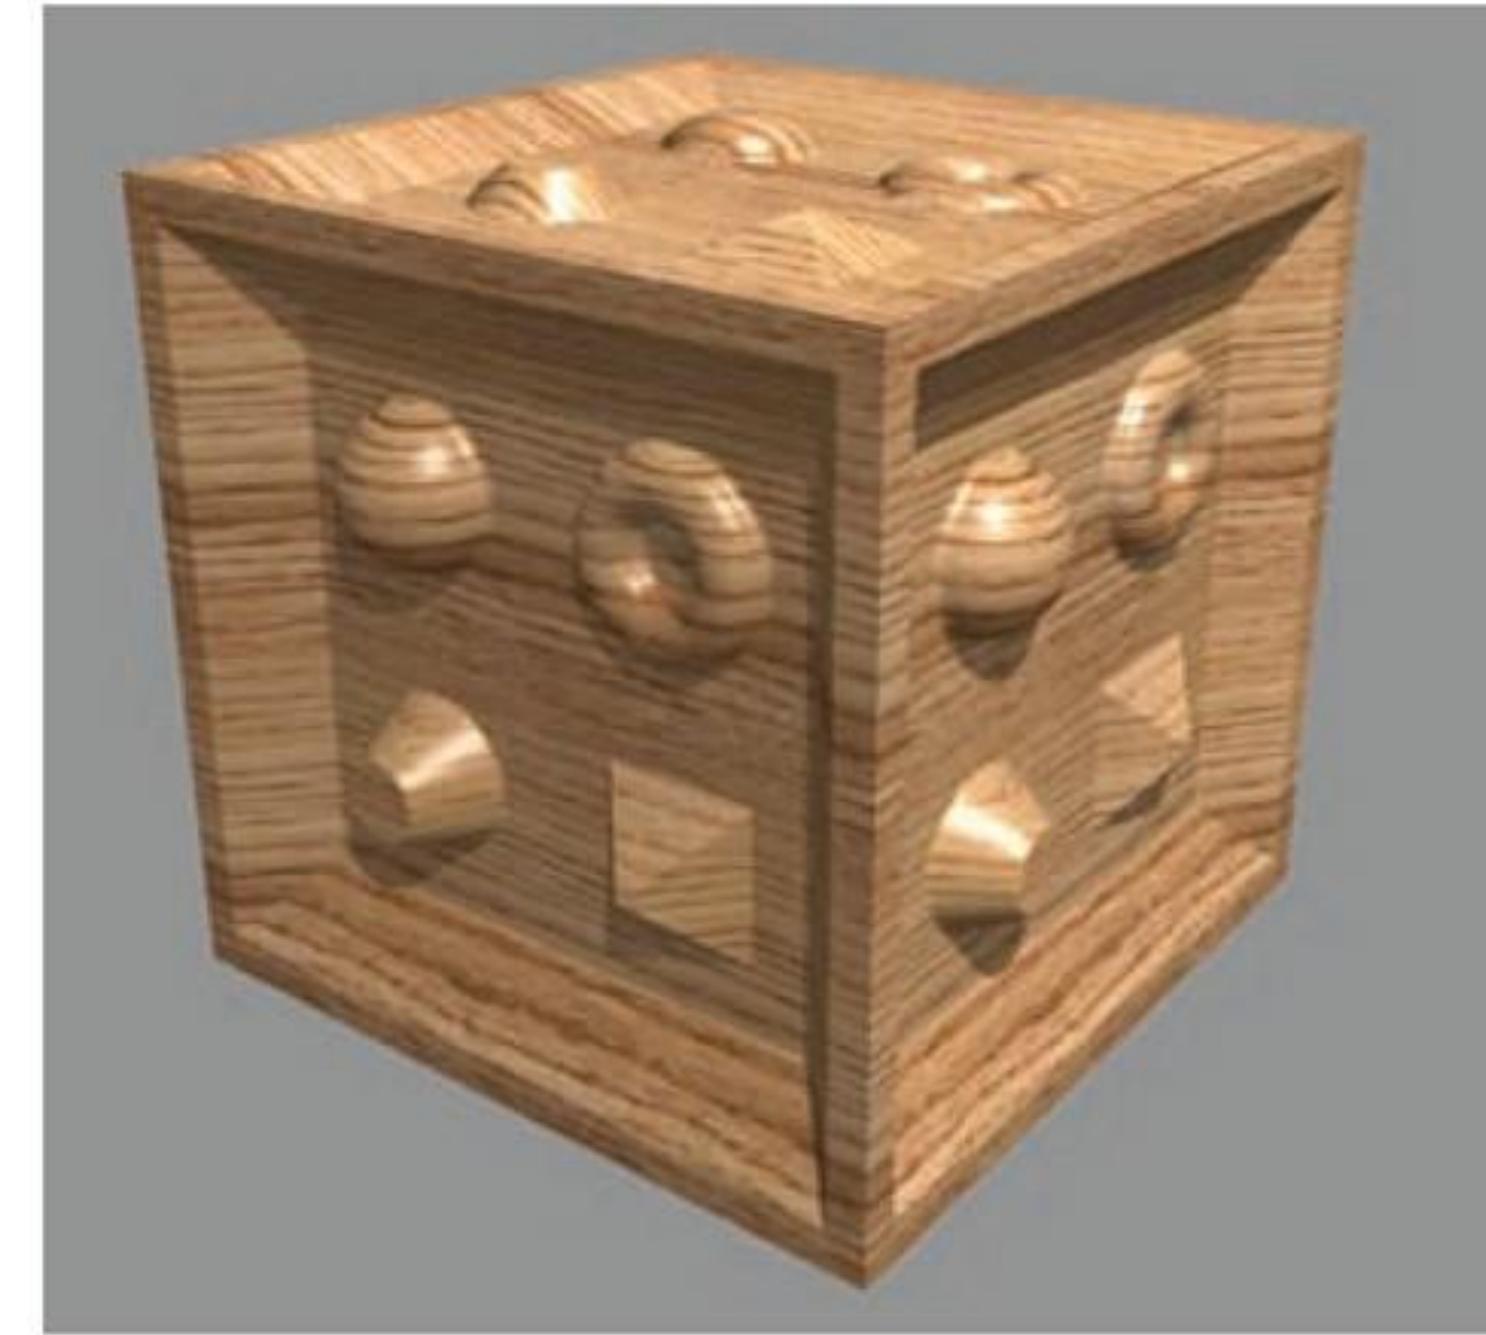


Sample along ray (green points)  
Lookup violet points (texture values)  
/\* Infer the black line shape \*/  
Compare green points with black points  
Find intersect between two conditions  
prev: green above black  
next: green below black





Parallax Mapping



Relief Mapping

<http://www.youtube.com/watch?v=5gorm90TXJM>

# Crysis, Crytek

